



JNAN VIKAS MANDAL'S

JVM'S DEGREE COLLEGE OF INFORMATION TECHNOLOGY, AIROLI,
NAAC RE-ACCREDITED Grade 'A+' (CGPA-3.31) NAVI MUMBAI – 400708

DEPARTMENT OF INFORMATION TECHNOLOGY

Practical Work Submitted in partial fulfilment of the
Requirements for the award of the Degree of
MASTER OF SCIENCE (INFORMATION TECHNOLOGY)
By
SONALI RUPCHAND THAKUR

Seat No: 1313626

ADVANCED AI (602)
MACHINE LEARNING (604)
SECURITY OPERATIONS CENTER (606b)

Under the esteemed guidance of

Mrs.Janhavi Kshirsagar

Mrs Bhagyashree Karnekar

Dr.Sanjivani Nalkar

ARTIFICIAL INTELLIGENCE



JNAN VIKAS MANDAL'S

Mohanlal Raichand Mehta College of Commerce
Diwali Maa College of Science
Amritlal Raichand Mehta College of Arts
Dr. R.T. Doshi College of Computer Science
NAAC Re-Accredited Grade 'A+' (CGPA : 3.31) (3rd Cycle)

DEPARTMENT OF INFORMATION TECHNOLOGY

CERTIFICATE

This is to certify that Sonali Rupchand Thakur bearing seat no. 1313626 has done the project work/journal work in the subject of Artificial intelligence of semester 3 Practical Examination during the academic year 2025-26 under the guidance of Mrs. Janhavi Kshirsagar being the partial requirement for the fulfilment of the curriculum of Degree in Master of Science in Information Technology under University of Mumbai.

Place: Airoli

Date:

Sign of Subject in- Charge

Sign of External Examiner

Sign of Coordinator

TABLE OF CONTENTS

Practical No	Practicals	Date	Sign
1.	Implementing advanced deep learning algorithms such as convolutional neural networks (CNNs) or recurrent neural networks (RNNs) using Python libraries like TensorFlow or PyTorch.		
2.	.Building a natural language processing (NLP) model for sentiment analysis or text classification.		
3.	Creating a chatbot using advanced techniques like transformer models.		
4.	Developing a recommendation system using collaborative filtering or deep learning approaches.		
5.	Implementing a computer vision project, such as object detection or image segmentation.		
6.	Training a generative adversarial network (GAN) for generating realistic images.		
7.	Applying reinforcement learning algorithms to solve complex decision-making problems.		
8.	Utilizing transfer learning to improve model performance on limited datasets.		
9.	Building a deep learning model for time series forecasting or anomaly detection.		
10.	Implementing a machine learning pipeline for automated feature engineering and model selection.		
11.	Using advanced optimization techniques like evolutionary algorithms or Bayesian optimization for hyperparameter tuning.		
12.	Use Python libraries such as GPT-2 or textgenrnn to train generative models on a corpus of text data and generate new text based on the patterns it has learned.		

Practical 1

Aim- Implementing advanced deep learning algorithms such as convolutional neural networks (CNNs) or recurrent neural networks (RNNs) using Python libraries like TensorFlow or PyTorch.

```
import tensorflow as tf
from tensorflow.keras import layers, models
(train_images, train_labels), (test_images, test_labels) = tf.keras.datasets.mnist.load_data()
train_images = train_images.reshape((train_images.shape[0], 28, 28, 1))
test_images = test_images.reshape((test_images.shape[0], 28, 28, 1))
train_images, test_images = train_images / 255.0, test_images / 255.0
model = models.Sequential([
    layers.Conv2D(32, (3, 3), activation='relu', input_shape=(28, 28, 1)),
    layers.MaxPooling2D((2, 2)), layers.Conv2D(64, (3, 3), activation='relu'),
    layers.MaxPooling2D((2, 2)), layers.Conv2D(64, (3, 3), activation='relu'), layers.Flatten(),
    layers.Dense(64, activation='relu'), layers.Dense(10, activation='softmax')])
model.compile(optimizer='adam', loss='sparse_categorical_crossentropy', metrics=['accuracy'])
model.fit(train_images, train_labels, epochs=5)
test_loss, test_acc = model.evaluate(test_images, test_labels)
print(f'Test accuracy: {test_acc}')
```

OUTPUT:



The screenshot shows a Jupyter Notebook terminal window with the title 'main'. The output displays the training progress of a neural network over 5 epochs. Each epoch shows 1875/1875 training steps taking 12s, with accuracy and loss values. The final test accuracy is printed at the end.

```
Epoch 1/5
1875/1875 12s 6ms/step - accuracy: 0.8990 - loss: 0.3412
Epoch 2/5
1875/1875 12s 6ms/step - accuracy: 0.9857 - loss: 0.0460
Epoch 3/5
1875/1875 12s 6ms/step - accuracy: 0.9895 - loss: 0.0327
Epoch 4/5
1875/1875 12s 6ms/step - accuracy: 0.9920 - loss: 0.0240
Epoch 5/5
1875/1875 12s 6ms/step - accuracy: 0.9940 - loss: 0.0183
313/313 1s 2ms/step - accuracy: 0.9877 - loss: 0.0473
Test accuracy: 0.9897000193595886

Process finished with exit code 0
```

Practical 2

Aim- Building a natural language processing (NLP) model for sentiment analysis or text classification.

```
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.feature_extraction.text import CountVectorizer
from sklearn.naive_bayes import MultinomialNB
from sklearn.metrics import accuracy_score, classification_report
data = { 'text': [
    'I love this product, it is amazing!',
    'This is the worst purchase I have ever made.',
    'I am very happy with the service.',
    'I will never buy from this store again.',
    'The quality is excellent!',
    'Terrible experience, very disappointed.',
    'Absolutely fantastic! Highly recommend it.',
    'Not good at all, I want a refund.',
    'Very satisfied with my order.',
    'The item broke after one use, awful.' ],
    'label': ['positive', 'negative', 'positive',
    'negative', 'positive', 'negative', 'positive',
    'negative', 'positive', 'negative']}
df = pd.DataFrame(data)
X_train, X_test, y_train, y_test = train_test_split( df['text'], df['label'], test_size=0.2,
random_state=42, stratify=df['label'])
vectorizer = CountVectorizer()
X_train_vectors = vectorizer.fit_transform(X_train)
X_test_vectors = vectorizer.transform(X_test)
model = MultinomialNB()
model.fit(X_train_vectors, y_train)
```

```

y_pred = model.predict(X_test_vectors)

accuracy = accuracy_score(y_test, y_pred)

print(f'Accuracy: {accuracy:.2f}')

print('Classification Report:')

print(classification_report(y_test, y_pred, zero_division=0))

custom_reviews = [
    "This product exceeded my expectations!",
    "Worst service ever, completely dissatisfied.",
    "The experience was okay, not great, not terrible.",
    "Highly recommend this to everyone!"

]custom_review_vectors = vectorizer.transform(custom_reviews)

custom_predictions = model.predict(custom_review_vectors)

for review, sentiment in zip(custom_reviews, custom_predictions):

    print(f"Review: \"{review}\" -> Sentiment: {sentiment}")

```

OUTPUT:

```

=====
RESTART: D:/Janu_mscIT_P2/AI-practs/(NLP) model.py =====
Accuracy: 1.00
Classification Report:
      precision    recall   f1-score   support
  negative      1.00      1.00      1.00       1
  positive      1.00      1.00      1.00       1

      accuracy           1.00       2
  macro avg      1.00      1.00      1.00       2
weighted avg      1.00      1.00      1.00       2

Review: "This product exceeded my expectations!" -> Sentiment: positive
Review: "Worst service ever, completely dissatisfied." -> Sentiment: negative
Review: "The experience was okay, not great, not terrible." -> Sentiment: negative
Review: "Highly recommend this to everyone!" -> Sentiment: positive

```

Practical 3

Aim- Creating a chatbot using advanced techniques like transformer models.

```
from transformers import GPT2LMHeadModel, GPT2Tokenizer, import torch, import os  
os.environ["HF_HOME"] = "./hf_cache"  
model_name = "microsoft/DialoGPT-medium"  
  
try:  
  
    model = GPT2LMHeadModel.from_pretrained(model_name)  
  
    tokenizer = GPT2Tokenizer.from_pretrained(model_name)  
  
except Exception as e:  
  
    print(f"Error loading model: {e}")  
  
    exit()  
  
def chat_with_bot(input_text):  
  
    try:  
  
        inputs = tokenizer.encode(input_text + tokenizer.eos_token, return_tensors="pt")  
  
        attention_mask = torch.ones(inputs.shape, dtype=torch.long) outputs = model.generate(  
inputs, max_length=100, pad_token_id=tokenizer.eos_token_id,  
attention_mask=attention_mask, no_repeat_ngram_size=2,  
temperature=0.7, top_p=0.9, top_k=50, do_sample=True, )  
  
        Decode the response back into text and filter it  
  
        response = tokenizer.decode(outputs[:, inputs.shape[-1]:][0], skip_special_tokens=True)  
  
        return filter_response(response)  
  
    except Exception as e:  
  
        return f"Error during response generation: {e}"  
  
def filter_response(response):  
  
    if len(response.strip()) < 5: # Adjusted to allow shorter valid responses  
  
        return "Could you say that again? I'm not sure I understood."  
  
    elif "?" in response: # If response is a question, it might be a valid conversational turn
```

```
return response return response

if __name__ == "__main__":
    print("Chatbot: Hello! Type 'exit' to end the conversation.")

    while True: user_input = input("You: ")

        if user_input.lower() == "exit":
            print("Chatbot: Goodbye!") break

        response = chat_with_bot(user_input)
        print(f"Chatbot: {response}")
```

OUTPUT:

```
Chatbot: Hello! Type 'exit' to end the conversation.
You: Hello, how are you?
Chatbot: Hi, I'm good thanks!
You: i am also fine
Chatbot: okay im in the room now
```

Practical 4

Aim- Developing a recommendation system using collaborative filtering or deep learning approaches.

```
import pandas as pd
import numpy as np
from sklearn.decomposition import TruncatedSVD
data = {
    "UserID": [1, 1, 2, 2, 3, 3, 4, 4, 5, 5],
    "MovieID": [101, 102, 101, 103, 102, 103, 101, 104, 103, 104],
    "Rating": [5, 4, 3, 4, 2, 5, 5, 3, 4, 2]}movies = {
    "MovieID": [101, 102, 103, 104],
    "Title": ["Dangal (2016)", "3 Idiots (2009)", "Chhichhore (2019)", "Lagaan (2001)"]}
ratings = pd.DataFrame(data)
movies = pd.DataFrame(movies)
ratings = pd.merge(ratings, movies, on="MovieID")
user_item_matrix = ratings.pivot_table(index="UserID", columns="Title",
values="Rating").fillna(0)
svd = TruncatedSVD(n_components=2) # Use a smaller number of components for
simplicity
user_factors = svd.fit_transform(user_item_matrix)
item_factors = svd.components_
predicted_ratings = np.dot(user_factors, item_factors)
min_rating, max_rating = predicted_ratings.min(), predicted_ratings.max()
predicted_ratings = 1 + 4 * (predicted_ratings - min_rating) / (max_rating - min_rating)
predicted_df = pd.DataFrame(predicted_ratings, index=user_item_matrix.index,
columns=user_item_matrix.columns)
predicted_df = predicted_df.round()
print("Predicted Ratings (Scaled):\n", predicted_df)
```

```

def recommend_movies(user_id, user_item_matrix, predicted_df, top_n=3, threshold=2):

    user_ratings = predicted_df.loc[user_id]

    rated_movies = user_item_matrix.loc[user_id] > 0

    recommendations = user_ratings[~rated_movies]

    print(f"Predicted ratings for User {user_id}:\n", user_ratings)

    recommendations = recommendations[recommendations >
threshold].sort_values(ascending=False)

    if recommendations.empty:

        print("No recommendations above threshold. Returning top N highest-rated movies.")

        recommendations = user_ratings.sort_values(ascending=False)

    return recommendations.head(top_n)

recommendations = recommend_movies(1, user_item_matrix, predicted_df, top_n=3,
threshold=2)

print("\nRecommended Movies for User 1 (High Ratings):\n", recommendations)

```

OUTPUT:

```

IDLE Shell 3.10.11
File Edit Shell Debug Options Window Help
Python 3.10.11 (tags/v3.10.11:7d4cc5a, Apr 5 2023, 00:38:17) [MSC v.1929 64 bit
(AMD64)] on win32
Type "help", "copyright", "credits" or "license()" for more information.
>>> ===== RESTART: C:/Users/Admin/Desktop/recommandation sys.py =====
Predicted Ratings (Scaled):
Title 3 Idiots (2009) Chhichhore (2019) Dangal (2016) Lagaan (2001)
UserID
1 2.0 1.0 5.0 2.0
2 2.0 4.0 3.0 2.0
3 2.0 5.0 1.0 1.0
4 2.0 1.0 5.0 2.0
5 1.0 4.0 1.0 1.0
Predicted ratings for User 1:
Title
3 Idiots (2009) 2.0
Chhichhore (2019) 1.0
Dangal (2016) 5.0
Lagaan (2001) 2.0
Name: 1, dtype: float64
No recommendations above threshold. Returning top N highest-rated movies.

Recommended Movies for User 1 (High Ratings):
Title
Dangal (2016) 5.0
3 Idiots (2009) 2.0
Lagaan (2001) 2.0
Name: 1, dtype: float64
>>> | Activate W
                                         Go to Settings


```

Practical 5

Aim- Implementing a computer vision project, such as object detection or image segmentation.

```
import cv2
from matplotlib import pyplot as plt
image_path = r"C:\Users\Janhavi\Downloads\stop.jpg"
cascade_path = r"C:\Users\Janhavi\Downloads\stop_data.xml"
img = cv2.imread(image_path)
if img is None:
    raise FileNotFoundError("Image file could not be loaded. Check the path.")
img_gray = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)
img_rgb = cv2.cvtColor(img, cv2.COLOR_BGR2RGB)
stop_data = cv2.CascadeClassifier(cascade_path)
if stop_data.empty():
    raise FileNotFoundError("Cascade file could not be loaded. Check the path.")
for (x, y, w, h) in stop_data.detectMultiScale(img_gray, minSize=(20, 20)):
    cv2.rectangle(img_rgb, (x, y), (x + h, y + w), (0, 255, 0), 5)
plt.imshow(img_rgb)
plt.show()
```

OUTPUT:



Practical 6

Aim- Training a generative adversarial network (GAN) for generating realistic images.

```
import torch

import torch.nn as nn

import torch.optim as optim

from torchvision import datasets, transforms

from torch.utils.data import DataLoader

import torchvision.utils as vutils

import matplotlib.pyplot as plt

class Generator(nn.Module):

    def __init__(self, input_size, output_size):

        super(Generator, self).__init__()

        self.model = nn.Sequential(nn.Linear(input_size, 128),

                                nn.ReLU(),nn.Linear(128, 256),nn.BatchNorm1d(256), nn.ReLU(),nn.Linear(256, 512), nn.BatchNorm1d(512), nn.ReLU(),nn.Linear(512, output_size), nn.Tanh()

                               )

    def forward(self, x):

        return self.model(x)

class Discriminator(nn.Module):

    def __init__(self, input_size):

        super(Discriminator, self).__init__()

        self.model = nn.Sequential( nn.Linear(input_size, 512), nn.LeakyReLU(0.2), nn.Linear(512, 256), nn.LeakyReLU(0.2), nn.Linear(256, 1), nn.Sigmoid()

                               )

    def forward(self, x):

        return self.model(x)

latent_size = 100 image_size = 28 * 28 batch_size = 64

epochs = 50 lr = 0.0002

transform = transforms.Compose([

    transforms.ToTensor(),

    transforms.Normalize(mean=[0.5], std=[0.5])])

dataset = datasets.MNIST(root=".data", train=True, transform=transform, download=True)
```

```

data_loader = DataLoader(dataset, batch_size=batch_size, shuffle=True)

generator = Generator(input_size=latent_size, output_size=image_size)

discriminator = Discriminator(input_size=image_size)

device = torch.device("cuda" if torch.cuda.is_available() else "cpu")

generator.to(device) discriminator.to(device)

criterion = nn.BCELoss()

optimizer_g = optim.Adam(generator.parameters(), lr=lr)

optimizer_d = optim.Adam(discriminator.parameters(), lr=lr)

for epoch in range(epochs):

    for batch_idx, (real_images, _) in enumerate(data_loader):

        real_images = real_images.view(-1, image_size).to(device) batch_size =
        real_images.size(0) real_labels = torch.ones(batch_size, 1).to(device) fake_labels =
        torch.zeros(batch_size, 1).to(device) outputs = discriminator(real_images) d_loss_real =
        criterion(outputs, real_labels) z = torch.randn(batch_size, latent_size).to(device) fake_images =
        generator(z) outputs = discriminator(fake_images.detach()) d_loss_fake = criterion(outputs,
        fake_labels) d_loss = d_loss_real + d_loss_fake
        optimizer_d.zero_grad() d_loss.backward() optimizer_d.step()

        z = torch.randn(batch_size, latent_size).to(device) fake_images = generator(z)

        outputs = discriminator(fake_images) g_loss = criterion(outputs, real_labels)

        optimizer_g.zero_grad() g_loss.backward() optimizer_g.step()

        print(f"Epoch [{epoch + 1}/{epochs}] D Loss: {d_loss.item():.4f} G Loss:
        {g_loss.item():.4f}")

    if (epoch + 1) % 10 == 0:

        with torch.no_grad():

            z = torch.randn(64, latent_size).to(device)

            fake_images = generator(z).view(-1, 1, 28, 28).cpu()

            print(f"Epoch {epoch + 1}: Generated Images Mean Pixel Value =
            {fake_images.mean().item():.4f}")

            grid = vutils.make_grid(fake_images, nrow=8, normalize=True)


            plt.figure(figsize=(8, 8))

```

```
plt.imshow(grid.permute(1, 2, 0).numpy(), cmap="gray")
plt.title(f"Generated Images at Epoch {epoch + 1}")
plt.axis("off")
filename = f"generated_images_epoch_{epoch + 1}.png"
plt.savefig(filename)
plt.close()

print(f"Images saved to {filename}")
```

OUTPUT:



Practical 7

Aim- Applying reinforcement learning algorithms to solve complex decision-making problems.


```
import numpy as np
import matplotlib.pyplot as plt
class GridWorld:
    def __init__(self, size, start, goal, obstacles):
        self.size = size
        self.start = start
        self.goal = goal
        self.obstacles = obstacles
        self.reset()
    def reset(self):
        self.state = self.start
        return self.state
    def step(self, action):
        x, y = self.state
        if action == 0: # Up
            next_state = (max(x - 1, 0), y)
        elif action == 1: # Down
            next_state = (min(x + 1, self.size - 1), y)
        elif action == 2: # Left
            next_state = (x, max(y - 1, 0))
        elif action == 3: # Right
            next_state = (x, min(y + 1, self.size - 1))
        if next_state in self.obstacles:
            next_state = self.state
            reward = -1
            done = False
        if next_state == self.goal:
            reward = 10
            done = True
        self.state = next_state
        return next_state, reward, done
start = (0, 0)
goal = (4, 4)
obstacles = [(2, 2), (3, 2)]
env = GridWorld(size, start, goal, obstacles)
```

```

actions = 4 q_table = np.zeros((size, size, actions)) episodes = 1000 alpha = 0.1
gamma = 0.9 epsilon = 1.0 epsilon_decay = 0.995 min_epsilon = 0.1
rewards = []
for episode in range(episodes): state = env.reset() total_reward = 0 while not done:
    x, y = state if np.random.uniform(0, 1) < epsilon: action = np.random.choice(actions) # Explore
    else: action = np.argmax(q_table[x, y])
    next_state, reward, done = env.step(action) next_x, next_y = next_state
    plt.plot(rewards) plt.xlabel('Episode') plt.ylabel('Total Reward')
    plt.title('Reward Progression') plt.show()
print("Learned Q-values:")
for i in range(size):
    for j in range(size):
        print(f"({i}, {j}): {q_table[i, j]}")

```

OUTPUT:



Practical 8

Aim- Utilizing transfer learning to improve model performance on limited datasets.

```
import tensorflow as tf

from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, GlobalAveragePooling2D, Dropout
from tensorflow.keras.preprocessing.image import ImageDataGenerator
from PIL import Image
import matplotlib.pyplot as plt

train_dir = r'C:\Users\Janhavi\Downloads\cats_dogs_light\train'
validation_dir = r'C:\Users\Janhavi\Downloads\cats_dogs_light\test'

print(f"Train directory exists: {os.path.isdir(train_dir)}")
print(f"Validation directory exists: {os.path.isdir(validation_dir)}")
print(f"Found {len(os.listdir(train_dir))} subdirectories in the train directory.")
print(f"Found {len(os.listdir(validation_dir))} subdirectories in the validation directory.")

def check_images(directory):
    print(f"Checking files in {directory}...")
    for subdir in os.listdir(directory):
        subdir_path = os.path.join(directory, subdir)
        if os.path.isdir(subdir_path):
            for file_name in os.listdir(subdir_path):
                file_path = os.path.join(subdir_path, file_name)

for layer in base_model.layers[:-10]: # Unfreeze all but the last 10 layers
    layer.trainable = False
model = Sequential([
    base_model,
    GlobalAveragePooling2D(),
    Dense(64, activation='relu'), # Reduced neurons for simplicity
    Dropout(0.5), # Dropout to prevent overfitting
    Dense(1, activation='sigmoid') # Binary classification
])
model.compile(
```

```

optimizer=tf.keras.optimizers.Adam(learning_rate=1e-5), # Use lower learning rate
loss='binary_crossentropy', metrics=['accuracy'])

train_generator = train_datagen.flow_from_directory(train_dir, target_size=(224, 224),
batch_size=32, class_mode='binary' # Binary classification (cats vs dogs)
history = model.fit(
    train_generator, validation_data=validation_generator, epochs=30,
    callbacks=[early_stopping, lr_scheduler]
plt.figure(figsize=(12, 4))


plt.subplot(1, 2, 1)
plt.plot(history.history['accuracy'], label='Training Accuracy')
plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
plt.ylabel('Accuracy')
plt.legend()
plt.title('Accuracy vs Epochs')

plt.subplot(1, 2, 2)
plt.plot(history.history['loss'], label='Training Loss')
plt.plot(history.history['val_loss'], label='Validation Loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()
plt.title('Loss vs Epochs')

plt.show()

```

OUTPUT:




Practical 9

Aim- Building a deep learning model for time series forecasting or anomaly detection.

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import LSTM, Dense, Dropout
from sklearn.preprocessing import MinMaxScaler
from sklearn.model_selection import train_test_split
def generate_sine_wave(length):
    x = np.linspace(0, 50, length) y = np.sin(x) return y
def create_dataset(data, look_back):
    X, y = [], [] for i in range(len(data) - look_back):X.append(data[i:i + look_back])
    y.append(data[i + look_back]) return np.array(X), np.array(y)
TIME_STEPS = 50 # Look-back window size EPOCHS = 20 BATCH_SIZE = 32
data = generate_sine_wave(1000) scaler = MinMaxScaler(feature_range=(0, 1))
scaled_data = scaler.fit_transform(data.reshape(-1, 1)).flatten()
X, y = create_dataset(scaled_data, TIME_STEPS)
X = np.expand_dims(X, axis=-1) # Add a channel dimension for LSTM
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
model = Sequential([LSTM(50, return_sequences=True, input_shape=(TIME_STEPS, 1)),
Dropout(0.2), LSTM(50), Dropout(0.2), Dense(1)])model.compile(optimizer='adam',
loss='mean_squared_error')
history = model.fit(X_train, y_train, epochs=EPOCHS, batch_size=BATCH_SIZE,
validation_data=(X_test, y_test))
train_loss = model.evaluate(X_train, y_train, verbose=0)
test_loss = model.evaluate(X_test, y_test, verbose=0)
print(f"Train Loss: {train_loss:.4f}, Test Loss: {test_loss:.4f}")
```

```
predicted = model.predict(X_test)
predicted = scaler.inverse_transform(predicted)
y_test_actual = scaler.inverse_transform(y_test.reshape(-1, 1))
plt.figure(figsize=(10, 6))
plt.plot(y_test_actual, label='Actual')
plt.plot(predicted, label='Predicted')
plt.legend()
plt.title("Time Series Forecasting with LSTM")
plt.show()
```

OUTPUT:



Practical 10

Aim- Implementing a machine learning pipeline for automated feature engineering and model selection.

```
import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split, cross_val_score
from sklearn.preprocessing import StandardScaler
from sklearn.ensemble import RandomForestClassifier
import xgboost as xgb
import lightgbm as lgb
import matplotlib.pyplot as plt
url = "https://raw.githubusercontent.com/datasciencedojo/datasets/master/titanic.csv"
data = pd.read_csv(url)
data = data.drop(columns=['Name', 'Ticket', 'Cabin']) # Drop irrelevant columns
data['Sex'] = data['Sex'].map({'male': 0, 'female': 1}) # Encode 'Sex' feature
data = pd.get_dummies(data, drop_first=True) # One-hot encoding for 'Embarked'
X = data.drop(columns='Survived')
y = data['Survived']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)
numeric_features = X.select_dtypes(include=['int64', 'float64']).columns
numeric_transformer = Pipeline(steps=[
    ('imputer', SimpleImputer(strategy='mean')), # Handle missing values
    ('scaler', StandardScaler()) # Feature scaling])
categorical_features = X.select_dtypes(include=['object']).columns
('encoder', 'passthrough') # Use 'passthrough' to handle categorical features in
pd.get_dummies outside the pipeline
])preprocessor = ColumnTransformer( transformers=[('num', numeric_transformer,
numeric_features), ('cat', categorical_transformer, categorical_features)])
```

```

models = [
    ('Logistic Regression', LogisticRegression(max_iter=1000)),
    ('Random Forest', RandomForestClassifier()),
    ('SVM', SVC()), ('XGBoost', xgb.XGBClassifier()), ('LightGBM',
    lgb.LGBMClassifier())]

cv_score = cross_val_score(pipeline, X_train, y_train, cv=5, scoring='accuracy')

mean_score = np.mean(cv_score)

if hasattr(best_pipeline.named_steps['classifier'], 'feature_importances_'):

    importances = best_pipeline.named_steps['classifier'].feature_importances_

    indices = np.argsort(importances)[::-1]

    plt.figure(figsize=(10, 6))

    plt.title(f"Feature Importances for {best_model}")

    plt.bar(range(X_train.shape[1]), importances[indices], align='center')

    plt.xticks(range(X_train.shape[1]), X_train.columns[indices], rotation=90)

    plt.show()

```

OUTPUT:

```

Training Logistic Regression...
Logistic Regression Mean Cross-Validation Score: 0.7913
Training Random Forest...
Random Forest Mean Cross-Validation Score: 0.8058
Training SVM...
SVM Mean Cross-Validation Score: 0.8170
Training XGBoost...
XGBoost Mean Cross-Validation Score: 0.7978
Training LightGBM...
[LightGBM] [Info] Number of positive: 184, number of negative: 314
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing was 0.000175 seconds.

```

Practical 11

Aim- Using advanced optimization techniques like evolutionary algorithms or Bayesian optimization for hyperparameter tuning.

```
import os  
  
os.environ["OPTUNA_NO_COLOR"] = "1" # Disable colored output  
  
import optuna  
  
from sklearn.datasets import load_iris  
  
from sklearn.ensemble import RandomForestClassifier  
  
from sklearn.model_selection import cross_val_score  
  
from sklearn.model_selection import train_test_split  
  
data = load_iris()  
  
X, y = data.data, data.target  
  
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)  
  
def objective(trial):  
  
    # Define hyperparameter search space  
  
    n_estimators = trial.suggest_int("n_estimators", 10, 200)  
  
    max_depth = trial.suggest_int("max_depth", 2, 20)  
  
    min_samples_split = trial.suggest_int("min_samples_split", 2, 20)  
  
    min_samples_leaf = trial.suggest_int("min_samples_leaf", 1, 10)  
  
    clf = RandomForestClassifier(  
  
        n_estimators=n_estimators,  
  
        max_depth=max_depth,  
  
        min_samples_split=min_samples_split,  
  
        min_samples_leaf=min_samples_leaf,  
  
        random_state=42, )  
  
    score = cross_val_score(clf, X_train, y_train, cv=5, scoring="accuracy").mean()  
  
    return score
```

```
study = optuna.create_study(direction="maximize")
study.optimize(objective, n_trials=50)
print("Best hyperparameters:", study.best_params)
best_params = study.best_params
best_model = RandomForestClassifier(**best_params, random_state=42)
best_model.fit(X_train, y_train)
accuracy = best_model.score(X_test, y_test)
print(f"Accuracy with best hyperparameters: {accuracy:.2f}")
```

OUTPUT:

```
meters: {'n_estimators': 154, 'max_depth': 6, 'min_samples_split': 5, 'min_samples_leaf': 3}. Best is trial 1 with value: 0.9583333333333334. [0m
[32m[I 2024-11-22 19:24:57,858] [0m Trial 49 finished with value: 0.9416666666666667 and parameters: {'n_estimators': 144, 'max_depth': 9, 'min_samples_split': 6, 'min_samples_leaf': 4}. Best is trial 1 with value: 0.9583333333333334. [0m
Best hyperparameters: {'n_estimators': 92, 'max_depth': 20, 'min_samples_split': 15, 'min_samples_leaf': 3}
Accuracy with best hyperparameters: 1.00
>>>
```

Practical 12

Aim- Use Python libraries such as GPT-2 or textgenrnn to train generative models on a corpus of text data and generate new text based on the patterns it has learned.

```
from transformers import GPT2LMHeadModel, GPT2Tokenizer  
  
model_name = "gpt2" # You can change to a larger version like 'gpt2-medium' if needed  
  
tokenizer = GPT2Tokenizer.from_pretrained(model_name)  
  
model = GPT2LMHeadModel.from_pretrained(model_name)  
  
tokenizer.pad_token = tokenizer.eos_token  
  
prompt = "Once upon a time in a faraway land"  
  
inputs = tokenizer.encode(prompt, return_tensors="pt")  
  
outputs = model.generate(inputs, max_length=50,  
    num_return_sequences=5, do_sample=True, top_k=50, top_p=0.95,  
    attention_mask=None, pad_token_id=tokenizer.eos_token_id) # Set pad_token_id to  
eos_token_id  
  
print("\nGenerated Texts:") for i, output in enumerate(outputs, start=1): print(f"{i}:\n{tokenizer.decode(output, skip_special_tokens=True)}")
```

OUTPUT:

```
Generated Texts:  
1: Once upon a time in a faraway land, when in my travels the wind and fire are  
both so keen they break down to pieces, yet in a vast place a stone is broken, t  
hough it cannot break. When in a faraway land,  
2: Once upon a time in a faraway land there is a mighty earthquake, and all that  
can be seen is that there is an enormous and massive earth that is completely c  
overed with molten lava, and there is still a great amount of energy left over,  
3: Once upon a time in a faraway land in the South, he stood on a ledge on which  
his body was resting. He saw in its depths a woman with no name or name of any  
kind with his wife. And he took her out to  
4: Once upon a time in a faraway land, a spirit spoke to the inhabitants and tol  
d them to stop. The inhabitants thought they had been tricked by the spirit and  
were prepared to fight it for their lives with the hope of having the spirit out  
of  
5: Once upon a time in a faraway land, there existed a civilization of humans an  
d elves. One of their basic needs was to produce food, so that humans could get  
by without relying on agriculture for their livelihood. Many of the human races  
of the  
>>> |
```

MACHINE LEARNING



JNAN VIKAS MANDAL'S

Mohanlal Raichand Mehta College of Commerce
Diwali Maa College of Science
Amritlal Raichand Mehta College of Arts
Dr. R.T. Doshi College of Computer Science
NAAC Re-Accredited Grade 'A+' (CGPA : 3.31) (3rd Cycle)

DEPARTMENT OF INFORMATION TECHNOLOGY

CERTIFICATE

This is to certify that Sonali Rupchand Thakur bearing seat no. 1313626 has done the project work/journal work in the subject of Machine Learning of semester III Practical Examination during the academic year 2025-26 under the guidance of **Mrs Bhagyashree Karnekar** being the partial requirement for the fulfilment of the curriculum of Degree in Master of Science in Information Technology under University of Mumbai

Place: Airoli

Date:

Sign of Subject in- Charge

Sign of External Examiner

Sign of Coordinator

Index

Sr.No	Practical	Date	Sign
1.	a)Load a CSV dataset. Handle missing values, inconsistent formatting, and outliers. b) Load a dataset, calculate descriptive summary statistics, create visualizations using different graphs, and identify potential features and target variables c) Create or Explore datasets to use all pre-processing routines like label encoding, scaling, and binarization.		
2.	Implement and demonstrate the Find-S algorithm for finding the most specific hypothesis based on a given set of training data samples. Read the training data from a .CSV file and generate the final specific hypothesis. (Create your dataset)		
3.	a)Fit a linear regression model on a dataset. Interpret coefficients, make predictions, and evaluate performance using metrics like R-squared and MSE b) Extend linear regression to multiple features. Handle feature selection and potential multicollinearity. c)Implement regression variants like LASSO and Ridge on any generated dataset.		
4.	A) Perform binary classification using logistic regression. Calculate accuracy, precision, recall, and understand the ROC curve. B) Implement and demonstrate k-nearest Neighbor algorithm. Read the training data from a .CSV file and build the model to classify a test sample. Print both correct and wrong predictions. C)Build a decision tree classifier or regressor. Control hyperparameters like tree depth to avoid overfitting. Visualize the tree. D)Implement a Support Vector Machine for any relevant dataset. E) Train a random forest ensemble. Experiment with the number of trees and feature sampling. Compare performance to a single decision tree. F)Implement a gradient boosting machine (e.g., XGBoost). Tune hyperparameters and explore feature importance.		
5.	A)Implement and demonstrate the working of a Naive Bayesian classifier using a sample data set. Build the model to classify a test sample B) Implement Hidden Markov Models using hmmlearn		
6.	A) Implement Bayesian Linear Regression to explore prior and posterior distribution B) Implement Gaussian Mixture Models for density estimation and unsupervised clustering		
7.	A) Implement cross-validation techniques (k-fold, stratified, etc.) for robust model Evaluation B) Systematically explore combinations of hyperparameters to optimize model performance. (Use grid and randomized search)		
8.	Implement Bayesian Learning using inferences		

PRACTICAL NO 1

Aim-Load a CSV dataset. Handle missing values, inconsistent formatting, and outliers.

```
import pandas as pd
import numpy as np
# Step 1: Load the dataset
file_path = 'C:/Users/admin/Downloads/enjoysport.csv'
df = pd.read_csv(file_path)
print("Original Dataset:\n", df.head())
# Step 2: Handle inconsistent formatting for string columns only
for col in df.select_dtypes(include=['object', 'string']):
    df[col] = df[col].map(lambda x: x.strip().lower() if isinstance(x, str) else x)
for col in df.columns:
    if df[col].dtype == 'object' or df[col].dtype.name == 'string':
        df[col] = df[col].fillna(df[col].mode(dropna=True)[0]) # Fill categorical with mode
    else:
        df[col] = df[col].fillna(df[col].mean())
print("\nCleaned Dataset:\n", df.head())
```

OUTPUT

```
(base) C:\Users\admin>python Desktop/MLPRAC/prac2.py
Original Dataset:
    sky airtemp humidity     wind water forcast enjoysport
0  sunny     warm   normal  strong  warm    same      yes
1  sunny     warm     high  strong  warm    same      yes
2  rainy    cold     high  strong  warm  change      no
3  sunny     warm     high  strong  cool  change      yes




Cleaned Dataset:
    sky airtemp humidity     wind water forcast enjoysport
0  sunny     warm   normal  strong  warm    same      yes
1  sunny     warm     high  strong  warm    same      yes
2  rainy    cold     high  strong  warm  change      no
3  sunny     warm     high  strong  cool  change      yes
```

B)Load a dataset, calculate descriptive summary statistics, create visualizations using different graphs, and identify potential features and target variables
Note: Explore Univariate and Bivariate graphs (Matplotlib) and Seaborn for visualization.

```
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.datasets import load_iris
iris = load_iris()
df = pd.DataFrame(iris.data, columns=iris.feature_names)
df['species'] = pd.Categorical.from_codes(iris.target, iris.target_names)
df.hist(figsize=(10, 6), edgecolor='black')
plt.suptitle("Univariate Histograms (Matplotlib)", fontsize=14)
plt.tight_layout()
plt.show()
plt.figure(figsize=(10, 6))
sns.boxplot(data=df.drop(columns='species'))
plt.title("Univariate Boxplots (Seaborn)", fontsize=14)
plt.xticks(rotation=45)
plt.show()
sns.pairplot(df, hue="species", corner=True, diag_kind='hist')
plt.suptitle("Bivariate Scatter Plots (Pairplot)", y=1.02, fontsize=14)
plt.show()
plt.figure(figsize=(8, 5))
sns.heatmap(df.drop('species', axis=1).corr(), annot=True, cmap="coolwarm", fmt=".2f")
plt.title("Feature Correlation Heatmap", fontsize=14)
plt.show()
features = df.columns[:-1]
target = 'species'
print("\nFeatures:", list(features))
print("Target:", target)
```

OUTPUT

```
(base) C:\Users\admin>python Desktop\MLPRAC\prac1b.py
Descriptive Statistics:
   sepal length (cm)  sepal width (cm)  petal length (cm)  petal width (cm)
count      150.000000      150.000000      150.000000      150.000000
mean       5.843333        3.057333       3.758000       1.199333
std        0.828066        0.435866       1.765298       0.762238
min        4.300000        2.000000       1.000000       0.100000
25%        5.100000        2.800000       1.600000       0.300000
50%        5.800000        3.000000       4.350000       1.300000
75%        6.400000        3.300000       5.100000       1.800000
max        7.900000        4.400000       6.900000       2.500000
```



C)Create or Explore datasets to use all pre-processing routines like label encoding, scaling, and binarization.

```
import pandas as pd
import numpy as np
from sklearn.preprocessing import LabelEncoder, OneHotEncoder, MinMaxScaler, StandardScaler, Binarizer
data = {
    'City': ['Mumbai', 'Delhi', 'Chennai', 'Delhi', 'Mumbai'],
    'Temperature': [33, 45, 30, 40, 35],
    'Humidity': [70, 60, 80, 75, 72],
    'Rainy': ['Yes', 'No', 'Yes', 'No', 'Yes']
}

df = pd.DataFrame(data)
print("Original Data:\n", df)
le = LabelEncoder()
df['City_Label'] = le.fit_transform(df['City'])
df['Rainy_Label'] = le.fit_transform(df['Rainy'])
ohe = OneHotEncoder(sparse_output=False)
city_encoded = ohe.fit_transform(df[['City']])
city_df = pd.DataFrame(city_encoded, columns=ohe.get_feature_names_out(['City']))
df = pd.concat([df, city_df], axis=1)
scaler_minmax = MinMaxScaler()
df[['Temperature_MinMax', 'Humidity_MinMax']] = scaler_minmax.fit_transform(df[['Temperature', 'Humidity']])
scaler_std = StandardScaler()
df[['Temperature_Std', 'Humidity_Std']] = scaler_std.fit_transform(df[['Temperature', 'Humidity']])
binarizer = Binarizer(threshold=70)
df['Humidity_Binarized'] = binarizer.fit_transform(df[['Humidity']])
print("\nProcessed Data:\n", df)
```

```
(base) C:\Users\admin>python Desktop\MLPAC/practic.py
Original Data:
   City  Temperature  Humidity Rainy
0  Mumbai        33       70    Yes
1    Delhi        45       68    No
2  Chennai        30       80   Yes
3    Delhi        40       75    No
4  Mumbai        35       72   Yes

Processed Data:
   City  Temperature  Humidity Rainy ...  Humidity_MinMax  Temperature_Std  Humidity_Std  Humidity_Binarized
0  Mumbai        33       70    Yes ...      0.50      -0.077439      -0.211443          0
1    Delhi        45       68    No ...      0.00      1.580691      -1.721748          0
2  Chennai        30       80   Yes ...      1.00     -1.241971      1.298863          1
3    Delhi        40       75    No ...      0.75      0.639803      0.543710          1
4  Mumbai        35       72   Yes ...      0.60      -0.381084      0.890618          1

[5 rows x 14 columns]
```

Practical no 2

Implement and demonstrate the FIND-S algorithm for finding the most specific hypothesis based on a given set of training data samples. Read the training data from a CSV file and generate the final specific hypothesis. (Create your dataset)

```
import pandas as pd
data = {
    'Sky': ['Sunny', 'Sunny', 'Rainy', 'Sunny', 'Sunny'],
    'AirTemp': ['Warm', 'Warm', 'Cold', 'Warm', 'Warm'],
    'Humidity': ['Normal', 'High', 'High', 'High', 'Normal'],
    'Wind': ['Strong', 'Strong', 'Strong', 'Strong', 'Strong'],
    'Water': ['Warm', 'Warm', 'Warm', 'Cool', 'Warm'],
    'Forecast': ['Same', 'Same', 'Change', 'Change', 'Same'],
    'EnjoySport': ['Yes', 'Yes', 'No', 'Yes', 'Yes']
}
df = pd.DataFrame(data)
df.to_csv("C:/Users/admin/Downloads/enjoysport.csv", index=False)
# Step 2: Load dataset
dataset = pd.read_csv("C:/Users/admin/Downloads/enjoysport.csv")
# Step 3: FIND-S algorithm
def find_s_algorithm(dataframe):
    # Extract only positive examples
    positive_examples = dataframe[dataframe['EnjoySport'] == 'Yes'].iloc[:, :-1].values
    hypothesis = positive_examples[0].copy()
    for example in positive_examples[1:]:
        for i in range(len(hypothesis)):
            if hypothesis[i] != example[i]:
                hypothesis[i] = '?'
    return hypothesis
final_hypothesis = find_s_algorithm(dataset)
print("Final Specific Hypothesis:", final_hypothesis)
```

OUTPUT

```
(base) C:\Users\admin>python Desktop/MLPRAC/prac2.py
Final Specific Hypothesis: ['Sunny' 'Warm' '?' 'Strong' '?' '?']
```

Practical no 3

A)Fit a linear regression model on a dataset. Interpret coefficients, make predictions, and evaluate performance using metrics like R-squared and MSE

```
import pandas as pd
import numpy as np
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score
from sklearn.model_selection import train_test_split
data = pd.read_csv("C:/Users/admin/Downloads/winequality_red.csv")
print(data.head())
data = data.dropna()
X = data[['fixed acidity', 'volatile acidity', 'citric acid']] # update as needed
y = data['quality']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
model = LinearRegression()
model.fit(X_train, y_train)
print("Intercept:", model.intercept_)
print("Coefficients:")
for feature, coef in zip(X.columns, model.coef_):
    print(f"{feature}: {coef}")
y_pred = model.predict(X_test)
mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)
print("Mean Squared Error (MSE):", mse)
print("R-squared (R2):", r2)
```

```
(base) C:\Users\admin>python Desktop\ML\prac3.py
   fixed acidity  volatile acidity  citric acid  residual sugar  chlorides ...  density  pH  sulphates  alcohol  quality
0            7.4            0.70          0.00         1.9        0.076 ...      0.9978  3.51     0.56     9.4      5
1            7.8            0.88          0.00         2.6        0.098 ...      0.9968  3.20     0.68     9.8      5
2            7.8            0.76          0.04         2.3        0.092 ...      0.9970  3.26     0.65     9.8      5
3           11.2            0.28          0.56         1.9        0.075 ...      0.9980  3.16     0.58     9.8      6
4            7.4            0.70          0.00         1.9        0.076 ...      0.9978  3.51     0.56     9.4      5

[5 rows x 12 columns]
Intercept: 6.378972118987667
Coefficients:
fixed acidity: 0.02053646472220784
volatile acidity: -1.7035172962776297
citric acid: -0.08136606250766007
Mean Squared Error (MSE): 0.533975815696693
R-squared (R2): 0.18290584845348323
```

B) Extend linear regression to multiple features. Handle feature selection and potential multicollinearity.

```
from sklearn.linear_model import LinearRegression, Lasso, Ridge
from sklearn.feature_selection import RFE
from sklearn.preprocessing import StandardScaler
from statsmodels.stats.outliers_influence import variance_inflation_factor
import pandas as pd
import numpy as np
df = pd.read_csv("C:/Users/admin/Downloads/winequality_red.csv")
X = df.drop(columns='quality') # target is 'quality'
y = df['quality']
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)
vif_data = pd.DataFrame()
vif_data['Feature'] = X.columns
vif_data['VIF'] = [variance_inflation_factor(X_scaled, i) for i in range(X.shape[1])]
print(vif_data)
model = LinearRegression()
rfe = RFE(model, n_features_to_select=5)
X_rfe = rfe.fit_transform(X_scaled, y)
selected_features = X.columns[rfe.support_]
print("Selected features:", list(selected_features))
ridge = Ridge(alpha=1.0)
ridge.fit(X_scaled, y)
```

Output

```
(base) C:\Users\admin>python Desktop\ML\prac3b.py
      Feature      VIF
0    fixed acidity  7.767512
1 volatile acidity  1.789390
2        citric acid  3.128022
3   residual sugar  1.702588
4       chlorides  1.481932
5  free sulfur dioxide  1.963019
6  total sulfur dioxide  2.186813
7          density  6.343760
8            pH  3.329732
9       sulphates  1.429434
10      alcohol  3.031160
Selected features: ['volatile acidity', 'chlorides', 'total sulfur dioxide', 'sulphates', 'alcohol']
```


C)Implement regression variants like LASSO and Ridge on any generated dataset.

```
import numpy as np
import matplotlib.pyplot as plt
from sklearn.datasets import make_regression
from sklearn.linear_model import LinearRegression, Lasso, Ridge
from sklearn.metrics import mean_squared_error
from sklearn.model_selection import train_test_split

X, y, coef = make_regression(
    n_samples=100,
    n_features=10,
    n_informative=5,
    noise=10,
    coef=True,
    random_state=42
)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
lr = LinearRegression()
lasso = Lasso(alpha=1.0)
ridge = Ridge(alpha=1.0)
lr.fit(X_train, y_train)
lasso.fit(X_train, y_train)
ridge.fit(X_train, y_train)
y_pred_lr = lr.predict(X_test)
y_pred_lasso = lasso.predict(X_test)
y_pred_ridge = ridge.predict(X_test)
print("Linear Regression MSE:", mean_squared_error(y_test, y_pred_lr))
print("Lasso Regression MSE:", mean_squared_error(y_test, y_pred_lasso))
print("Ridge Regression MSE:", mean_squared_error(y_test, y_pred_ridge))
plt.figure(figsize=(10,6))
plt.plot(lr.coef_, label='Linear', marker='o')
plt.plot(lasso.coef_, label='Lasso', marker='x')
plt.plot(ridge.coef_, label='Ridge', marker='s')
plt.title("Comparison of Coefficients")
plt.xlabel("Feature Index")
plt.ylabel("Coefficient Value")
plt.legend()
plt.grid(True)
plt.show()
```

Output

```
(base) C:\Users\admin>python Desktop\ML\prac3c.py
Linear Regression MSE: 126.39524754125014
Lasso Regression MSE: 154.83952646022993
Ridge Regression MSE: 141.31120885254484
```



Practical no 4

A) Perform binary classification using logistic regression. Calculate accuracy, precision, recall, and understand the ROC curve.

```
import numpy as np
import matplotlib.pyplot as plt
from sklearn.datasets import make_classification
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import (
    accuracy_score, precision_score, recall_score,
    roc_curve, roc_auc_score, confusion_matrix
)
from sklearn.model_selection import train_test_split
# Create binary classification dataset
X, y = make_classification(
    n_samples=1000,
    n_features=10,
    n_informative=5,
    n_redundant=2,
    n_classes=2,
    random_state=42
)
# Split into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42
)
model = LogisticRegression()
model.fit(X_train, y_train)

# Predict probabilities and classes
y_pred = model.predict(X_test)
y_prob = model.predict_proba(X_test)[:, 1] # probability of class 1
# Accuracy, Precision, Recall
accuracy = accuracy_score(y_test, y_pred)
precision = precision_score(y_test, y_pred)
recall = recall_score(y_test, y_pred)

print(f"Accuracy : {accuracy:.2f}")
print(f"Precision: {precision:.2f}")
print(f"Recall   : {recall:.2f}")
# Compute ROC curve and AUC
```

```


fpr, tpr, thresholds = roc_curve(y_test, y_prob)
auc_score = roc_auc_score(y_test, y_prob)

# Plot ROC
plt.figure(figsize=(8,6))
plt.plot(fpr, tpr, label=f'Logistic Regression (AUC = {auc_score:.2f})', color='darkorange')
plt.plot([0, 1], [0, 1], linestyle='--', color='gray') # baseline
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve')
plt.legend()
plt.grid(True)
plt.show()

from sklearn.metrics import ConfusionMatrixDisplay
ConfusionMatrixDisplay.from_estimator(model, X_test, y_test, cmap='Blues')
plt.title("Confusion Matrix")
plt.show()

```

Output



```
(base) C:\Users\admin>python Desktop\ML\prac4.py
Accuracy : 0.83
Precision: 0.81
Recall    : 0.82
```

B) Implement and demonstrate k-nearest Neighbor algorithm. Read the training data from a .CSV file and build the model to classify a test sample. Print both correct and wrong predictions.

```
import csv
import math
from collections import Counter
from sklearn.model_selection import train_test_split

# Function to read CSV data
def read_csv(filename):
    data = []
    with open(filename, 'r') as file:
        csv_reader = csv.reader(file)
        header = next(csv_reader) # skip header
        for row in csv_reader:
            features = list(map(float, row[:-1])) # convert features to float
            label = row[-1]
            data.append((features, label))
    return data

# Euclidean distance
def euclidean_distance(a, b):
    return math.sqrt(sum((x - y) ** 2 for x, y in zip(a, b)))

# k-NN algorithm
def knn_predict(train_data, test_sample, k=3):
    distances = []
    for features, label in train_data:
        distance = euclidean_distance(features, test_sample)
        distances.append((distance, label))
    distances.sort(key=lambda x: x[0])
    k_nearest = [label for _, label in distances[:k]]
    most_common = Counter(k_nearest).most_common(1)
    return most_common[0][0]

# Evaluation function
def evaluate_knn(train_data, test_data, k=3):
    correct = 0
    wrong = 0
    for features, actual_label in test_data:
        predicted = knn_predict(train_data, features, k)
        if predicted == actual_label:
```

```

    correct += 1
    print(f" Correct: Predicted = {predicted}, Actual = {actual_label}")
else:
    wrong += 1
    print(f" Wrong: Predicted = {predicted}, Actual = {actual_label}")
print(f"\nSummary: Correct = {correct}, Wrong = {wrong}, Accuracy = {correct / (correct + wrong) * 100:.2f}%")

# Main execution
if __name__ == "__main__":
    filename = "C:/Users/admin/Downloads/winequality_red.csv" # Replace with your file
    dataset = read_csv(filename)
    train_set, test_set = train_test_split(dataset, test_size=0.3, random_state=42)
    evaluate_knn(train_set, test_set, k=3)

```

Output-

```

Correct: Predicted = 6, Actual = 6
Wrong: Predicted = 5, Actual = 6
Correct: Predicted = 7, Actual = 7
Wrong: Predicted = 6, Actual = 4
Wrong: Predicted = 4, Actual = 6
Correct: Predicted = 5, Actual = 5
Correct: Predicted = 6, Actual = 6
Wrong: Predicted = 7, Actual = 6
Correct: Predicted = 7, Actual = 7
Correct: Predicted = 5, Actual = 5
Correct: Predicted = 6, Actual = 6
Wrong: Predicted = 6, Actual = 5
Correct: Predicted = 5, Actual = 5
Correct: Predicted = 6, Actual = 6
Correct: Predicted = 7, Actual = 7
Wrong: Predicted = 6, Actual = 5
Correct: Predicted = 5, Actual = 5
Correct: Predicted = 6, Actual = 6
Wrong: Predicted = 6, Actual = 5
Correct: Predicted = 5, Actual = 5
Correct: Predicted = 6, Actual = 6
Correct: Predicted = 6, Actual = 5
Correct: Predicted = 6, Actual = 6
Wrong: Predicted = 6, Actual = 5
Correct: Predicted = 5, Actual = 5
Correct: Predicted = 6, Actual = 6
Correct: Predicted = 6, Actual = 6
Wrong: Predicted = 6, Actual = 4
Wrong: Predicted = 7, Actual = 5
Correct: Predicted = 6, Actual = 6
Correct: Predicted = 5, Actual = 5
Correct: Predicted = 7, Actual = 7
Wrong: Predicted = 6, Actual = 8
Correct: Predicted = 6, Actual = 6
Wrong: Predicted = 8, Actual = 7
Wrong: Predicted = 6, Actual = 4

Summary: Correct = 229, Wrong = 251, Accuracy = 47.71%

```

C)Build a decision tree classifier or regressor. Control hyperparameters like tree depth to avoid overfitting. Visualize the tree.

```
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeClassifier, plot_tree
from sklearn.metrics import accuracy_score, classification_report
import matplotlib.pyplot as plt
data = pd.read_csv("C:/Users/admin/Downloads/winequality_red.csv")
X = data.iloc[:, :-1]
y = data.iloc[:, -1]
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)
clf = DecisionTreeClassifier(max_depth=3, random_state=42)
clf.fit(X_train, y_train)
y_pred = clf.predict(X_test)
print("\nAccuracy:", accuracy_score(y_test, y_pred))
print("\nClassification Report:\n", classification_report(y_test, y_pred, zero_division=0))
plt.figure(figsize=(20, 10))
plot_tree(clf, feature_names=X.columns, class_names=[str(c) for c in clf.classes_], filled=True)
plt.title("Decision Tree Visualization (max_depth=3)")
plt.show()
```

Output

```
(base) C:\Users\admin>python Desktop\ML\prac4c.py
Accuracy: 0.5166666666666667
Classification Report:
precision    recall    f1-score    support
          3       0.00      0.00      0.00        1
          4       0.00      0.00      0.00       17
          5       0.53      0.88      0.66      195
          6       0.52      0.30      0.38      200
          7       0.42      0.25      0.31       61
          8       0.00      0.00      0.00        6
   accuracy                           0.52      480
  macro avg       0.24      0.24      0.23      480
weighted avg     0.48      0.52      0.47      480

(base) C:\Users\admin>
```





D)Implement a Support Vector Machine for any relevant dataset.

```
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn import datasets
from sklearn.model_selection import train_test_split
from sklearn.svm import SVC
from sklearn.metrics import classification_report, accuracy_score, confusion_matrix
iris = datasets.load_iris()
X = iris.data[:, :2] # use only first 2 features for visualization
y = iris.target
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)
model = SVC(kernel='linear', C=1.0)
model.fit(X_train, y_train)
y_pred = model.predict(X_test)
print("Accuracy:", accuracy_score(y_test, y_pred))
print("Classification Report:\n", classification_report(y_test, y_pred))
cm = confusion_matrix(y_test, y_pred)
sns.heatmap(cm, annot=True, fmt='d', xticklabels=iris.target_names, yticklabels=iris.target_names)
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.title("Confusion Matrix")
plt.show()
def plot_decision_boundary(X, y, model):
    x_min, x_max = X[:, 0].min() - 1, X[:, 0].max() + 1
    y_min, y_max = X[:, 1].min() - 1, X[:, 1].max() + 1
    xx, yy = np.meshgrid(np.linspace(x_min, x_max, 200),
                         np.linspace(y_min, y_max, 200))
    Z = model.predict(np.c_[xx.ravel(), yy.ravel()])
    Z = Z.reshape(xx.shape)
    plt.contourf(xx, yy, Z, alpha=0.3)
    plt.scatter(X[:, 0], X[:, 1], c=y, cmap=plt.cm.Set1)
    plt.xlabel(iris.feature_names[0])
    plt.ylabel(iris.feature_names[1])
    plt.title("SVM Decision Boundary")
    plt.show()
plot_decision_boundary(X, y, model)
```

Output

```
(base) C:\Users\admin>python Desktop\ML\prac4d.py
Accuracy: 0.8
Classification Report:
precision    recall    f1-score   support
          0       1.00     1.00      1.00      19
          1       0.70     0.54      0.61      13
          2       0.62     0.77      0.69      13
   accuracy                           0.80      45
  macro avg       0.78     0.77      0.77      45
weighted avg     0.81     0.80      0.80      45
```




E) Train a random forest ensemble. Experiment with the number of trees and feature sampling.

Compare performance to a single decision tree.

```
import numpy as np
import matplotlib.pyplot as plt
from sklearn.datasets import load_breast_cancer
from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score
data = load_breast_cancer()
X = data.data
y = data.target
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)
tree_clf = DecisionTreeClassifier(random_state=42)
tree_clf.fit(X_train, y_train)
tree_preds = tree_clf.predict(X_test)
tree_acc = accuracy_score(y_test, tree_preds)
n_trees_list = [1, 5, 10, 50, 100, 200]
results = []
for n_trees in n_trees_list:
    rf_clf = RandomForestClassifier(n_estimators=n_trees, max_features='sqrt', random_state=42)
    rf_clf.fit(X_train, y_train)
    rf_preds = rf_clf.predict(X_test)
    rf_acc = accuracy_score(y_test, rf_preds)
    results.append((n_trees, rf_acc))
print(f"Single Decision Tree Accuracy: {tree_acc:.4f}")
print("Random Forest Accuracies:")
for n_trees, acc in results:
    print(f" Trees: {n_trees}, Accuracy: {acc:.4f}")
n_vals, acc_vals = zip(*results)
plt.figure(figsize=(8, 4))
plt.plot(n_vals, acc_vals, marker='o', label='Random Forest')
plt.axhline(y=tree_acc, color='r', linestyle='--', label='Decision Tree')
plt.xlabel("Number of Trees")
plt.ylabel("Accuracy")
plt.title("Random Forest vs Decision Tree")
plt.legend()
plt.grid(True)
plt.show()
```

Output

```
Single Decision Tree Accuracy: 0.9415
Random Forest Accuracies:
Trees: 1, Accuracy: 0.9474
Trees: 5, Accuracy: 0.9649
Trees: 10, Accuracy: 0.9649
Trees: 50, Accuracy: 0.9708
Trees: 100, Accuracy: 0.9708
Trees: 200, Accuracy: 0.9708
```



F) Implement a gradient boosting machine (e.g., XGBoost). Tune hyperparameters and explore feature importance.

```
import xgboost as xgb

from sklearn.datasets import load_breast_cancer

from sklearn.model_selection import train_test_split, GridSearchCV

from sklearn.metrics import accuracy_score, classification_report

import matplotlib.pyplot as plt

data = load_breast_cancer()

X = data.data

y = data.target

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)

xgb_model = xgb.XGBClassifier(eval_metric='logloss', use_label_encoder=False, random_state=42)

xgb_model.fit(X_train, y_train)

y_pred = xgb_model.predict(X_test)

print("Basic XGBoost Accuracy:", accuracy_score(y_test, y_pred))

param_grid = {

    'n_estimators': [50, 100],

    'max_depth': [3, 5],

    'learning_rate': [0.01, 0.1],

    'subsample': [0.8, 1.0]

}

grid = GridSearchCV(

    xgb.XGBClassifier(eval_metric='logloss', use_label_encoder=False, random_state=42),

    param_grid,

    cv=3,

    scoring='accuracy',

    verbose=1

)
```

```

grid.fit(X_train, y_train)

best_model = grid.best_estimator_

y_best_pred = best_model.predict(X_test)

print("\nBest Parameters:", grid.best_params_)

print("Tuned XGBoost Accuracy:", accuracy_score(y_test, y_best_pred))

print("\nClassification Report:\n", classification_report(y_test, y_best_pred))

xgb.plot_importance(best_model, importance_type='gain', title='Feature Importance',
max_num_features=10)

plt.tight_layout()

plt.show()

```


```

Best Parameters: {'learning_rate': 0.1, 'max_depth': 3, 'n_estimators': 100, 'subsample': 0.8}
Tuned XGBoost Accuracy: 0.9645122807017544

Classification Report:
precision    recall    f1-score   support
          0       0.95      0.95      0.95      63
          1       0.97      0.97      0.97     108

   accuracy                           0.96      171
  macro avg       0.96      0.96      0.96      171
weighted avg       0.96      0.96      0.96      171

```



Practical no 5

Implement and demonstrate the working of a Naive Bayesian classifier using a sample data set.

Build the model to classify a test sample.

```
import pandas as pd  
from sklearn.preprocessing import LabelEncoder  
from sklearn.naive_bayes import CategoricalNB  
  
data = {  
  
    'Outlook': [  
        'Sunny', 'Sunny', 'Overcast', 'Rain', 'Rain', 'Rain', 'Overcast',  
        'Sunny', 'Sunny', 'Rain', 'Sunny', 'Overcast', 'Overcast', 'Rain'  
    ],  
  
    'Temp': [  
        'Hot', 'Hot', 'Hot', 'Mild', 'Cool', 'Cool', 'Cool',  
        'Mild', 'Cool', 'Mild', 'Mild', 'Mild', 'Hot', 'Mild'  
    ],  
  
    'Humidity': [  
        'High', 'High', 'High', 'High', 'Normal', 'Normal', 'Normal',  
        'High', 'Normal', 'Normal', 'Normal', 'High', 'Normal', 'High'  
    ],  
  
    'Wind': [  
        'Weak', 'Strong', 'Weak', 'Weak', 'Weak', 'Strong', 'Strong',  
        'Weak', 'Weak', 'Weak', 'Strong', 'Strong', 'Weak', 'Strong'  
    ],  
  
    'Play': [  
        'No', 'No', 'Yes', 'Yes', 'Yes', 'No', 'Yes',  
        'No', 'Yes', 'Yes', 'Yes', 'Yes', 'Yes', 'Yes', 'No'  
    ]  
}
```

```

df = pd.DataFrame(data)

encoders = {}

for col in df.columns:

    le = LabelEncoder()

    df[col] = le.fit_transform(df[col])

    encoders[col] = le # Store encoder for future decoding

X = df[['Outlook', 'Temp', 'Humidity', 'Wind']]

y = df['Play']

model = CategoricalNB()

model.fit(X, y)

test_sample_dict = {

    'Outlook': encoders['Outlook'].transform(['Sunny'])[0],

    'Temp': encoders['Temp'].transform(['Cool'])[0],

    'Humidity': encoders['Humidity'].transform(['High'])[0],

    'Wind': encoders['Wind'].transform(['Strong'])[0]

}

test_sample = pd.DataFrame([test_sample_dict])

prediction = model.predict(test_sample)

decoded_prediction = encoders['Play'].inverse_transform(prediction)

print("Prediction:", decoded_prediction[0])

```

Output

```

base) C:\Users\admin>python Desktop\MLPRAC\prac5a.py
prediction: No

```

B) Implement Hidden Markov Models using hmmlearn

```
import numpy as np
from hmmlearn import hmm
states = ["Rainy", "Sunny"]
n_states = len(states)
observations = ["Walk", "Shop", "Clean"]
n_observations = len(observations)
transition_prob = np.array([
    [0.7, 0.3], # Rainy -> (Rainy, Sunny)
    [0.4, 0.6] # Sunny -> (Rainy, Sunny)
])
emission_prob = np.array([
    [0.1, 0.4, 0.5], # Rainy -> (Walk, Shop, Clean)
    [0.6, 0.3, 0.1] # Sunny -> (Walk, Shop, Clean)
])
start_prob = np.array([0.6, 0.4])
model = hmm.CategoricalHMM(n_components=n_states, n_iter=100, init_params="")
model.startprob_ = start_prob
model.transmat_ = transition_prob
model.emissionprob_ = emission_prob
obs_sequence = np.array([[0], [1], [2]]) # Walk -> Shop -> Clean
logprob, hidden_states = model.decode(obs_sequence, algorithm="viterbi")
print("Observation sequence:", [observations[i[0]] for i in obs_sequence])
print("Most likely hidden states:", [states[i] for i in hidden_states])
Output
```

```
(base) C:\Users\admin>python Desktop\MLPRAC\prac5b.py
Observation sequence: ['Walk', 'Shop', 'Clean']
Most likely hidden states: ['Sunny', 'Rainy', 'Rainy']
```

Practical no 6

A) Implement Bayesian Linear Regression to explore prior and posterior distribution

```
import numpy as np
import matplotlib.pyplot as plt
from numpy.linalg import inv
np.random.seed(42)
n = 40
x = np.random.uniform(-3, 3, size=n)
X = np.column_stack([np.ones(n), x]) # Add intercept term
true_w = np.array([1.5, -2.0])
sigma = 0.5 # Known noise standard deviation
y = X @ true_w + np.random.normal(0, sigma, size=n)
m0 = np.array([0.0, 0.0]) # Prior mean
S0 = np.diag([5.0**2, 5.0**2]) # Prior covariance (broad)
SN = inv(inv(S0) + (X.T @ X) / (sigma**2))
mN = SN @ (inv(S0) @ m0 + (X.T @ y) / (sigma**2))
print("True weights:", true_w)
print("Prior mean:", m0)
print("Posterior mean:", mN)
print("Posterior covariance:\n", SN)
x_grid = np.linspace(-3.5, 3.5, 200)
Xg = np.column_stack([np.ones_like(x_grid), x_grid])
pred_mean = Xg @ mN
pred_var = sigma**2 + np.einsum('ij,jk,ik->i', Xg, SN, Xg) # Predictive variance
pred_sd = np.sqrt(pred_var)
plt.figure(figsize=(7, 5))
plt.scatter(x, y, label="Data", color="black")
plt.plot(x_grid, pred_mean, label="Posterior mean", color="blue")
plt.fill_between(
    x_grid,
    pred_mean - 1.96 * pred_sd,
    pred_mean + 1.96 * pred_sd,
    color="blue", alpha=0.2,
    label="95% predictive interval"
)
plt.xlabel("x")
plt.ylabel("y")
plt.title("Bayesian Linear Regression: Posterior Prediction")
plt.legend()
plt.grid(True)
plt.tight_layout()
plt.show()
```

```

def gaussian_contour(ax, mean, cov, color, label):
    w0 = np.linspace(mean[0] - 4 * np.sqrt(cov[0, 0]), mean[0] + 4 * np.sqrt(cov[0, 0]), 200)
    w1 = np.linspace(mean[1] - 4 * np.sqrt(cov[1, 1]), mean[1] + 4 * np.sqrt(cov[1, 1]), 200)
    W0, W1 = np.meshgrid(w0, w1)
    W = np.stack([W0 - mean[0], W1 - mean[1]], axis=-1)
    icov = inv(cov)
    M = (W[..., 0]*icov[0, 0] + W[..., 1]*icov[1, 0])*W[..., 0] + \
        (W[..., 0]*icov[0, 1] + W[..., 1]*icov[1, 1])*W[..., 1]
    Z = np.exp(-0.5 * M)
    ax.contour(W0, W1, Z, levels=5, colors=color, alpha=0.7)
    ax.plot([], [], color=color, label=label) # For legend
plt.figure(figsize=(6, 6))
ax = plt.gca()
gaussian_contour(ax, m0, S0, "red", "Prior")
gaussian_contour(ax, mN, SN, "blue", "Posterior")
ax.scatter(*true_w, marker="x", s=100, color="black", label="True weights")
ax.set_xlabel("w0 (intercept)")
ax.set_ylabel("w1 (slope)")
ax.set_title("Prior vs Posterior in Parameter Space")
ax.legend()
plt.grid(True)
plt.tight_layout()
plt.show()



```

Output

```

(base) [C:\Users\admin>python Desktop\MLPRAC\ml6a.py
True weights: [ 1.5 -2. ]
Prior mean: [0. 0.]
Posterior mean: [ 1.4633892  -1.98871086]
Posterior covariance:
[[0.00638859 0.00053425]
 [0.00053425 0.00203646]]



```



B) Implement Gaussian Mixture Models for density estimation and unsupervised clustering

```
import numpy as np
import matplotlib.pyplot as plt
from sklearn.mixture import GaussianMixture
from sklearn.datasets import make_blobs
X, y_true = make_blobs(
    n_samples=400,
    centers=3,
    cluster_std=[1.0, 1.5, 0.5],
    random_state=42
)
plt.scatter(X[:, 0], X[:, 1], s=30, cmap="viridis")
plt.title("Synthetic Data")
plt.show()
gmm = GaussianMixture(
    n_components=3,
    covariance_type='full',
    random_state=42
)
gmm.fit(X)
labels = gmm.predict(X)
plt.scatter(X[:, 0], X[:, 1], c=labels, s=30, cmap="viridis")
plt.scatter(
    gmm.means_[:, 0], gmm.means_[:, 1],
    c="red", s=100, marker="X", label="Cluster Means"
)
plt.title("GMM Clustering")
plt.legend()
plt.show()
x = np.linspace(np.min(X[:, 0]) - 1, np.max(X[:, 0]) + 1, 200)
y = np.linspace(np.min(X[:, 1]) - 1, np.max(X[:, 1]) + 1, 200)
Xgrid, Ygrid = np.meshgrid(x, y)
XX = np.column_stack([Xgrid.ravel(), Ygrid.ravel()])
Z = np.exp(gmm.score_samples(XX))
Z = Z.reshape(Xgrid.shape)
plt.contourf(Xgrid, Ygrid, Z, levels=20, cmap="viridis")
plt.scatter(X[:, 0], X[:, 1], s=10, c="white", alpha=0.6)
plt.title("GMM Density Estimation")
plt.show()
print("Cluster means:\n", gmm.means_)
print("\nCovariances:\n", gmm.covariances_)
print("\nWeights (mixture proportions):\n", gmm.weights_)
```

Output



Practical no 7

A) Implement cross-validation techniques (k-fold, stratified, etc.) for robust model Evaluation

```
from sklearn.datasets import load_iris
from sklearn.model_selection import KFold, StratifiedKFold, cross_val_score
from sklearn.ensemble import RandomForestClassifier
from sklearn.svm import SVC
import numpy as np
data = load_iris()
X = data.data
y = data.target
model = RandomForestClassifier(random_state=42)
kf = KFold(n_splits=5, shuffle=True, random_state=42)
scores = cross_val_score(model, X, y, cv=kf, scoring='accuracy')
print("K-Fold CV Accuracy Scores:", scores)
print("Mean K-Fold CV Accuracy:", np.mean(scores))
# Stratified K-Fold ensures class proportions are preserved
skf = StratifiedKFold(n_splits=5, shuffle=True, random_state=42)
stratified_scores = cross_val_score(model, X, y, cv=skf, scoring='accuracy')
print("\nStratified K-Fold CV Accuracy Scores:", stratified_scores)
print("Mean Stratified K-Fold CV Accuracy:", np.mean(stratified_scores))
from sklearn.model_selection import LeaveOneOut
loo = LeaveOneOut()
loo_scores = cross_val_score(model, X, y, cv=loo, scoring='accuracy')
print("\nLeave-One-Out CV Accuracy Mean:", np.mean(loo_scores))
```

Output

```
(base) C:\Users\admin>python D:/ML/mlprac7a.py
K-Fold CV Accuracy Scores: [1.          0.96666667 0.93333333 0.93333333 0.96666667]
Mean K-Fold CV Accuracy: 0.9600000000000002

Stratified K-Fold CV Accuracy Scores: [0.96666667 0.96666667 0.93333333 0.96666667 0.9       ]
Mean Stratified K-Fold CV Accuracy: 0.9466666666666667

Leave-One-Out CV Accuracy Mean: 0.953333333333334

(base) C:\Users\admin>
```

B) Systematically explore combinations of hyperparameters to optimize model performance. (Use grid and randomized search)

```
from sklearn.datasets import load_iris
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import GridSearchCV, RandomizedSearchCV, train_test_split
from sklearn.metrics import accuracy_score
from scipy.stats import randint
import numpy as np
data = load_iris()
X = data.data
y = data.target
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42
)
model = RandomForestClassifier(random_state=42)
param_grid = {
    'n_estimators': [50, 100, 200],
    'max_depth': [None, 5, 10],
    'min_samples_split': [2, 5],
    'min_samples_leaf': [1, 2]
}
grid_search = GridSearchCV(
    estimator=model,
    param_grid=param_grid,
    cv=5,
    scoring='accuracy',
    n_jobs=1 # You can increase this (e.g., to 2 or 4) if stable
)
grid_search.fit(X_train, y_train)
print("Best Grid Search Parameters:", grid_search.best_params_)
print("Best Grid Search CV Score:", grid_search.best_score_)
y_pred_grid = grid_search.best_estimator_.predict(X_test)
print("Test Set Accuracy (Grid Search):", accuracy_score(y_test, y_pred_grid))
param_dist = {
    'n_estimators': randint(50, 300),
    'max_depth': [None, 5, 10, 20],
    'min_samples_split': randint(2, 10),
    'min_samples_leaf': randint(1, 5)
}
random_search = RandomizedSearchCV(
    estimator=model,
    param_distributions=param_dist,
    n_iter=20,
    cv=5,
    scoring='accuracy',
    n_jobs=1,
    random_state=42
)
random_search.fit(X_train, y_train)
```

```
print("\nBest Randomized Search Parameters:", random_search.best_params_)
print("Best Randomized Search CV Score:", random_search.best_score_)
y_pred_random = random_search.best_estimator_.predict(X_test)
print("Test Set Accuracy (Randomized Search):", accuracy_score(y_test, y_pred_random))
```

Output:

```
(base) C:\Users\admin>python D:/ML/mlprac7b.py
Best Grid Search Parameters: {'max_depth': None, 'min_samples_leaf': 2, 'min_samples_split': 2, 'n_estimators': 200}
Best Grid Search CV Score: 0.958333333333334
Test Set Accuracy (Grid Search): 1.0

Best Randomized Search Parameters: {'max_depth': 10, 'min_samples_leaf': 3, 'min_samples_split': 9, 'n_estimators': 166}
Best Randomized Search CV Score: 0.958333333333334
Test Set Accuracy (Randomized Search): 1.0
```

Practical no 8

Implement Bayesian Learning using inferences

```
import numpy as np
import matplotlib.pyplot as plt
from sklearn.linear_model import BayesianRidge
from sklearn.model_selection import train_test_split
np.random.seed(42)
X = 2 * np.random.rand(100, 1)
true_coef = 3.5
true_intercept = 1.5
y = true_coef * X.flatten() + true_intercept + np.random.randn(100) * 0.5
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
from sklearn.linear_model import BayesianRidge
model = BayesianRidge()
model.fit(X_train, y_train)
y_pred, y_std = model.predict(X_test, return_std=True)
plt.scatter(X_train, y_train, color='blue', label='Training Data')
X_line = np.linspace(0, 2, 100).reshape(-1, 1)
y_true = true_coef * X_line.flatten() + true_intercept
plt.plot(X_line, y_true, 'g--', label='True Function')
y_mean, y_std_line = model.predict(X_line, return_std=True)
plt.plot(X_line, y_mean, 'r-', label='Predicted Mean')
plt.fill_between(
    X_line.flatten(),
    y_mean - 2 * y_std_line,
    y_mean + 2 * y_std_line,
    color='pink',
    alpha=0.5,
    label='Uncertainty ( $\pm 2$  std)'
)
plt.legend()
```


```

plt.xlabel('X')
plt.ylabel('y')
plt.title('Bayesian Linear Regression with Posterior Inference')
plt.show()

print("Learned Coefficients Mean:", model.coef_)
print("Learned Intercept:", model.intercept_)
print("Alpha (Precision of noise):", model.alpha_)
print("Lambda (Precision of weights prior):", model.lambda_)

```

Output



```

(base) C:\Users\admin>python D:/ML/mlprac8.py
Learned Coefficients Mean: [3.39733606]
Learned Intercept: 1.5736340376397093
Alpha (Precision of noise): 4.6597831289883205
Lambda (Precision of weights prior): 0.08658179117729597

```

Security Operations Center



JNAN VIKAS MANDAL'S

Mohanlal Raichand Mehta College of Commerce
Diwali Maa College of Science
Amritlal Raichand Mehta College of Arts
Dr. R.T. Doshi College of Computer Science
NAAC Re-Accredited Grade 'A+' (CGPA : 3.31) (3rd Cycle)

DEPARTMENT OF INFORMATION TECHNOLOGY

CERTIFICATE

This is to certify that Sonali Rupchand Thakur bearing seat no. 1313626 has done the project work/journal work in the subject of Security Operations Center of semester 3 Practical Examination during the academic year 2025-26 under the guidance of **Dr.Sanjivani Nalkar** being the partial requirement for the fulfilment of the curriculum of Degree in Master of Science in Information Technology under University of Mumbai.

Place: Airoli

Date:

Sign of Subject in- Charge

Sign of External Examiner

Sign of Coordinator

TABLE OF CONTENTS

Practical No	Practicals	Date	Sign
1.	a. Encrypting and Decrypting Data Using a Hacker Tool b. Encrypting and Decrypting Data Using OpenSSL c. Hashing a Text File with OpenSSL and Verifying Hashes		
2.	a. Examining Telnet and SSH in Wireshark		
3.	a. Demonstrate the use of Snort and Firewall Rules b. Demonstrate Extract an Executable from a PCAP c. Demonstrate a practical for Exploring DNS Traffic		
4.	b. Exploring Processes, Threads, Handles, and Windows Registry		
5.	Perform a practical to Attack on a mySQL Database by using PCAP File.		
6.	Create your own syslog Server		
7.	Configure your Linux system to send syslog messages to a syslog server and read them		
8.	Install and Run Splunk on Linux		
9.	Install and Configure ELK on Linux		

Practical 1

a. Encrypting and Decrypting Data Using a Hacker Tool

Part 1: Create and Encrypt Files

```
cd ~  
mkdir Zip-Files  
cd Zip-Files  
echo This is a sample text file > sample-1.txt  
echo This is a sample text file > sample-2.txt  
echo This is a sample text file > sample-3.txt  ls  
Step 2: Zip and encrypt the text files.  
zip -e file-1.zip sample*  
enter a one-character password = B  
zip -e file-2.zip sample*  
enter a one-character password = R2  
zip -e file-3.zip sample*  
enter a one-character password = 0B1  
zip -e file-4.zip sample*  
enter a one-character password = Y0Da  
zip -e file-5.zip sample*  
enter a one-character password = C-3P0  ls -l f*
```

Attempt to open a zip using an incorrect password
unzip file-1.zip enter wrong password


Part 2: Recover Encrypted Zip File Passwords

Step 1: Introduction to fcrackzip

```
fcrackzip -h
```

Step 2: Recovering Passwords using fcrackzip

```
fcrackzip -vul 1-4 file-1.zip  fcrackzip -vul 1-4 file-2.zip  fcrackzip -vul 1-4 file-3.zip  
fcrackzip -vul 1-4 file-4.zip  fcrackzip -vul 1-5 file-5.zip
```



```
Terminal - analyst@secOps:~/Zip-Files  
File Edit View Terminal Tabs Help  
[analyst@secOps Zip-Files]$ fcrackzip -vul 1-4 file-2.zip  
found file 'sample-1.txt', (size cp/uc      39/     27, flags 9, chk 728f)  
found file 'sample-2.txt', (size cp/uc      39/     27, flags 9, chk 729c)  
found file 'sample-3.txt', (size cp/uc      40/     28, flags 9, chk 72a6)  
  
PASSWORD FOUND!!!!: pw == R2  
[analyst@secOps Zip-Files]$
```


b. Encrypting and Decrypting Data Using OpenSSL

Create a text file named input.txt and add some text or data in it. Encrypt the file with file command. After encrypting the file, a .enc file will be generated. And decrypt it into another file and name it output.txt.

```
C:\SOCPractical>openssl enc -aes-256-ecb -a -salt -in input.txt -out cipher.enc
enter AES-256-ECB encryption password:
Verifying - enter AES-256-ECB encryption password:
*** WARNING : deprecated key derivation used.
Using -iter or -pbkdf2 would be better.

C:\SOCPractical>openssl enc -d -aes-256-ecb -a -salt -in cipher.enc -out output.txt
enter AES-256-ECB decryption password:
*** WARNING : deprecated key derivation used.
Using -iter or -pbkdf2 would be better.

C:\SOCPractical>
```



c. Hashing a Text File with OpenSSL and Verifying Hashes

Part 1: Hashing a Text File with OpenSSL

```
cd /home/sec_admin/lab.support.files/  
cat letter_to_grandma.txt  
openssl sha256 letter_to_grandma.txt  
nano letter_to_grandma.txt  
change the first sentence from 'Hi Grandma' to 'Hi Grandpa'  
After the change = <CONTROL+X> , 'Y' , <Enter>  
openssl sha256 letter_to_grandma.txt  
  
openssl sha512 letter_to_grandma.txt
```

Part 2: Verifying Hashes

```
cat sample.img_SHA256.sig  
sha256sum sample.img
```

```
[sec_admin@secOps lab.support.files]$ cat sample.img_SHA256.sig  
c56c4724c26eb0157963c0d62b76422116be31804a39c82fd44ddf0ca5013e6a
```

```
[sec_admin@secOps lab.support.files]$ sha256sum sample.img  
c56c4724c26eb0157963c0d62b76422116be31804a39c82fd44ddf0ca5013e6a
```

Practical 2

a. Examining Telnet and SSH in Wireshark

Examining Telnet and SSH in Wireshark

Part 1: Examining a Telnet Session with Wireshark

- a. Start the CyberOps Workstation VM and log in with username analyst and password cyberops.
- b. Open a terminal window and start Wireshark. Press OK to continue after reading the warning message.
- c. Start a Wireshark capture on the Loopback: lo interface.

Step 1: Capture data.

- a. Start the CyberOps Workstation VM and log in with username analyst and password cyberops.
- b. Open a terminal window and start Wireshark. Press OK to continue after reading the warning message.
- c. Start a Wireshark capture on the Loopback: lo interface.

- d. Open another terminal window. Start a Telnet session to the localhost. Enter username analyst and password cyberops when prompted. Note that it may take several minutes for the “connected to localhost” and login prompt to appear.

write the following cmd:

```
[analyst@secOps ~]$ sudo systemctl start telnet.socket
```

```
[analyst@secOps ~]$ telnet localhost
```

- e. Stop the Wireshark capture after you have provided the user credentials.

Step 2: Examine the Telnet session.

- a. Apply a filter that only displays Telnet-related traffic. Enter Telnet in the filter field and click Apply.

- b. Right-click one of the Telnet lines in the Packet list section of Wireshark, and from the drop-down list, select Follow TCP Stream.

- Step 1: Capture data.
- a. Start the CyberOps Workstation VM and log in with username analyst and password cyberops.
 - b. Open a terminal window and start Wireshark. Press OK to continue after reading the warning message.

```
[analyst@secOps analyst]$ sudo wireshark-gtk
```

```
[sudo] password for analyst: cyberops
```

- c. Start a Wireshark capture on the Loopback: lo interface.

- d. Open another terminal window. Start a Telnet session to the localhost. Enter username analyst and password cyberops when prompted. Note that it may take several minutes for the “connected to localhost” and login prompt to appear.

```
[analyst@secOps ~]$ telnet localhost
```

```
secOps login: analyst
```

Password:

```
Last login: Fri Apr 28 10:50:52 from localhost.localdomain
```

```
[analyst@secOps ~]$
```

e. Stop the Wireshark capture after you have provided the user credentials.

Step 2: Examine the Telnet session.

a. Apply a filter that only displays Telnet-related traffic. Enter Telnet in the filter field and click Apply.

b. Right-click one of the Telnet lines in the Packet list section of Wireshark, and from the drop-down list, select Follow TCP Stream.

c. The Follow TCP Stream window displays the data for your Telnet session with the CyberOps Workstation VM. The entire session is displayed in plaintext, including your password. Notice that the username that you entered is displayed with duplicate characters. This is caused by the echo setting in Telnet to allow you to view the characters that you type on the screen.

d. After you have finished reviewing your Telnet session in the Follow TCP Stream window, click Close.

e. Type exit at the terminal to exit the Telnet session.

```
[analyst@secOps ~]$ exit
```

Part 2: Examine an SSH Session with Wireshark

same steps for ssh also just replace telnet with ssh localhost

The screenshot shows the 'Follow TCP Stream' dialog in Wireshark. The title bar says 'Follow TCP Stream (tcp.stream eq 2)'. The main pane displays the raw ASCII data of an SSH session. The data starts with the SSH protocol version 'SSH-2.0-OpenSSH_7.4' and continues with the key exchange process, including various cryptographic parameters and salts. At the bottom of the pane, it says 'Entire conversation (5237 bytes)'. Below the pane are several buttons: 'Find', 'Save As', 'Print', 'ASCII', 'EBCDIC', 'Hex Dump', 'C Arrays', 'Raw', 'Help', 'Filter Out This Stream', and 'Close'. The 'Raw' button is currently selected.

```
SSH-2.0-OpenSSH_7.4
SSH-2.0-OpenSSH_7.4
...4..x5....t.?g.B....#....curve25519-sha256,curve25519-sha256@libssh.org,ecdh-sha2-
nistp256,ecdh-sha2-nistp384,ecdh-sha2-nistp521,diffie-hellman-group-exchange-
sha256,diffie-hellman-group16-sha512,diffie-hellman-group18-sha512,diffie-hellman-
group14-sha256,diffie-hellman-group14-sha1...Assh-rsa,rsa-sha2-512,rsa-sha2-256,ecdsa-
sha2-nistp256,ssh-ed25519...lchacha20-poly1305@openssh.com,aes128-ctr,aes192-ctr,aes256-
ctr,aes128-gcm@openssh.com,aes256-gcm@openssh.com...lchacha20-
poly1305@openssh.com,aes128-ctr,aes192-ctr,aes256-ctr,aes128-gcm@openssh.com,aes256-
gcm@openssh.com...umac-64-etm@openssh.com,umac-128-etm@openssh.com,hmac-sha2-256-
etm@openssh.com,hmac-sha2-512-etm@openssh.com,hmac-sha1-
etm@openssh.com,umac-64@openssh.com,umac-128@openssh.com,hmac-sha2-256,hmac-
sha2-512,hmac-sha1...umac-64-etm@openssh.com,umac-128-etm@openssh.com,hmac-sha2-256-
etm@openssh.com,hmac-sha2-512-etm@openssh.com,hmac-sha1-
etm@openssh.com,umac-64@openssh.com,umac-128@openssh.com,hmac-sha2-256,hmac-
sha2-512,hmac-
sha1...none,zlib@openssh.com.none,zlib@openssh.com.....i....
g.!....m....0curve25519-sha256,curve25519-sha256@libssh.org,ecdh-sha2-nistp256,ecdh-
sha2-nistp384,ecdh-sha2-nistp521,diffie-hellman-group-exchange-sha256,diffie-hellman-
group16-sha512,diffie-hellman-group18-sha512,diffie-hellman-group-exchange-sha1,diffie-
hellman-group14-sha256,diffie-hellman-group14-sha1,ext-info-c..."ecdsa-sha2-nistp256-
cert-v01@openssh.com,ecdsa-sha2-nistp384-cert-v01@openssh.com,ecdsa-sha2-nistp521-cert-
v01@openssh.com,ecdsa-sha2-nistp256,ecdsa-sha2-nistp384,ecdsa-sha2-nistp521,ssh-ed25519-
cert-v01@openssh.com,ssh-rsa-cert-v01@openssh.com,ssh-ed25519,rsa-sha2-512,rsa-
sha2-256,ssh-rsa...chacha20-poly1305@openssh.com,aes128-ctr,aes192-ctr,aes256-
```



Practical 3

a. Demonstrate the use of Snort and Firewall Rules


Step-1-> Download snort application from <https://snort.org/downloads> from binaries section [Snort 2.9.20 Installer.x64.exe](#)

Step-2-> Download rules tar file from the registered section [snortrules-snapshot-29200.tar.gz](#).


Step-3-> Install snort.exe file




Step-4-> Unzip and extract the .tar file



Step-5-> Copy rules and preproc_rules folder from the extracted folder and paste it to the Snort folder in C:\Snort thus replacing the existing folder in the same.



Step-6-> Download and install notepad++ from <https://notepad-plus-plus.org/downloads/> and open the snort.conf file from C:\Snort\etc in the same.



Step-7-> Search for blacklist from the snort folder in C:\Snort, open in notepad++, do save as and name it to whitelist.rules now you will have files with name blacklist.rules and whitelist.rules in the folder.

Step-8-> Edit the config file in the following lines

```
44      # Setup the network addresses you i
45      ipvar HOME_NET 192.168.0.0/24
46
47      # Set up the external network address
48      ipvar EXTERNAL_NET !$HOME_NET
49
104     var RULE_PATH C:\Snort\rules|
105     # var SO_RULE_PATH ../so_rules
106     var PREPROC_RULE_PATH C:\Snort\preproc_rules
107
112     # Set the absolute path appropriately
113     var WHITE_LIST_PATH C:\Snort\rules
114     var BLACK_LIST_PATH C:\Snort\rules
115
```

```

185      #
186      config logdir: C:\Snort\log
187
188
189      # path to dynamic preprocessor libraries
190      dynamicpreprocessor directory C:\Snort\lib\snort_dynamicpreprocessor
191
192      # path to base preprocessor engine
193      dynamicengine C:\Snort\lib\snort_dynamicengine\sf_engine.dll
194
195      # path to dynamic rules libraries
196      # dynamicdetection directory /usr/local/lib/snort_dynamicrules
197
198
199      # Does nothing in IDS mode
200      # preprocessor normalize_ip4
201      # preprocessor normalize_tcp: ips ecn stream
202      # preprocessor normalize_icmp4
203      # preprocessor normalize_ip6
204      # preprocessor normalize_icmp6
205
206
207
208
209
210
211
212
213
214
215
216
217
218
219
220
221
222
223
224
225
226
227
228
229
230
231
232
233
234
235
236
237
238
239
240
241
242
243
244
245
246
247
248
249
250
251
252
253
254
255
256
257
258
259
260
261
262
263
264
265
266
267
268
269
270

```

Step-9-> From line no. 546 to 651 replace forward slash to backslash of the path given.

Step-10-> Open local.rules from C:\Snort\rules in notepad++ and add the following lines of code in it.

```

21
22  alert icmp any any -> any any (msg:"Testing ICMP"; sid:1000001;)
23
24  alert udp any any -> any any (msg:"Testing UDP"; sid:1000002;)
25
26  alert tcp any any -> any any (msg:"Testing TCP"; sid:1000003;)

```

Step-11-> Open command line in the path C:\Snort\bin and run the command snort -W

```
C:\Snort\bin>snort -W
--> Snort! <-
o...`-> Version 2.9.20-WIN64 GRE (Build 82)
By The Snort Team & The Snort Project. http://www.snort.org/contact#team
Copyright (C) 2014-2022 Cisco and/or its affiliates. All rights reserved.
Copyright (C) 1998-2013 Sourcefire, Inc., et al.
Using PCRE version: 8.10 2010-06-25
Using ZLIB version: 1.2.11

Index Physical Address          IP Address       Device Name   Description
----- -----
 1  00:00:00:00:00:00           disabled        \Device\NPF_{70B0A322-CB1B-4D42-AAE1-6053BA230C02}    WAN Miniport (Network Monitor)
 2  00:00:00:00:00:00           disabled        \Device\NPF_{2D41C9C4-C45F-4414-9246-4E57EGAF3779}  WAN Miniport (IP)
 3  00:00:00:00:00:00           disabled        \Device\NPF_{A02C2DF4-6A26-47A8-841F-EFA013BD4B4D}  WAN Miniport (IP)
 4  E8:D0:FC:E1:EF:58          169.254.39.155  \Device\NPF_{5F24DE34-CF9E-468C-91A3-7FCF44A08EGF}  Bluetooth Device (Personal Area Network)
 5  E8:D0:FC:E1:EF:4F          192.168.0.102   \Device\NPF_{BDD3EE5E-90F0-45E9-B649-C54C390F5069}  Qualcomm Atheros QCA9377 Wireless Network Adapter
 6  FA:D0:FC:E1:EF:4F          169.254.184.44  \Device\NPF_{B857CDB5-AC63-4D2C-B689-73EB8461E90B}  Microsoft Wi-Fi Direct Virtual Adapter #4
 7  EA:D0:FC:E1:EF:4F          169.254.7.234   \Device\NPF_{FFD03A86-D792-4D13-A9AF-599631CBFBAA}  Microsoft Wi-Fi Direct Virtual Adapter #3
 8  00:00:00:00:00:00           0000:0000:0000:0000:0000:0000 \Device\NPF_Loopback   Adapter for loopback traffic capture


C:\Snort\bin>

```

Step-12-> Run the command: snort -i 5 -c C:\Snort\etc\snort.conf -A console

b. Demonstrate Extract an Executable from a PCAP


Step-1-> Download packet capture file from :
<https://www.cloudshark.org/captures/a9472fbe700a>




Step-2-> Open the packet capture file in Wireshark by dragging and dropping the capture file into the Wireshark interface.



Step-3-> Click on File->Export Objects-> HTTP




Step-4-> Select image/png and click on Save. Select Save Location.



c. Demonstrate a practical for Exploring DNS Traffic


Step-1->Open Wireshark, select Wi-Fi and click on Capture to capture the DNS Traffic




Learn
User's Guide · Wiki · Questions and Answers · Mailing Lists · Sharkfest · Wireshark Discord · Donate
You are running Wireshark 4.0.4 (<4.0.4-gea14d46dd9ca). You receive automatic updates.

Ready to load or capture

Step-2->Click start



Step-3->Following window will open



Step-4->Search for dns in the search bar

Step-5->Search for some websites and check them on Wireshark, it will give you the following output Step-6->Double click on that and open the DNS arrow and analyse it

Practical 4

b. Exploring Processes, Threads, Handles, and Windows Registry

Download Windows SysInternals Suite.

<https://technet.microsoft.com/en-us/sysinternals/bb842062.aspx>

After the download is completed, extract the files from the folder.

Open folder -> Open proexp.exe.

The Microsoft Edge process can be terminated in the Process Explorer. Right-click the selected process and select Kill Process. Click OK to continue.

Open a Command Prompt, Drag the Find Window's Process icon into the Command Prompt, then seen in Process Explorer.

In cmd type

ping -t www.cisco.com

these ping.exe display in process explorer.

to check malicious content, right-click conhost.exe -> Check VirusTotal -> Yes -> Ok

Step 1: Exploring Threads and Handles

right-click conhost.exe -> Properties -> Threads = view active thread

Step 2: Explore handles.

View > select Lower Pane View > Handles = view handle associated with conhost.exe process

Exploring Windows Registry

click Start -> Registry Editor -> Yes

HKEY_CURRENT_USER> Software> Sysinternals> Process Explorer.

Scroll down to locate the key EulaAccepted = currently value is (1)

double click and change value to 0 (indicates that the EULA was not accepted)

open proexp.exe

Result - EULA licence seen

Name	Type	Data
DbgHelpPath	REG_SZ	C:\Windows\SYSTEM32\dbghelp.dll
DefaultDllPropP...	REG_DWORD	0x00000000 (0)
DefaultProcProp...	REG_DWORD	0x00000006 (6)
DefaultSysInfoP...	REG_DWORD	0x00000000 (0)
Divider	REG_BINARY	ae 47 c1 7a 14 ae ef 3f
DllColumnCount	REG_DWORD	0x00000004 (4)
DllPropWindow...	REG_BINARY	2c 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00
DllSortColumn...	REG_DWORD	0x00000000 (0)
DllSortDirection	REG_DWORD	0x00000001 (1)
ETWStandardUs...	REG_DWORD	0x00000000 (0)
EulaAccepted	REG_DWORD	0x00000000 (0)
FindWindowpla...	REG_BINARY	2c 00 00 00 00 00 00 01 00 00 00 ff ff ff ff ff
FormatlloBytes	REG_DWORD	0x00000001 (1)
GpuNodeUsage...	REG_DWORD	0x00000001 (1)
GpuNodeUsage...	REG_DWORD	0x00000000 (0)
HandleColumn...	REG_DWORD	0x00000002 (2)
HandleSortColu...	REG_DWORD	0x00000000 (0)
HandleSortDirec...	REG_DWORD	0x00000001 (1)
HideWhenMinim...	REG_DWORD	0x00000000 (0)
HighlightDelProc	REG_DWORD	0x00000001 (1)
HighlightDuration	REG_DWORD	0x0000003e8 (1000)
HighlightImmer...	REG_DWORD	0x00000001 (1)
HighlightJobs	REG_DWORD	0x00000000 (0)
HighlightNetPro...	REG_DWORD	0x00000000 (0)
HighlightNewPr...	REG_DWORD	0x00000001 (1)
HighlightOwnPr...	REG_DWORD	0x00000001 (1)

Practical 5

Perform a practical to Attack on a mySQL Database by using PCAP File.

Part 1: Open Wireshark and load the PCAP file.

Start the CyberOps Workstation VM.

Click Applications > CyberOPS > Wireshark

next go to File >open> /home/analyst/ > lab.support.files > SQL_Lab.pcap

Part 2: View the SQL Injection Attack.

right-click line 13 and select Follow > HTTP Stream.

Find field, enter 1=1. Click Find Next.

Clear display filter

Part 3: The SQL Injection Attack continues...

right-click line 19 and select Follow > HTTP Stream.

Find field, enter 1=1. Click Find Next.

Clear display filter

Part 4: The SQL Injection Attack provides system information.

right-click line 22 and select Follow > HTTP Stream.

Find field, enter 1=1. Click Find Next.

Clear display filter

Part 5: The SQL Injection Attack and Table Information.

right-click line 25 and select Follow > HTTP Stream.

Find field, enter 1=1. Click Find Next.

Clear display filter

Part 6: The SQL Injection Attack Concludes.

right-click line 28 and select Follow > HTTP Stream.

Find field, enter 1=1. Click Find Next.

Which user has the password hash of 8d3533d75ae2c3966d7e0d4fcc69216b?

first search password on find after these visible then go to

<https://crackstation.net/> and enter the password and hash the password

Free Password Hash Cracker

Enter up to 20 non-salted hashes, one per line:

8d3533d75ae2c3966d7e0d4fcc69216b

I'm not a robot

Crack Hashes

Supports: LM, NTLM, md2, md4, md5, md5(md5_hex), md5-half, sha1, sha224, sha256, sha384, sha512, ripeMD160, whirlpool, MySQL 4.1+ (sha1(sh1_bin)), QubesV3.1BackupDefaults

Hash	Type	Result
8d3533d75ae2c3966d7e0d4fcc69216b	md5	charley

Color Codes: Exact match, Partial match, Not found.

Practical 6

Create your own syslog Server

Step 1->Run sudo apt-get install openssh-server command to install ssh

```
ubuntu@ubuntu:~$ sudo apt-get install openssh-server
[sudo] password for ubuntu:
Reading package lists... Done
Building dependency tree
Reading state information... Done
openssh-server is already the newest version (1:8.2p1-0.5).
The following package was automatically installed and
longer required:
  libfprint-2-tod1
Use 'sudo apt autoremove' to remove it.
0 upgraded, 0 newly installed, 0 to remove and 32 not
d.
```

Step 2->Run sudo apt-get install net-tools to install net-tools


```
ubuntu@ubuntu:~$ sudo apt-get install net-tools
Reading package lists... Done
Building dependency tree
Reading state information... Done
net-tools is already the newest version (1.60+gitd88e-1ubuntu1).
The following package was automatically installed and
longer required:
  libfprint-2-tod1
Use 'sudo apt autoremove' to remove it.
0 upgraded, 0 newly installed, 0 to remove and 3
d
```

Step 3->Run sudo service rsyslog restart to restart rsyslog service

```
ubuntu@ubuntu:/etc$ sudo service rsyslog restart
ubuntu@ubuntu:/etc$
```

Step 4->Edit rsyslog config using sudo nano /etc/rsyslog.conf command and uncomment the lines - module(load="imudp") input(type="imudp" port="514") module(load="imtcp") input(type="imtcp" port="514") Press Ctrl+X & Enter to save file

```
ubuntu@ubuntu:/etc$ sudo nano /etc/rsyslog.conf
```



```
GNU nano 4.8      /etc/rsyslog.conf
# /etc/rsyslog.conf configuration file for rsyslog
#
# For more information install rsyslog-doc and see
# /usr/share/doc/rsyslog-doc/html/configuration/index.html
#
# Default logging rules can be found in /etc/rsyslog.d/*.conf

#####
##### MODULES #####
#####

module(load="imuxsock") # provides support for local log files
#module(load="immark")  # provides --MARK-- message capability

# provides UDP syslog reception
module(load="imudp")
input(type="imudp" port="514")

# provides TCP syslog reception
[ Read 61 lines ]
^G Get Help  ^O Write Out  ^W Where Is  ^K Cut Text
^X Exit      ^R Read File  ^\ Replace   ^U Paste Text
```

Step 5-> Restart rsyslog service with **sudo service rsyslog restart** command

```
ubuntu@ubuntu:/etc$ sudo nano /etc/rsyslog.conf
ubuntu@ubuntu:/etc$ sudo service rsyslog restart
```

Step 6-> Check status of rsyslog service with **sudo service rsyslog status** command

```
ubuntu@ubuntu:/etc$ sudo service rsyslog status
● rsyslog.service - System Logging Service
  Loaded: loaded (/lib/systemd/system/rsyslog.service; enabled; reloaded)
  Active: active (running) since Fri 2023-04-21 04:16:20 UTC
    Docs: man:rsyslogd(8)
           https://www.rsyslog.com/doc/
   Main PID: 14060 (rsyslogd)
     Tasks: 9 (limit: 3449)
    Memory: 1.2M
   CGroup: /system.slice/rsyslog.service
           └─14060 /usr/sbin/rsyslogd -n -iNONE

Apr 21 04:16:20 ubuntu systemd[1]: Starting System Logging Service...
Apr 21 04:16:20 ubuntu systemd[1]: Started System Logging Service.
Apr 21 04:16:20 ubuntu rsyslogd[14060]: imuxsock: Acq...
Apr 21 04:16:20 ubuntu rsyslogd[14060]: rsyslogd's gr...
Apr 21 04:16:20 ubuntu rsyslogd[14060]: rsyslogd's us...
Apr 21 04:16:20 ubuntu rsyslogd[14060]: [origin softw...
lines 1-18/18 (END)
```

Practical 7

**Configure your Linux system to send syslog messages to a syslog server
and read them**

Step 1: Configuring the Syslog Server (server-syslog)

```
sudo nano /etc/sysconfig/syslog
sudo apt update
sudo apt install rsyslog
sudo nano /etc/rsyslog.conf
remove # following line
module(load="imudp") # UDP listener
input(type="imudp" port="514")
add line = *.debug    /var/log/messages
sudo systemctl restart rsyslog
sudo netstat -anu | grep 514
```

Step 2: Configuring the Syslog Client (oer-host)

```
sudo nano /etc/hosts
add line = 10.10.10.1 server-syslog.domain.com server-syslog loghost
sudo nano /etc/rsyslog.conf
add line = *.debug @loghost
sudo systemctl restart rsyslog
Step 3: Testing the Configuration
logger -p user.debug "Test message from client"
sudo cat /var/log/messages | grep "Test message from client"
```

Practical 8

Install and Run Splunk on Linux

- 1) sudo aptget update
- 2) sudo aptget upgrade
- 3) sudo apt-get install net-tools wget openssh-server
- 4) sudo systemctl restart ssh
- 5) sudo nano /etc/rsyslog.conf
- 6) module(load="imuxsock")
- 7) module(load="imklog")
- 8) module(load="imudp") input(type="imudp" port="514")
- 9) module(load="imtcp") input(type="imtcp" port="514")
- 11) sudo systemctl restart rsyslog

if we get

```
$ sudo systemctl restart rsyslog
```

Failed to restart rsyslog.service: Unit rsyslog.service not found.

then use following command:

```
$ sudo apt-get purge rsyslog  
$ sudo apt-get install rsyslog  
$ sudo systemctl restart rsyslog
```

10)sudo systemctl status rsyslog

11)sudo groupadd splunkdemo

12)sudo useradd -d /opt/splunk -m -g splunk splunkdemo

13) Now go to <https://www.splunk.com/> . Register and Click on Free Splunk at top right of

page. Click on Free Trials and Downloads page link. On the next screen Click on Get My Free Trial

under Splunk Enterprise. Next Choose Linux and click on Download for .tgz. Stop the download and

click

14)Stop the download and click on Dowload via Command Line (wget) and copy the command.

15)su -

Password:

16)su - splunkdemo

if error come like su: user splunkdemo does not exist or the user entry does not contain all the required fields

command

```
cat /etc/passwd | grep splunkdemo  
sudo groupadd splunk  
sudo useradd -d /opt/splunk -m -g splunk -s /bin/bash splunkdemo  
sudo passwd splunkdemo  
cat /etc/passwd | grep splunkdemo  
cat /etc/group | grep splunk
```

```
cat /etc/passwd | grep splunkdemo  
sudo groupadd splunk  
sudo useradd -d /opt/splunk -m -g splunk -s /bin/bash splunkdemo  
sudo passwd splunkdemo  
cat /etc/passwd | grep splunkdemo  
cat /etc/group | grep splunk
```

17)su - splunkdemo

18) paste the command of tgz file(cmd that start wget step 14)

19) tar xvzf splunk-9.4.0-6b4ebe426ca6-linux-amd64.tgz -C /opt (edit these command with your tgz file version)

20) access root using su – (write root password (kali))

21) chown -R splunkdemo: /opt/splunk/

22) cd /opt/splunk/bin


23) ./splunk start

24) ./splunk start --accept-license

after these we get "The Splunk web interface is at http://kali:8000"

25) interface on http://kali:8000.

go these link if your step is correct then get login screen else any error in your steps.



```
File Actions Edit View Help  
[root@kali ~]# sudo apt-get install net-tools wget openssh-server  
Reading package lists... Done  
Building dependency tree... Done  
Reading state information... Done  
net-tools is already the newest version (2.10-0.1).  
net-tools is not manually installed.  
wget is already the newest version (1.21.3-1+b2).  
openssh-server is already the newest version (1:9.2p1-2).  
The following packages were automatically installed and are no longer required:  
bluez-firmware catfish cryptsetup-run faraday-client fastjar  
firmware-atheros firmware-brcm80211 firmware-intel-sound firmware-iwlwifi  
firmware-libertas firmware-realtek firmware-sof-signed  
firmware-alienware-firmware libfirmware-2.12.1-firmware-robotnslab  
gir1.2-xfconf-0 gstreamer1.0-pulseaudio jarwrapper kali-wallpapers-2021.4  
libarmadillo10 libatk1.0-data libavfilter7 libavformat58 libavresample4  
libcboro libcfitsio9 libcharls2 libdap27 libdapclient6v5 libepsilon1  
libevey libexporter-tiny-perl libflac8 libfmt7 libgdal28  
libgdk-pixbuf-xlib-2.0-0 libgdk-pixbuf2.0-0 libgeos-3.9.0 libgs9-common  
libhttp-server-simple-perl libidn11 libigdgmm11 liblist-moreutils-perl  
libjansson4 libkeyringsafe1 liblouis0 liblouisdb1 liblouisdb3  
libnetcdf18 libnginx-mod-http-geoip libnginx-mod-http-image-filter  
libnginx-mod-stream-geoip libnfts-3g883 libodbc1 libodbcocr2 libopts25  
libperl5.32 libpppler102 libpostproc55 libproj19 libprotobuf23  
libpython3.9 libpython3.9-dev libpython3.9-minimal libpython3.9-stdlib  
libqhull18.0 librest-0.7-0 libsrt1.4-gnutls libsswscale5 libtbb2 libtiff5  
libwputil12 libxara4 maltego-nginx-common nmap-ncscope odbcinst  
odbcinstdebian2 openjdk-11-jre perl-modules-5.32  
python-mpltoolkits.basemap-data python-pastedeploy-tpl  
python3-deprecation python3-editor python3-exif python3-fastjsonschema  
python3-gevent python3-gevent-websocket python3-iputils  
python3-ipython-genutils python3-jupyter-core python3-m2crypto  
python3-matplotlib python3-nbformat python3-ntlm-auth python3-parameterized  
python3-pbr python3-plytify python3-pylint python3-pyproj python3-pyshp  
python3-requests-ntlm python3-singledispatch python3-speaklater
```

Practical 9

Install and Configure ELK on Linux

- 1.sudo apt-get update
password :kali
- 2.sudo apt-get upgrade
- 3.sudo apt-get install default-jdk default-jre
- 4.java -version
- 5.update-alternatives --config java
- 6.sudo nano /etc/environment
set the path :
JAVA_HOME="/usr/lib/jvm/java-23-openjdk-amd64/bin/java"
#in case if the above cmd isn't working then to temporary set the path use this command because
sudo nano isn't working :
export JAVA_HOME=/usr/lib/jvm/java-23-openjdk-amd64/bin/java
- 7.source /etc/environment
- 8.sudo apt-get install apt-transport-https
- 9.wget -qO - https://artifacts.elastic.co/GPG-KEY-elasticsearch | sudo gpg --dearmor -o
/usr/share/keyrings/elasticsearch-keyring.gpg
- 10.echo "deb [signed-by=/usr/share/keyrings/elasticsearch-keyring.gpg]
https://artifacts.elastic.co/packages/8.x/apt stable main" | sudo tee /etc/apt/sources.list.d/elastic-
8.x.list
- 11.sudo apt-get update
- 12.sudo apt-get install elasticsearch
- 13.sudo systemctl daemon-reload
- 14.sudo systemctl enable elasticsearch.service
- 15.sudo systemctl start elasticsearch.service
- 16.sudo systemctl status elasticsearch.service
- 17.sudo nano /etc/elasticsearch/elasticsearch.yml
network.host: 0.0.0.0
discovery.seed_hosts:[]
xpack.security.enabled: false
- 18.sudo systemctl restart elasticsearch.service
- 19.sudo nano /etc/elasticsearch/jvm.options
edit and change to:
-Xms512m
-Xmx512m
- 20.curl -X GET "localhost:9200"
- 21.ifconfig
copy the inet 10.0.2.15 append :9200 to it (10.0.2.15:9200) run this in the browser
- 22.sudo apt-get install logstash

```
sudo systemctl start logstash  
sudo systemctl enable logstash  
sudo systemctl status logstash
```

23.sudo apt-get install kibana
sudo systemctl start kibana
sudo systemctl enable kibana
sudo systemctl status kibana
24.sudo nano /etc/kibana/kibana.yml
make changes as:
server.port:5601
server.host: "0.0.0.0"
elasticsearch.hosts: ["http://localhost:9200"]
25.sudo systemctl restart kibana

26.sudo apt-get install filebeat
27.sudo nano /etc/filebeat/filebeat.yml
make changes as shown in the screen shots
28.sudo filebeat setup --index-management -E output.logstash.enabled=false -E 'output.elasticsearch.hosts=["0.0.0.0:9200"]'
29.sudo systemctl start filebeat
30.sudo systemctl enable filebeat
31.curl -X GET http://10.0.2.15:9200/_cat/indices?v
32.In the browser run the following address:
http://10.0.2.15:9200/_cat/indices?v
<http://10.0.2.15:5601>

