

**BIG  
DATA  
ANALYTICS**



JNAN VIKAS MANDAL'S (**Linguistic Minority**)  
Mohanlal Raichand Mehta College of Commerce  
Diwali Maa College of Science  
Dr. R.T. Doshi College of Computer Science  
Plot No.9, Sector -19, Airoli, Navi Mumbai  
Permanently Unaided, Affiliated to University of Mumbai  
NAAC Re-Accredited Grade "A" Grade (3.33 )

### **DEPARTMENT OF M.Sc. (INFORMATION TECHNOLOGY)**

#### **CERTIFICATE**

This is to certify that \_\_\_\_\_ bearing seat no. \_\_\_\_\_ has done the practical work in the subject **BIG DATA ANALYSIS** of Semester **II** Practical Examination during the Academic Year 2024-25 under the guidance of **Dr. Sunita Joshi** being the requirement for the fulfilment of the curriculum of Degree of Master of Science in Information Technology Part I under University of Mumbai.

**Dr. Sunita Joshi**  
**Subject In-Charge**  
**M.Sc. IT Co Ordinator**

**External Examiner**  
**Signature**

**DATE:**

**PLACE: AIROLI**

Sr No	Practicals	Date	Sign
1	<b>Practical No 1</b>		
A.	K-Means Clustering: Clustering algorithms for unsupervised classification. Read a datafile grades_km_input.csv and apply k-means clustering. Plot the cluster data using R visualizations.		
B.	Apriori Algorithm (PBL): Implement Apriori Algorithm Recommending grocery items to a customer that is most frequently bought together, given a data set of transactions by customers of a store, using Ariory algorithm using Market_Basket_Optimisation.csv file		
2	<b>Practical No 2</b>		
A.	Regression Model: Import data from web storage – binary.csv. Name the dataset and do Logistic Regression to find out relation between variables that are affecting the admission of a student in an institute based on his or her GRE score, GPA obtained and rank of the student. Also check the model is fit or not.		
B.	MULTIPLE REGRESSION MODEL: Apply multiple regressions, if data have a continuous independent variable. Apply on above dataset – binary.csv.		
3.	<b>Practical No.3</b>		
A.	Decision Tree: Implement Decision Tree classification technique using Social_Network_Ads.csv dataset.		
B.	SVM Classification: Implement SVM Classification technique using Social_Network_Ads.csv dataset. Evaluate the performance of classifier.		
4	<b>Practical No.4</b>		
A.	Naïve Bayes Classification: Implement Naïve Bayes Classification technique using Social_Network_Ads.csv dataset. Evaluate the performance of classifier.		
B.	Text Analysis (PBL): Find the confusion matrix to find restaurant review based of sentiment analysis of Natural Language processing. Use Resaurent reviews.tsv file for your study.		

<b>5</b>	<b>Practical No.5</b>		
	Comparative Study of various machine learning models (Newly added): Take the inbuilt data file: iris and perform classification on that data using various classification models – Decision Tree, K Nearest Neighbour and Support Vector Machine. Find the confusion matrix for all three models and evaluate them by finding their accuracy. Find the algorithm which performs best on the given data file, out of all these three models		
<b>6</b>	<b>Practical No.6</b>		
	Install, configure and run Hadoop and HDFS and explore HDFS on Windows		
<b>7</b>	<b>Practical No.7</b>		
	Implement word count / frequency programs using MapReduce.		
<b>8</b>	<b>Practical No.8</b>		
	Implement an application that stores big data in Hbase / MongoDB and manipulate it using R / Python		

# Practical No. 1

A) Aim: To implement K-Means Clustering: Clustering algorithms for unsupervised classification. Read a datafile grades\_km\_input.csv and apply k-means clustering. Plot the cluster data using R visualizations.

## Source Code:

```
install.packages("plyr")
install.packages("ggplot2")
install.packages("cluster")
install.packages("lattice")
install.packages("grid")
install.packages("gridExtra")
library(plyr)
library(ggplot2)
library(cluster)
library(lattice)
library(grid)
library(gridExtra)

grade_input <- read.csv("C:/Users/HP/Downloads/grades_km_input.csv")

kmdata <- as.matrix(grade_input[, c("English", "Math", "Science")])

head(kmdata, 10)

wss <- numeric(15)
for (k in 1:15) {
  set.seed(123) # for reproducibility
  wss[k] <- sum(kmeans(kmdata, centers = k, nstart = 25)$withinss)
}

plot(1:15, wss, type = "b", pch = 19, frame = FALSE,
     xlab = "Number of Clusters (k)",
     ylab = "Within-Cluster Sum of Squares (WSS)",
     main = "Elbow Method for Optimal k")

set.seed(123)
km <- kmeans(kmdata, centers = 3, nstart = 25)

print(km)

df <- grade_input
df$cluster <- as.factor(km$cluster)

centers <- as.data.frame(km$centers)
centers$cluster <- as.factor(1:3)

g1 <- ggplot(data = df, aes(x = English, y = Math, color = cluster)) +
  geom_point(size = 3) +
  geom_point(data = centers, aes(x = English, y = Math), color = "black", size = 10, alpha = 0.3) +
  gtitle("English vs Math") + theme_minimal()
```

```

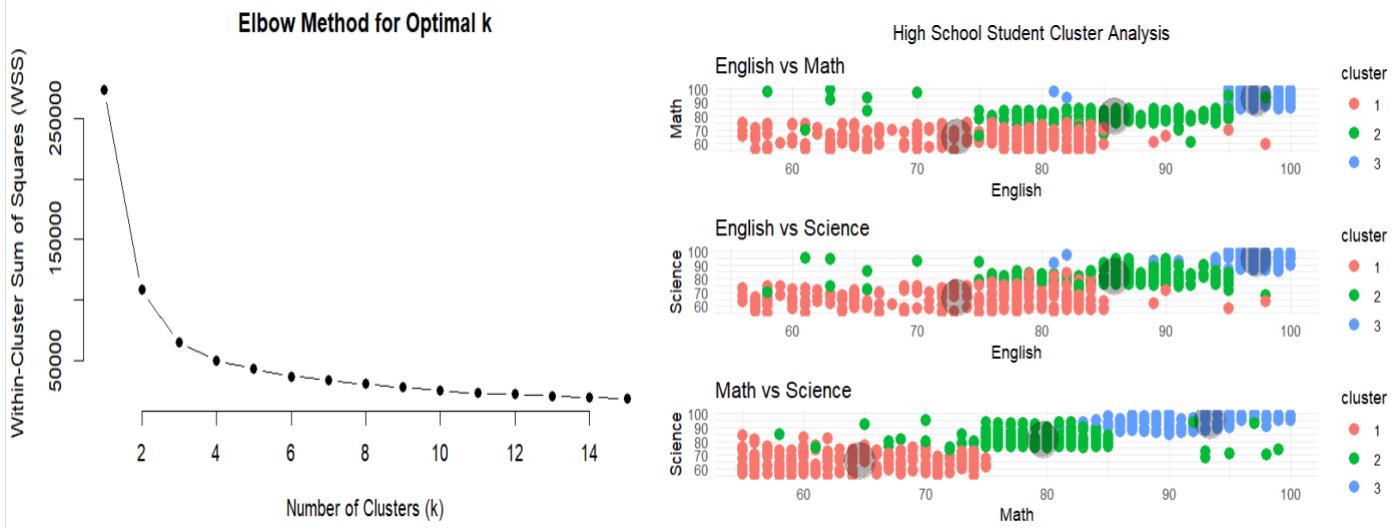
g2 <- ggplot(data = df, aes(x = English, y = Science, color = cluster)) +
  geom_point(size = 3) +
  geom_point(data = centers, aes(x = English, y = Science), color = "black", size = 10, alpha = 0.3) +
  ggttitle("English vs Science") + theme_minimal()

g3 <- ggplot(data = df, aes(x = Math, y = Science, color = cluster)) +
  geom_point(size = 3) +
  geom_point(data = centers, aes(x = Math, y = Science), color = "black", size = 10, alpha = 0.3) +
  ggttitle("Math vs Science") + theme_minimal()

grid.arrange(g1, g2, g3, ncol = 1, top = "High School Student Cluster Analysis")

```

### Output:



## B) Aim: Implement Apriori Algorithm

**Implement Apriori Algorithm** Recommending grocery items to a customer that is most frequently bought together, given a data set of transactions by customers of a store, using Ariory algorithm using Market\_Basket\_Optimisation.csv file.

### Source Code:

```
install.packages("arules")
install.packages("arulesViz")
install.packages("RColorBrewer")

library(arules)
library(arulesViz)
library(RColorBrewer)

data(Groceries)

Groceries
summary(Groceries)
class(Groceries)

rules <- apriori(Groceries, parameter = list(supp = 0.02, conf = 0.2))

summary(rules)

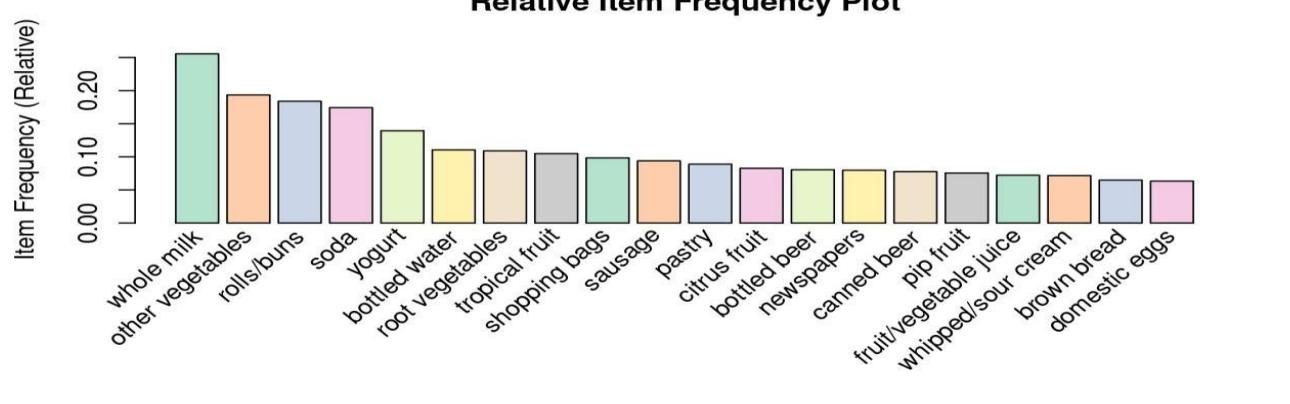
inspect(rules[1:10])

itemFrequencyPlot(Groceries, topN = 20,
                  col = brewer.pal(8, 'Pastel2'),
                  main = 'Relative Item Frequency Plot',
                  type = "relative",
                  ylab = "Item Frequency (Relative)")

itemsets <- apriori(Groceries, parameter = list(minlen = 2, maxlen = 2, support = 0.02, target =
"frequent itemsets"))
summary(itemsets)
inspect(itemsets[1:10])

itemsets_3 <- apriori(Groceries, parameter = list(minlen = 3, maxlen = 3, support = 0.02, target =
"frequent itemsets"))
summary(itemsets_3)
inspect(itemsets_3[1:10])
```

## Output:



```
> inspect(itemsets[1:10])
```

	items	support	count
[1]	{whole milk, frozen vegetables}	0.02043721	201
[2]	{beef, whole milk}	0.02125064	209
[3]	{whole milk, curd}	0.02613116	257
[4]	{pork, other vegetables}	0.02165735	213
[5]	{pork, whole milk}	0.02216573	218
[6]	{frankfurter, whole milk}	0.02053889	202
[7]	{whole milk, bottled beer}	0.02043721	201
[8]	{whole milk, brown bread}	0.02521607	248
[9]	{whole milk, margarine}	0.02419929	238
[10]	{other vegetables, butter}	0.02003050	197

## Practical No. 2

### A) Aim: To implement Regression Model

Import data from web storage – binary.csv. Name the dataset and do Logistic Regression to find out relation between variables that are affecting the admission of a student in an institute based on his or her GRE score, GPA obtained and rank of the student. Also check the model is fit or not.

### Source Code:

```
college <- read.csv("C:/Users/HP/Documents/binary.csv")
head(college)
nrow(college)

install.packages("caTools")
library(caTools)

set.seed(123) # for reproducibility
split <- sample.split(college$admit, SplitRatio = 0.75)
training_reg <- subset(college, split == TRUE)
test_reg <- subset(college, split == FALSE)

fit_logistic_model <- glm(admit ~ ., data = training_reg, family = "binomial")

predict_reg <- predict(fit_logistic_model, test_reg, type = "response")
predict_reg

cdplot(as.factor(admit) ~ gpa, data = college)
cdplot(as.factor(admit) ~ gre, data = college)
cdplot(as.factor(admit) ~ rank, data = college)

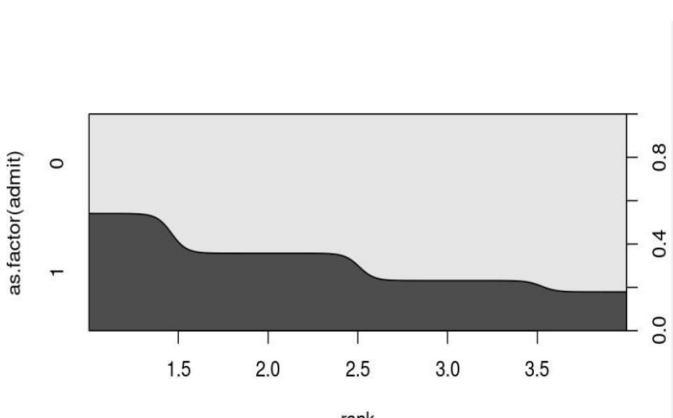
predict_reg <- ifelse(predict_reg > 0.5, 1, 0)
predict_reg

table(test_reg$admit, predict_reg)
```

### Output:

```
> table(test_reg$admit, predict_reg)

predict_reg
  0  1
0 53 12
1 26  9
>
```



## B) Aim: MULTIPLE REGRESSION MODEL:

Apply multiple regressions, if data have a continuous independent variable. Apply on above dataset – binary.csv.

Source Code:

```
college <- read.csv("C:/Users/HP/OneDrive/Documents/binary.csv")
head(college)
nrow(college)

# Load library for splitting data
install.packages("caTools")
library(caTools)

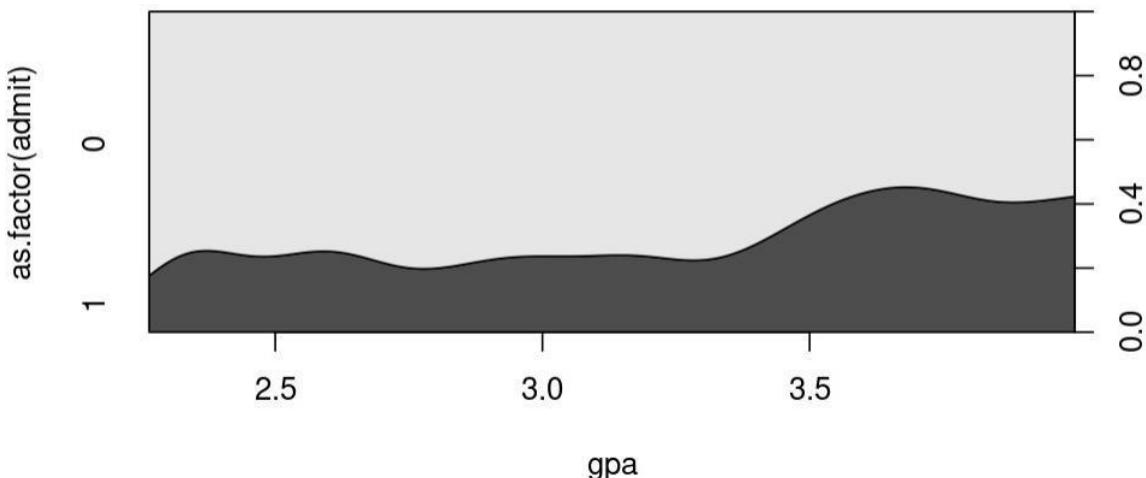
set.seed(123)
split <- sample.split(college$admit, SplitRatio = 0.75)
training_reg <- subset(college, split == TRUE)
test_reg <- subset(college, split == FALSE)

fit_MRegressor_model <- lm(formula = admit ~ gre + gpa + rank, data = training_reg)
summary(fit_MRegressor_model)

predict_reg <- predict(fit_MRegressor_model, newdata = test_reg)
print(predict_reg)

cdplot(as.factor(admit) ~ gpa, data = college)
cdplot(as.factor(admit) ~ gre, data = college)
cdplot(as.factor(admit) ~ rank, data = college)
```

Output:



# Practical No. 3

## A) Aim: Decision Tree: Implement Decision Tree classification technique using Social\_Network\_Ads.csv dataset.

```
# Importing the dataset
dataset <- read.csv('C:\\\\Users\\\\HP\\\\Downloads\\\\Social_Network_Ads.csv')
dataset <- dataset[3:5]

# Encoding target variable as factor
dataset$Purchased <- factor(dataset$Purchased, levels = c(0, 1))

# Splitting dataset into Training and Test set
install.packages('caTools')
library(caTools)

set.seed(123)
split <- sample.split(dataset$Purchased, SplitRatio = 0.75)
training_set <- subset(dataset, split == TRUE)
test_set <- subset(dataset, split == FALSE)

# Feature Scaling
training_set[, 1:2] <- scale(training_set[, 1:2])
test_set[, 1:2] <- scale(test_set[, 1:2])

install.packages('rpart')
library(rpart)

classifier <- rpart(formula = Purchased ~ ., data = training_set)

y_pred <- predict(classifier, newdata = test_set[, 1:2], type = 'class')

cm <- table(test_set[, 3], y_pred)
print(cm)

install.packages("ggplot2")
library(ggplot2)

visualize <- function(set, title) {
  X1 <- seq(min(set[, 1]) - 1, max(set[, 1]) + 1, by = 0.01)
  X2 <- seq(min(set[, 2]) - 1, max(set[, 2]) + 1, by = 0.01)
  grid_set <- expand.grid(Age = X1, EstimatedSalary = X2)
  y_grid <- predict(classifier, newdata = grid_set, type = 'class')
  plot(set[, -3],
       main = title,
       xlab = 'Age', ylab = 'Estimated Salary',
       xlim = range(X1), ylim = range(X2))
  contour(X1, X2, matrix(as.numeric(y_grid), length(X1), length(X2)), add = TRUE)
  points(grid_set, pch = '.', col = ifelse(y_grid == 1, 'springgreen3', 'tomato'))
  points(set, pch = 21, bg = ifelse(set[, 3] == 1, 'green4', 'red3'))
}
```

```

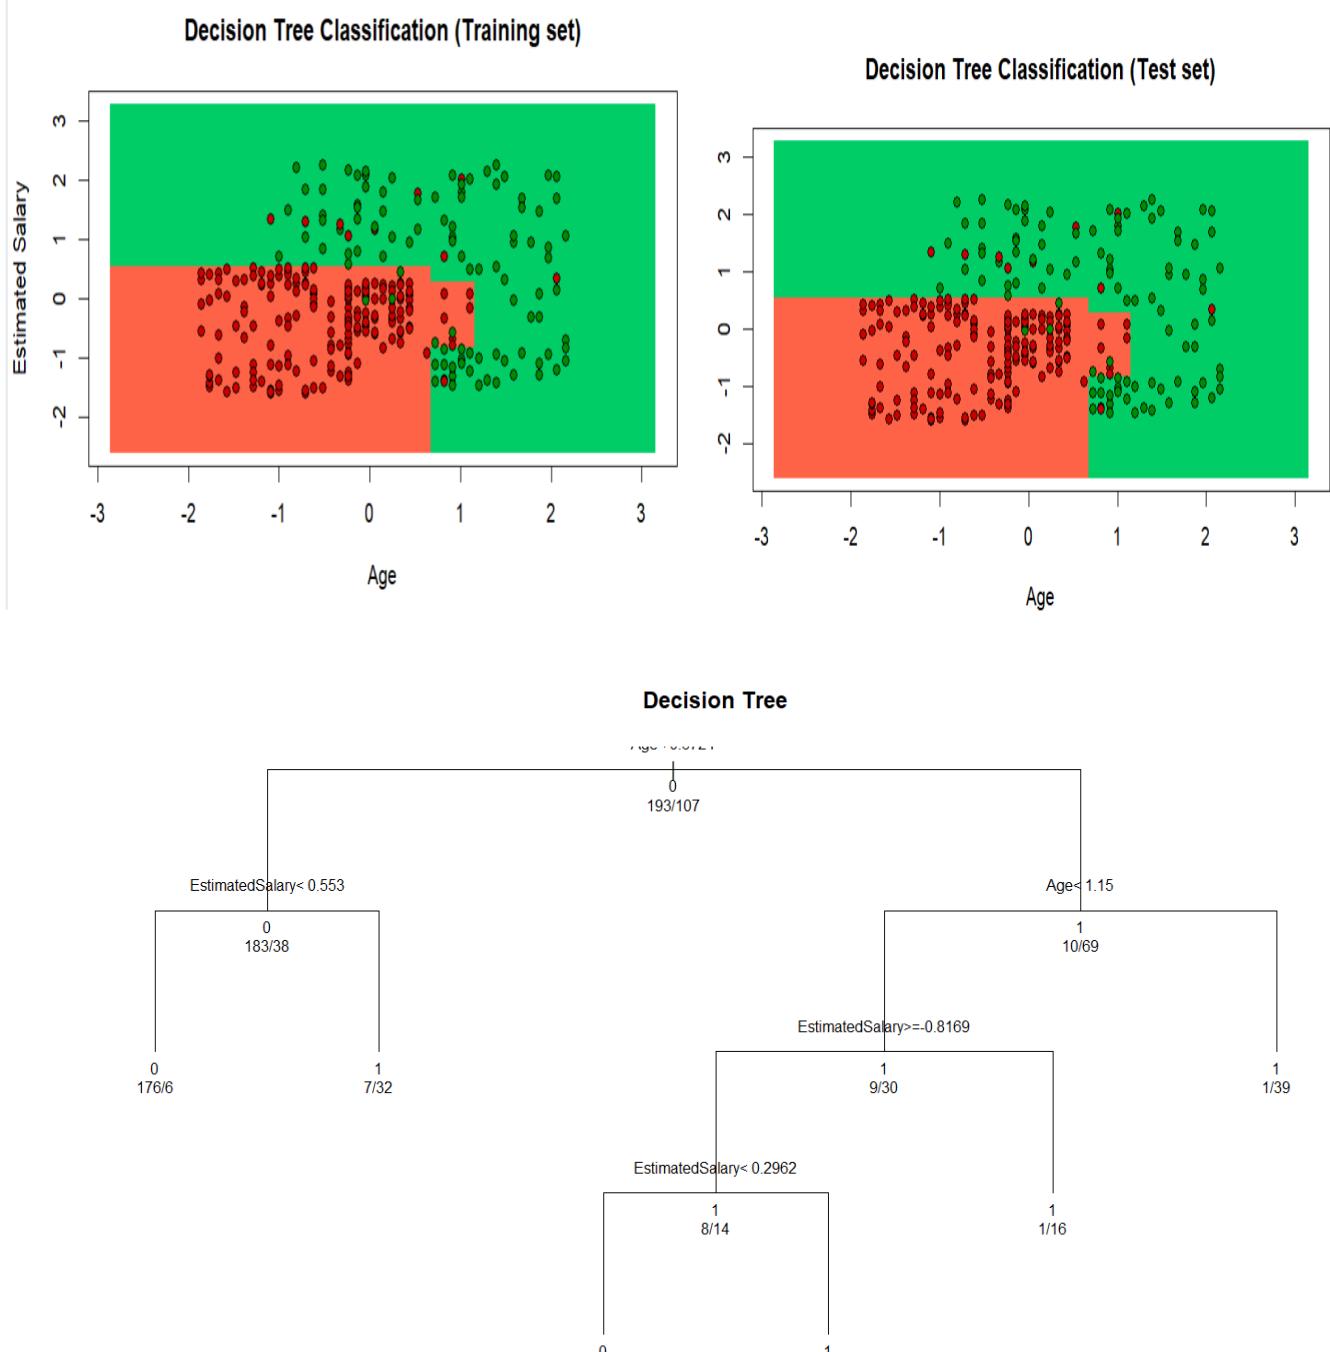
# Visualizing Training set
visualize(training_set, 'Decision Tree Classification (Training set)')

# Visualizing Test set
visualize(test_set, 'Decision Tree Classification (Test set)')

plot(classifier, uniform = TRUE, main = "Decision Tree")
text(classifier, use.n = TRUE, all = TRUE, cex = 0.8)

```

## OUTPUT



**B) Aim: SVM Classification: Implement SVM Classification technique using Social\_Network\_Ads.csv dataset. Evaluate the performance of classifier**

**Source Code:**

```
dataset <- read.csv('C:\\\\Users\\\\HP\\\\Downloads\\\\Social_Network_Ads.csv')
dataset <- dataset[3:5]
dataset$Purchased <- factor(dataset$Purchased, levels = c(0, 1))

install.packages("caTools")
library(caTools)
set.seed(123)
split <- sample.split(dataset$Purchased, SplitRatio = 0.75)
training_set <- subset(dataset, split == TRUE)
test_set <- subset(dataset, split == FALSE)

training_set[, 1:2] <- scale(training_set[, 1:2])
test_set[, 1:2] <- scale(test_set[, 1:2])

install.packages("e1071")
library(e1071)
classifier <- svm(formula = Purchased ~.,
                    data = training_set,
                    type = 'C-classification',
                    kernel = 'linear')

y_pred <- predict(classifier, newdata = test_set[, 1:2])
cm <- table(test_set[, 3], y_pred)
print(cm)

install.packages("ggplot2")
library(ggplot2)

# Visualization function
visualize <- function(set, title) {
  X1 <- seq(min(set[, 1]) - 1, max(set[, 1]) + 1, by = 0.01)
  X2 <- seq(min(set[, 2]) - 1, max(set[, 2]) + 1, by = 0.01)
  grid_set <- expand.grid(Age = X1, EstimatedSalary = X2)
  y_grid <- predict(classifier, newdata = grid_set)

  plot(set[, -3],
       main = title,
       xlab = 'Age', ylab = 'Estimated Salary',
       xlim = range(X1), ylim = range(X2))

  contour(X1, X2, matrix(as.numeric(y_grid), length(X1), length(X2)), add = TRUE)
  points(grid_set, pch = '.', col = ifelse(y_grid == 1, 'springgreen3', 'tomato'))
  points(set, pch = 21, bg = ifelse(set[, 3] == 1, 'green4', 'red3'))
}

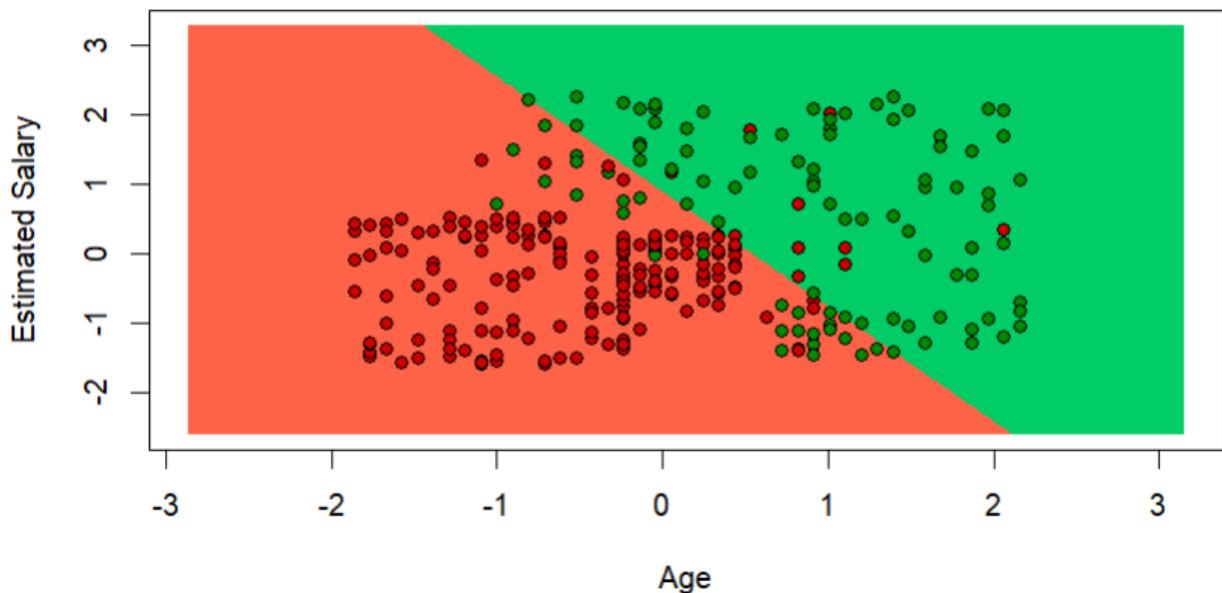
visualize(training_set, 'SVM (Training Set)')

visualize(test_set, 'SVM (Test Set)')
```

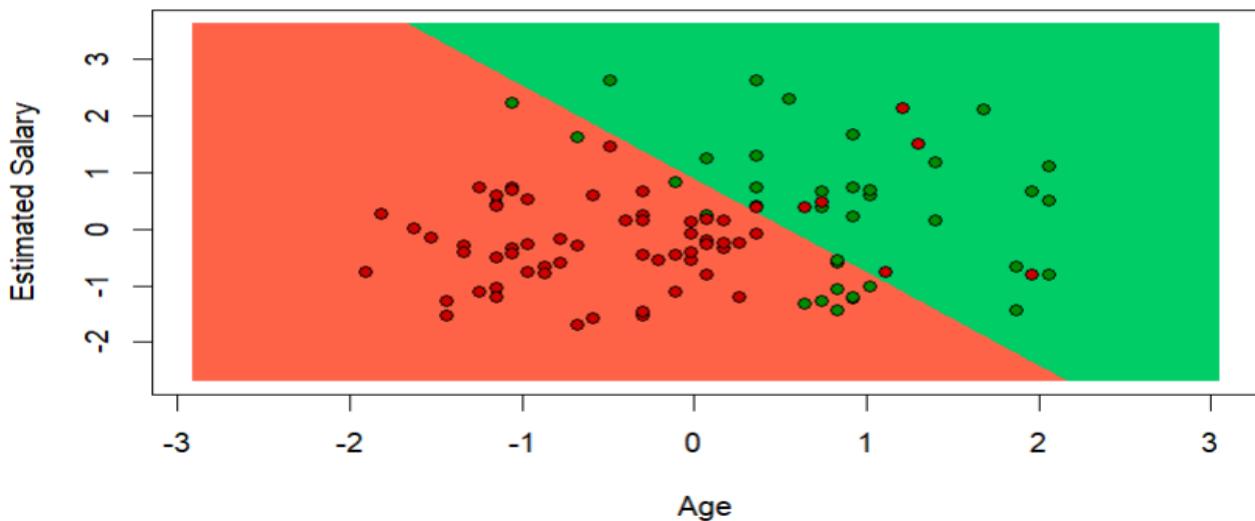
**Output:**

```
> # -----
> # 5. Predict on Test Set
> # -----
> y_pred <- predict(classifier, newdata = test_set[, 1:2])
> cm <- table(test_set[, 3], y_pred)
> print(cm)
y_pred
  0  1
0 57  7
1 13 23
```

**SVM (Training Set)**



**SVM (Test Set)**



## Practical No. 4

**A) Aim: Naïve Bayes Classification: Implement Naïve Bayes Classification technique using Social\_Network\_Ads.csv dataset. Evaluate the performance of classifier.**

### Source Code:

```
# Install latest package
install.packages("naivebayes")
library(naivebayes)
dataset <- read.csv('C:\\\\Users\\\\HP\\\\Downloads\\\\Social_Network_Ads.csv')
dataset <- dataset[3:5]
dataset$Purchased <- factor(dataset$Purchased, levels = c(0, 1))
install.packages("caTools")
library(caTools)
set.seed(123)
split <- sample.split(dataset$Purchased, SplitRatio = 0.75)
training_set <- subset(dataset, split == TRUE)
test_set <- subset(dataset, split == FALSE)
training_set[, 1:2] <- scale(training_set[, 1:2])
test_set[, 1:2] <- scale(test_set[, 1:2])

classifier <- naive_bayes(Purchased ~ ., data = training_set)

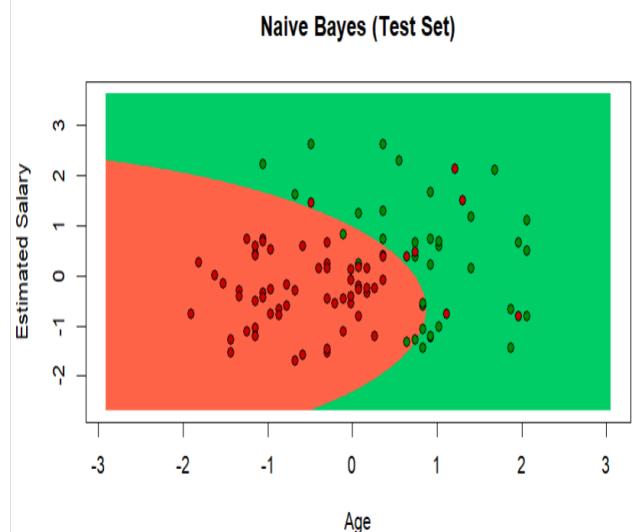
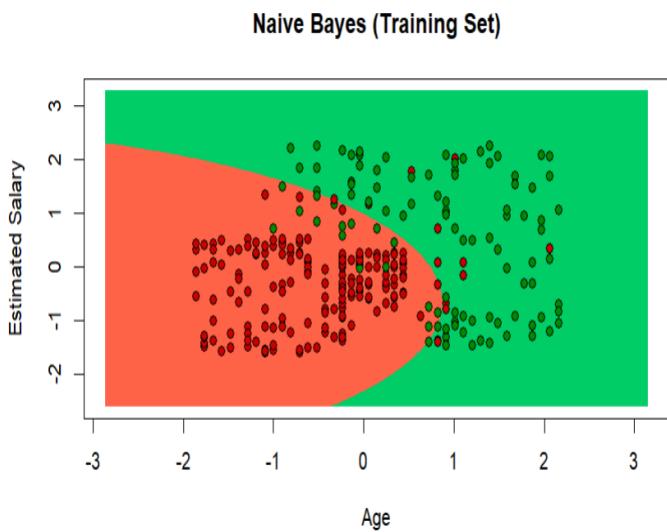
y_pred <- predict(classifier, newdata = test_set)
cm <- table(test_set[, 3], y_pred)
print(cm)
visualize <- function(set, title) {
  X1 <- seq(min(set[, 1]) - 1, max(set[, 1]) + 1, by = 0.01)
  X2 <- seq(min(set[, 2]) - 1, max(set[, 2]) + 1, by = 0.01)
  grid_set <- expand.grid(Age = X1, EstimatedSalary = X2)
  y_grid <- predict(classifier, newdata = grid_set)
  plot(set[, -3],
       main = title,
       xlab = 'Age', ylab = 'Estimated Salary',
       xlim = range(X1), ylim = range(X2))

  contour(X1, X2, matrix(as.numeric(y_grid), length(X1), length(X2)), add = TRUE)
  points(grid_set, pch = '.', col = ifelse(y_grid == 1, 'springgreen3', 'tomato'))
  points(set, pch = 21, bg = ifelse(set[, 3] == 1, 'green4', 'red3'))
}
```

```
visualize(training_set, 'Naive Bayes (Training Set)')  
visualize(test_set, 'Naive Bayes (Test Set)')
```

**Output:**

```
> cm = table(test_  
> print(cm)  
y_pred  
  0  1  
 0 57  7  
 1  7 29
```



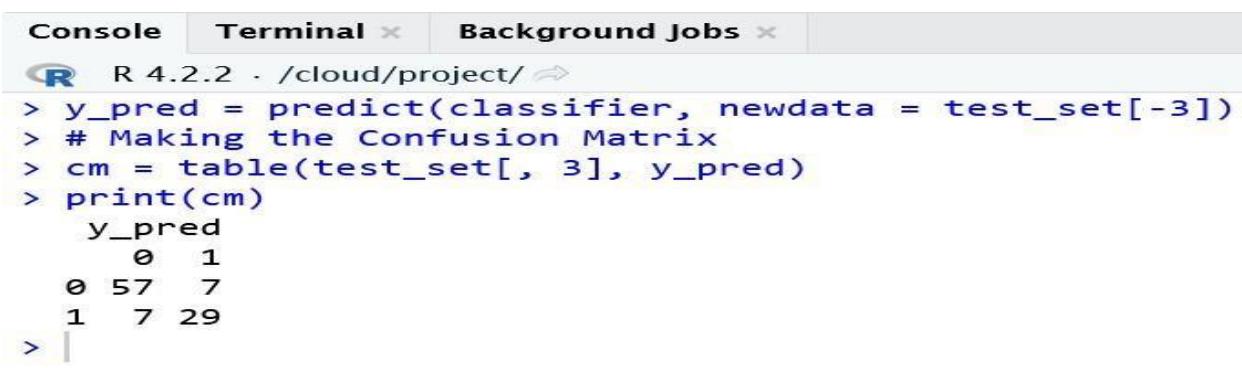
**B) Aim: Text Analysis (PBL): Find the confusion matrix to find restaurant review based of sentiment analysis of Natural Language processing. Use Restaurant reviews.tsv file for your study.**

**Source Code:**

```
dataset_original = read.delim('C:\\\\Users\\\\HP\\\\Downloads\\\\Restaurant_Reviews.tsv', quote = ",  
stringsAsFactors = FALSE)
```

```
install.packages('tm')  
install.packages('SnowballC')  
library(tm)  
library(SnowballC)  
corpus = VCorpus(VectorSource(dataset_original$Review))  
corpus = tm_map(corpus, content_transformer(tolower))  
corpus = tm_map(corpus, removeNumbers)  
corpus = tm_map(corpus, removePunctuation)  
corpus = tm_map(corpus, removeWords, stopwords())  
corpus = tm_map(corpus, stemDocument)  
corpus = tm_map(corpus, stripWhitespace)  
dtm = DocumentTermMatrix(corpus)  
dtm = removeSparseTerms(dtm, 0.999)  
dataset = as.data.frame(as.matrix(dtm))  
dataset$Liked = dataset_original$Liked  
print(dataset$Liked)  
dataset$Liked = factor(dataset$Liked, levels = c(0, 1))  
install.packages('caTools')  
library(caTools)  
set.seed(123)  
split = sample.split(dataset$Liked, SplitRatio = 0.8)  
training_set = subset(dataset, split == TRUE)  
test_set = subset(dataset, split == FALSE)  
install.packages('randomForest')  
library(randomForest)  
classifier = randomForest(x = training_set[-692],  
                           y = training_set$Liked,  
                           ntree = 10)  
y_pred = predict(classifier, newdata = test_set[-692])  
cm = table(test_set[, 692], y_pred)  
print(cm)
```

**Output:**



```
Console Terminal × Background Jobs ×  
R 4.2.2 · /cloud/project/ ↗  
> y_pred = predict(classifier, newdata = test_set[-3])  
> # Making the Confusion Matrix  
> cm = table(test_set[, 3], y_pred)  
> print(cm)  
y_pred  
  0   1  
0  57   7  
1   7  29  
>
```



## Practical No. 5

**Aim:** Comparative Study of various machine learning models (Newly added): Take the inbuilt data file: iris and perform classification on that data using various classification models – Decision Tree, K Nearest Neighbour and Support Vector Machine. Find the confusion matrix for all three models and evaluate them by finding their accuracy. Find the algorithm which performs best on the given data file, out of all these three models.

### Source Code:

```
install.packages("rpart")
install.packages("rpart.plot")
install.packages("caTools")
install.packages("e1071")
install.packages("class") # for kNN

library(rpart)
library(rpart.plot)
library(caTools)
library(e1071)
library(class)
data(iris)
summary(iris)

# Normalize continuous variables (only columns 1:4)
temp <- as.data.frame(scale(iris[, 1:4]))
temp$Species <- iris$Species

# Train-test split
set.seed(123)
split <- sample.split(temp$Species, SplitRatio = 0.75)
train <- subset(temp, split == TRUE)
test <- subset(temp, split == FALSE)
dt_classifier <- rpart(formula = Species ~ ., data = train)
rpart.plot(dt_classifier)

# Predict on test set
dt_y_pred <- predict(dt_classifier, newdata = test, type = "class")
dt_cm <- table(test$Species, dt_y_pred)
print(dt_cm)
dt_accuracy <- (sum(diag(dt_cm)) / sum(dt_cm)) * 100
print(paste("Decision Tree Accuracy:", round(dt_accuracy, 2), "%"))
cl <- train$Species
set.seed(1234)
```

```

knn_y_pred <- knn(train[, 1:4], test[, 1:4], cl, k = 5)

# Confusion matrix
knn_cm <- table(test$Species, knn_y_pred)
print(knn_cm)

# Accuracy calculation
knn_accuracy <- (sum(diag(knn_cm)) / sum(knn_cm)) * 100
print(paste("k-NN Accuracy:", round(knn_accuracy, 2), "%"))

svm_classifier <- svm(Species ~ ., data = train)
svm_y_pred <- predict(svm_classifier, newdata = test)

# Confusion matrix
svm_cm <- table(test$Species, svm_y_pred)
print(svm_cm)

# Accuracy calculation
svm_accuracy <- (sum(diag(svm_cm)) / sum(svm_cm)) * 100
print(paste("SVM Accuracy:", round(svm_accuracy, 2), "%"))

models <- data.frame(
  Technique = c("Decision Tree", "k-NN", "SVM"),
  Accuracy_Percentage = c(round(dt_accuracy, 2), round(knn_accuracy, 2), round(svm_accuracy, 2))
)
print(models)
cat("Decision Tree vs kNN: ", which(dt_y_pred != knn_y_pred), "\n")
cat("Decision Tree vs SVM: ", which(dt_y_pred != svm_y_pred), "\n")
cat("SVM vs kNN: ", which(svm_y_pred != knn_y_pred), "\n")

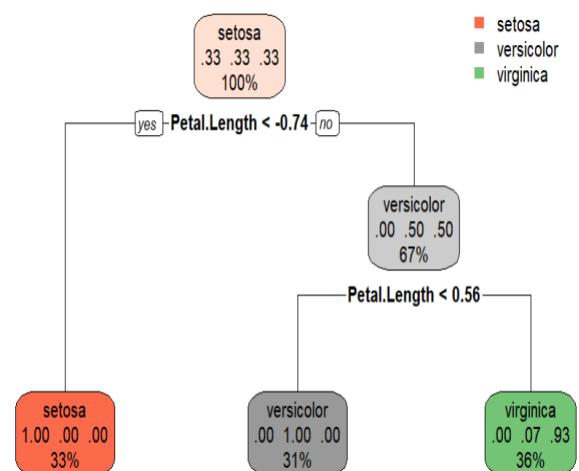
```

## Output:

```

> print(dt_cm)
      dt_y_pred
      setosa versicolor virginica
setosa     12       0       0
versicolor    0       9       3
virginica     0       1      11
> print(models)
      Technique Accuracy_Percentage
1  Decision Tree             88.89
2        k-NN                94.44
3         SVM                94.44

```



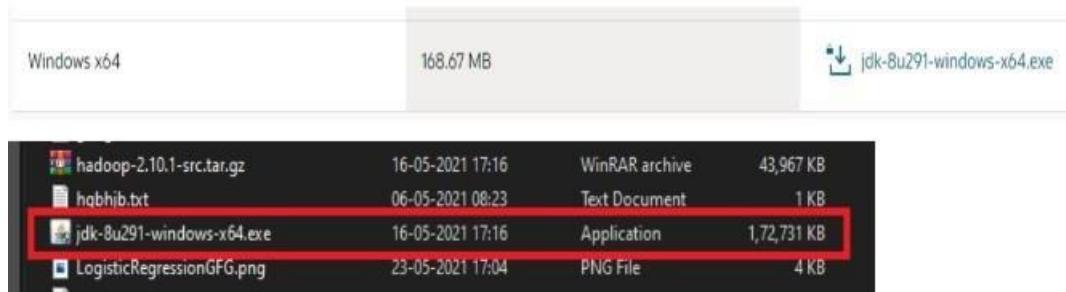
# Practical No. 6

**Aim: Install, configure and run Hadoop and HDFS ad explore HDFS.**

Description:

Hadoop Installation.

Step 1: download java jdk first .the package size 168.67MB



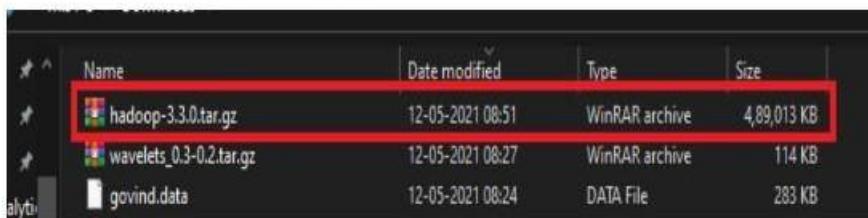
Step 2: download Hadoop binaries from the official website. The binary package size is about 342 MB.

Download

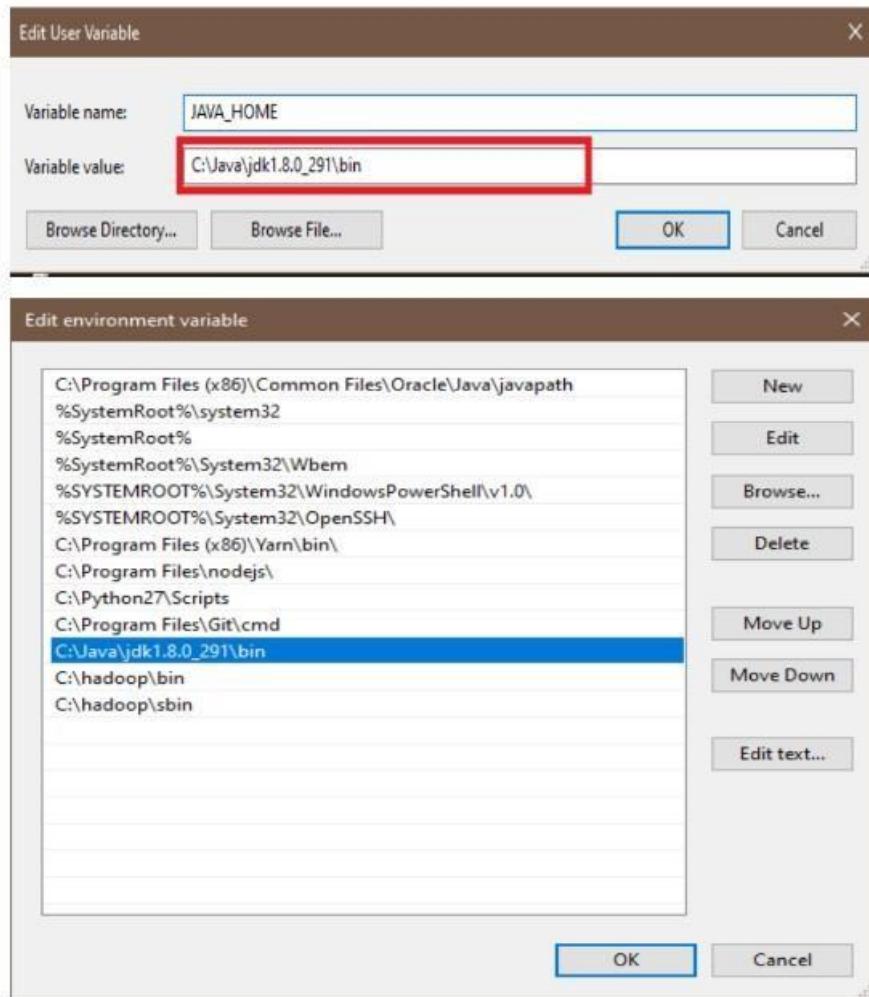
Hadoop is released as source code tarballs with corresponding binary tarballs for convenience. The downloads are distributed via mirror sites and should be checked for tampering using GPG or SHA-512.

Version	Release date	Source download	Binary download	Release notes
3.2.2	2021 Jan 9	source (checksum signature)	binary (checksum signature)	Announcement
2.10.1	2020 Sep 21	source (checksum signature)	binary (checksum signature)	Announcement
3.1.4	2020 Aug 3	source (checksum signature)	binary (checksum signature)	Announcement
3.3.0	2020 Jul 14	source (checksum signature)	binary (checksum signature) binary-aarch64 (checksum signature)	Announcement

Step 3: After finishing the file download, we should unpack the package using 7zip in two steps. First, we should extract the hadoop-3.2.1.tar.gz library, and then, we should unpack the extracted tar file:



Step 4: When the “Advanced system settings” dialog appears, go to the “Advanced” tab and click on the “Environment variables” button located on the bottom of the dialog.



### Step 5: Check the version of java

```
on C:\Windows\system32\cmd.exe
Microsoft Windows [Version 10.0.19041.928]
(c) Microsoft Corporation. All rights reserved.

C:\Users\hp>javac
Usage: javac <options> <source files>
where possible options include:
  -g                         Generate all debugging info
  -g:none                     Generate no debugging info
  -g:{lines,vars,source}       Generate only some debugging info
  -nowarn                     Generate no warnings
  -verbose                    Output messages about what the compiler is doing
  -deprecation               Output source locations where deprecated APIs are used
  -classpath <path>          Specify where to find user class files and annotation process
  -cp <path>                  Specify where to find user class files and annotation proces
  -sourcepath <path>          Specify where to find input source files
  -bootclasspath <path>      Override location of bootstrap class files
  -extdirs <dirs>             Override location of installed extensions
  -endorseddirs <dirs>        Override location of endorsed standards path
  -proc:{none,only}            Control whether annotation processing and/or compilation is o
  -processor <class1>[,<class2>,<class3>...] Names of the annotation processors to run; by
cs

C:\Users\hp>java -version
java version "1.8.0_291"
Java(TM) SE Runtime Environment (build 1.8.0_291-b10)
Java HotSpot(TM) 64-Bit Server VM (build 25.291-b10, mixed mode)
```

## Step 6: Configuration core-site.xml



```
C: > hadoop > etc > hadoop > core-site.xml
1  <?xml version="1.0" encoding="UTF-8"?>
2  <?xml-stylesheet type="text/xsl" href="configuration.xsl"?>
3
4  <configuration>
5
6    <property>
7      <name>fs.defaultFS</name>
8      <value>hdfs://localhost:9000</value>
9    <property>
10   </configuration>
```

## Step 7: Configuration core-site.xml

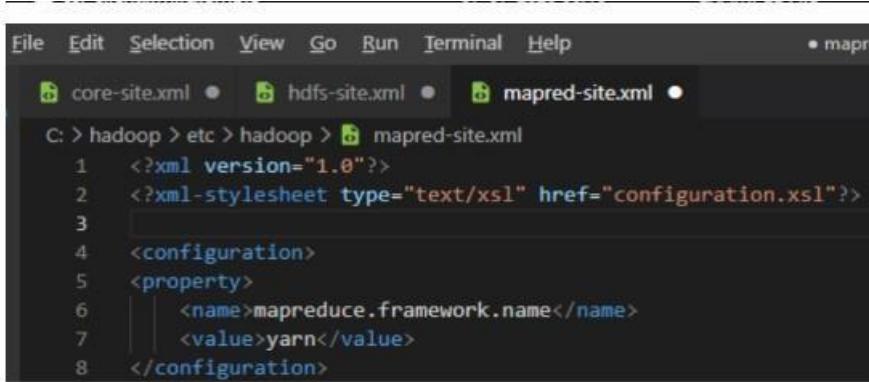
hdfs-rbf-site.xml	07-07-2020 00:26	XML File
hdfs-site.xml	19-05-2021 17:58	XML File
https-env.sh	07-07-2020 00:25	Shell Script



```
C: > hadoop > etc > hadoop > hdfs-site.xml
1  <?xml version="1.0" encoding="UTF-8"?>
2  <?xml-stylesheet type="text/xsl" href="configuration.xsl"?>
3
4  <configuration>
5    <property>
6      <name>dfs.replication</name>
7      <value>1</value>
8    </property>
9    <property>
10      <name>dfs.namenode.name.dir</name>
11      <value>C:\hadoop\data\namenode</value>
12
13  </property>
14  <property>
15      <name>dfs.namenode.data.dir</name>
16      <value>C:\hadoop\data\datanode</value>
17  </property>
18  </configuration>
```

## Step 8: Configuration core-site.xml

mapred-queues.xml.template	07-07-2020 01:04	TEMPLATE File
mapred-site.xml	19-05-2021 17:58	XML File
ssl-client.xml.example	07-07-2020 00:16	EXAMPLE File

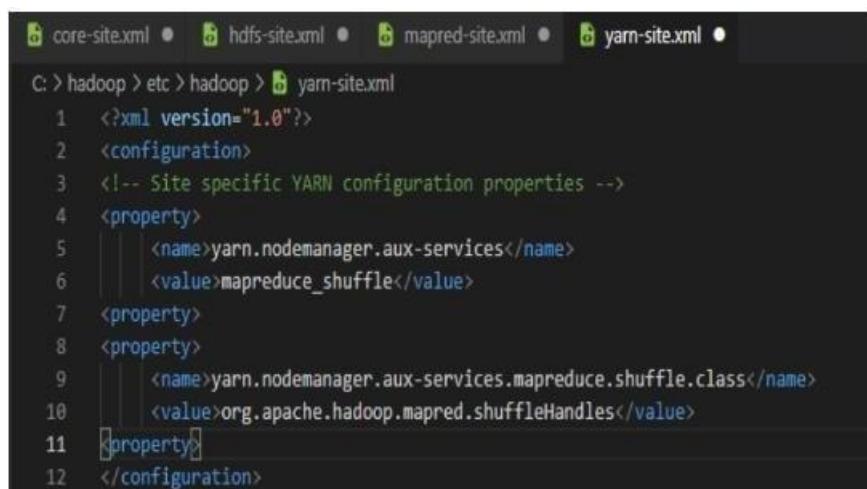


```
C: > hadoop > etc > hadoop > mapred-site.xml
1  <?xml version="1.0"?>
2  <?xml-stylesheet type="text/xsl" href="configuration.xsl"?>
3
4  <configuration>
5    <property>
6      <name>mapreduce.framework.name</name>
7      <value>yarn</value>
8    </configuration>
```

Step 9: Configuration core-site.xml

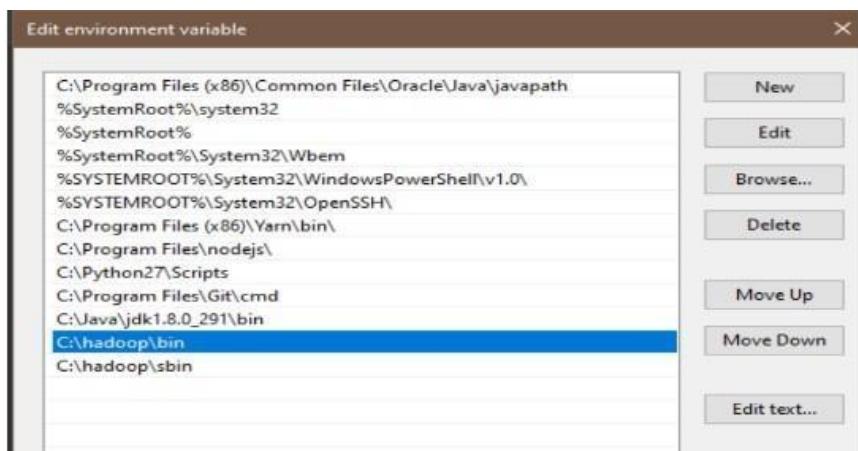


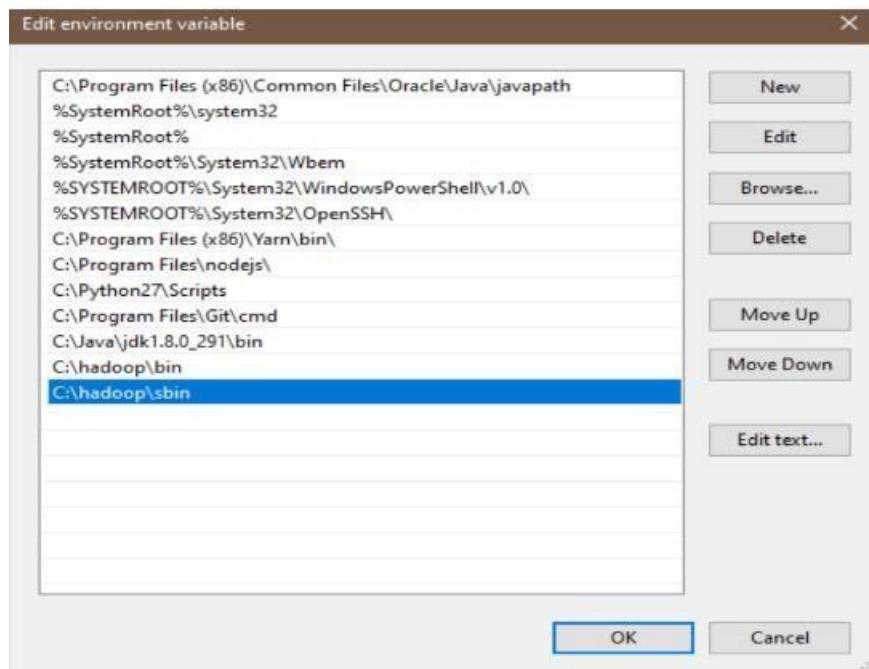
yamservice-log4j.properties	07-07-2020 01:03	PROPERTIES File
<b>yarn-site.xml</b>	19-05-2021 17:58	XML File



```
C: > hadoop > etc > hadoop > yarn-site.xml
1  <?xml version="1.0"?>
2  <configuration>
3    <!-- Site specific YARN configuration properties -->
4  <property>
5    <name>yarn.nodemanager.aux-services</name>
6    <value>mapreduce_shuffle</value>
7  <property>
8  <property>
9    <name>yarn.nodemanager.aux-services.mapreduce.shuffle.class</name>
10   <value>org.apache.hadoop.mapred.shuffleHandles</value>
11 </property>
12 </configuration>
```

Step 10: When the “Advanced system settings” dialog appears, go to the “Advanced” tab and click on the “Environment variables” button located on the bottom of the dialog.





Step 11: let's check Hadoop install Successfully

```
C:\Windows\system32\cmd.exe
Java(TM) SE Runtime Environment (build 1.8.0_291-b10)
Java HotSpot(TM) 64-Bit Server VM (build 25.291-b10, mixed mode)

C:\Users\hp>hdfs namenode -format
2021-05-23 17:17:11,111 INFO namenode.NameNode: STARTUP_MSG:
*****STARTUP_MSG: Starting NameNode
STARTUP_MSG:   host = DESKTOP-VUUFK2Q/192.168.0.104
STARTUP_MSG:   args = [-format]
STARTUP_MSG:   version = 3.3.0
STARTUP_MSG:   classpath = C:\hadoop\etc\hadoop;C:\hadoop\share\hadoop\common;C:\hs-smart-1.2.jar;C:\hadoop\share\hadoop\common\lib\animal-sniffer-annotations-1.17.asm-5.0.4.jar;C:\hadoop\share\hadoop\common\lib\audience-annotations-0.5.0.jar;C:\7.7.jar;C:\hadoop\share\hadoop\common\lib\checker-qual-2.5.2.jar;C:\hadoop\share\hadoop\share\hadoop\common\lib\commons-cli-1.2.jar;C:\hadoop\share\hadoop\share\hadoop\common\lib\commons-collections-3.2.2.jar;C:\hadoop\share\hadoop\share\hadoop\common\lib\commons-configuration2-2.1.1.jar;C:\hadoop\share\hadoop\share\hadoop\common\lib\commons-io-2.5.jar;C:\hadoop\share\hadoop\share\hadoop\common\lib\commons-logging-1.1.3.jar;C:\hadoop\share\hadoop\share\hadoop\common\lib\commons-net-3.6.jar;C:\hadoop\share\hadoop\common\lib\curator-client-4.2.0.jar;C:\hadoop\share\hadoop\common\lib\curator-common\lib\curator-recipes-4.2.0.jar;C:\hadoop\share\hadoop\common\lib\dr\failureaccess-1.0.jar;C:\hadoop\share\hadoop\common\lib\gson-2.2.4.jar;va-27.0-jre.jar;C:\hadoop\share\hadoop\common\lib\hadoop-annotations-3.3.0.jar;C:\auth-3.3.0.jar;C:\hadoop\share\hadoop\common\lib\hadoop-shaded-protobuf_3_7-1.0.0.htrace-core4-4.1.0-incubating.jar;C:\hadoop\share\hadoop\common\lib\httpclient-4.5.1b\httpcore-4.4.10.jar;C:\hadoop\share\hadoop\common\lib\j2objc-annotations-1.1.ja
```

```
Apache Hadoop Distribution
DEPRECATED: Use of this script to execute hdfs command is deprecated.
Instead use the hdfs command for it.
2021-05-23 17:19:33,116 INFO namenode.NameNode: STARTUP_MSG:
*****
STARTUP_MSG: Starting NameNode
STARTUP_MSG: host = DESKTOP-VUUFK2Q/192.168.0.104
STARTUP_MSG: args = []
STARTUP_MSG: version = 3.3.0
STARTUP_MSG: classpath = C:\hadoop\etc\hadoop;C:\hadoop\share\hadoop\common;C:\hadoop\share\hadoop\common\lib\accessor-s-smart-1.2.jar;C:\hadoop\share\hadoop\common\lib\animal-sniffer-annotations-1.17.jar;C:\hadoop\share\hadoop\common\lib\asm-5.0.4.jar;C:\hadoop\share\hadoop\common\lib\audience-annotations-0.5.0.jar;C:\hadoop\share\hadoop\common\lib\avro-1.7.7.jar;C:\hadoop\share\hadoop\common\lib\checker-qual-2.5.2.jar;C:\hadoop\share\hadoop\common\lib\commons-beanutils-1.9.4.jar;C:\hadoop\share\hadoop\common\lib\commons-cli-1.2.jar;C:\hadoop\share\hadoop\common\lib\commons-codec-1.11.jar;C:\hadoop\share\hadoop\common\lib\commons-collections-3.2.2.jar;C:\hadoop\share\hadoop\common\lib\commons-compress-1.19.jar;C:\hadoop\share\hadoop\common\lib\commons-configuration2-2.1.1.jar;C:\hadoop\share\hadoop\common\lib\commons-daemon-1.0.13.jar;C:\hadoop\share\hadoop\common\lib\commons-io-2.5.jar;C:\hadoop\share\hadoop\common\lib\commons-lang3-3.7.jar;C:\hadoop\share\hadoop\common\lib\commons-logging-1.1.3.jar;C:\hadoop\share\hadoop\common\lib\commons-math3-3.1.1.jar;C:\hadoop\share\hadoop\common\lib\commons-net-3.6.jar;C:\hadoop\share\hadoop\common\lib\commons-text-1.4.jar;C:\hadoop\share\hadoop\common\lib\curator-client-4.2.0.jar;C:\hadoop\share\hadoop\common\lib\curator-framework-4.2.0.jar;C:\hadoop\share\hadoop\common\lib\curator-recipes-4.2.0.jar;C:\hadoop\share\hadoop\common\lib\dnsjava-2.1.7.jar;C:\hadoop\share\hadoop
```

```
Apache Hadoop Distribution
at com.ctc.wstx.sr.StreamScanner.throwParseError(StreamScanner.java:491)
at com.ctc.wstx.sr.StreamScanner.throwParseError(StreamScanner.java:475)
at com.ctc.wstx.sr.BasicStreamReader.reportWrongEndElement(BasicStreamReader.java:3365)
at com.ctc.wstx.sr.BasicStreamReader.readEndElement(BasicStreamReader.java:3292)
at com.ctc.wstx.sr.BasicStreamReader.nextFromTree(BasicStreamReader.java:2911)
at com.ctc.wstx.sr.BasicStreamReader.next(BasicStreamReader.java:1123)
at org.apache.hadoop.conf.Configuration$Parser.parseNext(Configuration.java:3347)
at org.apache.hadoop.conf.Configuration$Parser.parse(Configuration.java:3141)
at org.apache.hadoop.conf.Configuration.loadResource(Configuration.java:3034)
... 9 more
```

## Step 12: Let check bin

```
C:\Windows\system32\cmd.exe
C:\Users\hp>cd C:\hadoop\sbin
C:\hadoop\sbin>start-all.cmd
This script is Deprecated. Instead use start-dfs.cmd and start-yarn.cmd
starting yarn daemons
C:\hadoop\sbin>
```

# Practical No. 7

**Aim.: Implement word count / frequency programs using MapReduce.**

**Source Code:**

```
Microsoft Windows [Version 10.0.18363.418] (c)
2019 Microsoft Corporation. All rights reserved.
C:\Windows\system32>cd..
```

```
C:\Windows>cd..
```

```
C:>cd C:\hadoop-3.2.2\hadoop-3.2.2\sbin
```

```
C:\hadoop-3.2.2\hadoop-3.2.2\sbin>start-dfs
'C:\Program' is not recognized as an internal or external command,
operable program or batch file.
```

```
C:\hadoop-3.2.2\hadoop-3.2.2\sbin>start-yarn starting
yarn daemons
'C:\Program' is not recognized as an internal or external command,
operable program or batch file.
```

```
C:\hadoop-3.2.2\hadoop-3.2.2\sbin>jps
12368 Jps
2752 DataNode
12228 NodeManager
8360 NameNode
9500 ResourceManager
```

```
C:\hadoop-3.2.2\hadoop-3.2.2\sbin>hadoop fs -ls /input/
'C:\Program' is not recognized as an internal or external command,
operable program or batch file.
```

```
Found 1 items
-rw-r--r-- 1 HP supergroup 196 2023-02-18 15:13 /input/data.txt
C:\hadoop-3.2.2\hadoop-3.2.2\sbin>hadoop dfs -cat /input/data.txt
'C:\Program' is not recognized as an internal or external command,
operable program or batch file.
```

DEPRECATED: Use of this script to execute hdfs command is deprecated.  
Instead use the hdfs command for it.

```
'C:\Program' is not recognized as an internal or external command,
operable program or batch file.
```

```
Peter Piper picked a peck of pickled peppers
A peck of pickled peppers Peter Piper picked
If Peter Piper picked a peck of pickled peppers
```

```
Where's the peck of pickled peppers Peter Piper picked?
```

```
C:\hadoop-3.2.2\hadoop-3.2.2\sbin>hadoop jar C:\hadoop-3.2.2\hadoop-
3.2.2\share\hadoop\mapreduce\hadoopmapreduce-examples-3.2.2.jar wordcount /input /out
```

'C:\Program' is not recognized as an internal or external command,  
operable program or batch file.

2023-02-18 15:48:10,371 INFO impl.MetricsConfig: Loaded properties from hadoop-metrics2.properties 2023-02-18 15:48:10,491 INFO impl.MetricsSystemImpl: Scheduled Metric snapshot period at 10 second(s).

2023-02-18 15:48:10,491 INFO impl.MetricsSystemImpl: JobTracker metrics system started

2023-02-18 15:48:11,307 INFO input.FileInputFormat: Total input files to process : 1

2023-02-18 15:48:11,347 INFO mapreduce.JobSubmitter: number of splits:1

2023-02-18 15:48:11,477 INFO mapreduce.JobSubmitter: Submitting tokens for job: job\_local1756910747\_0001

2023-02-18 15:48:11,478 INFO mapreduce.JobSubmitter: Executing with tokens: []

2023-02-18 15:48:11,735 INFO mapreduce.Job: The url to track the job: http://localhost:8080/

2023-02-18 15:48:11,737 INFO mapreduce.Job: Running job: job\_local1756910747\_0001

2023-02-18 15:48:11,742 INFO mapred.LocalJobRunner: OutputCommitter set in config null

2023-02-18 15:48:11,769 INFO output.FileOutputCommitter: File Output Committer Algorithm version is 2

2023-02-18 15:48:11,769 INFO output.FileOutputCommitter: FileOutputCommitter skip cleanup \_temporary folders under output directory:false, ignore cleanup failures: false

2023-02-18 15:48:11,772 INFO mapred.LocalJobRunner: OutputCommitter is org.apache.hadoop.mapreduce.lib.output.FileOutputCommitter

2023-02-18 15:48:11,979 INFO mapred.LocalJobRunner: Waiting for map tasks 2023-02-18 15:48:11,986 INFO mapred.LocalJobRunner: Starting task: attempt\_local1756910747\_0001\_m\_000000\_0

2023-02-18 15:48:12,091 INFO output.FileOutputCommitter: File Output Committer Algorithm version is 2

2023-02-18 15:48:12,091 INFO output.FileOutputCommitter: FileOutputCommitter skip cleanup \_temporary folders under output directory:false, ignore cleanup failures: false

2023-02-18 15:48:12,111 INFO util.ProcfsBasedProcessTree: ProcfsBasedProcessTree currently is supported only on Linux.

2023-02-18 15:48:12,154 INFO mapred.Task: Using ResourceCalculatorProcessTree :

A 1  
If 1  
Peter 4  
Piper 4 Where's 1  
a 2  
of 4  
peck 4  
peppers 4  
picked 3  
picked? 1  
pickled 4 the  
1

C:\hadoop-3.2.2\hadoop-3.2.2\sbin>

## Output:

Block information – Block 0

Block ID: 1073741834  
Block Pool ID: BP-870543631-10.0.0.7-1676388394632  
Generation Stamp: 1010  
Size: 96  
Availability:

- DESKTOP-I2BH387

File contents

```
A 1
If 1
Peter 4
Piper 4
Where's 1
a 2
of 4
peck 4
```

**Close**

## Practical No. 8

**Aim: Implement an application that stores big data in Hbase / MongoDB and manipulate it using R / Python**

**Requirements :**

- a. PyMongo
- b. Mongo Database

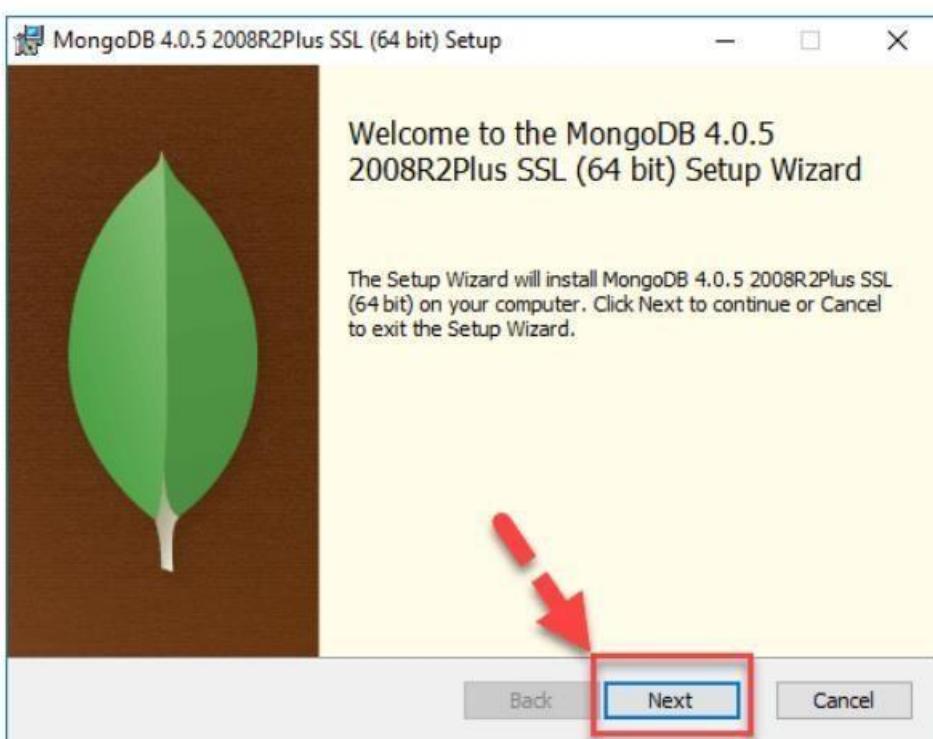
### Step A: Install Mongo database

Step 1) Go to (<https://www.mongodb.com/download-center/community>) and Download MongoDB Community Server. We will install the 64-bit version for Windows.

Select the server you would like to run:

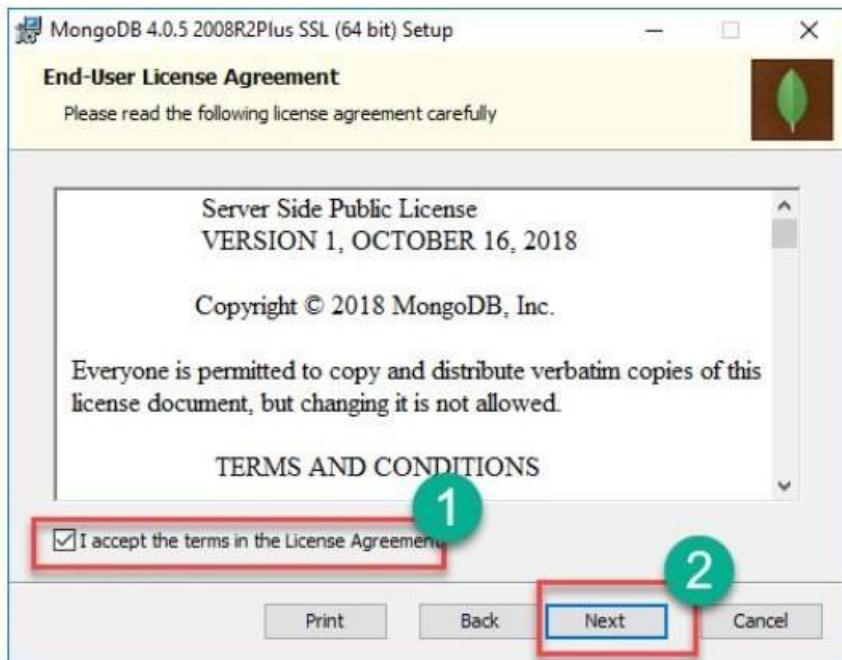


Step 2) Once download is complete open the msi file. Click Next in the start up screen

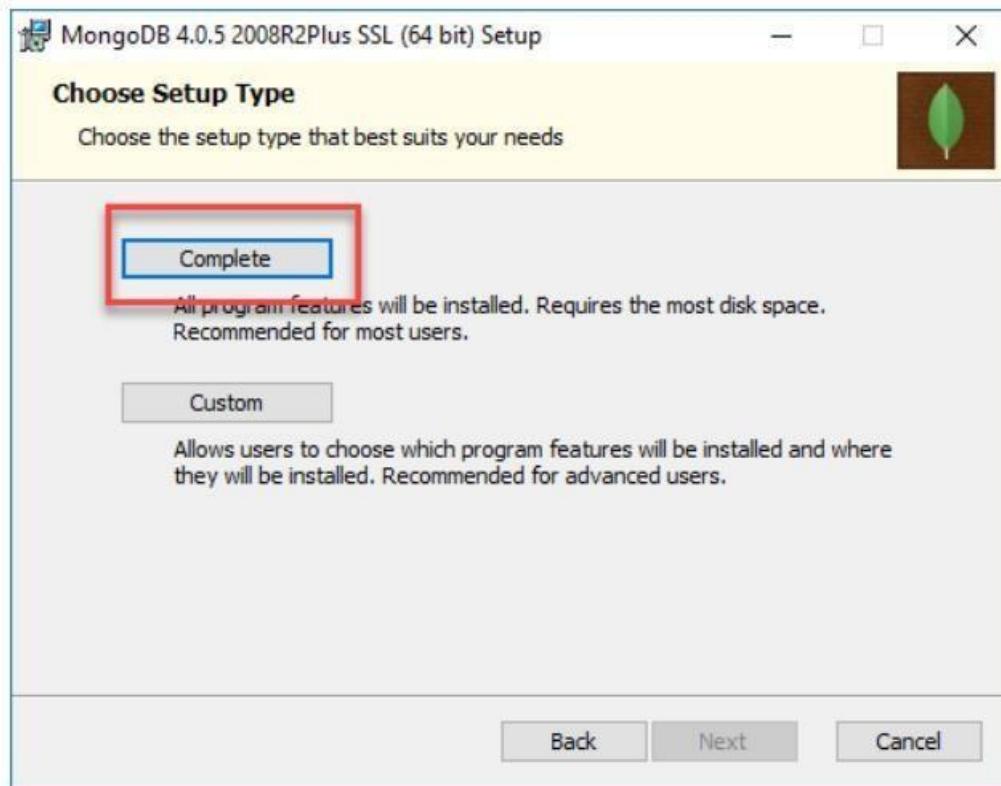


Step 3)

1. Accept the End-User License Agreement
2. Click Next



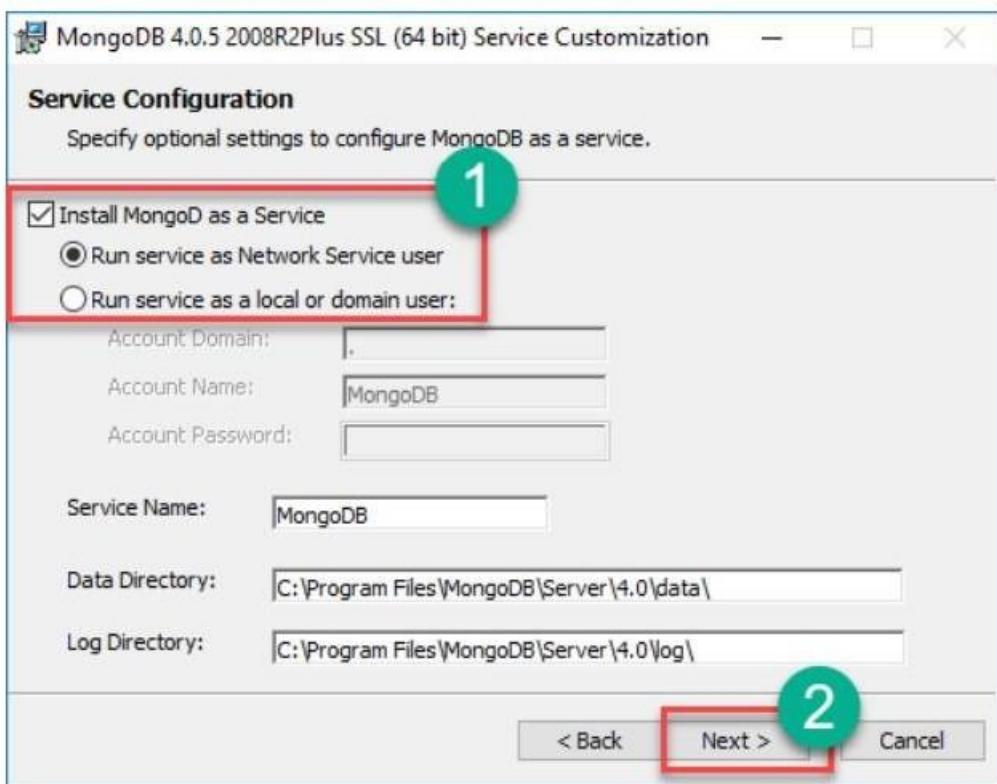
Step 4) Click on the "complete" button to install all of the components. The custom option can be used to install selective components or if you want to change the location of the installation.



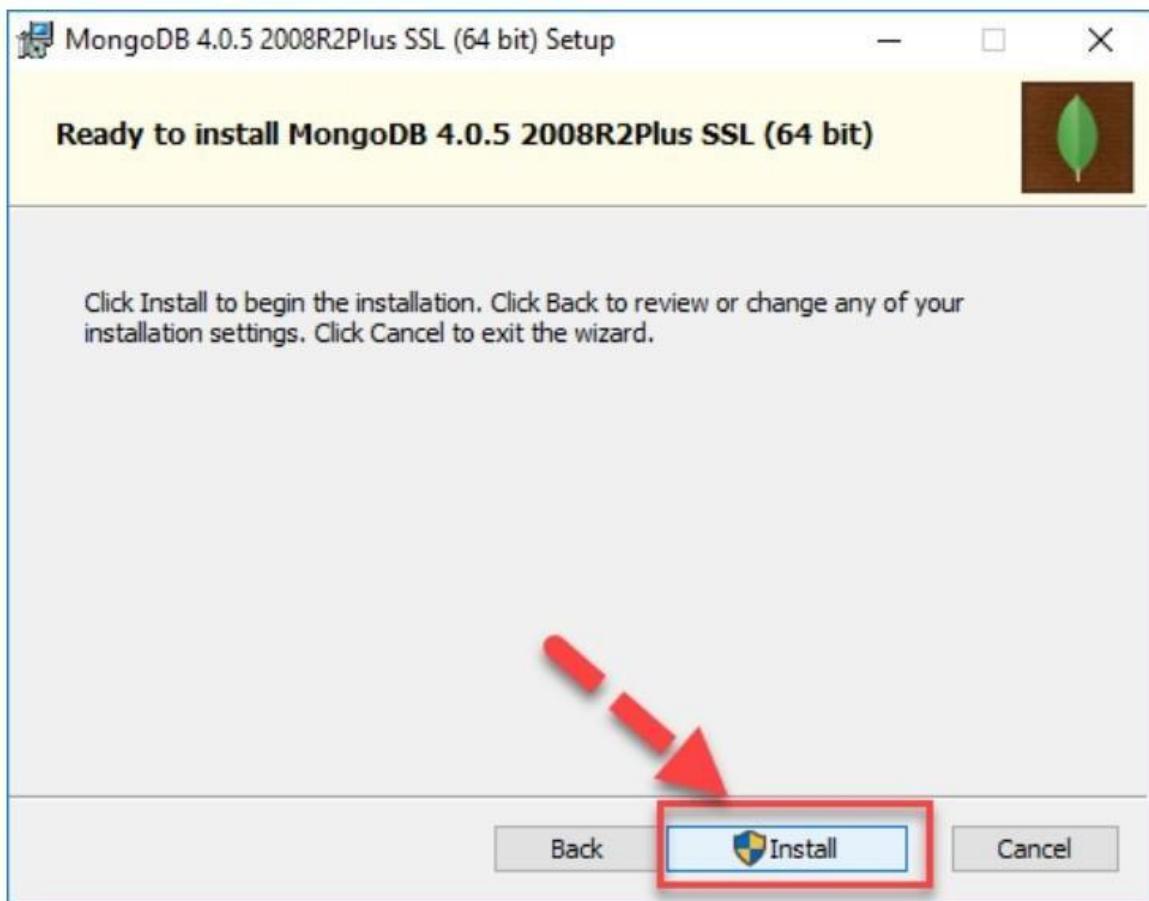
Step 5)

1. Select “Run service as Network Service user”. make a note of the data directory, we’ll need this later.

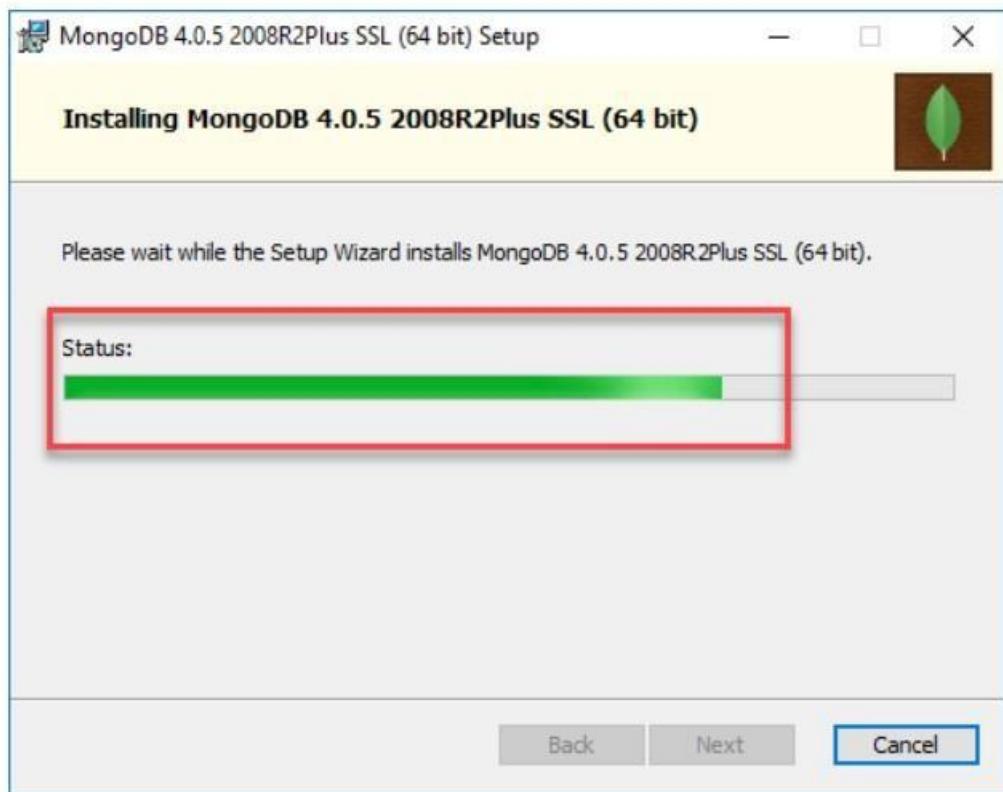
2. Click Next



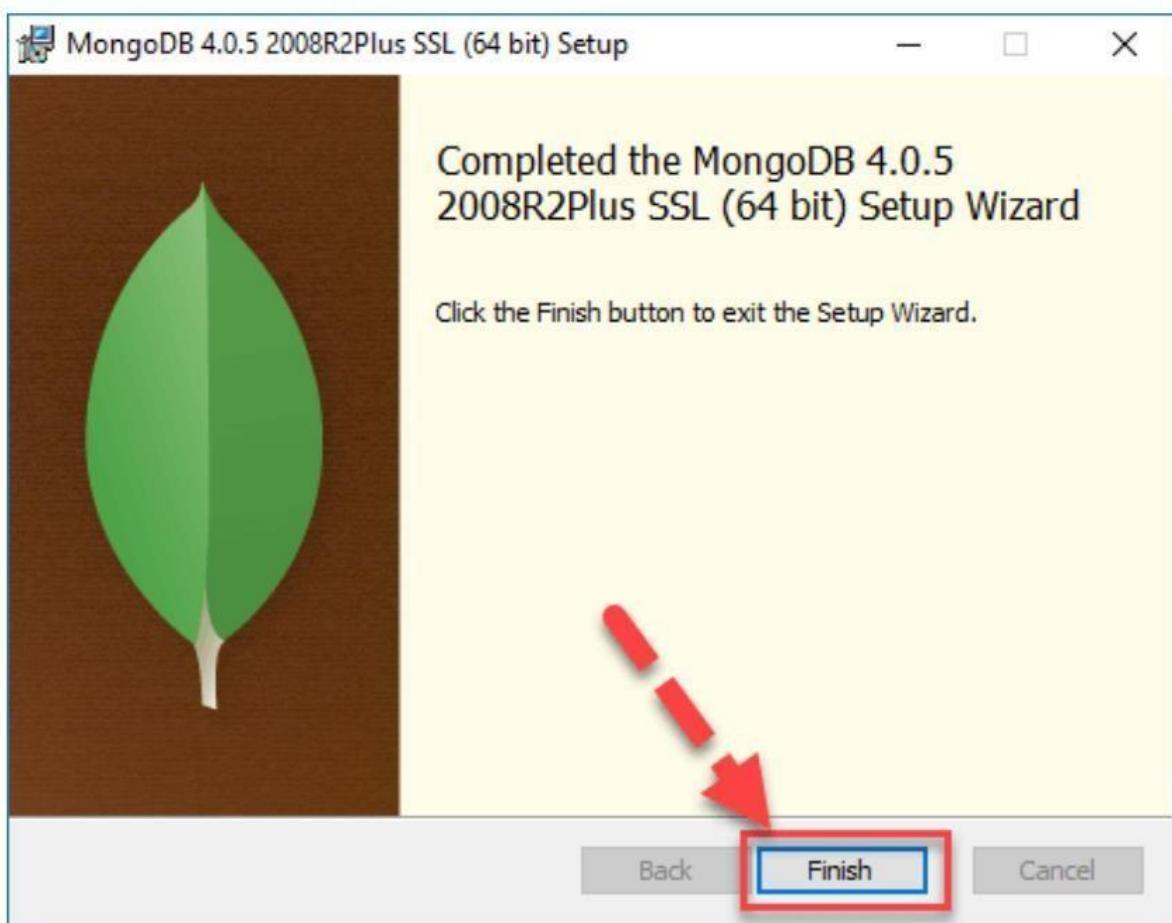
Step 6) Click on the Install button to start the installation.



Step 7) Installation begins. Click Next once completed.



Step 8) Click on the Finish button to complete the installation.



## Test Mongodb

**Step 1)** Go to " C:\Program Files\MongoDB\Server\4.0\bin" and double click on **mongo.exe**. Alternatively, you can also click on the MongoDB desktop icon.

### □ Create the directory where MongoDB will store its files.

Open command prompt window and apply following commands

```
C:\users\admin> cd\  
C:>md data\db
```

### Step 2) Execute mongod

Open another command prompt window.

```
C:>> cd C:\Program Files\MongoDB\Server\4.0\bin  
C:\Program Files\MongoDB\Server\4.0\bin> mongod
```

*In case if it gives an error then run the following command:*

```
C:\Program Files\MongoDB\Server\4.0\bin> mongod --repair
```

```
{ v: 2, key: { lastUse: 1 }, name: "lsidTTLIndex", ns: "config.system.sessions", expireAfterSeconds: 1800 }  
2023-03-03T09:46:21.011+0530 I INDEX [LogicalSessionCacheRefresh] building index using bulk method; build may temporarily use up to 500 megabytes of RAM  
2023-03-03T09:46:21.044+0530 I INDEX [LogicalSessionCacheRefresh] build index done. scanned 0 total records. 0 secs  
2023-03-03T09:46:21.045+0530 I COMMAND [LogicalSessionCacheRefresh] command config.$cmd command: createIndexes { createIndexes: "system.sessions", indexes: [ { key: { lastUse: 1 }, name: "lsidTTLIndex", expireAfterSeconds: 1800 } ], $db: "config" } numYields:0 reslen:114 locks:{ Global: { acquireCount: { r: 2, w: 2 } }, Database: { acquireCount: { w: 2, W: 1 } }, Collection: { acquireCount: { w: 2 } } } protocol:op_msg 254ms
```

### Step 3) Connect to MongoDB using the Mongo shell

Let the MongoDB daemon to run.

Open another command prompt window and run the following commands:

```
C:\users\admin> cd C:\Program Files\MongoDB\Server\4.0\bin  
C:\Program Files\MongoDB\Server\4.0\bin> mongo
```

```
The monitoring data will be available on a MongoDB website with a unique URL accessible to you and anyone you share the URL with. MongoDB may use this information to make product improvements and to suggest MongoDB products and deployment options to you.  
  
To enable free monitoring, run the following command: db.enableFreeMonitoring()  
To permanently disable this reminder, run the following command: db.disableFreeMonitoring()  
---  
  
> show dbs  
admin      0.000GB  
config     0.000GB  
local      0.000GB  
mybigdata  0.000GB  
> use mybigdata
```

### Step 4) Install PyMongo

Open another command prompt window and run the following commands:

Check the python version on your desktop / laptop and copy that path from window explorer

```
C:\users\admin>cd C:\Program Files\Python311\Scripts  
C:\Program Files\<Python38>\Scripts > python -m pip install pymongo
```

```
C:\Program Files\Python38\Scripts>python -m pip install pymongo  
Defaulting to user installation because normal site-packages is not writeable  
Collecting pymongo  
  Downloading pymongo-4.3.3-cp38-cp38-win_amd64.whl (382 kB)  
    |██████████| 382 kB 3.3 MB/s  
Collecting dnspython<3.0.0,>=1.16.0  
  Downloading dnspython-2.3.0-py3-none-any.whl (283 kB)  
    |██████████| 283 kB ...  
Installing collected packages: dnspython, pymongo  
Successfully installed dnspython-2.3.0 pymongo-4.3.3  
WARNING: You are using pip version 20.2.1; however, version 23.0.1 is available.  
You should consider upgrading via the 'C:\Program Files\Python38\python.exe -m pip install --upgrade pip' command.
```

Note: # **-m** option is for <module-name>

Now you have downloaded and installed a mongoDB driver.

### Step 5) Test PyMongo

Run the following command from python command prompt

```
import pymongo
```

Now, either create a file in Python IDLE or run all commands one by one in sequence on Python cell

#### Program 1: Creating a Database: `create_dp.py`

```
import pymongo
myclient = pymongo.MongoClient("mongodb://localhost:27017/")
mydb      = myclient["mybigdata"]
print(myclient.list_database_names())
['admin', 'config', 'local']
```

#### Program 2: Creating a Collection: `create_collection.py`

```
import pymongo
myclient = pymongo.MongoClient("mongodb://localhost:27017/")
mydb      = myclient["mybigdata"]    mycol= mydb["student"]
print(mydb.list_collection_names())
[]
```

#### Program 3: Insert into Collection: `insert_into_collection.py`

```
import pymongo
myclient = pymongo.MongoClient("mongodb://localhost:27017/")
mydb = myclient["mybigdata"] mycol= mydb["student"]
mydict={"name":"Beena", "address":"Mumbai"}
x= mycol.insert_one(mydict) # insert_one(containing the name(s) and value(s) of each field
```

#### Program 4: Insert Multiple data into Collection: `insert_many.py`

```
import pymongo
myclient = pymongo.MongoClient("mongodb://localhost:27017/")
mydb = myclient["mybigdata"] mycol= mydb["student"]
mylist=[ {"name":"Khyati", "address":"Mumbai"}, {"name":"Kruti", "address":"Mumbai"}, {"name":"Nidhi", "address":"Pune"}, {"name":"Komal", "address":"Pune"},] x= mycol.insert_many(mylist)
```

### Step 6) Test in Mongodb to check database and data inserted in collection

- If you want to check your database list, use the command `show dbs` in mongo command prompt

```
> show dbs
```

admin	0.000GB
config	0.000GB
local	0.000GB
mybigdata	0.000GB

- If you want to use a database with name `mybigdata`, then use `use` database statement would be as follow:

```
> use mybigdata
```

```
|switched to db mybigdata
```

- c. If you want to check collection in mongodb use the command show collections > show collections

```
|student
```

- d. If you want to display the first row from collection:  
db.collection\_name.find() > db.student.findOne()

```
> db.student.findOne()
{
    "_id" : ObjectId("640178face663db608cef72f"),
    "name" : "Beena",
    "address" : "Mumbai"
}
```

- e. If you want to display all the data from collection: db.collection\_name.find()

```
> db.student.find()
> db.student.find()
{
    "_id" : ObjectId("640178face663db608cef72f"), "name" : "Beena", "address" : "Mumbai" }
{
    "_id" : ObjectId("640179336ce317082c266dc1"), "name" : "Khyati", "address" : "Mumbai" }
{
    "_id" : ObjectId("640179336ce317082c266dc2"), "name" : "Kruti", "address" : "Mumbai" }
{
    "_id" : ObjectId("640179336ce317082c266dc3"), "name" : "Nidhi", "address" : "Pune" }
{
    "_id" : ObjectId("640179336ce317082c266dc4"), "name" : "Komal", "address" : "Pune" }
```

- f. count number of rows in a collection > db.student.count()

```
5
```

**Site for R packages documentation:**

[https://cran.r-project.org/web/packages/available\\_packages\\_by\\_name.html](https://cran.r-project.org/web/packages/available_packages_by_name.html)

# **MODERN NETWORKING**



## JNAN VIKAS MANDAL'S

Mohanlal Raichand Mehta College of Commerce

Diwali Maa College of Science

Amritlal Raichand Mehta College of Arts

Dr. R.T. Doshi College of Computer Science

NAAC Re-Accredited Grade 'A+' (CGPA : 3.31) (3rd Cycle)

### DEPARTMENT OF INFORMATION TECHNOLOGY CERTIFICATE

This is to certify that \_\_\_\_\_ bearing seat no. \_\_\_\_\_

has Completed practical work of **Semester II** Practical Examination during the Academic Year 2024-25 under the guidance of **Ass.Prof. Shraddha Kenjale** being the requirement for the fulfilment of the curriculum of Degree of Master of Science in Information Technology Part I under University of Mumbai.

Date:

Place: Airoli

Ass Prof Shraddha Kenjale  
Subject – In Charge

External Examiner  
Signature

Dr. Sunita Joshi  
Msc-IT Co Ordinator

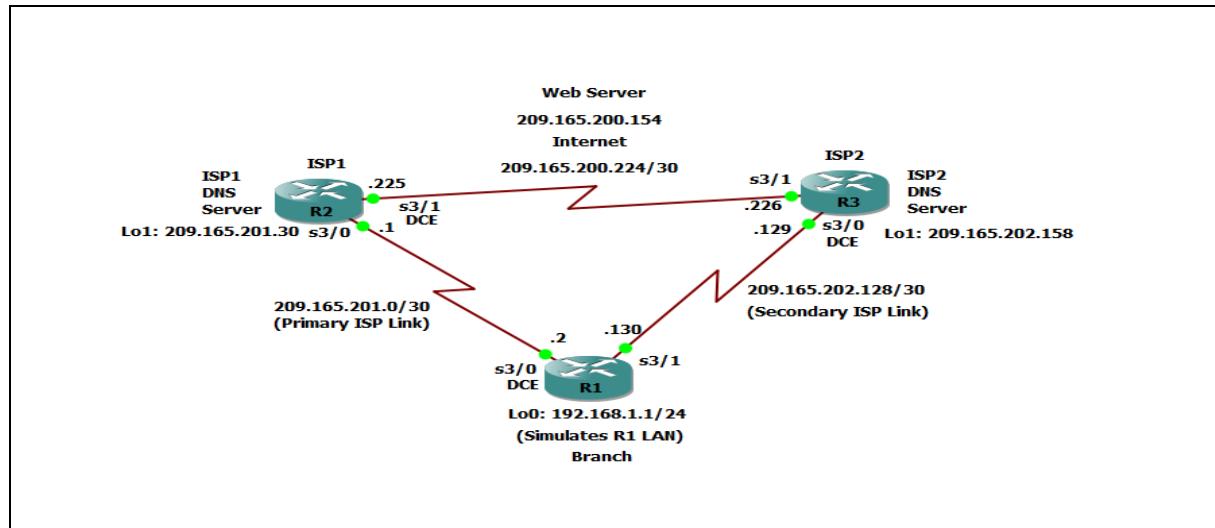
# INDEX

SR.NO	TITLE	PAGE NO	DATE	SIGN
1	Configure IP SLA Tracking and Path Control.	1		
2	Using the AS_PATH Attribute.	12		
3	Configuring IBGP and EBGP Sessions, Local Preference and MED.	19		
4	Secure the Management Plane.	39		
5	Configure and Verify Path Control Using PBR.	48		
6	Inter VLAN Routing.	60		
7	Simulating MPLS environment.	71		

## Practical no 1

Aim:- Configure IP SLA Tracking and Path Control.

### Topology



### Step 1: Prepare the routers and configure the router hostname and interface addresses.

- Cable the network as shown in the topology diagram. Erase the startup configuration and reload each router to clear the previous configurations. Using the addressing scheme in the diagram, create the loopback interfaces and apply IP addresses to them as well as the serial interfaces on R1, ISP1, and ISP2.  
You can copy and paste the following configurations into your routers to begin.

#### Router R1

```
R1#conf t
Enter configuration commands, one per line. End with CNTL/Z.
R1(config)#hostname R1
R1(config)#interface Loopback 0
R1(config-if)#ip address 192.168.1.1 255.255.255.0
R1(config-if)#exit
R1(config)#[
```

```
R1(config)#interface S3/0
R1(config-if)#ip address 209.165.201.2 255.255.255.252
R1(config-if)#clock rate 128000
R1(config-if)#bandwidth 128
R1(config-if)#no shutdown
R1(config-if)#exit
R1(config)#[
```

```
R1(config)#interface S3/1
R1(config-if)#ip address 209.165.202.130 255.255.255.252
R1(config-if)#bandwidth 128
R1(config-if)#no shutdown
R1(config-if)#exit
R1(config)#[
```

### **Router ISP1 (R2)**

```
R2#conf t
Enter configuration commands, one per line. End with CNTL/Z.
R2(config)#hostname ISP1
ISP1(config)#interface Loopback0
ISP1(config-if)#ip address 209.165.200.254 255.255.255.255
ISP1(config-if)#exit
ISP1(config)#[
```

```
ISP1(config)#interface Loopback1
ISP1(config-if)#ip address 209.165.201.30 255.255.255.255
ISP1(config-if)#interface S3/0
ISP1(config-if)#ip address 209.165.201.1 255.255.255.252
ISP1(config-if)#bandwidth 128
ISP1(config-if)#no shutdown
ISP1(config-if)#exit
```

```
ISP1(config)#interface S3/1
ISP1(config-if)#ip address 209.165.200.225 255.255.255.252
ISP1(config-if)#clock rate 128000
ISP1(config-if)#bandwidth 128
ISP1(config-if)#no shutdown
ISP1(config-if)#exit
```

### **Router ISP2 (R3)**

```
R3#conf t
Enter configuration commands, one per line. End with CNTL/Z.
R3(config)#hostname ISP2
ISP2(config)#interface Loopback0
ISP2(config-if)#ip address 209.165.200.254 255.255.255.255
ISP2(config-if)#exit
ISP2(config)#[
```

```
ISP2(config)#interface Loopback1
ISP2(config-if)#ip address 209.165.202.158 255.255.255.255
ISP2(config-if)#exit
```

```
ISP2(config)#interface S3/0
ISP2(config-if)#ip address 209.165.202.129 255.255.255.252
ISP2(config-if)#clock rate 128000
ISP2(config-if)#bandwidth 128
ISP2(config-if)#no shutdown
ISP2(config-if)#exit
```

```
ISP2(config)#interface S3/1
ISP2(config-if)#ip address 209.165.200.226 255.255.255.252
ISP2(config-if)#bandwidth 128
ISP2(config-if)#no shutdown
ISP2(config-if)#exit
```

- b. Verify the configuration by using the show interfaces description command. The output from router R1 is shown here as an example.

```
R1#show interface description
Interface          Status      Protocol Description
Fa0/0              admin down  down
Gi1/0              admin down  down
Gi2/0              admin down  down
Se3/0              up         up
Se3/1              up         up
Se3/2              admin down  down
Se3/3              admin down  down
Se4/0              admin down  down
Se4/1              admin down  down
Se4/2              admin down  down
Se4/3              admin down  down
Lo0               up         up
R1#
```

All three interfaces should be active. Troubleshoot if necessary.

c. The current routing policy in the topology is as follows:

- Router R1 establishes connectivity to the Internet through ISP1 using a default static route.
- ISP1 and ISP2 have dynamic routing enabled between them, advertising their respective public address pools.
- ISP1 and ISP2 both have static routes back to the ISP LAN.

### **Router R1**

```
R1#conf t
Enter configuration commands, one per line. End with CNTL/Z.
R1(config)#ip route 0.0.0.0 0.0.0.0 209.165.201.1
R1(config)#
```

### **Router ISP1 (R2)**

```
ISP1(config)#router eigrp 1
ISP1(config-router)#network 209.165.200.224 0.0.0.3
ISP1(config-router)#network 209.165.201.0 0.0.0.31
ISP1(config-router)#no auto-summary
ISP1(config-router)#exit
ISP1(config)#ip route 192.168.1.0 255.255.255.0 209.165.201.2
ISP1(config)#
```

### **Router ISP2 (R3)**

```
ISP2(config)#router eigrp 1
ISP2(config-router)#network 209.165.200.224 0.0.0.3
ISP2(config-router)#network 209.165.202.128 0.0.0.31
ISP2(config-router)#no auto-summary
ISP2(config-router)#exit
ISP2(config)#ip route 192.168.1.0 255.255.255.0 209.165.202.130
ISP2(config)#exit
```

### **Step 2: Verify server reachability.**

- a. Before implementing the Cisco IOS SLA feature, you must verify reachability to the Internet servers. From router R1, ping the web server, ISP1 DNS server, and ISP2 DNS server to verify connectivity. You can copy the following Tcl script and paste it into R1.

```
R1(tcl)#foreach address {
+>(tcl) #209.165.200.254
+>(tcl) #209.165.201.30
+>(tcl) #209.165.202.158
+>(tcl) #} {
+>(tcl) #ping $address source 192.168.1.1
+>(tcl) #
Type escape sequence to abort.
Sending 5, 100-byte ICMP Echos to 209.165.200.254, timeout is 2 seconds:
Packet sent with a source address of 192.168.1.1
!!!!!
Success rate is 100 percent (5/5), round-trip min/avg/max = 12/32/52 ms
Type escape sequence to abort.
Sending 5, 100-byte ICMP Echos to 209.165.201.30, timeout is 2 seconds:
Packet sent with a source address of 192.168.1.1
!!!!!
Success rate is 100 percent (5/5), round-trip min/avg/max = 16/39/68 ms
Type escape sequence to abort.
Sending 5, 100-byte ICMP Echos to 209.165.202.158, timeout is 2 seconds:
Packet sent with a source address of 192.168.1.1
!!!!!
Success rate is 100 percent (5/5), round-trip min/avg/max = 16/61/120 ms
R1(tcl) #
```

- b. Trace the path taken to the web server, ISP1 DNS server, and ISP2 DNS server. You can copy the following Tcl script and paste it into R1.

```
R1(tcl)#foreach address {
+>(tcl) #209.165.200.254
+>(tcl) #209.165.201.30
+>(tcl) #209.165.202.158
+>(tcl) #} {
+>(tcl) #trace $address source 192.168.1.1
+>(tcl) #
Type escape sequence to abort.
Tracing the route to 209.165.200.254
VRF info: (vrf in name/id, vrf out name/id)
 1 209.165.201.1 60 msec 48 msec 28 msec
Type escape sequence to abort.
Tracing the route to 209.165.201.30
VRF info: (vrf in name/id, vrf out name/id)
 1 209.165.201.1 44 msec 28 msec 36 msec
Type escape sequence to abort.
Tracing the route to 209.165.202.158
VRF info: (vrf in name/id, vrf out name/id)
 1 209.165.201.1 32 msec 32 msec 44 msec
 2 209.165.200.226 72 msec 68 msec 44 msec
R1(tcl) #
```

### Step 3: Configure IP SLA probes.

- a. Create an ICMP echo probe on R1 to the primary DNS server on ISP1 using the **ip sla** command.

```
R1#conf t
Enter configuration commands, one per line. End with CNTL/Z.
R1(config)#ip sla 11
R1(config-ip-sla)#icmp-echo 209.165.201.30
R1(config-ip-sla-echo)#frequency 10
R1(config-ip-sla-echo)#exit
R1(config)#ip sla schedule 11 life forever start-time now
R1(config) #
```

The operation number of 11 is only locally significant to the router. The frequency 10 command schedules the connectivity test to repeat every 10 seconds. The probe is scheduled to start now and to run forever.

- b. Verify the IP SLAs configuration of operation 11 using the **show ip sla configuration 11** command.

```
R1#show ip sla configuration 11
IP SLAs Infrastructure Engine-III
Entry number: 11
Owner:
Tag:
Operation timeout (milliseconds): 5000
Type of operation to perform: icmp-echo
Target address/Source address: 209.165.201.30/0.0.0.0
Type Of Service parameter: 0x0
Request size (ARR data portion): 28
Verify data: No
Vrf Name:
Schedule:
  Operation frequency (seconds): 10  (not considered if randomly scheduled)
  Next Scheduled Start Time: Start Time already passed
  Group Scheduled : FALSE
  Randomly Scheduled : FALSE
  Life (seconds): Forever
  Entry Ageout (seconds): never
  Recurring (Starting Everyday): FALSE
  Status of entry (SNMP RowStatus): Active
Threshold (milliseconds): 5000
Distribution Statistics:
  Number of statistic hours kept: 2
  Number of statistic distribution buckets kept: 1
  Statistic distribution interval (milliseconds): 20
Enhanced History:
History Statistics:
  Number of history Lives kept: 0
  Number of history Buckets kept: 15
  History Filter Type: None
```

The output lists the details of the configuration of operation 11. The operation is an ICMP echo to 209.165.201.30, with a frequency of 10 seconds, and it has already started (the start time has already passed).

- c. Issue the **show ip sla statistics** command to display the number of successes, failures, and results of the latest operations.

```
R1#show ip sla statistics
IPSLAs Latest Operation Statistics

IPSLA operation id: 11
  Latest RTT: 48 milliseconds
Latest operation start time: 07:31:32 UTC Sat Jun 17 2023
Latest operation return code: OK
Number of successes: 20
Number of failures: 0
Operation time to live: Forever
```

You can see that operation 11 has already succeeded five times, has had no failures, and the last operation returned an OK result.

- d. Although not actually required because IP SLA session 11 alone could provide the desired fault tolerance, create a second probe, 22, to test connectivity to the second DNS server located on router ISP2. You can copy and paste the following commands on R1.

```
R1#conf t
Enter configuration commands, one per line. End with CNTL/Z.
R1(config)#ip sla 22
R1(config-ip-sla)#icmp-echo 209.165.202.158
R1(config-ip-sla-echo)#frequency 10
R1(config-ip-sla-echo)#exit
R1(config)#ip sla schedule 22 life forever start-time now
R1(config) #
```

- e. Verify the new probe using the **show ip sla configuration** and **show ip sla statistics** commands.

```
R1#show ip sla configuration 22
IP SLAs Infrastructure Engine-III
Entry number: 22
Owner:
Tag:
Operation timeout (milliseconds): 5000
Type of operation to perform: icmp-echo
Target address/Source address: 209.165.202.158/0.0.0.0
Type Of Service parameter: 0x0
Request size (ARR data portion): 28
Verify data: No
Vrf Name:
Schedule:
  Operation frequency (seconds): 10 (not considered if randomly scheduled)
  Next Scheduled Start Time: Start Time already passed
  Group Scheduled : FALSE
  Randomly Scheduled : FALSE
  Life (seconds): Forever
  Entry Ageout (seconds): never
  Recurring (Starting Everyday): FALSE
  Status of entry (SNMP RowStatus): Active
Threshold (milliseconds): 5000
Distribution Statistics:
  Number of statistic hours kept: 2
  Number of statistic distribution buckets kept: 1
  Statistic distribution interval (milliseconds): 20
Enhanced History:
History Statistics:
  Number of history Lives kept: 0
  Number of history Buckets kept: 15
  History Filter Type: None
```

```
R1#show ip sla statistics 22
IPSLAs Latest Operation Statistics

IPSLA operation id: 22
    Latest RTT: 56 milliseconds
Latest operation start time: 07:37:01 UTC Sat Jun 17 2023
Latest operation return code: OK
Number of successes: 19
Number of failures: 0
Operation time to live: Forever
```

The output lists the details of the configuration of operation 22. The operation is an ICMP echo to 209.165.202.158, with a frequency of 10 seconds, and it has already started (the start time has already passed). The statistics also prove that operation 22 is active.

#### **Step 4: Configure tracking options.**

Although PBR could be used, you will configure a floating static route that appears or disappears depending on the success or failure of the IP SLA.

- Remove the current default route on R1, and replace it with a floating static route having an administrative distance of 5.

```
R1#conf t
Enter configuration commands, one per line. End with CNTL/Z.
R1(config)#no ip route 0.0.0.0 0.0.0.0 209.165.201.1
R1(config)#ip route 0.0.0.0 0.0.0.0 209.165.201.1 5
R1(config)#exit
R1#
```

- Verify the routing table.

```
R1#show ip route
Codes: L - local, C - connected, S - static, R - RIP, M - mobile, B - BGP
      D - EIGRP, EX - EIGRP external, O - OSPF, IA - OSPF inter area
      N1 - OSPF NSSA external type 1, N2 - OSPF NSSA external type 2
      E1 - OSPF external type 1, E2 - OSPF external type 2
      i - IS-IS, su - IS-IS summary, L1 - IS-IS level-1, L2 - IS-IS level-2
      ia - IS-IS inter area, * - candidate default, U - per-user static route
      o - ODR, P - periodic downloaded static route, H - NHRP, 1 - LISP
      + - replicated route, % - next hop override

Gateway of last resort is 209.165.201.1 to network 0.0.0.0

S*   0.0.0.0/0 [5/0] via 209.165.201.1
      192.168.1.0/24 is variably subnetted, 2 subnets, 2 masks
C     192.168.1.0/24 is directly connected, Loopback0
L     192.168.1.1/32 is directly connected, Loopback0
      209.165.201.0/24 is variably subnetted, 2 subnets, 2 masks
C     209.165.201.0/30 is directly connected, Serial3/0
L     209.165.201.2/32 is directly connected, Serial3/0
      209.165.202.0/24 is variably subnetted, 2 subnets, 2 masks
C     209.165.202.128/30 is directly connected, Serial3/1
L     209.165.202.130/32 is directly connected, Serial3/1
R1#
```

- Use the **track 1 ip sla 11 reachability** command to enter the config-track subconfiguration mode.

```
R1#conf t  
Enter configuration commands, one per line. End with CNTL/Z.  
R1(config)#track 1 ip sla 11 reachability  
R1(config-track) #
```

- d. Specify the level of sensitivity to changes of tracked objects to 10 seconds of down delay and 1 second of up delay using the **delay down 10 up 1** command. The delay helps to alleviate the effect of flapping objects—objects that are going down and up rapidly. In this situation, if the DNS server fails momentarily and comes back up within 10 seconds, there is no impact.

```
R1(config-track)#delay down 10 up 1  
R1(config-track)#exit  
R1(config) #
```

- e. Configure the floating static route that will be implemented when tracking object 1 is active. To view routing table changes as they happen, first enable the **debug ip routing** command. Next, use the **ip route 0.0.0.0 0.0.0.0 209.165.201.1 2 track 1** command to create a floating static default route via 209.165.201.1 (ISP1). Notice that this command references the tracking object number 1, which in turn references IP SLA operation number 11.

```
R1#debug ip routing  
IP routing debugging is on  
R1#
```

```
R1#conf t  
Enter configuration commands, one per line. End with CNTL/Z.  
R1(config)#ip route 0.0.0.0 0.0.0.0 209.165.201.1 2 track 1  
R1(config) #  
*Jun 17 07:44:51.595: RT: updating static 0.0.0.0/0 (0x0):  
    via 209.165.201.1 1048578  
  
*Jun 17 07:44:51.599: RT: closer admin distance for 0.0.0.0, flushing 1 routes  
*Jun 17 07:44:51.603: RT: add 0.0.0.0/0 via 209.165.201.1, static metric [2/0]  
*Jun 17 07:44:51.603: RT: updating static 0.0.0.0/0 (0x0):  
    via 209.165.201.1 1048578  
  
*Jun 17 07:44:51.607: RT: rib update return code: 17  
*Jun 17 07:44:51.611: RT: updating static 0.0.0.0/0 (0x0):  
    via 209.165.201.1 1048578  
  
*Jun 17 07:44:51.615: RT: rib update return code: 17  
R1(config) #
```

- f. Repeat the steps for operation 22, track number 2, and assign the static route an admin distance higher than track 1 and lower than 5. On R1, copy the following configuration, which sets an admin distance of 3.

```
R1(config)#track 2 ip sla 22 reachability  
R1(config-track)#delay down 10 up 1  
R1(config-track)#exit  
R1(config) #
```

```
R1(config)#ip route 0.0.0.0 0.0.0.0 209.165.202.129 3 track 2
R1(config)#
*Jun 17 07:48:39.263: RT: updating static 0.0.0.0/0 (0x0):
  via 209.165.201.1  1048578

*Jun 17 07:48:39.267: RT: updating static 0.0.0.0/0 (0x0):
  via 209.165.201.1  1048578

*Jun 17 07:48:39.271: RT: rib update return code: 17
*Jun 17 07:48:39.279: RT: updating static 0.0.0.0/0 (0x0):
  via 209.165.202.129  1048578

*Jun 17 07:48:39.279: RT: rib update return code: 17
R1(config)#

```

g. Verify the routing table again.

```
R1#show ip route
Codes: L - local, C - connected, S - static, R - RIP, M - mobile, B - BGP
      D - EIGRP, EX - EIGRP external, O - OSPF, IA - OSPF inter area
      N1 - OSPF NSSA external type 1, N2 - OSPF NSSA external type 2
      E1 - OSPF external type 1, E2 - OSPF external type 2
      i - IS-IS, su - IS-IS summary, L1 - IS-IS level-1, L2 - IS-IS level-2
      ia - IS-IS inter area, * - candidate default, U - per-user static route
      o - ODR, P - periodic downloaded static route, H - NHRP, l - LIS
      + - replicated route, % - next hop override

Gateway of last resort is 209.165.201.1 to network 0.0.0.0

S*   0.0.0.0/0 [2/0] via 209.165.201.1
      192.168.1.0/24 is variably subnetted, 2 subnets, 2 masks
C     192.168.1.0/24 is directly connected, Loopback0
L     192.168.1.1/32 is directly connected, Loopback0
      209.165.201.0/24 is variably subnetted, 2 subnets, 2 masks
C     209.165.201.0/30 is directly connected, Serial3/0
L     209.165.201.2/32 is directly connected, Serial3/0
      209.165.202.0/24 is variably subnetted, 2 subnets, 2 masks
C     209.165.202.128/30 is directly connected, Serial3/1
L     209.165.202.130/32 is directly connected, Serial3/1
R1#
```

Although a new default route was entered, its administrative distance is not better than 2. Therefore, it does not replace the previously entered default route.

### **Step 5: Verify IP SLA operation.**

In this step you observe and verify the dynamic operations and routing changes when tracked objects fail. The following summarizes the process:

- Disable the DNS loopback interface on ISP1 (R2).
- Observe the output of the debug command on R1.
- Verify the static route entries in the routing table and the IP SLA statistics of R1.
- Re-enable the loopback interface on ISP1 (R2) and again observe the operation of the IP SLA tracking feature.

```

ISP1(config)#interface loopback 1
ISP1(config-if)#shutdown
ISP1(config-if)#
*Jun 17 07:51:45.643: %LINEPROTO-5-UPDOWN: Line protocol on Interface Loopback1,
  changed state to down
*Jun 17 07:51:45.647: %LINK-5-CHANGED: Interface Loopback1, changed state to adm
inistratively down
ISP1(config-if)#

```

- a. Shortly after the loopback interface is administratively down, observe the debug output being generated on R1.

```

R1#
*Jun 17 07:51:57.099: %TRACKING-5-STATE: 1 ip sla 11 reachability Up->Down
*Jun 17 07:51:57.099: RT: del 0.0.0.0 via 209.165.201.1, static metric [2/0]
*Jun 17 07:51:57.099: RT: delete network route to 0.0.0.0/0
*Jun 17 07:51:57.099: RT: default path has been cleared
*Jun 17 07:51:57.099: RT: updating static 0.0.0.0/0 (0x0):
  via 209.165.202.129  1048578

*Jun 17 07:51:57.099: RT: add 0.0.0.0/0 via 209.165.202.129, static metric [3/0]
*Jun 17 07:51:57.099: RT: default path is now 0.0.0.0 via 209.165.202.129
*Jun 17 07:51:57.099: RT: updating static 0.0.0.0/0 (0x0):
  via 209.165.201.1  1048578

*Jun 17 07:51:57.099: RT: rib update return code: 17
*Jun 17 07:51:57.115: RT: updating static 0.0.0.0/0 (0x0):
  via 209.165.202.129  1048578

R1#
*Jun 17 07:51:57.123: RT: updating static 0.0.0.0/0 (0x0):
  via 209.165.201.1  1048578

*Jun 17 07:51:57.127: RT: rib update return code: 17
R1#

```

The tracking state of track 1 changes from up to down. This is the object that tracked reachability for IP SLA object 11, with an ICMP echo to the ISP1 DNS server at 209.165.201.30.

R1 then proceeds to delete the default route with the administrative distance of 2 and installs the next highest default route to ISP2 with the administrative distance of 3.

- b. On R1, verify the routing table.

```

R1#show ip route
Codes: L - local, C - connected, S - static, R - RIP, M - mobile, B - BGP
      D - EIGRP, EX - EIGRP external, O - OSPF, IA - OSPF inter area
      N1 - OSPF NSSA external type 1, N2 - OSPF NSSA external type 2
      E1 - OSPF external type 1, E2 - OSPF external type 2
      i - IS-IS, su - IS-IS summary, L1 - IS-IS level-1, L2 - IS-IS level-2
      ia - IS-IS inter area, * - candidate default, U - per-user static route
      o - ODR, P - periodic downloaded static route, H - NHRP, l - LISP
      + - replicated route, % - next hop override

Gateway of last resort is 209.165.202.129 to network 0.0.0.0

S*    0.0.0.0/0 [3/0] via 209.165.202.129
      192.168.1.0/24 is variably subnetted, 2 subnets, 2 masks
C      192.168.1.0/24 is directly connected, Loopback0
L      192.168.1.1/32 is directly connected, Loopback0
      209.165.201.0/24 is variably subnetted, 2 subnets, 2 masks
C      209.165.201.0/30 is directly connected, Serial3/0
L      209.165.201.2/32 is directly connected, Serial3/0
      209.165.202.0/24 is variably subnetted, 2 subnets, 2 masks
C      209.165.202.128/30 is directly connected, Serial3/1
L      209.165.202.130/32 is directly connected, Serial3/1
R1#

```

The new static route has an administrative distance of 3 and is being forwarded to ISP2 as it should.

- c. Verify the IP SLA statistics.

```
R1#show ip sla statistics
IPSLAs Latest Operation Statistics

IPSLA operation id: 11
    Latest RTT: NoConnection/Busy/Timeout
Latest operation start time: 07:54:32 UTC Sat Jun 17 2023
Latest operation return code: Timeout
Number of successes: 138
Number of failures: 20
Operation time to live: Forever


IPSLA operation id: 22
    Latest RTT: 8 milliseconds
Latest operation start time: 07:54:31 UTC Sat Jun 17 2023
Latest operation return code: OK
Number of successes: 122
Number of failures: 2
Operation time to live: Forever
```

- d. On R1, initiate a trace to the web server from the internal LAN IP address.

```
R1#trace 209.165.200.254 source 192.168.1.1
Type escape sequence to abort.
Tracing the route to 209.165.200.254
VRF info: (vrf in name/id, vrf out name/id)
  1 209.165.202.129 76 msec 56 msec 48 msec
R1#
```

This confirms that traffic is leaving router R1 and being forwarded to the ISP2 router.

- e. On ISP1, re-enable the DNS address by issuing the no shutdown command on the loopback 1 interface to examine the routing behavior when connectivity to the ISP1 DNS is restored.

```
ISP1(config-if)#no shutdown
ISP1(config-if)#
*Jun 17 07:56:17.731: %LINEPROTO-5-UPDOWN: Line protocol on Interface Loopback1, changed state to up
ISP1(config-if)#
*Jun 17 07:56:17.735: %LINK-3-UPDOWN: Interface Loopback1, changed state to up
ISP1(config-if) #
```

Notice the output of the **debug ip routing** command on R1.

```

R1#
*Jun 17 07:56:28.099: %TRACKING-5-STATE: 1 ip sla 11 reachability Down->Up
*Jun 17 07:56:28.099: RT: updating static 0.0.0.0/0 (0x0):
    via 209.165.201.1  1048578

*Jun 17 07:56:28.099: RT: closer admin distance for 0.0.0.0, flushing 1 routes
*Jun 17 07:56:28.103: RT: add 0.0.0.0/0 via 209.165.201.1, static metric [2/0]
*Jun 17 07:56:28.103: RT: updating static 0.0.0.0/0 (0x0):
    via 209.165.202.129  1048578

*Jun 17 07:56:28.107: RT: rib update return code: 17
*Jun 17 07:56:28.111: RT: updating static 0.0.0.0/0 (0x0):
    via 209.165.202.129  1048578

*Jun 17 07:56:28.115: RT: rib update return code: 17
*Jun 17 07:56:28.115: RT: updating static 0.0.0.0/0 (0x0):
    via 209.165.201.1  1048578

*Jun 17 07:56:28.115: RT:
R1#rib update return code: 17
R1#

```

Now the IP SLA 11 operation transitions back to an up state and reestablishes the default static route to ISP1 with an administrative distance of 2.

f. Again examine the IP SLA statistics.

```

R1#show ip sla statistics
IPSLAs Latest Operation Statistics

IPSLA operation id: 11
    Latest RTT: 24 milliseconds
Latest operation start time: 07:58:02 UTC Sat Jun 17 2023
Latest operation return code: OK
Number of successes: 149
Number of failures: 30
Operation time to live: Forever


IPSLA operation id: 22
    Latest RTT: 56 milliseconds
Latest operation start time: 07:58:01 UTC Sat Jun 17 2023
Latest operation return code: OK
Number of successes: 143
Number of failures: 2
Operation time to live: Forever

```

The IP SLA 11 operation is active again, as indicated by the OK return code, and the number of successes is incrementing.

g. Verify the routing table.

```
R1#show ip route
Codes: L - local, C - connected, S - static, R - RIP, M - mobile, B - BGP
      D - EIGRP, EX - EIGRP external, O - OSPF, IA - OSPF inter area
      N1 - OSPF NSSA external type 1, N2 - OSPF NSSA external type 2
      E1 - OSPF external type 1, E2 - OSPF external type 2
      i - IS-IS, su - IS-IS summary, L1 - IS-IS level-1, L2 - IS-IS level-2
      ia - IS-IS inter area, * - candidate default, U - per-user static route
      o - ODR, P - periodic downloaded static route, H - NHRP, 1 - LISP
      + - replicated route, % - next hop override

Gateway of last resort is 209.165.201.1 to network 0.0.0.0

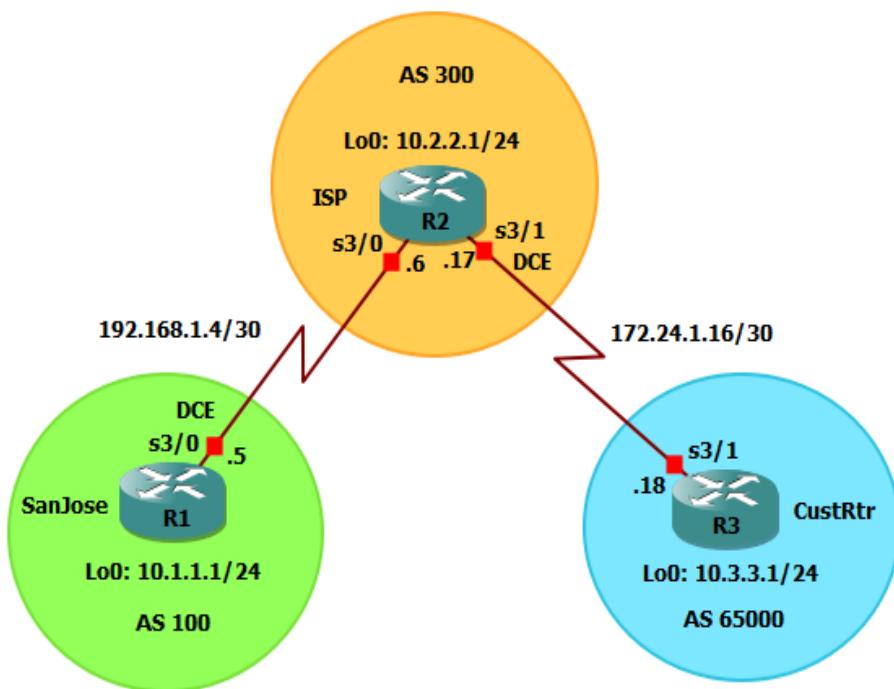
S*   0.0.0.0/0 [2/0] via 209.165.201.1
      192.168.1.0/24 is variably subnetted, 2 subnets, 2 masks
C     192.168.1.0/24 is directly connected, Loopback0
L     192.168.1.1/32 is directly connected, Loopback0
      209.165.201.0/24 is variably subnetted, 2 subnets, 2 masks
C     209.165.201.0/30 is directly connected, Serial3/0
L     209.165.201.2/32 is directly connected, Serial3/0
      209.165.202.0/24 is variably subnetted, 2 subnets, 2 masks
C     209.165.202.128/30 is directly connected, Serial3/1
L     209.165.202.130/32 is directly connected, Serial3/1
R1#
```

The default static through ISP1 with an administrative distance of 2 is restablished.

## Practical no 2

Aim:- Using the AS\_PATH Attribute.

### Topology



### Step 1: Prepare the routers for the lab.

Cable the network as shown in the topology diagram. Erase the startup configuration and reload each router to clear previous configurations.

### Step 2: Configure the hostname and interface addresses.

- You can copy and paste the following configurations into your routers to begin.

#### Router R1 (hostname SanJose)

```
R1#conf t
Enter configuration commands, one per line. End with CNTL/Z.
R1(config)#hostname SanJose
SanJose(config)#interface Loopback0
SanJose(config-if)#ip address 10.1.1.1 255.255.255.0
SanJose(config-if)#interface S3/0
SanJose(config-if)#ip address 192.168.1.5 255.255.255.252
SanJose(config-if)#clock rate 128000
SanJose(config-if)#no shutdown
```

#### Router R2 (hostname ISP)

```

R2#conf t
Enter configuration commands, one per line. End with CNTL/Z.
R2(config)#hostname ISP
ISP(config)#interface Loopback0
ISP(config-if)#
*May 16 19:58:30.931: %LINEPROTO-5-UPDOWN: Line protocol on Interface Loopback0,
  changed state to up
ISP(config-if)#ip address 10.2.2.1 255.255.255.0
ISP(config-if)#interface S3/0
ISP(config-if)#ip address 192.168.1.6 255.255.255.252
ISP(config-if)#no shutdown
ISP(config-if)#interface
*May 16 19:59:37.587: %LINK-3-UPDOWN: Interface Serial3/0, changed state to up
ISP(config-if)#interface
*May 16 19:59:38.595: %LINEPROTO-5-UPDOWN: Line protocol on Interface Serial3/0,
  changed state to up
ISP(config-if)#interface S3/1
ISP(config-if)#ip address 172.24.1.17 255.255.255.252
ISP(config-if)#clock rate 128000
ISP(config-if)#no shutdown

```

### **Router R3 (hostname CustRtr)**

```

R3#conf t
Enter configuration commands, one per line. End with CNTL/Z.
R3(config)#hostname CustRtr
CustRtr(config)#interface Loopback0
CustRtr(config-if)#
*May 16 20:01:20.763: %LINEPROTO-5-UPDOWN: Line protocol on Interface Loopback0,
  changed state to up
CustRtr(config-if)#ip address 10.3.3.1 255.255.255.0
CustRtr(config-if)#interface S3/1
CustRtr(config-if)#ip address 172.24.1.18 255.255.255.252
CustRtr(config-if)#no shutdown
CustRtr(config-if)#

```

### **Step 3: Configure BGP.**

- Configure BGP for normal operation. Enter the appropriate BGP commands on each router so that they identify their BGP neighbors and advertise their loopback networks.

```

SanJose(config)#router bgp 100
SanJose(config-router)#neighbor 192.168.1.6 remote-as 300
SanJose(config-router)#network 10.1.1.0 mask 255.255.255.0
SanJose(config-router)#

```

```

ISP(config)#router bgp 300
ISP(config-router)#neighbor 192.168.1.5 remote-as 100
ISP(config-router)#neighbor 172.24.1.18 remote-as 65000
ISP(config-router)#network 10.2.2.0 mask 255.255.255.0
ISP(config-router)#

```

```

CustRtr(config)#router bgp 65000
CustRtr(config-router)#neighbor 172.24.1.17 remote-as 300
CustRtr(config-router)#network 10.3.3.0 mask 255.255.255.0

```

- b. Verify that these routers have established the appropriate neighbor relationships by issuing the show ip bgp neighbors command on each router.

```
ISP#show ip bgp neighbors
BGP neighbor is 172.24.1.18, remote AS 65000, external link
  BGP version 4, remote router ID 10.3.3.1
  BGP state = Established, up for 00:01:53
```

<output omitted>

```
BGP neighbor is 192.168.1.5, remote AS 100, external link
  BGP version 4, remote router ID 10.1.1.1
  BGP state = Established, up for 00:04:14
```

<output omitted>

#### **Step 4: Remove the private AS.**

- a. Display the SanJose routing table using the show ip route command. SanJose should have a route to both 10.2.2.0 and 10.3.3.0. Troubleshoot if necessary.

```
SanJose#show ip route
Codes: L - local, C - connected, S - static, R - RIP, M - mobile, B - BGP
      D - EIGRP, EX - EIGRP external, O - OSPF, IA - OSPF inter area
      N1 - OSPF NSSA external type 1, N2 - OSPF NSSA external type 2
      E1 - OSPF external type 1, E2 - OSPF external type 2
      i - IS-IS, su - IS-IS summary, L1 - IS-IS level-1, L2 - IS-IS level-2
      ia - IS-IS inter area, * - candidate default, U - per-user static route
      o - ODR, P - periodic downloaded static route, H - NHRP, l - LISP
      + - replicated route, % - next hop override

Gateway of last resort is not set

  10.0.0.0/8 is variably subnetted, 4 subnets, 2 masks
C        10.1.1.0/24 is directly connected, Loopback0
L        10.1.1.1/32 is directly connected, Loopback0
B        10.2.2.0/24 [20/0] via 192.168.1.6, 00:05:00
B        10.3.3.0/24 [20/0] via 192.168.1.6, 00:03:19
  192.168.1.0/24 is variably subnetted, 2 subnets, 2 masks
C        192.168.1.4/30 is directly connected, Serial3/0
L        192.168.1.5/32 is directly connected, Serial3/0
```

- b. Ping again, this time as an extended ping, sourcing from the Loopback0 interface address.

```
SanJose#ping
Protocol [ip]:
Target IP address: 10.3.3.1
Repeat count [5]:
Datagram size [100]:
Timeout in seconds [2]:
Extended commands [n]: y
Source address or interface: 10.1.1.1
Type of service [0]:
Set DF bit in IP header? [no]:
Validate reply data? [no]:
Data pattern [0xABCD]:
Loose, Strict, Record, Timestamp, Verbose[none]:
Sweep range of sizes [n]:
Type escape sequence to abort.
Sending 5, 100-byte ICMP Echos to 10.3.3.1, timeout is 2 seconds:
Packet sent with a source address of 10.1.1.1
!!!!!
Success rate is 100 percent (5/5), round-trip min/avg/max = 60/102/140 ms
```

- c. Check the BGP table from SanJose by using the show ip bgp command. Note the AS path for the 10.3.3.0 network. The AS 65000 should be listed in the path to 10.3.3.0.

```
SanJose#show ip bgp
BGP table version is 4, local router ID is 10.1.1.1
Status codes: s suppressed, d damped, h history, * valid, > best, i - internal,
               r RIB-failure, S Stale, m multipath, b backup-path, f RT-Filter,
               x best-external, a additional-path, c RIB-compressed,
Origin codes: i - IGP, e - EGP, ? - incomplete
RPKI validation codes: V valid, I invalid, N Not found

      Network          Next Hop            Metric LocPrf Weight Path
* 10.1.1.0/24        0.0.0.0              0        32768 i
* 10.2.2.0/24        192.168.1.6          0        0 300 i
* 10.3.3.0/24        192.168.1.6          0        0 300 65000 i
```

- d. Configure ISP to strip the private AS numbers from BGP routes exchanged with SanJose using the following commands.

```
ISP#conf t
Enter configuration commands, one per line. End with CNTL/Z.
ISP(config)#router bgp 300
ISP(config-router)#neighbor 192.168.1.5 remove-private-as
```

- e. SanJose should be able to ping 10.3.3.1 using its loopback 0 interface as the source of the ping.

```
SanJose#ping 10.3.3.1 source Lo0
Type escape sequence to abort.
Sending 5, 100-byte ICMP Echos to 10.3.3.1, timeout is 2 seconds:
Packet sent with a source address of 10.1.1.1
!!!!!
Success rate is 100 percent (5/5), round-trip min/avg/max = 44/63/88 ms
```

- f. Now check the BGP table on SanJose. The AS\_PATH to the 10.3.3.0 network should be AS 300. It no longer has the private AS in the path.

```
SanJose#show ip bgp
BGP table version is 5, local router ID is 10.1.1.1
Status codes: s suppressed, d damped, h history, * valid, > best, i - internal,
               r RIB-failure, S Stale, m multipath, b backup-path, f RT-Filter,
               x best-external, a additional-path, c RIB-compressed,
Origin codes: i - IGP, e - EGP, ? - incomplete
RPKI validation codes: V valid, I invalid, N Not found

      Network          Next Hop            Metric LocPrf Weight Path
* 10.1.1.0/24        0.0.0.0              0        32768 i
* 10.2.2.0/24        192.168.1.6          0        0 300 i
* 10.3.3.0/24        192.168.1.6          0        0 300 i
SanJose#
```

## Step 5: Use the AS\_PATH attribute to filter routes.

- a. Configure a special kind of access list to match BGP routes with an AS\_PATH attribute that both begins and ends with the number 100. Enter the following commands on ISP.

```
ISP#conf t
Enter configuration commands, one per line. End with CNTL/Z.
ISP(config)#ip as-path access-list 1 deny ^100$ 
ISP(config)#ip as-path access-list 1 permit .*
```

- b. Apply the configured access list using the neighbor command with the filter-list option.

```
ISP(config)#router bgp 300
ISP(config-router)#neighbor 172.24.1.18 filter-list 1 out
```

- c. Use the clear ip bgp \* command to reset the routing information. Wait several seconds and then check the routing table for ISP. The route to 10.1.1.0 should be in the routing table.

```
ISP#show ip route
Codes: L - local, C - connected, S - static, R - RIP, M - mobile, B - BGP
      D - EIGRP, EX - EIGRP external, O - OSPF, IA - OSPF inter area
      N1 - OSPF NSSA external type 1, N2 - OSPF NSSA external type 2
      E1 - OSPF external type 1, E2 - OSPF external type 2
      i - IS-IS, su - IS-IS summary, L1 - IS-IS level-1, L2 - IS-IS level-2
      ia - IS-IS inter area, * - candidate default, U - per-user static route
      o - ODR, P - periodic downloaded static route, H - NHRP, l - LISP
      + - replicated route, % - next hop override

Gateway of last resort is not set

      10.0.0.0/8 is variably subnetted, 4 subnets, 2 masks
B        10.1.1.0/24 [20/0] via 192.168.1.5, 00:15:20
C        10.2.2.0/24 is directly connected, Loopback0
L        10.2.2.1/32 is directly connected, Loopback0
B        10.3.3.0/24 [20/0] via 172.24.1.18, 00:12:03
      172.24.0.0/16 is variably subnetted, 2 subnets, 2 masks
C        172.24.1.16/30 is directly connected, Serial3/1
L        172.24.1.17/32 is directly connected, Serial3/1
      192.168.1.0/24 is variably subnetted, 2 subnets, 2 masks
C        192.168.1.4/30 is directly connected, Serial3/0
L        192.168.1.6/32 is directly connected, Serial3/0
```

- d. Check the routing table for CustRtr. It should not have a route to 10.1.1.0 in its routing table.

```
CustRtr#show ip route
Codes: L - local, C - connected, S - static, R - RIP, M - mobile, B - BGP
      D - EIGRP, EX - EIGRP external, O - OSPF, IA - OSPF inter area
      N1 - OSPF NSSA external type 1, N2 - OSPF NSSA external type 2
      E1 - OSPF external type 1, E2 - OSPF external type 2
      i - IS-IS, su - IS-IS summary, L1 - IS-IS level-1, L2 - IS-IS level-2
      ia - IS-IS inter area, * - candidate default, U - per-user static route
      o - ODR, P - periodic downloaded static route, H - NHRP, l - LISP
      + - replicated route, % - next hop override

Gateway of last resort is not set

      10.0.0.0/8 is variably subnetted, 3 subnets, 2 masks
B        10.2.2.0/24 [20/0] via 172.24.1.17, 00:13:43
C        10.3.3.0/24 is directly connected, Loopback0
L        10.3.3.1/32 is directly connected, Loopback0
      172.24.0.0/16 is variably subnetted, 2 subnets, 2 masks
C        172.24.1.16/30 is directly connected, Serial3/1
L        172.24.1.18/32 is directly connected, Serial3/1
CustRtr#
```

- e. Return to ISP and verify that the filter is working as intended. Issue the show ip bgp regexp ^100\$ command.

```
ISP#show ip bgp regexp ^100$
BGP table version is 4, local router ID is 10.2.2.1
Status codes: s suppressed, d damped, h history, * valid, > best, i - internal,
               r RIB-failure, S Stale, m multipath, b backup-path, f RT-Filter,
               x best-external, a additional-path, c RIB-compressed,
Origin codes: i - IGP, e - EGP, ? - incomplete
RPKI validation codes: V valid, I invalid, N Not found

      Network          Next Hop            Metric LocPrf Weight Path
*>  10.1.1.0/24      192.168.1.5        0          0 100 i
```

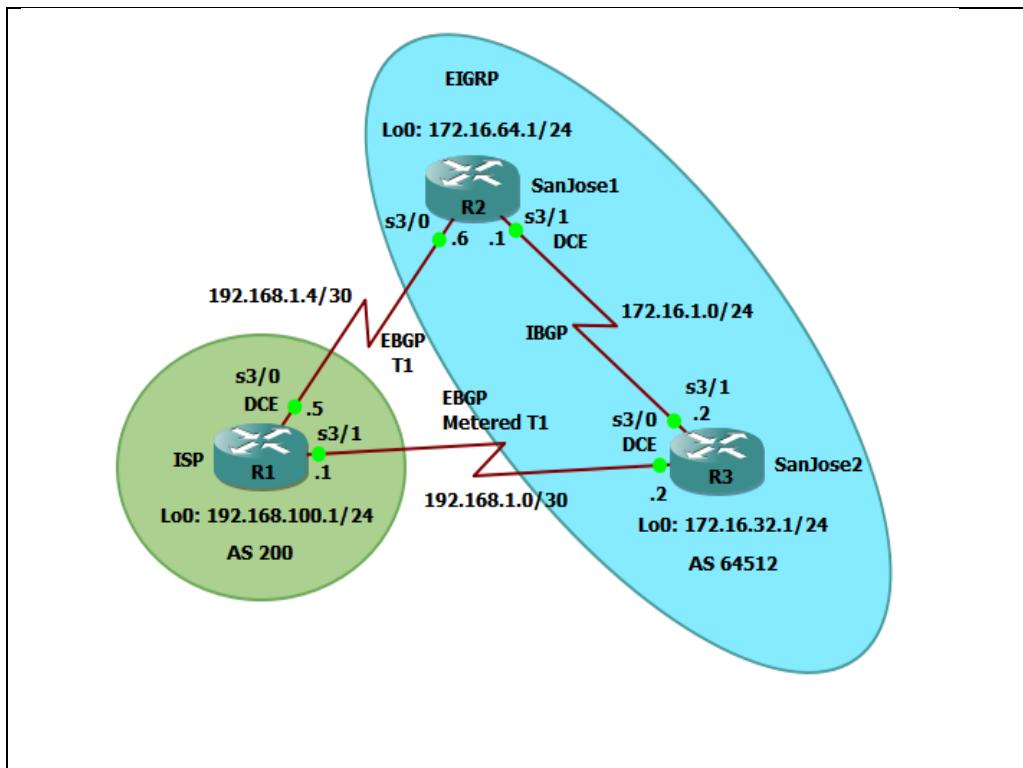
- f. Run the following Tcl script on all routers to verify whether there is connectivity. All pings from ISP should be successful. SanJose should not be able to ping the CustRtr loopback 10.3.3.1 or the WAN link 172.24.1.16/30. CustRtr should not be able to ping the SanJose loopback 10.1.1.1 or the WAN link 192.168.1.4/30.

```
ISP#tclsh
ISP(tcl)#foreach address {
+>10.1.1.1
+>10.2.2.1
+>10.3.3.1
+>192.168.1.5
+>192.168.1.6
+>172.24.1.17
+>172.24.1.18
+>} {
+>ping 192.168.1.5 }
Type escape sequence to abort.
Sending 5, 100-byte ICMP Echos to 192.168.1.5, timeout is 2 seconds:
!!!!!
Success rate is 100 percent (5/5), round-trip min/avg/max = 44/83/184 ms
Type escape sequence to abort.
Sending 5, 100-byte ICMP Echos to 192.168.1.5, timeout is 2 seconds:
!!!!!
Success rate is 100 percent (5/5), round-trip min/avg/max = 24/55/88 ms
Type escape sequence to abort.
Sending 5, 100-byte ICMP Echos to 192.168.1.5, timeout is 2 seconds:
!!!!!
Success rate is 100 percent (5/5), round-trip min/avg/max = 20/34/52 ms
Type escape sequence to abort.
Sending 5, 100-byte ICMP Echos to 192.168.1.5, timeout is 2 seconds:
!!!!!
Success rate is 100 percent (5/5), round-trip min/avg/max = 16/29/52 ms
Type escape sequence to abort.
Sending 5, 100-byte ICMP Echos to 192.168.1.5, timeout is 2 seconds:
!!!!!
Success rate is 100 percent (5/5), round-trip min/avg/max = 16/30/48 ms
Type escape sequence to abort.
Sending 5, 100-byte ICMP Echos to 192.168.1.5, timeout is 2 seconds:
!!!!!
Success rate is 100 percent (5/5), round-trip min/avg/max = 20/35/56 ms
Type escape sequence to abort.
Sending 5, 100-byte ICMP Echos to 192.168.1.5, timeout is 2 seconds:
!!!!!
Success rate is 100 percent (5/5), round-trip min/avg/max = 20/33/48 ms
ISP(tcl)#[
```

### Practical no 3

Aim:- Configuring IBGP and EBGP Sessions, Local Preference and MED

#### Topology



#### Step 0: Suggested starting configurations.

- Apply the following configuration to each router along with the appropriate **hostname**. The **exec-timeout 0 0** command should only be used in a lab environment.

```
R1#conf t
Enter configuration commands, one per line. End with CNTL/Z.
R1(config)#no ip domain-lookup
R1(config)#line con 0
R1(config-line)#logging synchronous
R1(config-line)#exec-timeout 0 0
R1(config-line)#exit
R1(config)#[
```

```
R2#conf t
Enter configuration commands, one per line. End with CNTL/Z.
R2(config)#no ip domain-lookup
R2(config)#line con 0
R2(config-line)#logging synchronous
R2(config-line)#exec-timeout 0 0
R2(config-line)#exit
R2(config)#[
```

```
R3#conf t
Enter configuration commands, one per line. End with CNTL/Z.
R3(config)#no ip domain-lookup
R3(config)#line con 0
R3(config-line)#logging synchronous
R3(config-line)#exec-timeout 0 0
R3(config-line)#exit
R3(config)#[
```

### Step 1: configure interface addresses.

- a. Using the addressing scheme in the diagram, create the loopback interfaces and apply IPv4 addresses to these and the serial interfaces on ISP (R1), SanJose1 (R2), and SanJose2 (R3).

#### Router R1 (hostname ISP)

```
R1(config)#hostname ISP
ISP(config)#interface Loopback0
ISP(config-if)#ip address 192.168.100.1 255.255.255.0
ISP(config-if)#exit
ISP(config)#interface S3/0
ISP(config-if)#ip address 192.168.1.5 255.255.255.252
ISP(config-if)#clock rate 128000
ISP(config-if)#no shutdown
ISP(config-if)#exit
```

```
[ISP(config)#interface S3/1
ISP(config-if)#ip address 192.168.1.1 255.255.255.252
ISP(config-if)#no shutdown
ISP(config-if)#end
```

#### Router R2 (hostname SanJose1)

```
R2(config)#hostname SanJosel
SanJosel(config)#interface Loopback0
SanJosel(config-if)#ip address 172.16.64.1 255.255.255.0
SanJosel(config-if)#exit
SanJosel(config)#interface S3/0
SanJosel(config-if)#ip address 192.168.1.6 255.255.255.252
SanJosel(config-if)#no shutdown
SanJosel(config-if)#exit
```

```
[SanJosel(config)#interface S3/1
SanJosel(config-if)#ip address 172.16.1.1 255.255.255.0
SanJosel(config-if)#clock rate 128000
SanJosel(config-if)#no shutdown
SanJosel(config-if)#end
```

#### Router R3 (hostname SanJose2)

```
R3(config)#hostname SanJose2
SanJose2(config)#interface Loopback0
SanJose2(config-if)#ip address 172.16.32.1 255.255.255.0
SanJose2(config-if)#exit
SanJose2(config)#interface S3/0
SanJose2(config-if)#ip address 192.168.1.2 255.255.255.252
SanJose2(config-if)#clock rate 128000
SanJose2(config-if)#no shutdown
SanJose2(config-if)#exit
```

```
SanJose2(config)#interface S3/1
SanJose2(config-if)#ip address 172.16.1.2 255.255.255.0
SanJose2(config-if)#no shutdown
SanJose2(config-if)#end
SanJose2#
```

- b. Use **ping** to test the connectivity between the directly connected routers. Both SanJose routers should be able to ping each other and their local ISP serial link IP address. The ISP router cannot reach the segment between SanJose1 and SanJose2.

## Step 2: Configure EIGRP.

Configure EIGRP between the SanJose1 and SanJose2 routers. (Note: If using an IOS prior to 15.0, use the no auto-summary router configuration command to disable automatic summarization. This command is the default beginning with IOS 15.)

```
SanJosel#conf t
Enter configuration commands, one per line.  End with CNTL/Z.
SanJosel(config)#router eigrp 1
SanJosel(config-router)#network 172.16.0.0
SanJosel(config-router)#
```

```
SanJose2#conf t
Enter configuration commands, one per line.  End with CNTL/Z.
SanJose2(config)#router eigrp 1
SanJose2(config-router)#network 172.16.0.0
SanJose2(config-router)#
```

## Step 3: Configure IBGP and verify BGP neighbors.

- a. Configure IBGP between the SanJose1 and SanJose2 routers. On the SanJose1 router, enter the following configuration.

```
SanJosel(config-router)#exit
SanJosel(config)#router bgp 64512
SanJosel(config-router)#neighbor 172.16.32.1 remote-as 64512
SanJosel(config-router)#neighbor 172.16.32.1 update-source lo0
SanJosel(config-router)#exit
SanJosel(config)#
```

If multiple pathways to the BGP neighbor exist, the router can use multiple IP interfaces to communicate with the neighbor. The source IP address therefore depends on the outgoing interface. The update-source lo0 command instructs the router to use the IP address of the interface Loopback0 as the source IP address for all BGP messages sent to that neighbor.

- b. Complete the IBGP configuration on SanJose2 using the following commands.

```
SanJose2(config-router)#exit
SanJose2(config)#router bgp 64512
SanJose2(config-router)#neighbor 172.16.64.1 remote-as 64512
SanJose2(config-router)#neighbor 172.16.64.1 update-source lo0
SanJose2(config-router)#exit
SanJose2(config)#exit
```

- C. Verify that SanJose1 and SanJose2 become BGP neighbors by issuing the **show ip bgp neighbors** command on SanJose1. View the following partial output. If the BGP state is not established, troubleshoot the connection.

```
SanJose2#show ip bgp neighbors
BGP neighbor is 172.16.64.1, remote AS 64512, internal link
  BGP version 4, remote router ID 172.16.64.1
  BGP state = Established, up for 00:01:43
  Last read 00:00:55, last write 00:00:45, hold time is 180, keepalive interval
is 60 seconds
  Neighbor sessions:
    1 active, is not multisession capable (disabled)
  Neighbor capabilities:
    Route refresh: advertised and received(new)
    Four-octets ASN Capability: advertised and received
    Address family IPv4 Unicast: advertised and received
    Enhanced Refresh Capability: advertised and received
    Multisession Capability:
      Stateful switchover support enabled: NO for session 1
  Message statistics:
    InQ depth is 0
    OutQ depth is 0

              Sent        Rcvd
  Opens:          1          1
  Notifications: 0          0
  Updates:        1          1
  Keepalives:     3          3
--More--
```

<Output Omitted>

The link between SanJose1 and SanJose2 should be identified as an internal link indicating an IBGP peering relationship, as shown in the output.

#### **Step 4: Configure EBGP and verify BGP neighbors.**

- Configure ISP to run EBGP with SanJose1 and SanJose2. Enter the following commands on ISP.

```
ISP#conf t
Enter configuration commands, one per line. End with CNTL/Z.
ISP(config)#router bgp 200
ISP(config-router)#neighbor 192.168.1.6 remote-as 64512
ISP(config-router)#neighbor 192.168.1.2 remote-as 64512
ISP(config-router)#network 192.168.100.0
```

Because EBGP sessions are almost always established over point-to-point links, there is no reason to use the **update-source** keyword in this configuration. Only one path exists between the peers. If this path goes down, alternative paths are not available.

- Configure a discard static route for the 172.16.0.0/16 network. Any packets that do not have a more specific match (longer match) for a 172.16.0.0 subnet will be dropped instead of sent to the ISP. Later in this lab we will configure a default route to the ISP.
- Configure SanJose1 as an EBGP peer to ISP.

```
SanJose1(config)#ip route 172.16.0.0 255.255.0.0 null0
SanJose1(config)#router bgp 64512
SanJose1(config-router)#neighbor 192.168.1.5 remote-as 200
SanJose1(config-router)#network 172.16.0.0
SanJose1(config-router)#end
SanJose1#
```

- d. Use the **show ip bgp neighbors** command to verify that SanJose1 and ISP have reached the established state. Troubleshoot if necessary.

```
SanJose1#show ip bgp neighbors
BGP neighbor is 172.16.32.1, remote AS 64512, internal link
  BGP version 4, remote router ID 172.16.32.1
  BGP state = Established, up for 00:07:35
  Last read 00:00:15, last write 00:00:12, hold time is 180, keepalive interval
is 60 seconds
  Neighbor sessions:
    1 active, is not multisession capable (disabled)
  Neighbor capabilities:
    Route refresh: advertised and received(new)
    Four-octets ASN Capability: advertised and received
    Address family IPv4 Unicast: advertised and received
    Enhanced Refresh Capability: advertised and received
    Multisession Capability:
      Stateful switchover support enabled: NO for session 1
  Message statistics:
    InQ depth is 0
    OutQ depth is 0

                                Sent          Rcvd
  Opens:                  1            1
  Notifications:         0            0
  Updates:                3            1
  Keepalives:             10           10
--More--
```

<Output Omitted>

- e. Configure a discard static route for 172.16.0.0/16 on SanJose2 and as an EBGP peer to ISP.

```
SanJose2(config)#ip route 172.16.0.0 255.255.0.0 null0
SanJose2(config)#router bgp 64512
SanJose2(config-router)#neighbor 192.168.1.1 remote-as 200
SanJose2(config-router)#network 172.16.0.0
SanJose2(config-router)#end
SanJose2#
*May 22 10:29:32.183: %SYS-5-CONFIG_I: Configured from console by console
SanJose2#
```

### Step 5: View BGP summary output.

In Step 4, the **show ip bgp neighbors** command was used to verify that SanJose1 and ISP had reached the established state. A useful alternative command is **show ip bgp summary**. The output should be similar to the following.

```

SanJose2#show ip bgp summary
BGP router identifier 172.16.32.1, local AS number 64512
BGP table version is 4, main routing table version 4
2 network entries using 288 bytes of memory
4 path entries using 320 bytes of memory
4/2 BGP path/bestpath attribute entries using 544 bytes of memory
1 BGP AS-PATH entries using 24 bytes of memory
0 BGP route-map cache entries using 0 bytes of memory
0 BGP filter-list cache entries using 0 bytes of memory
BGP using 1176 total bytes of memory
BGP activity 2/0 prefixes, 4/0 paths, scan interval 60 secs

Neighbor          V        AS MsgRcvd MsgSent   TblVer  InQ OutQ Up/Down  State
/PfxRcd
172.16.64.1      4        64512    21       20        4     0     0 00:14:14
  2
192.168.1.1      4        200      8       7        4     0     0 00:01:38
  1
SanJose2#

```

### **Step 6: Verify which path the traffic takes.**

- f. Clear the IP BGP conversation with the **clear ip bgp \*** command on ISP. Wait for the conversations to reestablish with each SanJose router.

```

ISP#clear ip bgp *
ISP#
*May 22 10:31:43.727: %BGP-5-ADJCHANGE: neighbor 192.168.1.2 Down User reset
*May 22 10:31:43.727: %BGP_SESSION-5-ADJCHANGE: neighbor 192.168.1.2 IPv4 Unicast topology base
removed from session User reset
*May 22 10:31:43.735: %BGP-5-ADJCHANGE: neighbor 192.168.1.6 Down User reset
*May 22 10:31:43.735: %BGP_SESSION-5-ADJCHANGE: neighbor 192.168.1.6 IPv4 Unicast topology base
removed from session User reset
*May 22 10:31:44.315: %BGP-5-ADJCHANGE: neighbor 192.168.1.6 Up
*May 22 10:31:44.315: %BGP-5-ADJCHANGE: neighbor 192.168.1.2 Up
ISP#

```

- g. Test whether ISP can ping the loopback 0 address of 172.16.64.1 on SanJose1 and the serial link between SanJose1 and SanJose2, 172.16.1.1

```

ISP#ping 172.16.64.1
Type escape sequence to abort.
Sending 5, 100-byte ICMP Echos to 172.16.64.1, timeout is 2 seconds:
.....
Success rate is 0 percent (0/5)
ISP#ping 172.16.1.1
Type escape sequence to abort.
Sending 5, 100-byte ICMP Echos to 172.16.1.1, timeout is 2 seconds:
.....
Success rate is 0 percent (0/5)
ISP#

```

- h. Now ping from ISP to the loopback 0 address of 172.16.32.1 on SanJose2 and the serial link between SanJose1 and SanJose2, 172.16.1.2.

```

ISP#ping 172.16.1.1
Type escape sequence to abort.
Sending 5, 100-byte ICMP Echos to 172.16.1.1, timeout is 2 seconds:
.....
Success rate is 0 percent (0/5)
ISP#ping 172.16.32.1
Type escape sequence to abort.
Sending 5, 100-byte ICMP Echos to 172.16.32.1, timeout is 2 seconds:
!!!!!
Success rate is 100 percent (5/5), round-trip min/avg/max = 36/43/60 ms
ISP#ping 172.16.1.2
Type escape sequence to abort.
Sending 5, 100-byte ICMP Echos to 172.16.1.2, timeout is 2 seconds:
!!!!!
Success rate is 100 percent (5/5), round-trip min/avg/max = 24/45/88 ms
ISP#

```

- i. Issue the **show ip bgp** command on ISP to verify BGP routes and metrics.

```

ISP#show ip bgp
BGP table version is 3, local router ID is 192.168.100.1
Status codes: s suppressed, d damped, h history, * valid, > best, i - internal,
               r RIB-failure, S Stale, m multipath, b backup-path, f RT-Filter,
               x best-external, a additional-path, c RIB-compressed,
Origin codes: i - IGP, e - EGP, ? - incomplete
RPKI validation codes: V valid, I invalid, N Not found

      Network          Next Hop            Metric LocPrf Weight Path
*>  172.16.0.0      192.168.1.2        0          0 64512 i
*   192.168.1.6      192.168.1.6        0          0 64512 i
*>  192.168.100.0   0.0.0.0          0          32768 i
ISP#

```

- j. At this point, the ISP router should be able to get to each network connected to SanJose1 and SanJose2 from the loopback address 192.168.100.1. Use the extended **ping** command and specify the source address of ISP Lo0 to test.

```

ISP#ping 172.16.1.1 source 192.168.100.1
Type escape sequence to abort.
Sending 5, 100-byte ICMP Echos to 172.16.1.1, timeout is 2 seconds:
Packet sent with a source address of 192.168.100.1
!!!!!
Success rate is 100 percent (5/5), round-trip min/avg/max = 92/108/128 ms
ISP#ping 172.16.32.1 source 192.168.100.1
Type escape sequence to abort.
Sending 5, 100-byte ICMP Echos to 172.16.32.1, timeout is 2 seconds:
Packet sent with a source address of 192.168.100.1
!!!!!
Success rate is 100 percent (5/5), round-trip min/avg/max = 44/84/140 ms
ISP#ping 172.16.1.2 source 192.168.100.1
Type escape sequence to abort.
Sending 5, 100-byte ICMP Echos to 172.16.1.2, timeout is 2 seconds:
Packet sent with a source address of 192.168.100.1
!!!!!
Success rate is 100 percent (5/5), round-trip min/avg/max = 44/76/116 ms
ISP#ping 172.16.64.1 source 192.168.100.1
Type escape sequence to abort.
Sending 5, 100-byte ICMP Echos to 172.16.64.1, timeout is 2 seconds:
Packet sent with a source address of 192.168.100.1
!!!!!
Success rate is 100 percent (5/5), round-trip min/avg/max = 72/104/136 ms

```

```

ISP#ping
Protocol [ip]:
Target IP address: 172.16.64.1
Repeat count [5]:
Datagram size [100]:
Timeout in seconds [2]:
Extended commands [n]: y
Source address or interface: 192.168.100.1
Type of service [0]:
Set DF bit in IP header? [no]:
Validate reply data? [no]:
Data pattern [0xABCD]:
Loose, Strict, Record, Timestamp, Verbose[none]:
Sweep range of sizes [n]:
Type escape sequence to abort.
Sending 5, 100-byte ICMP Echos to 172.16.64.1, timeout is 2 seconds:
Packet sent with a source address of 192.168.100.1
!!!!!
Success rate is 100 percent (5/5), round-trip min/avg/max = 80/120/140 ms
ISP#

```

Complete reachability has been demonstrated between the ISP router and both SanJose1 and SanJose2.

### **Step 7: Configure the BGP next-hop-self feature.**

SanJose1 is unaware of the link between ISP and SanJose2, and SanJose2 is unaware of the link between ISP and SanJose1. Before ISP can successfully ping all the internal serial interfaces of AS 64512, these serial links should be advertised via BGP on the ISP router. This can also be resolved via EIGRP on each SanJose router. One method is for ISP to advertise these links.

- Issue the following commands on the ISP router.

```

ISP#conf t
Enter configuration commands, one per line.  End with CNTL/Z.
ISP(config)#router bgp 200
ISP(config-router)#network 192.168.1.0 mask 255.255.255.252
ISP(config-router)#network 192.168.1.4 mask 255.255.255.252
ISP(config-router)#end
ISP#

```

- Issue the **show ip bgp** command to verify that the ISP is correctly injecting its own WAN links into BGP.

```

ISP#show ip bgp
BGP table version is 5, local router ID is 192.168.100.1
Status codes: s suppressed, d damped, h history, * valid, > best, i - internal,
              r RIB-failure, S Stale, m multipath, b backup-path, f RT-Filter,
              x best-external, a additional-path, c RIB-compressed,
Origin codes: i - IGP, e - EGP, ? - incomplete
RPKI validation codes: V valid, I invalid, N Not found

      Network          Next Hop            Metric LocPrf Weight Path
*->  172.16.0.0      192.168.1.2          0        0 64512 i
*          192.168.1.6          0        0 64512 i
*>  192.168.1.0/30   0.0.0.0          0        32768 i
*>  192.168.1.4/30   0.0.0.0          0        32768 i
*>  192.168.100.0    0.0.0.0          0        32768 i
ISP#

```

- c. Verify on SanJose1 and SanJose2 that the opposite WAN link is included in the routing table. The output from SanJose2 is as follows.

```
SanJose2#show ip route
Codes: L - local, C - connected, S - static, R - RIP, M - mobile, B - BGP
      D - EIGRP, EX - EIGRP external, O - OSPF, IA - OSPF inter area
      N1 - OSPF NSSA external type 1, N2 - OSPF NSSA external type 2
      E1 - OSPF external type 1, E2 - OSPF external type 2
      i - IS-IS, su - IS-IS summary, L1 - IS-IS level-1, L2 - IS-IS level-2
      ia - IS-IS inter area, * - candidate default, U - per-user static route
      o - ODR, P - periodic downloaded static route, H - NHRP, 1 - LIS
      + - replicated route, % - next hop override

Gateway of last resort is not set

      172.16.0.0/16 is variably subnetted, 6 subnets, 3 masks
S        172.16.0.0/16 is directly connected, Null0
C        172.16.1.0/24 is directly connected, Serial3/1
L        172.16.1.2/32 is directly connected, Serial3/1
C        172.16.32.0/24 is directly connected, Loopback0
L        172.16.32.1/32 is directly connected, Loopback0
D        172.16.64.0/24 [90/2297856] via 172.16.1.1, 00:30:54, Serial3/1
      192.168.1.0/24 is variably subnetted, 3 subnets, 2 masks
C        192.168.1.0/30 is directly connected, Serial3/0
L        192.168.1.2/32 is directly connected, Serial3/0
B        192.168.1.4/30 [20/0] via 192.168.1.1, 00:01:22
B        192.168.100.0/24 [20/0] via 192.168.1.1, 00:11:11
SanJose2#
```

The next issue to consider is BGP policy routing between autonomous systems. The next-hop attribute of a route in a different AS is set to the IP address of the border router in the next AS toward the destination, and this attribute is not modified by default when advertising this route through IBGP. Therefore, for all IBGP peers, it is either necessary to know the route to that border router (in a different neighboring AS), or our own border router needs to advertise the foreign routes using the next-hop-self feature, overriding the next-hop address with its own IP address. The SanJose2 router is passing a policy to SanJose1 and vice versa. The policy for routing from AS 64512 to AS 200 is to forward packets to the 192.168.1.1 interface. SanJose1 has a similar yet opposite policy: it forwards requests to the 192.168.1.5 interface. If either WAN link fails, it is critical that the opposite router become a valid gateway. This is achieved if the **next-hop-self** command is configured on SanJose1 and SanJose2.

- d. To better understand the **next-hop-self** command we will remove ISP advertising its two WAN links and shutdown the WAN link between ISP and SanJose2. The only possible path from SanJose2 to ISP's 192.168.100.0/24 is through SanJose1.

```
ISP#conf t
Enter configuration commands, one per line. End with CNTL/Z.
ISP(config)#router bgp 200
ISP(config-router)#no network 192.168.1.0 mask 255.255.255.252
ISP(config-router)#no network 192.168.1.4 mask 255.255.255.252
ISP(config-router)#exit
ISP(config)#interface S3/1
ISP(config-if)#shutdown
ISP(config-if)#
```

- e. Display SanJose2's BGP table using the **show ip bgp** command and the IPv4 routing table with **show ip route**.

```

SanJose2#show ip bgp
BGP table version is 13, local router ID is 172.16.32.1
Status codes: s suppressed, d damped, h history, * valid, > best, i - internal,
               r RIB-failure, S Stale, m multipath, b backup-path, f RT-Filter,
               x best-external, a additional-path, c RIB-compressed,
Origin codes: i - IGP, e - EGP, ? - incomplete
RPKI validation codes: V valid, I invalid, N Not found

      Network          Next Hop            Metric LocPrf Weight Path
* > 172.16.0.0        0.0.0.0              0        32768  i
* i                  172.16.64.1           0     100      0  i
* i 192.168.100.0    192.168.1.5          0     100      0 200 i
SanJose2#
SanJose2#show ip route
Codes: L - local, C - connected, S - static, R - RIP, M - mobile, B - BGP
       D - EIGRP, EX - EIGRP external, O - OSPF, IA - OSPF inter area
       N1 - OSPF NSSA external type 1, N2 - OSPF NSSA external type 2
       E1 - OSPF external type 1, E2 - OSPF external type 2
       i - IS-IS, su - IS-IS summary, L1 - IS-IS level-1, L2 - IS-IS level-2
       ia - IS-IS inter area, * - candidate default, U - per-user static route
       o - ODR, P - periodic downloaded static route, H - NHRP, l - LISP
       + - replicated route, % - next hop override

Gateway of last resort is not set

      172.16.0.0/16 is variably subnetted, 6 subnets, 3 masks
S       172.16.0.0/16 is directly connected, Null0
C       172.16.1.0/24 is directly connected, Serial3/1
L       172.16.1.2/32 is directly connected, Serial3/1
C       172.16.32.0/24 is directly connected, Loopback0
L       172.16.32.1/32 is directly connected, Loopback0
D       172.16.64.0/24 [90/2297856] via 172.16.1.1, 00:36:56, Serial3/1
SanJose2#

```

Notice that SanJose2 has 192.168.100.0 in its BGP table but not in its routing table. The BGP table shows the next hop to 192.168.100.0 as 192.168.1.5. Because SanJose2 does not have a route to this next hop address of 192.168.1.5 in its routing table, it will not install the 192.168.100.0 network into the routing table. It won't install a route if it doesn't know how to get to the next hop.

EBGP next hop addresses are carried into IBGP unchanged. As we saw previously, we could advertise the WAN link using BGP, but this is not always desirable. It means advertising additional routes when we are usually trying to minimize the size of the routing table. Another option is to have the routers within the IGP domain advertise themselves as the next hop router using the **next-hop-self** command.

- f. Issue the **next-hop-self** command on SanJose1 and SanJose2 to advertise themselves as the next hop to their IBGP peer.

```

SanJosel#conf t
Enter configuration commands, one per line. End with CNTL/Z.
SanJosel(config)#router bgp 64512
SanJosel(config-router)#neighbor 172.16.32.1 next-hop-self
SanJosel(config-router)#end
SanJosel#
SanJose2#conf t
Enter configuration commands, one per line. End with CNTL/Z.
SanJose2(config)#router bgp 64512
SanJose2(config-router)#neighbor 172.16.64.1 next-hop-self
SanJose2(config-router)#end
SanJose2#

```

g. Reset BGP operation on either router with the **clear ip bgp \*** command.

```
SanJose1#clear ip bgp *
SanJose1#
*May 22 10:53:41.855: %BGP-5-ADJCHANGE: neighbor 172.16.32.1 Down User reset
*May 22 10:53:41.855: %BGP_SESSION-5-ADJCHANGE: neighbor 172.16.32.1 IPv4 Unicast topology base removed from session User reset
*May 22 10:53:41.863: %BGP-5-ADJCHANGE: neighbor 192.168.1.5 Down User reset
*May 22 10:53:41.863: %BGP_SESSION-5-ADJCHANGE: neighbor 192.168.1.5 IPv4 Unicast topology base removed from session User reset
*May 22 10:53:42.759: %BGP-5-ADJCHANGE: neighbor 192.168.1.5 Up
*May 22 10:53:42.759: %BGP-5-ADJCHANGE: neighbor 172.16.32.1 Up
SanJose1# [REDACTED]
SanJose2#clear ip bgp *
SanJose2#
*May 22 10:54:42.771: %BGP-5-ADJCHANGE: neighbor 172.16.64.1 Down User reset
*May 22 10:54:42.771: %BGP_SESSION-5-ADJCHANGE: neighbor 172.16.64.1 IPv4 Unicast topology base removed from session User reset
*May 22 10:54:43.575: %BGP-5-ADJCHANGE: neighbor 172.16.64.1 Up
SanJose2# [REDACTED]
```

h. After the routers have returned to established BGP speakers, issue the **show ip bgp** command on SanJose2 and notice that the next hop is now SanJose1 instead of ISP.

```
SanJose2#show ip bgp
BGP table version is 3, local router ID is 172.16.32.1
Status codes: s suppressed, d damped, h history, * valid, > best, i - internal,
               r RIB-failure, S Stale, m multipath, b backup-path, f RT-Filter,
               x best-external, a additional-path, c RIB-compressed,
Origin codes: i - IGP, e - EGP, ? - incomplete
RPKI validation codes: V valid, I invalid, N Not found

      Network          Next Hop            Metric LocPrf Weight Path
* > 172.16.0.0        0.0.0.0              0        32768  i
*   i                  172.16.64.1           0       100      0  i
*>i 192.168.100.0    172.16.64.1           0       100      0  200  i
SanJose2# [REDACTED]
```

i. The **show ip route** command on SanJose2 now displays the 192.168.100.0/24 network because SanJose1 is the next hop, 172.16.64.1, which is reachable from SanJose2.

```
SanJose2#show ip route
Codes: L - local, C - connected, S - static, R - RIP, M - mobile, B - BGP
      D - EIGRP, EX - EIGRP external, O - OSPF, IA - OSPF inter area
      N1 - OSPF NSSA external type 1, N2 - OSPF NSSA external type 2
      E1 - OSPF external type 1, E2 - OSPF external type 2
      i - IS-IS, su - IS-IS summary, L1 - IS-IS level-1, L2 - IS-IS level-2
      ia - IS-IS inter area, * - candidate default, U - per-user static route
      o - ODR, P - periodic downloaded static route, H - NHRP, l - LIS
      + - replicated route, % - next hop override

Gateway of last resort is not set

      172.16.0.0/16 is variably subnetted, 6 subnets, 3 masks
S        172.16.0.0/16 is directly connected, Null0
C        172.16.1.0/24 is directly connected, Serial3/1
L        172.16.1.2/32 is directly connected, Serial3/1
C        172.16.32.0/24 is directly connected, Loopback0
L        172.16.32.1/32 is directly connected, Loopback0
D        172.16.64.0/24 [90/2297856] via 172.16.1.1, 00:44:41, Serial3/1
B        192.168.100.0/24 [200/0] via 172.16.64.1, 00:01:20
SanJose2# [REDACTED]
```

j. Before configuring the next BGP attribute, restore the WAN link between ISP and SanJose3. This will change the BGP table and routing table on both routers. For example, SanJose2's routing table shows 192.168.100.0/24 will now have a better path through ISP.

```
ISP(config)#interface S3/1
ISP(config-if)#no shutdown
ISP(config-if)#
SanJose2#show ip route
Codes: L - local, C - connected, S - static, R - RIP, M - mobile, B - BGP
      D - EIGRP, EX - EIGRP external, O - OSPF, IA - OSPF inter area
      N1 - OSPF NSSA external type 1, N2 - OSPF NSSA external type 2
      E1 - OSPF external type 1, E2 - OSPF external type 2
      i - IS-IS, su - IS-IS summary, L1 - IS-IS level-1, L2 - IS-IS level-2
      ia - IS-IS inter area, * - candidate default, U - per-user static route
      o - ODR, P - periodic downloaded static route, H - NHRP, l - LISP
      + - replicated route, % - next hop override

Gateway of last resort is not set

  172.16.0.0/16 is variably subnetted, 6 subnets, 3 masks
S        172.16.0.0/16 is directly connected, Null0
C        172.16.1.0/24 is directly connected, Serial3/1
L        172.16.1.2/32 is directly connected, Serial3/1
C        172.16.32.0/24 is directly connected, Loopback0
L        172.16.32.1/32 is directly connected, Loopback0
D        172.16.64.0/24 [90/2297856] via 172.16.1.1, 00:46:46, Serial3/1
  192.168.1.0/24 is variably subnetted, 2 subnets, 2 masks
C        192.168.1.0/30 is directly connected, Serial3/0
L        192.168.1.2/32 is directly connected, Serial3/0
B        192.168.100.0/24 [20/0] via 192.168.1.1, 00:00:51
SanJose2#
```

### Step 8: Set BGP local preference.

At this point, everything looks good, with the exception of default routes, the outbound flow of data, and inbound packet flow.

- a. Because the local preference value is shared between IBGP neighbors, configure a simple route map that references the local preference value on SanJose1 and SanJose2. This policy adjusts outbound traffic to prefer the link off the SanJose1 router instead of the metered T1 off SanJose2.

```
SanJosel#conf t
Enter configuration commands, one per line.  End with CNTL/Z.
SanJosel(config)#route-map PRIMARY_T1_IN permit 10
SanJosel(config-route-map)#set local-preference 150
SanJosel(config-route-map)#exit
SanJosel(config)#router bgp 64512
SanJosel(config-router)#neighbor 192.168.1.5 route-map PRIMARY_T1_IN in
SanJosel(config-router)#exit
SanJosel(config)#

```

```

SanJose2#conf t
Enter configuration commands, one per line. End with CNTL/Z.
SanJose2(config)#route-map SECONDARY_T1_IN permit 10
SanJose2(config-route-map)#set local-preference 125
SanJose2(config-route-map)#exit
SanJose2(config)#router bgp 64512
SanJose2(config-router)#neighbor 192.168.1.1 route-map SECONDARY_T1_IN in
SanJose2(config-router)#end
SanJose2#

```

- b. Use the **clear ip bgp \* soft** command after configuring this new policy. When the conversations have been reestablished, issue the **show ip bgp** command on SanJose1 and SanJose2.

```

SanJosel#clear ip bgp * soft
SanJosel#show ip bgp
BGP table version is 5, local router ID is 172.16.64.1
Status codes: s suppressed, d damped, h history, * valid, > best, i - internal,
              r RIB-failure, S Stale, m multipath, b backup-path, f RT-Filter,
              x best-external, a additional-path, c RIB-compressed,
Origin codes: i - IGP, e - EGP, ? - incomplete
RPKI validation codes: V valid, I invalid, N Not found

      Network          Next Hop            Metric LocPrf Weight Path
* i 172.16.0.0      172.16.32.1        0       100      0 i
*>                  0.0.0.0           0             32768 i
*> 192.168.100.0   192.168.1.5        0       150      0 200 i
SanJosel#

```

```

SanJose2#clear ip bgp * soft
SanJose2#show ip bgp
BGP table version is 5, local router ID is 172.16.32.1
Status codes: s suppressed, d damped, h history, * valid, > best, i - internal,
              r RIB-failure, S Stale, m multipath, b backup-path, f RT-Filter,
              x best-external, a additional-path, c RIB-compressed,
Origin codes: i - IGP, e - EGP, ? - incomplete
RPKI validation codes: V valid, I invalid, N Not found

      Network          Next Hop            Metric LocPrf Weight Path
*> 172.16.0.0      0.0.0.0           0             32768 i
* i                 172.16.64.1        0       100      0 i
* 192.168.100.0   192.168.1.1        0       125      0 200 i
*>i                172.16.64.1        0       150      0 200 i
SanJose2#

```

This now indicates that routing to the loopback segment for ISP 192.168.100.0 /24 can be reached only through the link common to SanJose1 and ISP. SanJose2's next hop to 192.168.100.0/24 is SanJose1 because both routers have been configured using the **next-hop-self** command.

### **Step 9: Set BGP MED.**

- a. In the previous step we saw that SanJose1 and SanJose2 will route traffic for 192.168.100.0/24 using the link between SanJose1 and ISP. Examine what the return path ISP takes to reach AS 64512. Notice that the return path is different from the original path. This is known as asymmetric routing and is not necessarily an unwanted trait.

```

ISP#show ip bgp
BGP table version is 11, local router ID is 192.168.100.1
Status codes: s suppressed, d damped, h history, * valid, > best, i - internal,
               r RIB-failure, S Stale, m multipath, b backup-path, f RT-Filter
               ,
               x best-external, a additional-path, c RIB-compressed,
Origin codes: i - IGP, e - EGP, ? - incomplete
RPKI validation codes: V valid, I invalid, N Not found

      Network          Next Hop            Metric LocPrf Weight Path
*   172.16.0.0        192.168.1.2        0          0 64512 i
*->                          192.168.1.6        0          0 64512 i
*>  192.168.100.0     0.0.0.0          0          32768 i
ISP#

```

```

ISP#show ip route
Codes: L - local, C - connected, S - static, R - RIP, M - mobile, B - BGP
       D - EIGRP, EX - EIGRP external, O - OSPF, IA - OSPF inter area
       N1 - OSPF NSSA external type 1, N2 - OSPF NSSA external type 2
       E1 - OSPF external type 1, E2 - OSPF external type 2
       i - IS-IS, su - IS-IS summary, L1 - IS-IS level-1, L2 - IS-IS level-2
       ia - IS-IS inter area, * - candidate default, U - per-user static route
       o - ODR, P - periodic downloaded static route, H - NHRP, l - LISP
       + - replicated route, % - next hop override

Gateway of last resort is not set

B    172.16.0.0/16 [20/0] via 192.168.1.6, 00:18:22
    192.168.1.0/24 is variably subnetted, 4 subnets, 2 masks
C      192.168.1.0/30 is directly connected, Serial3/1
L      192.168.1.1/32 is directly connected, Serial3/1
C      192.168.1.4/30 is directly connected, Serial3/0
L      192.168.1.5/32 is directly connected, Serial3/0
    192.168.100.0/24 is variably subnetted, 2 subnets, 2 masks
C      192.168.100.0/24 is directly connected, Loopback0
L      192.168.100.1/32 is directly connected, Loopback0
ISP#

```

- a. Use an extended **ping** command to verify this situation. Specify the **record** option and compare your output to the following. Notice the return path using the exit interface 192.168.1.1 to SanJose2.

```

SanJose2#ping
Protocol [ip]:
Target IP address: 192.168.100.1
Repeat count [5]:
Datagram size [100]:
Timeout in seconds [2]:
Extended commands [n]: y
Source address or interface: 172.16.32.1
Type of service [0]:
Set DF bit in IP header? [no]:
Validate reply data? [no]:
Data pattern [0xABCD]:
Loose, Strict, Record, Timestamp, Verbose[none]:
Sweep range of sizes [n]:
Type escape sequence to abort.
Sending 5, 100-byte ICMP Echos to 192.168.100.1, timeout is 2 seconds:
Packet sent with a source address of 172.16.32.1
!!!!!
Success rate is 100 percent (5/5), round-trip min/avg/max = 100/125/176 ms
SanJose2#

```

If you are unfamiliar with the **record** option, the important thing to note is that each IP address in brackets is an outgoing interface. The output can be interpreted as follows:

1. A ping that is sourced from 172.16.32.1 exits SanJose2 through s0/0/1, 172.16.1.2. It then arrives at the s0/0/1 interface for SanJose1.
2. SanJose1 S0/0/0, 192.168.1.6, routes the packet out to arrive at the S0/0/0 interface of ISP.
3. The target of 192.168.100.1 is reached: 192.168.100.1.
4. The packet is next forwarded out the S0/0/1, 192.168.1.1 interface for ISP and arrives at the S0/0/0 interface for SanJose2.
5. SanJose2 then forwards the packet out the last interface, loopback 0, 172.16.32.1.

Although the unlimited use of the T1 from SanJose1 is preferred here, ISP currently takes the link from SanJose2 for all return traffic.

- b. Create a new policy to force the ISP router to return all traffic via SanJose1. Create a second route map utilizing the MED (metric) that is shared between EBGP neighbors.

```
SanJose1#conf t
Enter configuration commands, one per line. End with CNTL/Z.
SanJose1(config)#route-map PRIMARY_T1_MED_OUT permit 10
SanJose1(config-route-map)#set Metric 50
SanJose1(config-route-map)#exit
SanJose1(config)#router bgp 64512
SanJose1(config-router)#neighbor 192.168.1.5 route-map PRIMARY_T1_MED_OUT out
SanJose1(config-router)#exit
SanJose1(config)#[
```

```
SanJose2#conf t
Enter configuration commands, one per line. End with CNTL/Z.
SanJose2(config)#route-map SECONDARY_T1_MED_OUT permit 10
SanJose2(config-route-map)#set Metric 75
SanJose2(config-route-map)#exit
SanJose2(config)#router bgp 64512
SanJose2(config-router)##$2.168.1.1 route-map SECONDARY_T1_MED_OUT out
SanJose2(config-router)#exit
SanJose2(config)#[
```

- c. Use the **clear ip bgp \* soft** command after issuing this new policy. Issuing the **show ip bgp** command as follows on SanJose1 or SanJose2 does not indicate anything about this newly defined policy.

```
SanJose1#clear ip bgp * soft
SanJose1#[
```

```
SanJose2#clear ip bgp * soft
SanJose2#[
```

```
SanJose1#show ip bgp
BGP table version is 5, local router ID is 172.16.64.1
Status codes: s suppressed, d damped, h history, * valid, > best, i - internal,
              r RIB-failure, S Stale, m multipath, b backup-path, f RT-Filter,
              x best-external, a additional-path, c RIB-compressed,
Origin codes: i - IGP, e - EGP, ? - incomplete
RPKI validation codes: V valid, I invalid, N Not found

      Network          Next Hop            Metric LocPrf Weight Path
* i 172.16.0.0        172.16.32.1        0     100      0 i
*>                  0.0.0.0             0           32768 i
*-> 192.168.100.0     192.168.1.5        0     150      0 200 i
SanJose1#[
```

```

SanJose2#show ip bgp
BGP table version is 5, local router ID is 172.16.32.1
Status codes: s suppressed, d damped, h history, * valid, > best, i - internal,
               r RIB-failure, S Stale, m multipath, b backup-path, f RT-Filter,
               x best-external, a additional-path, c RIB-compressed,
Origin codes: i - IGP, e - EGP, ? - incomplete
RPKI validation codes: V valid, I invalid, N Not found

      Network          Next Hop            Metric LocPrf Weight Path
* > 172.16.0.0        0.0.0.0             0          32768  i
* i                  172.16.64.1          0       100      0 i
*   192.168.100.0     192.168.1.1         0       125      0 200 i
*>i                172.16.64.1          0       150      0 200 i
SanJose2#

```

d. Reissue an extended **ping** command with the **record** command. Notice the change in return path using the exit interface 192.168.1.5 to SanJose1.

```

SanJose2#ping
Protocol [ip]:
Target IP address: 192.168.100.1
Repeat count [5]:
Datagram size [100]:
Timeout in seconds [2]:
Extended commands [n]: y
Source address or interface: 172.16.32.1
Type of service [0]:
Set DF bit in IP header? [no]:
Validate reply data? [no]:
Data pattern [0xABCD]:
Loose, Strict, Record, Timestamp, Verbose[none]:
Sweep range of sizes [n]:
Type escape sequence to abort.
Sending 5, 100-byte ICMP Echos to 192.168.100.1, timeout is 2 seconds:
Packet sent with a source address of 172.16.32.1
!!!!!
Success rate is 100 percent (5/5), round-trip min/avg/max = 108/150/180 ms
SanJose2#

```

The newly configured policy MED shows that the lower MED value is considered best. The ISP now prefers the route with the lower MED value of 50 to AS 64512. This is just opposite from the local-preference command configured earlier.

```

ISP#show ip bgp
BGP table version is 13, local router ID is 192.168.100.1
Status codes: s suppressed, d damped, h history, * valid, > best, i - internal,
               r RIB-failure, S Stale, m multipath, b backup-path, f RT-Filter,
               x best-external, a additional-path, c RIB-compressed,
Origin codes: i - IGP, e - EGP, ? - incomplete
RPKI validation codes: V valid, I invalid, N Not found

      Network          Next Hop            Metric LocPrf Weight Path
*    172.16.0.0        192.168.1.2          75          0 64512 i
* >   192.168.100.0     0.0.0.0             50          0 64512 i
*>   192.168.100.0     0.0.0.0             0          32768  i
ISP#

```

## Step 10: Establish a default route.

The final step is to establish a default route that uses a policy statement that adjusts to changes in the network.

- a. Configure ISP to inject a default route to both SanJose1 and SanJose2 using BGP using the **default-originate** command. This command does not require the presence of 0.0.0.0 in the ISP router. Configure the 10.0.0.0/8 network which will not be advertised using BGP. This network will be used to test the default route on SanJose1 and SanJose2.

```
ISP(config)#router bgp 200
ISP(config-router)#neighbor 192.168.1.6 default-originate
ISP(config-router)#neighbor 192.168.1.2 default-originate
ISP(config-router)#exit
ISP(config)#interface loopback 10
ISP(config-if)#ip address 10.0.0.1 255.255.255.0
ISP(config-if)#[
```

- b. Verify that both routers have received the default route by examining the routing tables on SanJose1 and SanJose2. Notice that both routers prefer the route between SanJose1 and ISP.

```
SanJosel#show ip bgp
BGP table version is 5, local router ID is 172.16.64.1
Status codes: s suppressed, d damped, h history, * valid, > best, i - internal,
              r RIB-failure, S Stale, m multipath, b backup-path, f RT-Filter,
              x best-external, a additional-path, c RIB-compressed,
Origin codes: i - IGP, e - EGP, ? - incomplete
RPKI validation codes: V valid, I invalid, N Not found

      Network          Next Hop            Metric LocPrf Weight Path
* i 172.16.0.0        172.16.32.1          0    100      0 i
 *->                   0.0.0.0             0          32768 i
 *> 192.168.100.0     192.168.1.5          0    150      0 200 i
SanJosel#show ip route
Codes: L - local, C - connected, S - static, R - RIP, M - mobile, B - BGP
       D - EIGRP, EX - EIGRP external, O - OSPF, IA - OSPF inter area
       N1 - OSPF NSSA external type 1, N2 - OSPF NSSA external type 2
       E1 - OSPF external type 1, E2 - OSPF external type 2
       i - IS-IS, su - IS-IS summary, L1 - IS-IS level-1, L2 - IS-IS level-2
       ia - IS-IS inter area, * - candidate default, U - per-user static route
       o - ODR, P - periodic downloaded static route, H - NHRP, l - LISP
       + - replicated route, % - next hop override

Gateway of last resort is 192.168.1.5 to network 0.0.0.0

B*   0.0.0.0/0 [20/0] via 192.168.1.5, 00:02:16
      172.16.0.0/16 is variably subnetted, 6 subnets, 3 masks
S     172.16.0.0/16 is directly connected, Null0
C     172.16.1.0/24 is directly connected, Serial3/1
L     172.16.1.1/32 is directly connected, Serial3/1
D     172.16.32.0/24 [90/2297856] via 172.16.1.2, 01:31:32, Serial3/1
C     172.16.64.0/24 is directly connected, Loopback0
L     172.16.64.1/32 is directly connected, Loopback0
      192.168.1.0/24 is variably subnetted, 2 subnets, 2 masks
C     192.168.1.4/30 is directly connected, Serial3/0
L     192.168.1.6/32 is directly connected, Serial3/0
B     192.168.100.0/24 [20/0] via 192.168.1.5, 00:36:46
SanJosel#[
```

```

SanJose2#show ip route
Codes: L - local, C - connected, S - static, R - RIP, M - mobile, B - BGP
      D - EIGRP, EX - EIGRP external, O - OSPF, IA - OSPF inter area
      N1 - OSPF NSSA external type 1, N2 - OSPF NSSA external type 2
      E1 - OSPF external type 1, E2 - OSPF external type 2
      i - IS-IS, su - IS-IS summary, L1 - IS-IS level-1, L2 - IS-IS level-2
      ia - IS-IS inter area, * - candidate default, U - per-user static route
      o - ODR, P - periodic downloaded static route, H - NHRP, l - LISPs
      + - replicated route, % - next hop override

Gateway of last resort is 172.16.64.1 to network 0.0.0.0

B*   0.0.0.0/0 [200/0] via 172.16.64.1, 00:03:22
      172.16.0.0/16 is variably subnetted, 6 subnets, 3 masks
S     172.16.0.0/16 is directly connected, Null0
C     172.16.1.0/24 is directly connected, Serial3/1
L     172.16.1.2/32 is directly connected, Serial3/1
C     172.16.32.0/24 is directly connected, Loopback0
L     172.16.32.1/32 is directly connected, Loopback0
D     172.16.64.0/24 [90/2297856] via 172.16.1.1, 01:32:38, Serial3/1
      192.168.1.0/24 is variably subnetted, 2 subnets, 2 masks
C     192.168.1.0/30 is directly connected, Serial3/0
L     192.168.1.2/32 is directly connected, Serial3/0
B     192.168.100.0/24 [200/0] via 172.16.64.1, 00:37:52
SanJose2#

```

c. The preferred default route is by way of SanJose1 because of the higher local preference attribute configured on SanJose1 earlier.

```

SanJose2#show ip bgp
BGP table version is 6, local router ID is 172.16.32.1
Status codes: s suppressed, d damped, h history, * valid, > best, i - internal,
               r RIB-failure, S Stale, m multipath, b backup-path, f RT-Filter,
               x best-external, a additional-path, c RIB-compressed,
Origin codes: i - IGP, e - EGP, ? - incomplete
RPKI validation codes: V valid, I invalid, N Not found

      Network          Next Hop            Metric LocPrf Weight Path
* 0.0.0.0           192.168.1.1        125    0 200 i
* >i               172.16.64.1          0      150    0 200 i
* > 172.16.0.0     0.0.0.0            0      32768 i
* i                172.16.64.1          0      100    0 i
* 192.168.100.0    192.168.1.1        0      125    0 200 i
* >i               172.16.64.1          0      150    0 200 i
SanJose2#

```

d. Using the traceroute command verify that packets to 10.0.0.1 is using the default route through SanJose1.

```

SanJose2#traceroute 10.0.0.1
Type escape sequence to abort.
Tracing the route to 10.0.0.1
VRF info: (vrf in name/id, vrf out name/id)
  1 172.16.1.1 92 msec 28 msec 88 msec
  2 192.168.1.5 [AS 200] 112 msec 156 msec 144 msec
SanJose2#

```

e. Next, test how BGP adapts to using a different default route when the path between SanJose1 and ISP goes down.

```

ISP(config)#interface S3/0
ISP(config-if)#shutdown
ISP(config-if)#
*May 22 11:48:07.007: %BGP-5-NBR_RESET: Neighbor 192.168.1.6 reset (Interface flap)
*May 22 11:48:07.031: %BGP-5-ADJCHANGE: neighbor 192.168.1.6 Down Interface flap
*May 22 11:48:07.031: %BGP_SESSION-5-ADJCHANGE: neighbor 192.168.1.6 IPv4 Unicast topology has been removed from session Interface flap
ISP(config-if)#
*May 22 11:48:08.975: %LINK-5-CHANGED: Interface Serial3/0, changed state to administratively down
*May 22 11:48:09.975: %LINEPROTO-5-UPDOWN: Line protocol on Interface Serial3/0, changed state to down
ISP(config-if)#

```

f. Verify that both routers are modified their routing tables with the default route using the path between SanJose2 and ISP.

```
SanJose1#show ip route
Codes: L - local, C - connected, S - static, R - RIP, M - mobile, B - BGP
      D - EIGRP, EX - EIGRP external, O - OSPF, IA - OSPF inter area
      N1 - OSPF NSSA external type 1, N2 - OSPF NSSA external type 2
      E1 - OSPF external type 1, E2 - OSPF external type 2
      i - IS-IS, su - IS-IS summary, L1 - IS-IS level-1, L2 - IS-IS level-2
      ia - IS-IS inter area, * - candidate default, U - per-user static route
      o - ODR, P - periodic downloaded static route, H - NHRP, l - LISP
      + - replicated route, % - next hop override

Gateway of last resort is 172.16.32.1 to network 0.0.0.0

B*   0.0.0.0/0 [200/0] via 172.16.32.1, 00:00:27
    172.16.0.0/16 is variably subnetted, 6 subnets, 3 masks
S     172.16.0.0/16 is directly connected, Null0
C     172.16.1.0/24 is directly connected, Serial3/1
L     172.16.1.1/32 is directly connected, Serial3/1
D     172.16.32.0/24 [90/2297856] via 172.16.1.2, 01:36:39, Serial3/1
C     172.16.64.0/24 is directly connected, Loopback0
L     172.16.64.1/32 is directly connected, Loopback0
B     192.168.100.0/24 [200/0] via 172.16.32.1, 00:00:27
SanJose1#
```

```
SanJose2#show ip route
Codes: L - local, C - connected, S - static, R - RIP, M - mobile, B - BGP
      D - EIGRP, EX - EIGRP external, O - OSPF, IA - OSPF inter area
      N1 - OSPF NSSA external type 1, N2 - OSPF NSSA external type 2
      E1 - OSPF external type 1, E2 - OSPF external type 2
      i - IS-IS, su - IS-IS summary, L1 - IS-IS level-1, L2 - IS-IS level-2
      ia - IS-IS inter area, * - candidate default, U - per-user static route
      o - ODR, P - periodic downloaded static route, H - NHRP, l - LISP
      + - replicated route, % - next hop override

Gateway of last resort is 192.168.1.1 to network 0.0.0.0

B*   0.0.0.0/0 [20/0] via 192.168.1.1, 00:01:10
    172.16.0.0/16 is variably subnetted, 6 subnets, 3 masks
S     172.16.0.0/16 is directly connected, Null0
C     172.16.1.0/24 is directly connected, Serial3/1
L     172.16.1.2/32 is directly connected, Serial3/1
C     172.16.32.0/24 is directly connected, Loopback0
L     172.16.32.1/32 is directly connected, Loopback0
D     172.16.64.0/24 [90/2297856] via 172.16.1.1, 01:37:22, Serial3/1
    192.168.1.0/24 is variably subnetted, 2 subnets, 2 masks
C     192.168.1.0/30 is directly connected, Serial3/0
L     192.168.1.2/32 is directly connected, Serial3/0
B     192.168.100.0/24 [20/0] via 192.168.1.1, 00:01:10
SanJose2#
```

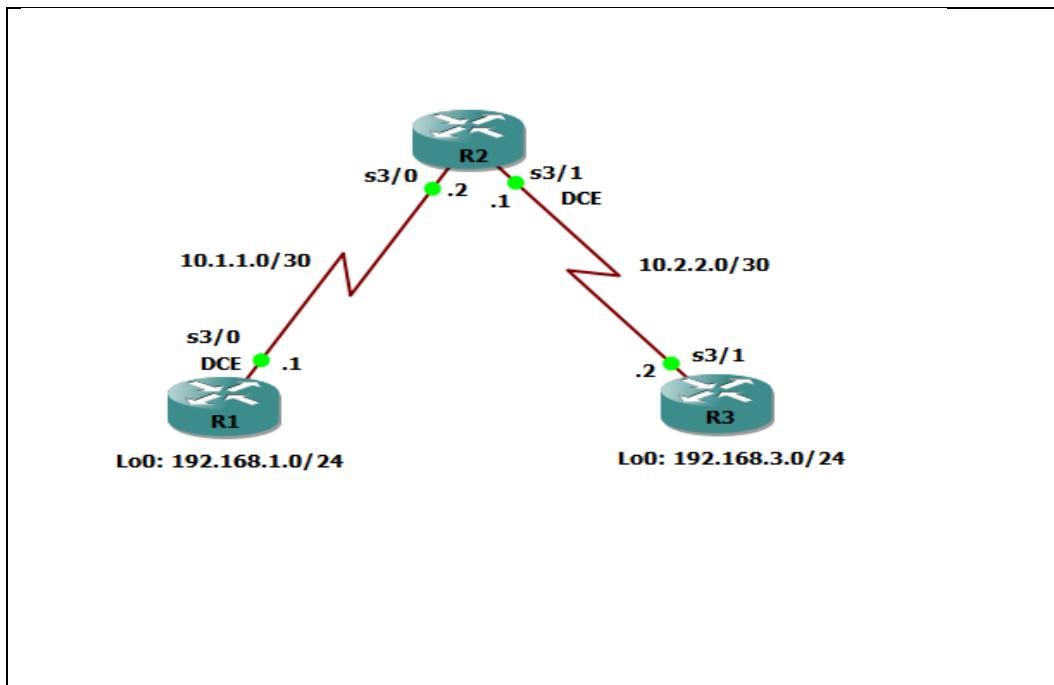
g. Verify the new path using the traceroute command to 10.0.0.1 from SanJose1. Notice the default route is now through SanJose2.

```
SanJose1#trace 10.0.0.1
Type escape sequence to abort.
Tracing the route to 10.0.0.1
VRF info: (vrf in name/id, vrf out name/id)
  1 172.16.1.2 56 msec 40 msec 28 msec
  2 192.168.1.1 [AS 200] 72 msec 76 msec 60 msec
SanJose1#
```

## Practical no 4

Aim:- Secure the Management Plane

### Topology



### Step 1: Configure loopbacks and assign addresses.

Cable the network as shown in the topology diagram. Erase the startup configuration and reload each router to clear previous configurations. Using the addressing scheme in the diagram, apply the IP addresses to the interfaces on the R1, R2, and R3 routers.

You can copy and paste the following configurations into your routers to begin.

Note: Depending on the router model, interfaces might be numbered differently than those listed. You might need to alter the designations accordingly.

R1

```
R1#conf t
Enter configuration commands, one per line.  End with CNTL/Z.
R1(config)#interface Loopback 0
R1(config-if)#ip address 192.168.1.1 255.255.255.0
R1(config-if)#exit
R1(config)#interface S3/0
R1(config-if)#ip address 10.1.1.1 255.255.255.252
R1(config-if)#clock rate 128000
R1(config-if)#no shutdown
R1(config-if)#exit
```

## R2

```
R2#conf t
Enter configuration commands, one per line. End with CNTL/Z.
R2(config)#hostname R2
R2(config)#interface S3/0
R2(config-if)#ip address 10.1.1.2 255.255.255.252
R2(config-if)#no shutdown
R2(config-if)#exit
```

```
R2(config)#interface S3/1
R2(config-if)#ip address 10.2.2.1 255.255.255.252
R2(config-if)#clock rate 128000
R2(config-if)#no shutdown
R2(config-if)#exit
R2(config)#end
```

## R3

```
R3#conf t
Enter configuration commands, one per line. End with CNTL/Z.
R3(config)#hostname R3
R3(config)#interface Loopback 0
R3(config-if)#
*May 22 12:20:27.183: %LINEPROTO-5-UPDOWN: Line protocol on Interface Loopback0,
 changed state to up
R3(config-if)#ip address 192.168.3.1 255.255.255.0
R3(config-if)#exit
R3(config)#interface S3/1
R3(config-if)#ip address 10.2.2.2 255.255.255.252
R3(config-if)#no shutdown
R3(config-if)#exit
R3(config)#end
R3#
```

### Step 2: Configure static routes.

- On R1, configure a default static route to ISP.

```
R1#conf t
Enter configuration commands, one per line. End with CNTL/Z.
R1(config)#ip route 0.0.0.0 0.0.0.0 10.1.1.2
R1(config)#
```

- On R2, configure two static route .

```
R2#conf t
Enter configuration commands, one per line. End with CNTL/Z.
R2(config)#ip route 192.168.1.0 255.255.255.0 10.1.1.1
R2(config)#ip route 192.168.3.0 255.255.255.0 10.2.2.2
R2(config)#
```

- On R3, configure a default static route to ISP.

```
R3#conf t
Enter configuration commands, one per line. End with CNTL/Z.
R3(config)#ip route 0.0.0.0 0.0.0.0 10.2.2.1
R3(config)#
```

- d. From the R1 router, run the following Tcl script to verify connectivity.

```
R1(tcl)#exit
R1#tclsh
R1(tcl)#foreach address {
+>(tcl) #192.168.1.1
+>(tcl) #10.1.1.1
+>(tcl) #10.1.1.2
+>(tcl) #10.2.2.1
+>(tcl) #10.2.2.2
+>(tcl) #192.168.3.1
+>(tcl) #} { ping 192.168.1.1 }
Type escape sequence to abort.
Sending 5, 100-byte ICMP Echos to 192.168.1.1, timeout is 2 seconds:
!!!!!
Success rate is 100 percent (5/5), round-trip min/avg/max = 4/7/8 ms
Type escape sequence to abort.
Sending 5, 100-byte ICMP Echos to 192.168.1.1, timeout is 2 seconds:
!!!!!
Success rate is 100 percent (5/5), round-trip min/avg/max = 1/1/1 ms
Type escape sequence to abort.
Sending 5, 100-byte ICMP Echos to 192.168.1.1, timeout is 2 seconds:
!!!!!
Success rate is 100 percent (5/5), round-trip min/avg/max = 1/2/8 ms
Type escape sequence to abort.
Sending 5, 100-byte ICMP Echos to 192.168.1.1, timeout is 2 seconds:
!!!!!
Success rate is 100 percent (5/5), round-trip min/avg/max = 1/1/1 ms
Type escape sequence to abort.
Sending 5, 100-byte ICMP Echos to 192.168.1.1, timeout is 2 seconds:
!!!!!
Success rate is 100 percent (5/5), round-trip min/avg/max = 1/2/4 ms
Type escape sequence to abort.
Sending 5, 100-byte ICMP Echos to 192.168.1.1, timeout is 2 seconds:
!!!!!
Success rate is 100 percent (5/5), round-trip min/avg/max = 1/1/1 ms
Go to Settings to activate Windows.
R1(tcl) #
```

### **Step 3: Secure management access.**

- On R1, use the security passwords command to set a minimum password length of 10 characters.
- Configure the enable secret encrypted password on both routers.

```
R1#conf t
Enter configuration commands, one per line. End with CNTL/Z.
R1(config)#security passwords min-length 10
R1(config)#enable secret class12345
```

- Configure a console password and enable login for routers. For additional security, the exec-timeout command causes the line to log out after 5 minutes of inactivity. The logging synchronous command prevents console messages from interrupting command entry.

```
R1(config)#line console 0
R1(config-line)#password ciscoconpass
R1(config-line)#exec-timeout 5 0
R1(config-line)#login
R1(config-line)#logging synchronous
R1(config-line)#exit
R1(config) #
```

Activate V  
Go to Settings

- d. Configure the password on the vty lines for router R1.

```
R1(config)#line vty 0 4
R1(config-line)#password ciscovtypass
R1(config-line)#exec-timeout 5 0
R1(config-line)#login
R1(config-line)#exit
```

- e. The aux port is a legacy port used to manage a router remotely using a modem and is hardly ever used. Therefore, disable the aux port.

```
R1(config)#line aux 0
R1(config-line)#no exec
R1(config-line)#end
R1#
```

- f. Enter privileged EXEC mode and issue the show run command. Can you read the enable secret password? Why or why not?
- g. Use the service password-encryption command to encrypt the line console and vty passwords.

```
R1#conf t
Enter configuration commands, one per line. End with CNTL/Z.
R1(config)#service password-encryption
R1(config) #
```

- h. Configure a warning to unauthorized users with a message-of-the-day (MOTD) banner using the banner motd command. When a user connects to one of the routers, the MOTD banner appears before the login prompt. In this example, the dollar sign (\$) is used to start and end the message.

```
R1#conf t
Enter configuration commands, one per line. End with CNTL/Z.
R1(config)#service password-encryption
R1(config)#banner motd $Unauthorized access strictly prohibited!$
R1(config)#exit
R1#
```

Activate Windows

- k. Repeat the configuration portion of steps 3a through 3k on router R3.

```
R3#conf t
Enter configuration commands, one per line. End with CNTL/Z.
R3(config)#security passwords min-length 10
R3(config)#enable secret class12345
```

```
R3(config)#line console 0
R3(config-line)#password ciscoconpass
R3(config-line)#exec-timeout 5 0
R3(config-line)#login
R3(config-line)#logging synchronous
R3(config-line)#exit
```

```
R3(config)#line vty 0 4
R3(config-line)#password ciscovtypass
R3(config-line)#exec-timeout 5 0
R3(config-line)#login
R3(config-line)#exit
R3(config)#line aux 0
R3(config-line)#no exec
R3(config-line)#end
R3#
```

#### Step 4: Configure enhanced username password security.

- To create local database entry encrypted to level 4 (SHA256), use the username name secret password global configuration command. In global configuration mode, enter the following command:

```
R1#conf t
Enter configuration commands, one per line. End with CNTL/Z.
R1(config)#username JR-ADMIN secret class12345 Activate Windows
R1(config)#username ADMIN secret class54321 Go to Settings to activate Windo
R1(config) #
```

- Set the console line to use the locally defined login accounts.

```
R1(config-line)#exit
R1(config)#line console 0
R1(config-line)#login local
R1(config-line)#exit
R1(config) #
```

- Set the vty lines to use the locally defined login accounts.

```
R1(config)#line vty 0 4
R1(config-line)#login local
R1(config-line)#end
R1#
```

- Repeat the steps 4a to 4c on R3.

```
R3#conf t
Enter configuration commands, one per line. End with CNTL/Z.
R3(config)#ip route 0.0.0.0 0.0.0.0 10.2.2.1
R3(config)#username JR-ADMIN secret class12345
R3(config)#username ADMIN secret class54321
```

```
R3(config)#line console 0
R3(config-line)#login local
R3(config-line)#exit
```

```
R3(config)#line vty 0 4
R3(config-line)#login local
R3(config-line)#end
R3#
```

- e. To verify the configuration, telnet to R3 from R1 and login using the ADMIN local database account

```
R1#telnet 10.2.2.2
Trying 10.2.2.2 ... Open

User Access Verification

Username: ADMIN
Password: [REDACTED]
R3>
```

#### Step 5: Enabling AAA RADIUS Authentication with Local User for Backup.

- Always have local database accounts created before enabling AAA. Since we created two local database accounts in the previous step, then we can proceed and enable AAA on R1.
- Configure the specifics for the first RADIUS server located at 192.168.1.101. Use RADIUS-1-pa55w0rd as the server password.

```
R1#conf t
Enter configuration commands, one per line. End with CNTL/Z.
R1(config)#aaa new-model
R1(config)#radius server RADIUS-1
R1(config-radius-server)#address ipv4 192.168.1.101
R1(config-radius-server)#key RADIUS-1-pa55w0rd
R1(config-radius-server)#exit
```

- Configure the specifics for the second RADIUS server located at 192.168.1.102. Use RADIUS-2-pa55w0rd as the server password.

```
R1(config)#radius server RADIUS-2
R1(config-radius-server)#address ipv4 192.168.1.102
R1(config-radius-server)#key RADIUS-2-pa55w0rd
R1(config-radius-server)#exit
```

- Assign both RADIUS servers to a server group.

```
R1(config)#aaa group server radius RADIUS-GROUP
R1(config-sg-radius)#server name RADIUS-1
R1(config-sg-radius)#server name RADIUS-2
R1(config-sg-radius)#exit
R1(config)#
Activate Windows  
Go to Settings to activate
```

- Enable the default AAA authentication login to attempt to validate against the server group. If they are not available, then authentication should be validated against the local database.

```
R1(config)#aaa authentication login default group RADIUS-GROUP local
R1(config)#
[REDACTED]
```

- f. Enable the default AAA authentication Telnet login to attempt to validate against the server group. If they are not available, then authentication should be validated against a case sensitive local database.

```
R1(config)#aaa authentication login TELNET-LOGIN group RADIUS-GROUP local-case
```

- g. Alter the VTY lines to use the TELNET-LGIN AAA authentiaitoOn method.

```
R1(config)#line vty 0 4
R1(config-line)#login authentication TELNET-LGIN
R1(config-line)#exit
R1(config)#[REDACTED]
```

- h. Repeat the steps 5a to 5g on R3.

```
R3#conf t
Enter configuration commands, one per line. End with CNTL/Z.
R3(config)#aaa new-model
R3(config)#[REDACTED]
```

```
R3(config)#radius server RADIUS-1
R3(config-radius-server)#address ipv4 192.168.1.101
R3(config-radius-server)#key RADIUS-1-pa55w0rd
R3(config-radius-server)#exit
R3(config)#[REDACTED]
```

```
R3(config)#radius server RADIUS-2
R3(config-radius-server)#address ipv4 192.168.1.102
R3(config-radius-server)#key RADIUS-2-pa55w0rd
R3(config-radius-server)#exit
R3(config)#[REDACTED]
```

```
R3(config)#aaa group server radius RADIUS-GROUP
R3(config-sg-radius)#server name RADIUS-1
R3(config-sg-radius)#server name RADIUS-2
R3(config-sg-radius)#exit
R3(config)#[REDACTED]
```

```
R3(config)#aaa authentication login default group RADIUS-GROUP local
R3(config)#[REDACTED]
```

```
R3(config)#aaa authentication login TELNET-LGIN group RADIUS-GROUP local-case
R3(config)#line vty 0 4
R3(config-line)#login authentication TELNET-LGIN
R3(config-line)#exit
R3(config)#[REDACTED]
```

- i. To verify the configuration, telnet to R3 from R1 and login using the ADMIN local database account.

```
R1#telnet 10.2.2.2
Trying 10.2.2.2 ... Open

User Access Verification

Username: admin
Password:

% Authentication failed

Username: ADMIN
Password:

R3>
```

#### Step 6: Enabling secure remote management using SSH.

- SSH requires that a device name and a domain name be configured. Since the router already has a name assigned, configure the domain name.

```
R1#conf t
Enter configuration commands, one per line. End with CNTL/Z.
R1(config)#ip domain-name ccnasecurity.com
```

- The router uses the RSA key pair for authentication and encryption of transmitted SSH data. Although optional it may be wise to erase any existing key pairs on the router.

```
R1(config)#crypto key zeroize rsa
% No Signature Keys found in configuration.
```

- Generate the RSA encryption key pair for the router. Configure the RSA keys with 1024 for the number of modulus bits. The default is 512, and the range is from 360 to 2048.

```
R1(config)#crypto key generate rsa general-keys modulus 1024
The name for the keys will be: R1.ccnasecurity.com

% The key modulus size is 1024 bits
% Generating 1024 bit RSA keys, keys will be non-exportable...
[OK] (elapsed time was 3 seconds)

R1(config)#
*May 22 13:54:08.811: %SSH-5-ENABLED: SSH 1.99 has been enabled
R1(config) #
```

- Cisco routers support two versions of SSH:

- SSH version 1 (SSHv1): Original version but has known vulnerabilities.
- SSH version 2 (SSHv2): Provides better security using the Diffie-Hellman key exchange and the strong integrity-checking message authentication code (MAC).

Configure SSH version 2 on R1.

```
R1(config)#ip ssh version 2
R1(config) #
```

- Configure the vty lines to use only SSH connections.

```
R1(config)#line vty 0 4
R1(config-line)#transport input ssh
R1(config-line)#end
R1#
```

f. Verify the SSH configuration using the show ip ssh command.

```
R1#show ip ssh
SSH Enabled - version 2.0
Authentication timeout: 120 secs; Authentication retries: 3
Minimum expected Diffie Hellman key size : 1024 bits
IOS Keys in SECSH format(ssh-rsa, base64 encoded):
ssh-rsa AAAAB3NzaC1yc2EAAAQABAAAAgQCdSdhVymeyiWrTKUAhgNKx1+Q3dthlsG9KDLt9nB
Aj
hqiJkkToVh5RhdFz3VAOn62hRjZ3nVDUS7Uaf2iFY2ZTuMD/rLrMUybeYP8Hf8E4VZMMo/AdqPqrRi
wB
bSWIzeJVhQie7dbb6H7bvbyzrx847jEl5+Vst1wiCrZbHOyXSQ==

R1#
```

g. Repeat the steps 6a to 6f on R3.

```
R3(config)#ip domain-name ccnasecurity.com
R3(config)#
```

```
R3(config)#crypto key zeroize rsa
% No Signature Keys found in configuration.
```

```
R3(config)#crypto key generate rsa general-keys modulus 1024
The name for the keys will be: R3.ccnasecurity.com

% The key modulus size is 1024 bits
% Generating 1024 bit RSA keys, keys will be non-exportable...
[OK] (elapsed time was 5 seconds)

R3(config)#
*May 22 14:02:03.515: %SSH-5-ENABLED: SSH 1.99 has been enabled
R3(config)#
```

```
R3(config)#ip ssh version 2
R3(config)#
```

```
R3(config)#line vty 0 4
R3(config-line)#transport input ssh
R3(config-line)#end
R3#
```

```
R3#show ip ssh
SSH Enabled - version 2.0
Authentication timeout: 120 secs; Authentication retries: 3
Minimum expected Diffie Hellman key size : 1024 bits
IOS Keys in SECSH format(ssh-rsa, base64 encoded):
ssh-rsa AAAAB3NzaC1yc2EAAAQABAAAAgQCXB0cO2yMVYpl3eT2GsD6vf4piyPnk/7BsggIIImFmP
Nht5jfiVx9V6NSi4AkSN0A2iG+iMfudRRkuY7Kmx7hRMpAuO23QGsghWaQFRLuMlneJHwlbUrjFhMnlm
vBMi5oTWhryRxTkbU9JiTkfkl/lpaXydJnBQbwUdl/tkxKdOfQw==
R3#
```

h. Although a user can SSH from a host using the SSH option of TeraTerm or PuTTY, a router can also SSH to another SSH enabled device. SSH to R3 from R1.

```
R1#ssh -l ADMIN 10.2.2.2
Password:

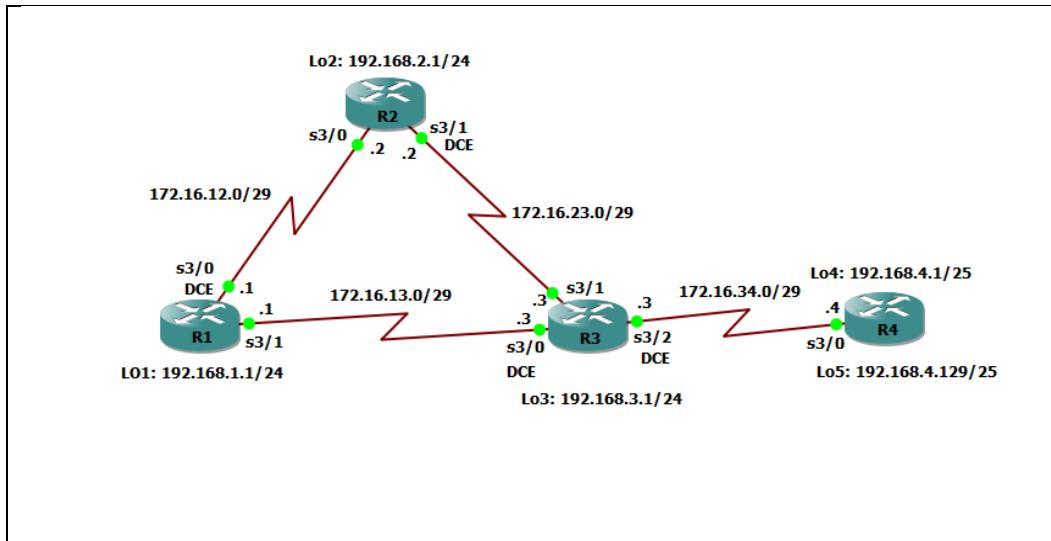
Password:

Password:
Unauthorized access strictly prohibited!R3>
R3>en
Password:
R3#
```

## Practical no 5

**Aim:-** Configure and Verify Path Control Using PBR.

### Topology



### Step 1: Prepare the routers for the lab.

Cable the network as shown in the topology diagram. Erase the startup configuration, and reload each router to clear previous configurations.

### Step 2: Configure router hostname and interface addresses.

- Using the addressing scheme in the diagram, create the loopback interfaces and apply IP addresses to these and the serial interfaces on R1, R2, R3, and R4. On the serial interfaces connecting R1 to R3 and R3 to R4, specify the bandwidth as 64 Kb/s and set a clock rate on the DCE using the clock rate 64000 command. On the serial interfaces connecting R1 to R2 and R2 to R3, specify the bandwidth as 128 Kb/s and set a clock rate on the DCE using the clock rate 128000 command.  
You can copy and paste the following configurations into your routers to begin.

#### Router R1

```
R1#conf t
Enter configuration commands, one per line. End with CNTL/Z.
R1(config)#hostname R1
R1(config)#interface Lo1
R1(config-if)#ip address 192.168.1.1 255.255.255.0
R1(config-if)#exit
```

```
R1(config)#interface S3/0
R1(config-if)#ip address 172.16.12.1 255.255.255.248
R1(config-if)#clock rate 128000
R1(config-if)#bandwidth 128
R1(config-if)#no shutdown
```

```
R1(config)#interface S3/1
R1(config-if)#ip address 172.16.13.1 255.255.255.248
R1(config-if)#bandwidth 64
R1(config-if)#no shutdown
R1(config-if)#end
```

## Router R2

```
R2#conf t
Enter configuration commands, one per line. End with CNTL/Z.
R2(config)#interface Lo2
R2(config-if)#ip address 192.168.2.1 255.255.255.0
R2(config-if)#exit
```

```
R2(config)#interface S3/0
R2(config-if)#ip address 172.16.12.2 255.255.255.248
R2(config-if)#bandwidth 128
R2(config-if)#no shutdown
R2(config-if)#exit
```

```
R2(config)#interface S3/1
R2(config-if)#ip address 172.16.23.2 255.255.255.248
R2(config-if)#clock rate 128000
R2(config-if)#bandwidth 128
R2(config-if)#no shutdown
R2(config-if)#end
```

## Router R3

```
R3#conf t
Enter configuration commands, one per line. End with CNTL/Z.
R3(config)#hostname R3
R3(config)#interface Lo3
R3(config-if)#ip address 192.168.3.1 255.255.255.0
R3(config-if)#exit
```

```
R3(config)#interface S3/0
R3(config-if)#ip address 172.16.13.3 255.255.255.248
R3(config-if)#clock rate 64000
R3(config-if)#bandwidth 64
R3(config-if)#no shutdown
R3(config-if)#exit
```

```
R3(config)#interface S3/1
R3(config-if)#ip address 172.16.23.3 255.255.255.248
R3(config-if)#bandwidth 128
R3(config-if)#no shutdown
R3(config-if)#exit
```

```
R3(config)#interface S3/2
R3(config-if)#ip address 172.16.34.3 255.255.255.248
R3(config-if)#clock rate 64000
R3(config-if)#bandwidth 64
R3(config-if)#no shutdown
R3(config-if)#end
```

## Router R4

```
R4#conf t
Enter configuration commands, one per line. End with CNTL/Z.
R4(config)#hostname R4
R4(config)#interface Lo4
R4(config-if)#ip address 192.168.4.1 255.255.255.128
```

```
R4(config)#interface Lo5
R4(config-if)#ip address 192.168.4.129 255.255.255.128
R4(config-if)#exit
```

```
R4(config)#interface S3/0
R4(config-if)#ip address 172.16.34.4 255.255.255.248
R4(config-if)#bandwidth 64
R4(config-if)#no shutdown
R4(config-if)#end
```

b. Verify the configuration with the show ip interface brief, show protocols, and show interfaces

description commands. The output from router R3 is shown here as an example.

```
R3#show ip interface brief
Interface          IP-Address      OK? Method Status          Protocol
FastEthernet0/0    unassigned      YES unset administratively down down
GigabitEthernet1/0 unassigned      YES unset administratively down down
GigabitEthernet2/0 unassigned      YES unset administratively down down
Serial3/0          172.16.13.3    YES manual up           up
Serial3/1          172.16.23.3    YES manual up           up
Serial3/2          172.16.34.3    YES manual up           up
Serial3/3          unassigned      YES unset administratively down down
Serial4/0          unassigned      YES unset administratively down down
Serial4/1          unassigned      YES unset administratively down down
Serial4/2          unassigned      YES unset administratively down down
Serial4/3          unassigned      YES unset administratively down down
Loopback3          192.168.3.1    YES manual up           up
R3#
```

```
R3#show protocols
Global values:
  Internet Protocol routing is enabled
FastEthernet0/0 is administratively down, line protocol is down
GigabitEthernet1/0 is administratively down, line protocol is down
GigabitEthernet2/0 is administratively down, line protocol is down
Serial3/0 is up, line protocol is up
  Internet address is 172.16.13.3/29
Serial3/1 is up, line protocol is up
  Internet address is 172.16.23.3/29
Serial3/2 is up, line protocol is up
  Internet address is 172.16.34.3/29
Serial3/3 is administratively down, line protocol is down
Serial4/0 is administratively down, line protocol is down
Serial4/1 is administratively down, line protocol is down
Serial4/2 is administratively down, line protocol is down
Serial4/3 is administratively down, line protocol is down
Loopback3 is up, line protocol is up
  Internet address is 192.168.3.1/24
R3#
```

```
R3#show interfaces description
Interface          Status      Protocol Description
Fa0/0              admin down  down
Gi1/0              admin down  down
Gi2/0              admin down  down
Se3/0              up          up
Se3/1              up          up
Se3/2              up          up
Se3/3              admin down  down
Se4/0              admin down  down
Se4/1              admin down  down
Se4/2              admin down  down
Se4/3              admin down  down
Lo3                up          up
R3#
```

### **Step 3: Configure basic EIGRP.**

- Implement EIGRP AS 1 over the serial and loopback interfaces as you have configured it for the other EIGRP labs.
- Advertise networks 172.16.12.0/29, 172.16.13.0/29, 172.16.23.0/29, 172.16.34.0/29, 192.168.1.0/24, 192.168.2.0/24, 192.168.3.0/24, and 192.168.4.0/24 from their respective routers.

You can copy and paste the following configurations into your routers.

#### **Router R1**

```
R1#conf t
Enter configuration commands, one per line. End with CNTL/Z.
R1(config)#router eigrp 1
R1(config-router)#network 192.168.1.0
R1(config-router)#network 172.16.12.0 0.0.0.7
R1(config-router)#network 172.16.13.0 0.0.0.7
R1(config-router)#no auto-summary
R1(config-router)#exit
R1(config) #
```

#### **Router R2**

```
R2#conf t
Enter configuration commands, one per line. End with CNTL/Z.
R2(config)#router eigrp 1
R2(config-router)#network 192.168.2.0
R2(config-router)#network 172.16.12.0 0.0.0.7
R2(config-router)#network 172.16.23.0 0.0.0.7
R2(config-router)#no auto-summary
R2(config-router)#exit
R2(config) #
```

## Router R3

```
R3(config)#router eigrp 1
R3(config-router)#network 192.168.3.0
R3(config-router)#network 172.16.13.0 0.0.0.7
R3(config-router)#network 172.16.23.0 0.0.0.7
R3(config-router)#network 172.16.34.0 0.0.0.7
R3(config-router)#no auto-summary
R3(config-router)#exit
R3(config) #
```

## Router R4

```
R4(config)#router eigrp 1
R4(config-router)#network 192.168.4.0
R4(config-router)#network 172.16.34.0 0.0.0.7
R4(config-router)#no auto-summary
R4(config-router)#exit
R4(config) #
```

You should see EIGRP neighbor relationship messages being generated.

### Step 4: Verify EIGRP connectivity.

- Verify the configuration by using the **show ip eigrp neighbors** command to check which routers have EIGRP adjacencies.

```
R1#show ip eigrp neighbors
EIGRP-IPv4 Neighbors for AS(1)
H   Address           Interface          Hold Uptime    SRTT     RTO   Q   Seq
   Address           Interface          (sec)   (ms)      Cnt Num
1   172.16.13.3       Se3/1            13 00:06:14  121  2340  0  13
0   172.16.12.2       Se3/0            14 00:08:36  73   1170  0  13
R1#
```

```
R2#show ip eigrp neighbors
EIGRP-IPv4 Neighbors for AS(1)
H   Address           Interface          Hold Uptime    SRTT     RTO   Q
   Address           Interface          (sec)   (ms)      Cnt
Seq
Num
1   172.16.23.3       Se3/1            11 00:06:26  75   1170  0
14
0   172.16.12.1       Se3/0            14 00:09:38  79   1170  0
14
R2#
```

```
R3#show ip eigrp neighbors
EIGRP-IPv4 Neighbors for AS(1)
H   Address           Interface          Hold Uptime    SRTT     RTO   Q   Seq
   Address           Interface          (sec)   (ms)      Cnt Num
2   172.16.34.4       Se3/2            11 00:06:16  78   2340  0  3
1   172.16.23.2       Se3/1            13 00:09:35  104  1170  0  12
0   172.16.13.1       Se3/0            10 00:10:25  87   2340  0  13
R3#
```

```
R4#show ip eigrp neighbors
EIGRP-IPv4 Neighbors for AS(1)
H   Address           Interface      Hold Uptime    SRTT     RTO   Q
Seq
                                         (sec)        (ms)       Cnt
Num
0   172.16.34.3          Se3/0          11  00:07:13   64  2340   0
15
R4#
```

- b. Run the following Tcl script on all routers to verify full connectivity.

```
R1#tclsh
R1(tcl)#foreach address {
+>(tcl)#172.16.12.1
+>(tcl)#172.16.12.2
+>(tcl)#172.16.13.1
+>(tcl)#172.16.13.3
+>(tcl)#172.16.23.2
+>(tcl)#172.16.23.3
+>(tcl)#172.16.34.3
+>(tcl)#172.16.34.4
+>(tcl)#192.168.1.1
+>(tcl)#192.168.2.1
+>(tcl)#192.168.3.1
+>(tcl)#192.168.4.1
+>(tcl)#192.168.4.129
+>(tcl)#{ ping 192.168.1.1 }
Type escape sequence to abort.
Sending 5, 100-byte ICMP Echos to 192.168.1.1, timeout is 2 seconds:
!!!!!
Success rate is 100 percent (5/5), round-trip min/avg/max = 4/5/8 ms
Type escape sequence to abort.
Sending 5, 100-byte ICMP Echos to 192.168.1.1, timeout is 2 seconds:
!!!!!
Success rate is 100 percent (5/5), round-trip min/avg/max = 4/6/8 ms
Type escape sequence to abort.
Sending 5, 100-byte ICMP Echos to 192.168.1.1, timeout is 2 seconds:
!!!!!
Success rate is 100 percent (5/5), round-trip min/avg/max = 4/6/8 ms
Type escape sequence to abort.
Sending 5, 100-byte ICMP Echos to 192.168.1.1, timeout is 2 seconds:
!!!!!
Success rate is 100 percent (5/5), round-trip min/avg/max = 8/8/8 ms
Type escape sequence to abort.
Sending 5, 100-byte ICMP Echos to 192.168.1.1, timeout is 2 seconds:
!!!!!
Success rate is 100 percent (5/5), round-trip min/avg/max = 4/5/8 ms
Type escape sequence to abort.
Sending 5, 100-byte ICMP Echos to 192.168.1.1, timeout is 2 seconds:
!!!!!
```

You should get ICMP echo replies for every address pinged. Make sure to run the Tcl script on each router.

#### **Step 5: Verify the current path.**

- a. On R1, use the show ip route command. Notice the next-hop IP address for all networks discovered by EIGRP.

```

R1#show ip route
Codes: L - local, C - connected, S - static, R - RIP, M - mobile, B - BGP
      D - EIGRP, EX - EIGRP external, O - OSPF, IA - OSPF inter area
      N1 - OSPF NSSA external type 1, N2 - OSPF NSSA external type 2
      E1 - OSPF external type 1, E2 - OSPF external type 2
      i - IS-IS, su - IS-IS summary, L1 - IS-IS level-1, L2 - IS-IS level-2
      ia - IS-IS inter area, * - candidate default, U - per-user static route
      o - ODR, P - periodic downloaded static route, H - NHRP, l - LISP
      + - replicated route, % - next hop override

Gateway of last resort is not set

  172.16.0.0/16 is variably subnetted, 6 subnets, 2 masks
C        172.16.12.0/29 is directly connected, Serial3/0
L        172.16.12.1/32 is directly connected, Serial3/0
C        172.16.13.0/29 is directly connected, Serial3/1
L        172.16.13.1/32 is directly connected, Serial3/1
D        172.16.23.0/29 [90/21024000] via 172.16.12.2, 00:14:57, Serial3/0
D        172.16.34.0/29 [90/41024000] via 172.16.13.3, 00:13:47, Serial3/1
  192.168.1.0/24 is variably subnetted, 2 subnets, 2 masks
C        192.168.1.0/24 is directly connected, Loopback1
L        192.168.1.1/32 is directly connected, Loopback1
D        192.168.2.0/24 [90/20640000] via 172.16.12.2, 00:14:56, Serial3/0
D        192.168.3.0/24 [90/21152000] via 172.16.12.2, 00:14:56, Serial3/0
  192.168.4.0/25 is subnetted, 2 subnets
D        192.168.4.0 [90/41152000] via 172.16.13.3, 00:11:37, Serial3/1
D        192.168.4.128 [90/41152000] via 172.16.13.3, 00:11:37, Serial3/1
R1#
R1#
```

- b. On R4, use the traceroute command to the R1 LAN address and source the ICMP packet from R4 LAN A and LAN B.

```

R4#traceroute 192.168.1.1 source 192.168.4.1
Type escape sequence to abort.
Tracing the route to 192.168.1.1
VRF info: (vrf in name/id, vrf out name/id)
  1  172.16.34.3  180 msec 120 msec 92 msec
  2  172.16.23.2  480 msec 160 msec 156 msec
  3  172.16.12.1  408 msec 212 msec 176 msec
R4#
```

```

R4#traceroute 192.168.1.1 source 192.168.4.129
Type escape sequence to abort.
Tracing the route to 192.168.1.1
VRF info: (vrf in name/id, vrf out name/id)
  1  172.16.34.3  96 msec 80 msec 80 msec
  2  172.16.23.2  132 msec 176 msec 180 msec
  3  172.16.12.1  184 msec 248 msec 224 msec
R4#
```

- c. On R3, use the show ip route command and note that the preferred route from R3 to R1 LAN

192.168.1.0/24 is via R2 using the R3 exit interface S3/1.

```

R3#show ip route
Codes: L - local, C - connected, S - static, R - RIP, M - mobile, B - BGP
      D - EIGRP, EX - EIGRP external, O - OSPF, IA - OSPF inter area
      N1 - OSPF NSSA external type 1, N2 - OSPF NSSA external type 2
      E1 - OSPF external type 1, E2 - OSPF external type 2
      i - IS-IS, su - IS-IS summary, L1 - IS-IS level-1, L2 - IS-IS level-2
      ia - IS-IS inter area, * - candidate default, U - per-user static route
      o - ODR, P - periodic downloaded static route, H - NHRP, l - LISP
      + - replicated route, % - next hop override

Gateway of last resort is not set

      172.16.0.0/16 is variably subnetted, 7 subnets, 2 masks
D        172.16.12.0/29 [90/21024000] via 172.16.23.2, 00:18:17, Serial3/1
C        172.16.13.0/29 is directly connected, Serial3/0
L        172.16.13.3/32 is directly connected, Serial3/0
C        172.16.23.0/29 is directly connected, Serial3/1
L        172.16.23.3/32 is directly connected, Serial3/1
C        172.16.34.0/29 is directly connected, Serial3/2
L        172.16.34.3/32 is directly connected, Serial3/2
D        192.168.1.0/24 [90/21152000] via 172.16.23.2, 00:18:17, Serial3/1
D        192.168.2.0/24 [90/20640000] via 172.16.23.2, 00:18:17, Serial3/1
      192.168.3.0/24 is variably subnetted, 2 subnets, 2 masks
C        192.168.3.0/24 is directly connected, Loopback3
L        192.168.3.1/32 is directly connected, Loopback3
      192.168.4.0/25 is subnetted, 2 subnets
D          192.168.4.0 [90/40640000] via 172.16.34.4, 00:14:58, Serial3/2
D          192.168.4.128 [90/40640000] via 172.16.34.4, 00:14:58, Serial3/2
R3#

```

d.On R3, use the show interfaces serial3/0 and show interfaces s3/1 commands.

```

R3#show interface S3/0
Serial3/0 is up, line protocol is up
  Hardware is M4T
  Internet address is 172.16.13.3/29
    MTU 1500 bytes, BW 64 Kbit/sec, DLY 20000 usec,
      reliability 255/255, txload 1/255, rxload 1/255
  Encapsulation HDLC, crc 16, loopback not set
  Keepalive set (10 sec)
  Restart-Delay is 0 secs
  Last input 00:00:03, output 00:00:02, output hang never
  Last clearing of "show interface" counters never
  Input queue: 0/75/0/0 (size/max/drops/flushes); Total output drops: 0
  Queueing strategy: weighted fair
  Output queue: 0/1000/64/0 (size/max total/threshold/drops)
    Conversations 0/1/256 (active/max active/max total)
    Reserved Conversations 0/0 (allocated/max allocated)
    Available Bandwidth 48 kilobits/sec
  5 minute input rate 0 bits/sec, 0 packets/sec
  5 minute output rate 0 bits/sec, 0 packets/sec
    630 packets input, 43784 bytes, 0 no buffer
    Received 292 broadcasts (0 IP multicasts)
    0 runts, 0 giants, 0 throttles
    0 input errors, 0 CRC, 0 frame, 0 overrun, 0 ignored, 0 abort
    568 packets output, 40361 bytes, 0 underruns
    0 output errors, 0 collisions, 2 interface resets
    0 unknown protocol drops
    0 output buffer failures, 0 output buffers swapped out
    2 carrier transitions      DCD=up  DSR=up  DTR=up  RTS=up  CTS=up

R3#

```

```

R3#show interfaces S3/1
Serial3/1 is up, line protocol is up
  Hardware is M4T
  Internet address is 172.16.23.3/29
  MTU 1500 bytes, BW 128 Kbit/sec, DLY 20000 usec,
    reliability 255/255, txload 1/255, rxload 1/255
  Encapsulation HDLC, crc 16, loopback not set
  Keepalive set (10 sec)
  Restart-Delay is 0 secs
  Last input 00:00:01, output 00:00:01, output hang never
  Last clearing of "show interface" counters never
  Input queue: 0/75/0/0 (size/max/drops/flushes); Total output drops: 0
  Queueing strategy: weighted fair
  Output queue: 0/1000/64/0 (size/max total/threshold/drops)
    Conversations 0/1/256 (active/max active/max total)
    Reserved Conversations 0/0 (allocated/max allocated)
    Available Bandwidth 96 kilobits/sec
  5 minute input rate 0 bits/sec, 0 packets/sec
  5 minute output rate 0 bits/sec, 0 packets/sec
    618 packets input, 43469 bytes, 0 no buffer
    Received 300 broadcasts (0 IP multicasts)
    0 runts, 0 giants, 0 throttles
    0 input errors, 0 CRC, 0 frame, 0 overrun, 0 ignored, 0 abort
    593 packets output, 41538 bytes, 0 underruns
    0 output errors, 0 collisions, 1 interface resets
    0 unknown protocol drops
    0 output buffer failures, 0 output buffers swapped out
    2 carrier transitions      DCD=up  DSR=up  DTR=up  RTS=up  CTS=up

R3#

```

e. Confirm that R3 has a valid route to reach R1 from its serial3/0 interface using the show ip eigrp topology 192.168.1.0 command.

```

R3#show ip eigrp topology 192.168.1.0
EIGRP-IPv4 Topology Entry for AS(1)/ID(192.168.3.1) for 192.168.1.0/24
  State is Passive, Query origin flag is 1, 1 Successor(s), FD is 21152000
  Descriptor Blocks:
  172.16.23.2 (Serial3/1), from 172.16.23.2, Send flag is 0x0
    Composite metric is (21152000/20640000), route is Internal
    Vector metric:
      Minimum bandwidth is 128 Kbit
      Total delay is 45000 microseconds
      Reliability is 255/255
      Load is 1/255
      Minimum MTU is 1500
      Hop count is 2
      Originating router is 192.168.1.1
  172.16.13.1 (Serial3/0), from 172.16.13.1, Send flag is 0x0
    Composite metric is (40640000/128256), route is Internal
    Vector metric:
      Minimum bandwidth is 64 Kbit
      Total delay is 25000 microseconds
      Reliability is 255/255
      Load is 1/255
      Minimum MTU is 1500
      Hop count is 1
      Originating router is 192.168.1.1

R3#

```

As indicated, R4 has two routes to reach 192.168.1.0. However, the metric for the route to R1 (172.16.13.1) is much higher (40640000) than the metric of the route to R2 (21152000), making the route through R2 the successor route.

### Step 6: Configure PBR to provide path control.

- On router R3, create a standard access list called PBR-ACL to identify the R4 LAN B network.

```
R3#conf t
Enter configuration commands, one per line. End with CNTL/Z.
R3(config)#ip access-list standard PBR-ACL
R3(config-std-nacl)#remark ACL matches R4 LAN B traffic
R3(config-std-nacl)#permit 192.168.4.128 0.0.0.127
R3(config-std-nacl)#exit
R3(config)#[
```

- Create a route map called R3-to-R1 that matches PBR-ACL and sets the next-hop interface to the R1 serial3/1 interface.

```
R3(config)#route-map R3-to-R1 permit
R3(config-route-map)#match ip address PER-ACL
R3(config-route-map)#set ip next-hop 172.16.13.1
R3(config-route-map)#exit
```

- Apply the R3-to-R1 route map to the serial interface on R3 that receives the traffic from R4. Use the ip policy route-map command on interface S3/2.

```
R3(config)#interface S3/2
R3(config-if)#ip policy route-map R3-to-R1
R3(config-if)#end
R3#[
```

- On R3, display the policy and matches using the **show route-map** command.

```
R3#show route-map
-
route-map R3-to-R1, permit, sequence 10
Match clauses:
  ip address (access-lists): PER-ACL
Set clauses:
  ip next-hop 172.16.13.1
Policy routing matches: 0 packets, 0 bytes
R3#[
```

### Step 7: Test the policy.

- On R3, create a standard ACL which identifies all of the R4 LANs.

```
R3#conf t
Enter configuration commands, one per line. End with CNTL/Z.
R3(config)#access-list 1 permit 192.168.4.0 0.0.0.255
R3(config)#exit
R3#[
```

- Enable PBR debugging only for traffic that matches the R4 LANs.

```
R3#debug ip policy ?
<1-199>  Access list
dynamic   dynamic PBR
early     Early PBR
<cr>
```

```
R3#debug ip policy 1
Policy routing debugging is on for access list 1
R3#
```

- c. Test the policy from R4 with the traceroute command, using R4 LAN A as the source network.

```
R4#traceroute 192.168.1.1 source 192.168.4.1
Type escape sequence to abort.
Tracing the route to 192.168.1.1
VRF info: (vrf in name/id, vrf out name/id)
 1 172.16.34.3 196 msec 52 msec 52 msec
 2 172.16.13.1 224 msec 248 msec 192 msec
R4#
```

Notice the path taken for the packet sourced from R4 LAN A is still going through R3 --> R2 -> R1.

As the traceroute was being executed, router R3 should be generating the following debug output.

```
R3#
*May 22 20:27:55.831: IP: s=192.168.4.1 (Serial3/2), d=192.168.1.1, len 28,
policy match
*May 22 20:27:55.835: IP: route map R3-to-R1, item 10, permit
*May 22 20:27:55.835: IP: s=192.168.4.1 (Serial3/2), d=192.168.1.1 (Serial3/
0), len 28, policy routed
*May 22 20:27:55.839: IP: Serial3/2 to Serial3/0 172.16.13.1
*May 22 20:27:55.927: IP: s=192.168.4.1 (Serial3/2), d=192.168.1.1, len 28,
policy match
*May 22 20:27:55.931: IP: route map R3-to-R1, item 10, permit
*May 22 20:27:55.931: IP: s=192.168.4.1 (Serial3/2), d=192.168.1.1 (Serial3/
0), len 28, policy routed
*May 22 20:27:55.931: IP: Serial3/2 to Serial3/0 172.16.13.1
*May 22 20:27:55.983: IP: s=192.168.4.1 (Serial3/2), d=192.168.1.1, len 28,
policy match
*May 22 20:27:55.983: IP: route map R3-to-R1, item 10, permit
*May 22 20:27:55.983: IP: s=192.168.4.1 (Serial3/2), d=192.168.1.1 (Serial3/
0), len 28, policy routed
*May 22 20:27:55.983: IP: Serial3/2 to Serial3/0 172.16.13.1
*May 22 20:27:56.051: IP: s=192.168.4.1 (Se
R3#rial3/2), d=192.168.1.1, len 28, FIB policy match
*May 22 20:27:56.051: IP: s=192.168.4.1 (Serial3/2), d=192.168.1.1, len 28,
PBR Counted
*May 22 20:27:56.051: IP: s=192.168.4.1 (Serial3/2), d=192.168.1.1, g=172.16
.13.1, len 28, FIB policy routed
*May 22 20:27:56.279: IP: s=192.168.4.1 (Serial3/2), d=192.168.1.1, len 28,
FIB policy match
*May 22 20:27:56.279: IP: s=192.168.4.1 (Serial3/2), d=192.168.1.1, len 28,
PBR Counted
*May 22 20:27:56.283: IP: s=192.168.4.1 (Serial3/2), d=192.168.1.1, g=172.16
.13.1, len 28, FIB policy routed
*May 22 20:27:56.567: IP: s=192.168.4.1 (Serial3/2), d=192.168.1.1, len 28,
FIB policy match
*May 22 20:27:56.567: IP: s=192.168.4.1 (Serial3/2), d=192.168.1.1, len 28,
PBR Counted
*May 22 20:27:56.567: IP: s=192.168.4.1 (Serial3/2), d=192.168.1.1, g=172.16
.13.1, len 28, FIB policy routed
R3#
```

- d. Test the policy from R4 with the traceroute command, using R4 LAN B as the source network.

```

R4#traceroute 192.168.1.1 source 192.168.4.129
Type escape sequence to abort.
Tracing the route to 192.168.1.1
VRF info: (vrf in name/id, vrf out name/id)
  1 172.16.34.3 48 msec 128 msec 116 msec
  2 172.16.13.1 184 msec 200 msec 204 msec
R4#  

R3#
*May 22 20:32:54.339: IP: s=192.168.4.129 (Serial3/2), d=192.168.1.1, len 28, policy match
*May 22 20:32:54.343: IP: route map R3-to-R1, item 10, permit
*May 22 20:32:54.343: IP: s=192.168.4.129 (Serial3/2), d=192.168.1.1 (Serial3/0), len 28, policy routed
*May 22 20:32:54.347: IP: Serial3/2 to Serial3/0 172.16.13.1
*May 22 20:32:54.419: IP: s=192.168.4.129 (Serial3/2), d=192.168.1.1, len 28, policy match
*May 22 20:32:54.419: IP: route map R3-to-R1, item 10, permit
*May 22 20:32:54.423: IP: s=192.168.4.129 (Serial3/2), d=192.168.1.1 (Serial3/0), len 28, policy routed
*May 22 20:32:54.427: IP: Serial3/2 to Serial3/0 172.16.13.1
*May 22 20:32:54.579: IP: s=192.168.4.129 (Serial3/2), d=192.168.1.1, len 28, policy match
*May 22 20:32:54.579: IP: route map R3-to-R1, item 10, permit
*May 22 20:32:54.579: IP: s=192.168.4.129 (Serial3/2), d=192.168.1.1 (Serial3/0), len 28, policy routed
*May 22 20:32:54.579: IP: Serial3/2 to Serial3/0 172.16.13.1
*May 22 20:32:54.663: IP: s=192
R3#.168.4.129 (Serial3/2), d=192.168.1.1, len 28, FIB policy match
*May 22 20:32:54.663: IP: s=192.168.4.129 (Serial3/2), d=192.168.1.1, len 28, PBR Counted
*May 22 20:32:54.663: IP: s=192.168.4.129 (Serial3/2), d=192.168.1.1, g=172.16.13.1, len 28, FIB policy routed
*May 22 20:32:54.863: IP: s=192.168.4.129 (Serial3/2), d=192.168.1.1, len 28, FIB policy match
*May 22 20:32:54.863: IP: s=192.168.4.129 (Serial3/2), d=192.168.1.1, len 28, PBR Counted
*May 22 20:32:54.863: IP: s=192.168.4.129 (Serial3/2), d=192.168.1.1, g=172.16.13.1, len 28, FIB policy routed
*May 22 20:32:55.115: IP: s=192.168.4.129 (Serial3/2), d=192.168.1.1, len 28, FIB policy match
*May 22 20:32:55.115: IP: s=192.168.4.129 (Serial3/2), d=192.168.1.1, len 28, PBR Counted
*May 22 20:32:55.115: IP: s=192.168.4.129 (Serial3/2), d=192.168.1.1, g=172.16.13.1, len 28, FIB policy routed
R3#
```

e.On R3, display the policy and matches using the show route-map command.

```

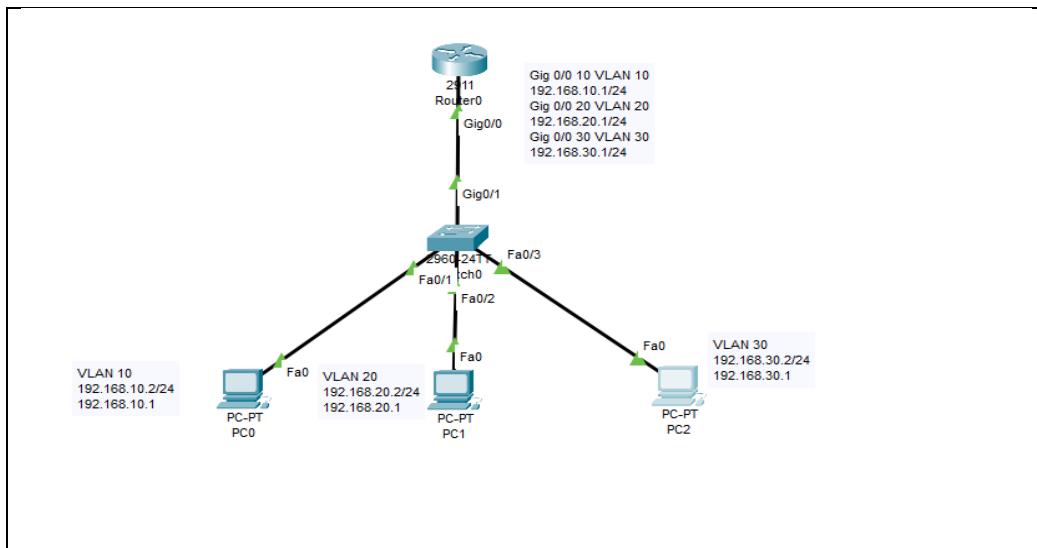
R3#show route-map
route-map R3-to-R1, permit, sequence 10
  Match clauses:
    ip address (access-lists): PER-ACL
  Set clauses:
    ip next-hop 172.16.13.1
Nexthop tracking current: 0.0.0.0
172.16.13.1, fib_nh:0,oce:0,status:0

  Policy routing matches: 12 packets, 384 bytes
R3#
```

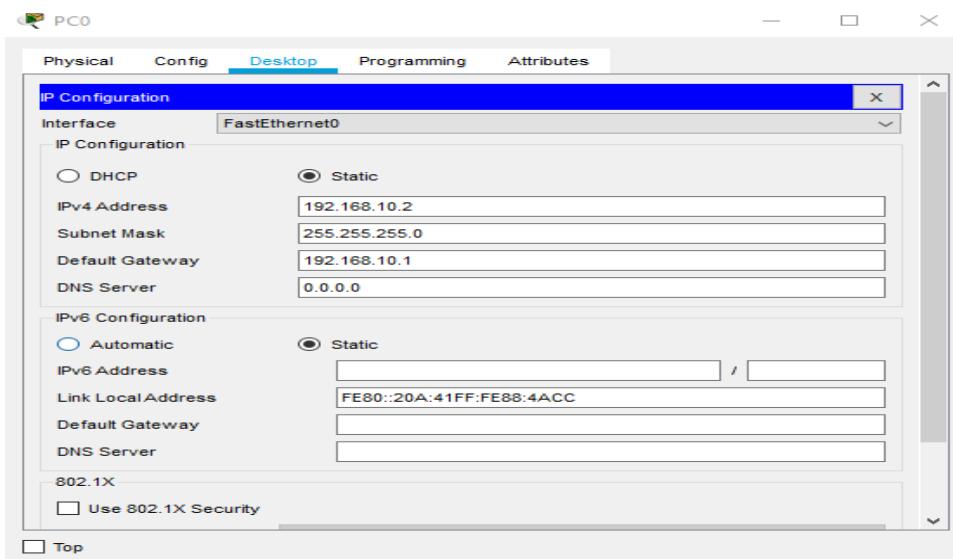
## Practical no 6

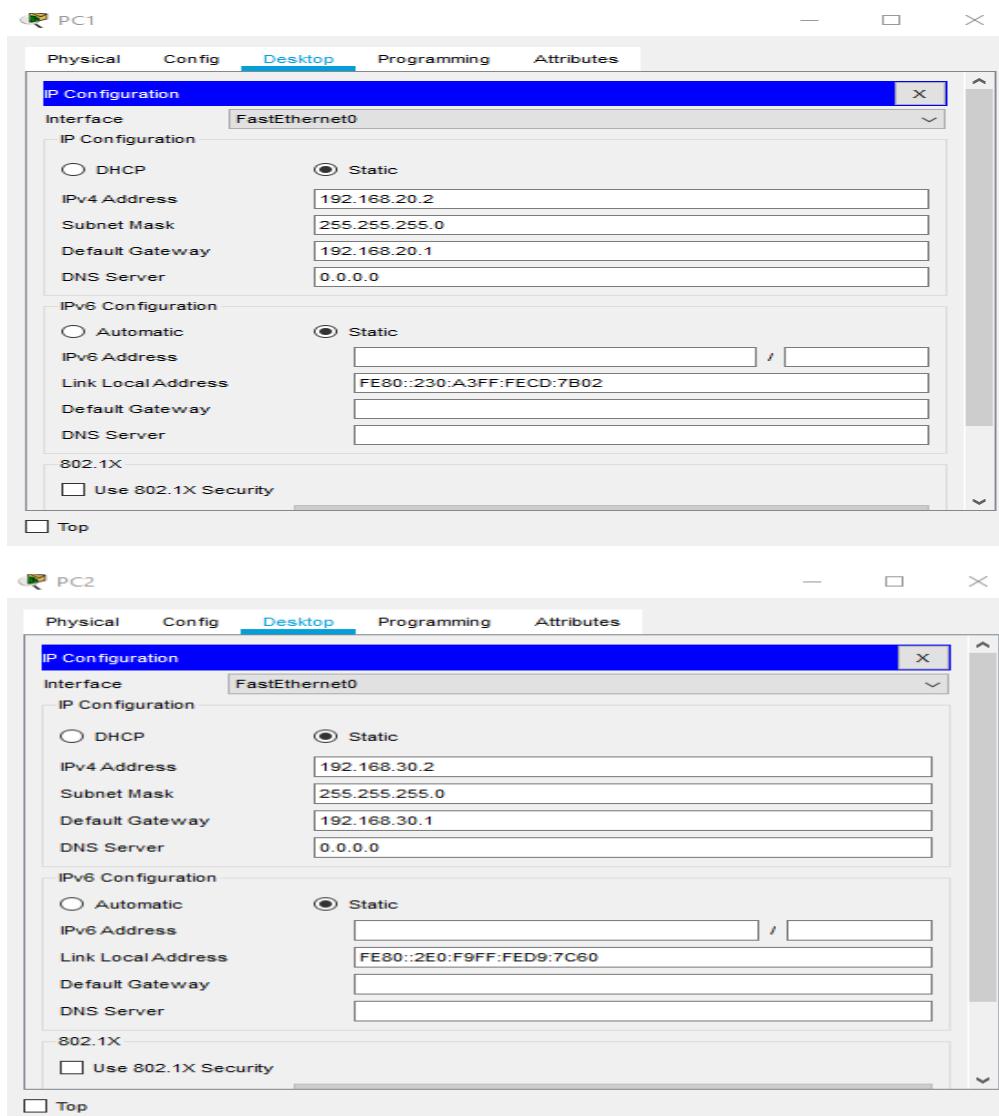
Aim:- Inter VLAN Routing.

### Topology



### Step 1: Assign the IP Addresses to PCs.





## Step 2: Configure within the switch.

- a. VLAN
- b. Assign the specified interfaces to the specific VLANs.
- c. Define access Ports and Trunk Ports.

```

Switch>
Switch>enable
Switch#configure terminal
Enter configuration commands, one per line. End with CNTL/Z.
Switch(config)#VLAN 10
Switch(config-vlan)#exit
Switch(config)#VLAN 20
Switch(config-vlan)#exit
Switch(config)#
Switch(config)#exit
Switch#
%SYS-5-CONFIG_I: Configured from console by console

```

```

Switch#show vlan brief

VLAN Name          Status    Ports
---- -----
1    default        active    Fa0/1, Fa0/2, Fa0/3, Fa0/4
                           Fa0/5, Fa0/6, Fa0/7, Fa0/8
                           Fa0/9, Fa0/10, Fa0/11, Fa0/12
                           Fa0/13, Fa0/14, Fa0/15, Fa0/16
                           Fa0/17, Fa0/18, Fa0/19, Fa0/20
                           Fa0/21, Fa0/22, Fa0/23, Fa0/24
                           Gig0/1, Gig0/2
10   VLAN0010       active
20   VLAN0020       active
1002 fddi-default  active
1003 token-ring-default  active
1004 fddinet-default  active
1005 trnet-default   active
Switch#

```

```

Switch#configure terminal
Enter configuration commands, one per line. End with CNTL/Z.
Switch(config)#interface fastethernet 0/1
Switch(config-if)#switchport access vlan 10
Switch(config-if)#switchport mode access
Switch(config-if)#exit
Switch(config)#

Switch(config)#interface fastethernet 0/2
Switch(config-if)#switchport access vlan 20
Switch(config-if)#switchport mode access
Switch(config-if)#exit
Switch(config)#

Switch(config)#interface fastethernet 0/3
Switch(config-if)#switchport access vlan 10
Switch(config-if)#switchport mode access
Switch(config-if)#end
Switch#
*SYS-5-CONFIG_I: Configured from console by console

```

```

Switch#show vlan brief

VLAN Name          Status    Ports
---- -----
1    default        active    Fa0/4, Fa0/5, Fa0/6, Fa0/7
                           Fa0/8, Fa0/9, Fa0/10, Fa0/11
                           Fa0/12, Fa0/13, Fa0/14, Fa0/15
                           Fa0/16, Fa0/17, Fa0/18, Fa0/19
                           Fa0/20, Fa0/21, Fa0/22, Fa0/23
                           Fa0/24, Gig0/1, Gig0/2
10   VLAN0010       active    Fa0/1, Fa0/3
20   VLAN0020       active    Fa0/2
1002 fddi-default  active
1003 token-ring-default  active
1004 fddinet-default  active
1005 trnet-default   active
Switch#

```

```
Switch#configure terminal
Enter configuration commands, one per line. End with CNTL/Z.
Switch(config)#interface gigabitethernet 0/1
Switch(config-if)#no shutdown
Switch(config-if)#
Switch(config-if)#switchport mode trunk
Switch(config-if)#exit
Switch(config)#exit
Switch#
%SYS-5-CONFIG_I: Configured from console by console

Switch#configure terminal
Enter configuration commands, one per line. End with CNTL/Z.
Switch(config)#interface gigabitethernet 0/1
Switch(config-if)#no shutdown
Switch(config-if)#
Switch(config-if)#switchport mode trunk
Switch(config-if)#exit
Switch(config)#exit
Switch#
%SYS-5-CONFIG_I: Configured from console by console

!
interface FastEthernet0/2
  switchport access vlan 20
  switchport mode access
!
interface FastEthernet0/3
  switchport access vlan 10
  switchport mode access
!
interface FastEthernet0/4
!
interface FastEthernet0/5
!
interface FastEthernet0/6
!
interface FastEthernet0/7
!
interface FastEthernet0/8
!
interface FastEthernet0/9
!
interface FastEthernet0/10
!
interface FastEthernet0/11
--More--
```

```

!
interface FastEthernet0/2
  switchport access vlan 20
  switchport mode access
!
interface FastEthernet0/3
  switchport access vlan 10
  switchport mode access
!
interface FastEthernet0/4
!
interface FastEthernet0/5
!
interface FastEthernet0/6
!
interface FastEthernet0/7
!
interface FastEthernet0/8
!
interface FastEthernet0/9
!
interface FastEthernet0/10
!
interface FastEthernet0/11
--More--

!
interface FastEthernet0/2
  switchport access vlan 20
  switchport mode access
!
interface FastEthernet0/3
  switchport access vlan 10
  switchport mode access
!
interface FastEthernet0/4
!
interface FastEthernet0/5
!
interface FastEthernet0/6
!
interface FastEthernet0/7
!
interface FastEthernet0/8
!
interface FastEthernet0/9
!
interface FastEthernet0/10
!
interface FastEthernet0/11
--More--

Switch#copy running-config startup-config
Destination filename [startup-config]?
Building configuration...
[OK]
Switch#

```

**Step 3: Configure router with dot1Q encapsulation by making each of sub-interfaces.**

Would you like to enter the initial configuration dialog? [yes/no]: no

Press RETURN to get started!

```
Router>enable
Router#configure terminal
Enter configuration commands, one per line. End with CNTL/Z.
Router(config)#interface gigabitethernet 0/0
Router(config-if)#no shutdown

Router(config-if)#
%LINK-5-CHANGED: Interface GigabitEthernet0/0, changed state to up

%LINEPROTO-5-UPDOWN: Line protocol on Interface GigabitEthernet0/0, changed state to up

Router(config-if)#exit
Router(config)#

Router(config)#interface gigabitethernet 0/0.10
Router(config-subif)#
%LINK-5-CHANGED: Interface GigabitEthernet0/0.10, changed state to up

%LINEPROTO-5-UPDOWN: Line protocol on Interface GigabitEthernet0/0.10, changed state to up

Router(config-subif)#encapsulation dot1Q 10
Router(config-subif)#ip address 192.168.10.1 255.255.255.0
Router(config-subif)#exit
Router(config)#

Router(config-subif)#encapsulation dot1Q 10
Router(config-subif)#ip address 192.168.10.1 255.255.255.0
Router(config-subif)#exit
Router(config)#

```

Fire	Last Status	Source	Destination	Type	Color	Time(sec)	Periodic	Num
	Successful	PC0	PC1	ICMP	Green	0.000	N	0
	Failed	PC1	PC2	ICMP	Brown	0.000	N	1

```
Router(config)#interface gigabitethernet 0/0.30
Router(config-subif)#
%LINK-5-CHANGED: Interface GigabitEthernet0/0.30, changed state to up

%LINEPROTO-5-UPDOWN: Line protocol on Interface GigabitEthernet0/0.30, changed state to up

Router(config-subif)#encapsulation dot1Q ?
<1-4094> IEEE 802.1Q VLAN ID
Router(config-subif)#encapsulation dot1Q 10

%Configuration of multiple subinterfaces of the same main
interface with the same VID (10) is not permitted.
This VID is already configured on GigabitEthernet0/0.10.

Router(config-subif)#shutdown

Router(config-subif)#
%LINK-5-CHANGED: Interface GigabitEthernet0/0.30, changed state to administratively down

%LINEPROTO-5-UPDOWN: Line protocol on Interface GigabitEthernet0/0.30, changed state to down

Router(config-subif)#exit
```

```
Router(config)#no interface gigabitethernet 0/0.30
Router(config)#end
Router#
%SYS-5-CONFIG_I: Configured from console by console

Router#show run
Building configuration...

Current configuration : 875 bytes
!
version 15.1
no service timestamps log datetime msec
no service timestamps debug datetime msec
no service password-encryption
!
hostname Router
!
!
!
!
!
!
!
ip cef
no ipv6 cef
!
!
!
!
license udi pid CISCO2911/K9 sn FTX1524B7CB-
!
!
!
!
!
!
!
!
!
!
!
!
!
!
!
!
spanning-tree mode pvst
```

```
!
!
!
!
!
interface GigabitEthernet0/0
no ip address
duplex auto
speed auto
!
interface GigabitEthernet0/0.10
encapsulation dot1Q 10
ip address 192.168.10.1 255.255.255.0
!
interface GigabitEthernet0/0.20
encapsulation dot1Q 20
ip address 192.168.20.1 255.255.255.0
!
interface GigabitEthernet0/1
no ip address
duplex auto
speed auto
shutdown
!
interface GigabitEthernet0/2
no ip address
duplex auto
speed auto
shutdown
!
interface Vlan1
no ip address
shutdown
!
ip classless
!
ip flow-export version 9
!
!
!
!
!
!
line con 0
!
line aux 0
!
line vty 0 4
login
!
!
!
end

Router#
```

```

Router#show ip route
Codes: L - local, C - connected, S - static, R - RIP, M - mobile, B - BGP
      D - EIGRP, EX - EIGRP external, O - OSPF, IA - OSPF inter area
      N1 - OSPF NSSA external type 1, N2 - OSPF NSSA external type 2
      E1 - OSPF external type 1, E2 - OSPF external type 2, E - EGP
      i - IS-IS, L1 - IS-IS level-1, L2 - IS-IS level-2, ia - IS-IS inter area
      * - candidate default, U - per-user static route, o - ODR
      P - periodic downloaded static route

Gateway of last resort is not set

      192.168.10.0/24 is variably subnetted, 2 subnets, 2 masks
C        192.168.10.0/24 is directly connected, GigabitEthernet0/0.10
L        192.168.10.1/32 is directly connected, GigabitEthernet0/0.10
      192.168.20.0/24 is variably subnetted, 2 subnets, 2 masks
C        192.168.20.0/24 is directly connected, GigabitEthernet0/0.20
L        192.168.20.1/32 is directly connected, GigabitEthernet0/0.20

Router#

Switch>enable
Switch#
Switch#show vlan brief

VLAN Name          Status     Ports
---- 
1    default        active     Fa0/4, Fa0/5, Fa0/6, Fa0/7
                           Fa0/8, Fa0/9, Fa0/10, Fa0/11
                           Fa0/12, Fa0/13, Fa0/14, Fa0/15
                           Fa0/16, Fa0/17, Fa0/18, Fa0/19
                           Fa0/20, Fa0/21, Fa0/22, Fa0/23
                           Fa0/24, Gig0/2
10   VLAN0010       active     Fa0/1, Fa0/3
20   VLAN0020       active     Fa0/2
1002 fddi-default   active
1003 token-ring-default   active
1004 fddinet-default   active
1005 trnet-default   active
Switch#

Switch#configure terminal
Enter configuration commands, one per line. End with CNTL/Z.
Switch(config)#vlan 30
Switch(config-vlan)#exit
Switch(config)#
Switch(config)#interface fastethernet 0/3
Switch(config-if)#switchport access vlan 30
Switch(config-if)#switchport mode access
Switch(config-if)#exit
Switch(config)#exit
Switch#
*SYS-5-CONFIG_I: Configured from console by console

```

```

Switch#show vlan brief

VLAN Name                               Status    Ports
---- -----
1   default                                active    Fa0/4, Fa0/5, Fa0/6, Fa0/7
                                                Fa0/8, Fa0/9, Fa0/10, Fa0/11
                                                Fa0/12, Fa0/13, Fa0/14, Fa0/15
                                                Fa0/16, Fa0/17, Fa0/18, Fa0/19
                                                Fa0/20, Fa0/21, Fa0/22, Fa0/23
                                                Fa0/24, Gig0/2
10  VLAN0010                             active    Fa0/1
20  VLAN0020                             active    Fa0/2
30  VLAN0030                             active    Fa0/3
1002 fddi-default                         active
1003 token-ring-default                  active
1004 fddinet-default                     active
1005 trnet-default                      active
Switch#

Router#configure terminal
Enter configuration commands, one per line. End with CNTL/Z.
Router(config)#interface gigabitethernet 0/0.30
Router(config-subif)#
%LINK-5-CHANGED: Interface GigabitEthernet0/0.30, changed state to up

%LINEPROTO-5-UPDOWN: Line protocol on Interface GigabitEthernet0/0.30, changed state to up

Router(config-subif)#encapsulation dot1Q 30
Router(config-subif)#ip address 192.168.30.1 255.255.255.0
Router(config-subif)#exit
Router(config)#exit
Router#
%SYS-5-CONFIG_I: Configured from console by console

Router#show ip route
Codes: L - local, C - connected, S - static, R - RIP, M - mobile, B - BGP
       D - EIGRP, EX - EIGRP external, O - OSPF, IA - OSPF inter area
       N1 - OSPF NSSA external type 1, N2 - OSPF NSSA external type 2
       E1 - OSPF external type 1, E2 - OSPF external type 2, E - EGP
       i - IS-IS, L1 - IS-IS level-1, L2 - IS-IS level-2, ia - IS-IS inter area
       * - candidate default, U - per-user static route, o - ODR
       P - periodic downloaded static route

Gateway of last resort is not set

  192.168.10.0/24 is variably subnetted, 2 subnets, 2 masks
C        192.168.10.0/24 is directly connected, GigabitEthernet0/0.10
L        192.168.10.1/32 is directly connected, GigabitEthernet0/0.10
  192.168.20.0/24 is variably subnetted, 2 subnets, 2 masks
C        192.168.20.0/24 is directly connected, GigabitEthernet0/0.20
L        192.168.20.1/32 is directly connected, GigabitEthernet0/0.20
  192.168.30.0/24 is variably subnetted, 2 subnets, 2 masks
C        192.168.30.0/24 is directly connected, GigabitEthernet0/0.30
L        192.168.30.1/32 is directly connected, GigabitEthernet0/0.30

```

Fire	Last Status	Source	Destination	Type	Color	Time(sec)	Periodic	Num
●	Successful	PC0	PC1	ICMP	■	0.000	N	0
●	Successful	PC1	PC2	ICMP	■	0.000	N	1
●	Successful	PC0	PC2	ICMP	■	0.000	N	2

```
Router#copy running-config startup-config
Destination filename [startup-config]?
Building configuration...
[OK]
Router#
```

```
Switch#copy running-config startup-config
Destination filename [startup-config]?
Building configuration...
[OK]
Switch#
```

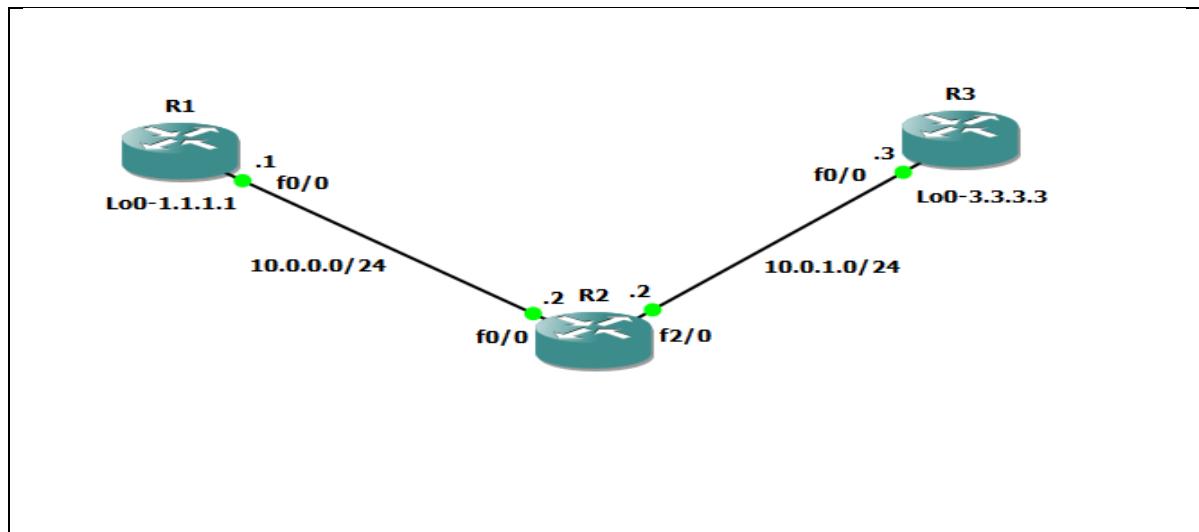
Fire	Last Status	Source	Destination	Type	Color	Time(sec)	Periodic	Num
●	Successful	PC0	PC1	ICMP	█	0.000	N	0
●	Successful	PC1	PC2	ICMP	█	0.000	N	1

Fire	Last Status	Source	Destination	Type	Color	Time(sec)	Periodic	Num
●	Successful	PC2	PC1	ICMP	█	0.000	N	0
●	Successful	PC2	PC0	ICMP	█	0.000	N	1

## Practical no 7

Aim:- Cisco MPLS Configuration.

### Topology



### Step 1 – IP addressing of MPLS Core and OSPF

First bring 3 routers into your topology R1, R2, R3 position them as below. We are going to address the routers and configure ospf to ensure loopback to loopback connectivity between R1 and R3.

R1

```
R1#conf t
Enter configuration commands, one per line. End with CNTL/Z.
R1(config)#hostname R1
R1(config)#int lo0
R1(config-if)#ip add 1.1.1.1 255.255.255.255
R1(config-if)#ip ospf 1 area 0
R1(config-if)#exit
```

```
R1(config)#int f0/0
R1(config-if)#ip add 10.0.0.1 255.255.255.0
R1(config-if)#no shut
R1(config-if)#ip ospf 1 area 0
R1(config-if)#
```

R2

```
R2#conf t
Enter configuration commands, one per line. End with CNTL/Z.
R2(config)#hostname R2
R2(config)#int lo0
R2(config-if)#ip add 2.2.2.2 255.255.255.255
R2(config-if)#ip ospf 1 area 0
R2(config-if)#exit
```

```
R2(config)#int f0/0
R2(config-if)#ip add 10.0.0.2 255.255.255.0
R2(config-if)#no shut
R2(config-if)#ip ospf 1 area 0
R2(config-if)#exit
```

```
R2(config)#int f2/0
R2(config-if)#ip add 10.0.1.2 255.255.255.0
R2(config-if)#no shut
R2(config-if)#ip ospf 1 area 0
R2(config-if)#exit
R2(config)#[
```

### R3

```
R3#conf t
Enter configuration commands, one per line. End with CNTL/Z.
R3(config)#hostname R3
R3(config)#int lo0
R3(config-if)#ip add 3.3.3.3 255.255.255.255
R3(config-if)#ip ospf 1 area 0
R3(config-if)#exit
```

```
R3(config)#int f0/0
R3(config-if)#ip add 10.0.1.3 255.255.255.0
R3(config-if)#no shut
R3(config-if)#ip ospf 1 area 0
R3(config-if)#exit
```

You should now have full ip connectivity between R1, R2, R3 to verify this we need to see if we can ping between the loopbacks of R1 and R3.

```
R1#ping 3.3.3.3 source lo0
Type escape sequence to abort.
Sending 5, 100-byte ICMP Echos to 3.3.3.3, timeout is 2 seconds:
Packet sent with a source address of 1.1.1.1
!!!!!
Success rate is 100 percent (5/5), round-trip min/avg/max = 96/216/300 ms
R1#[
```

You could show the routing table here, but the fact that you can ping between the loopbacks is verification enough and it is safe to move on.

### Step 2 – Configure LDP on all the interfaces in the MPLS Core

In order to run MPLS you need to enable it, there are two ways to do this.

- At each interface enter the **mpls ip** command
- Under the ospf process use the **mpls ldp autoconfig** command.

## R1

```
R1#conf t
Enter configuration commands, one per line. End with CNTL/Z.
R1(config)#router ospf 1
R1(config-router)#mpls ldp autoconfig
R1(config-router)#
*May 23 17:18:43.131: %PARSE_RC-3-PRC_INVALID_BLOCK_PTR:
R1(config-router)#
*May 23 17:19:26.811: %LDP-5-NBRCHG: LDP Neighbor 2.2.2.2:0 (1) is UP
R1(config-router)#[
```

## R2

```
R2(config)#router ospf 1
R2(config-router)#mpls ldp autoconfig
R2(config-router)#
*May 23 17:19:24.991: %PARSE_RC-3-PRC_INVALID_BLOCK_PTR:
*May 23 17:19:25.051: %PARSE_RC-3-PRC_INVALID_BLOCK_PTR:
R2(config-router)#[
```

## R3

```
R3(config)#router ospf 1
R3(config-router)#mpls ldp autoconfig
R3(config-router)#
*May 23 17:19:52.823: %PARSE_RC-3-PRC_INVALID_BLOCK_PTR:
R3(config-router)#
*May 23 17:19:53.935: %LDP-5-NBRCHG: LDP Neighbor 2.2.2.2:0 (1) is UP
R3(config-router)#[
```

You should see log messages coming up showing the LDP neighbors are up.

```
R2(config-router)#
*May 23 17:19:26.555: %LDP-5-NBRCHG: LDP Neighbor 1.1.1.1:0 (1) is UP
R2(config-router)#
*May 23 17:19:54.679: %LDP-5-NBRCHG: LDP Neighbor 3.3.3.3:0 (2) is UP
R2(config-router)#[
```

To verify the mpls interfaces the command is very simple – **sh mpls interface**  
This is done on R2 and you can see that both interfaces are running mpls and using LDP.

```
R2#sh mpls interface
Interface          IP           Tunnel   BGP  Static Operational
FastEthernet0/0    Yes (ldp)     No      No   No    Yes
FastEthernet2/0    Yes (ldp)     No      No   No    Yes
R2#[
```

You can also verify the LDP neighbors with the **sh mpls ldp neighbors** command.

```

R2#sh mpls ldp neigh
  Peer LDP Ident: 1.1.1.1:0; Local LDP Ident 2.2.2.2:0
    TCP connection: 1.1.1.1.646 - 2.2.2.2.39856
    State: Oper; Msgs sent/rcvd: 13/13; Downstream
    Up time: 00:05:05
    LDP discovery sources:
      FastEthernet0/0, Src IP addr: 10.0.0.1
      Addresses bound to peer LDP Ident:
        10.0.0.1          1.1.1.1
  Peer LDP Ident: 3.3.3.3:0; Local LDP Ident 2.2.2.2:0
    TCP connection: 3.3.3.3.17291 - 2.2.2.2.646
    State: Oper; Msgs sent/rcvd: 12/13; Downstream
    Up time: 00:04:37
    LDP discovery sources:
      FastEthernet2/0, Src IP addr: 10.0.1.3
      Addresses bound to peer LDP Ident:
        10.0.1.3          3.3.3.3
R2#

```

One more verification to confirm LDP is running ok is to do a trace between R1 and R3 and verify if you get MPLS Labels show up in the trace.

```

R1#trace 3.3.3.3
Type escape sequence to abort.
Tracing the route to 3.3.3.3
VRF info: (vrf in name/id, vrf out name/id)
  1 10.0.0.2 [MPLS: Label 17 Exp 0] 68 msec 52 msec 92 msec
  2 10.0.1.3 120 msec 172 msec 168 msec
R1#

```

As you can see the trace to R2 used an MPLS Label in the path, as this is a very small MPLS core only one label was used as R3 was the final hop.

So to review we have now configured IP addresses on the MPLS core, enabled OSPF and full IP connectivity between all routers and finally enabled mpls on all the interfaces in the core and have established ldp neighbors between all routers.

The next step is to configure MP-BGP between R1 and R3.

This is when you start to see the layer 3 vpn configuration come to life.

### **Step 3 – MPLS BGP Configuration between R1 and R3**

We need to establish a Multi Protocol BGP session between R1 and R3 this is done by configuring the vpng4 address family as below.

```

R1#conf t
Enter configuration commands, one per line.  End with CNTL/Z.
R1(config)#router bgp 1
R1(config-router)#neighbor 3.3.3.3 remote-as 1
R1(config-router)#neighbor 3.3.3.3 update-source Loopback0
R1(config-router)#no auto-summary
R1(config-router)#address-family vpng4
R1(config-router-af)#neighbor 3.3.3.3 activate
R1(config-router-af)#

```

```

R3#conf t
Enter configuration commands, one per line. End with CNTL/Z.
R3(config)#router bgp 1
R3(config-router)#neighbor 1.1.1.1 remote-as 1
R3(config-router)#neighbor 1.1.1.1 update-source Loopback0
R3(config-router)#no auto-summary
R3(config-router)#address-family vpnv4
R3(config-router-af)#neighbor 1.1.1.1 activate
R3(config-router-af)#
*May 23 17:30:40.751: %BGP-5-NBR_RESET: Neighbor 1.1.1.1 reset (Capability changed)
*May 23 17:30:40.767: %BGP-5-ADJCHANGE: neighbor 1.1.1.1 Down Capability changed
*May 23 17:30:40.771: %BGP_SESSION-5-ADJCHANGE: neighbor 1.1.1.1 IPv4 Unicast topology base removed from session Capability changed
R3(config-router-af)#
*May 23 17:30:45.291: %BGP-5-ADJCHANGE: neighbor 1.1.1.1 Up
R3(config-router-af)#

```

You should see log messages showing the BGP sessions coming up.

To verify the BGP session between R1 and R3 issue the command **sh bgp vpnv4 unicast all summary**.

```

R1#sh bgp vpnv4 unicast all summary
BGP router identifier 1.1.1.1, local AS number 1
BGP table version is 1, main routing table version 1

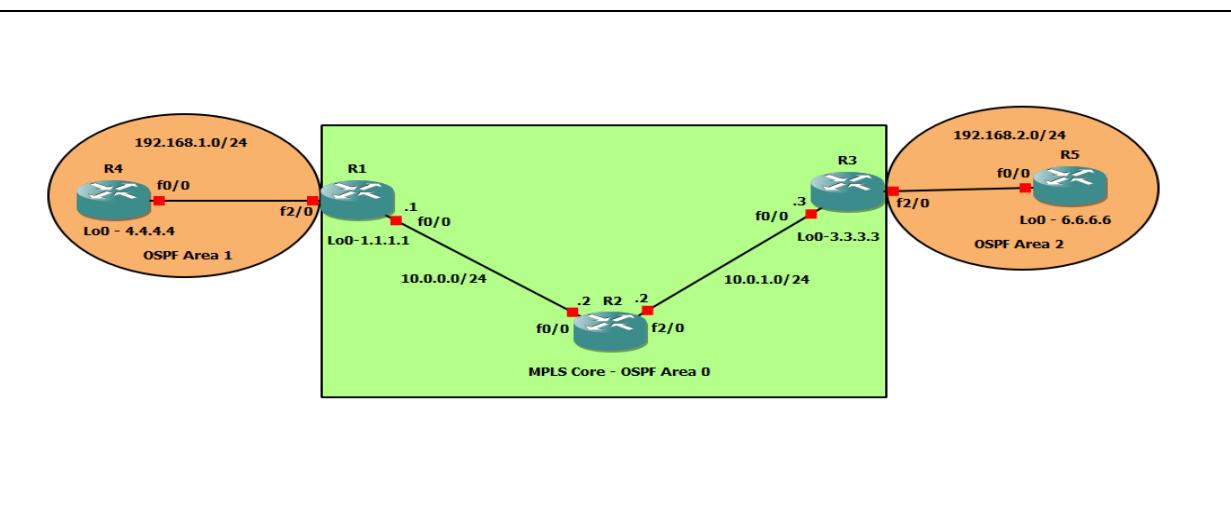
Neighbor          V      AS MsgRcvd MsgSent   TblVer  InQ OutQ Up/Down  State
/PfxRcd
3.3.3.3           4      1       6       6       1     0     0 00:01:02
  0
R1#

```

You can see here that we do have a bgp vpnv4 peering to R3 – looking at the PfxRcd you can see it says 0 this is because we have not got any routes in BGP. We are now going to add two more routers to the topology. These will be the customer sites connected to R1 and R3. We will then create a VRF on each router and put the interfaces connected to each site router into that VRF.

#### **Step 4 – Add two more routers, create VRFs.**

We will add two more routers into the topology so it now looks like the final topology



Router 4 will peer OSPF using process number 2 to a VRF configured on R1. It will use the local site addressing of 192.168.1.0/24.

## R4

```
R4#conf t
Enter configuration commands, one per line. End with CNTL/Z.
R4(config)#int lo0
R4(config-if)#ip add 4.4.4.4 255.255.255.255
R4(config-if)#ip ospf 2 area 2
R4(config-if)#int f0/0
R4(config-if)#ip add 192.168.1.4 255.255.255.0
R4(config-if)#ip ospf 2 area 2
R4(config-if)#no shut
R4(config-if)#
*May 23 17:34:00.139: %LINK-3-UPDOWN: Interface FastEthernet0/0, changed state to up
*May 23 17:34:01.139: %LINEPROTO-5-UPDOWN: Line protocol on Interface FastEthernet0/0, changed state to up
R4(config-if)#[
```

## R1

```
R1#conf t
Enter configuration commands, one per line. End with CNTL/Z.
R1(config)#int f2/0
R1(config-if)#no shut
R1(config-if)#ip add 192.168.1.1 255.255.255.0
R1(config-if)#[
```

Now at this point we have R4 peering to R1 but in the global routing table of R1 which is not what we want.

### We are now going to start using VRF's.

For this mpls tutorial I will be using VRF RED.

```
R1(config)#ip vrf RED
R1(config-vrf)#rd 4:4
R1(config-vrf)#route-target both 4:4
R1(config-vrf)#[
```

So now we have configured the VRF on R1 we need to move the interface F0/1 into that VRF.

```
R1(config)#int f2/0
R1(config-if)#ip vrf forwarding RED
% Interface FastEthernet2/0 IPv4 disabled and address(es) removed due to enabling VRF RED
R1(config-if)#[
```

You just need to re-apply it

```
R1(config)#int f2/0
R1(config-if)#ip address 192.168.1.1 255.255.255.0
R1(config-if)#end
[
```

Now if we view the config on R1 int f2/0 you can see the VRF configured.

```
R1#sh run int f2/0
Building configuration...

Current configuration : 107 bytes
!
interface FastEthernet2/0
  ip vrf forwarding RED
  ip address 192.168.1.1 255.255.255.0
  duplex full
end

R1#
```

Now we can start to look int VRF's and how they operate – you need to understand now that there are 2 routing tables within R1

- The Global Routing Table
- The Routing Table for VRF RED

If you issue the command **sh ip route** this shows the routes in the global table and you will notice that you do not see 192.168.1.0/24

```
R1#sh ip route
Codes: L - local, C - connected, S - static, R - RIP, M - mobile, B - BGP
      D - EIGRP, EX - EIGRP external, O - OSPF, IA - OSPF inter area
      N1 - OSPF NSSA external type 1, N2 - OSPF NSSA external type 2
      E1 - OSPF external type 1, E2 - OSPF external type 2
      i - IS-IS, su - IS-IS summary, L1 - IS-IS level-1, L2 - IS-IS level-2
      ia - IS-IS inter area, * - candidate default, U - per-user static route
      o - ODR, P - periodic downloaded static route, H - NHRP, l - LISP
      + - replicated route, % - next hop override

Gateway of last resort is not set

      1.0.0.0/32 is subnetted, 1 subnets
C        1.1.1.1 is directly connected, Loopback0
      2.0.0.0/32 is subnetted, 1 subnets
O        2.2.2.2 [110/2] via 10.0.0.2, 00:36:16, FastEthernet0/0
      3.0.0.0/32 is subnetted, 1 subnets
O        3.3.3.3 [110/3] via 10.0.0.2, 00:29:53, FastEthernet0/0
      10.0.0.0/8 is variably subnetted, 3 subnets, 2 masks
C        10.0.0.0/24 is directly connected, FastEthernet0/0
L        10.0.0.1/32 is directly connected, FastEthernet0/0
O        10.0.1.0/24 [110/2] via 10.0.0.2, 00:35:08, FastEthernet0/0
R1#
```

If you now issue the command **sh ip route vrf red** – this will show the routes in the routing table for VRF RED.

```
R1#sh ip route vrf red
% IP routing table vrf red does not exist
R1#
```

```
R1#sh ip route vrf RED

Routing Table: RED
Codes: L - local, C - connected, S - static, R - RIP, M - mobile, B - BGP
      D - EIGRP, EX - EIGRP external, O - OSPF, IA - OSPF inter area
      N1 - OSPF NSSA external type 1, N2 - OSPF NSSA external type 2
      E1 - OSPF external type 1, E2 - OSPF external type 2
      i - IS-IS, su - IS-IS summary, L1 - IS-IS level-1, L2 - IS-IS level-2
      ia - IS-IS inter area, * - candidate default, U - per-user static route
      o - ODR, P - periodic downloaded static route, H - NHRP, l - LISP
      + - replicated route, % - next hop override

Gateway of last resort is not set

      192.168.1.0/24 is variably subnetted, 2 subnets, 2 masks
C            192.168.1.0/24 is directly connected, FastEthernet2/0
L            192.168.1.1/32 is directly connected, FastEthernet2/0
R1#
```

We just need to enable OSPF on this interface and get the loopback address for R4 in the VRF RED routing table before proceeding.

```
R1#conf t
Enter configuration commands, one per line. End with CNTL/Z.
R1(config)#int f2/0
R1(config-if)#ip ospf 2 area 2
R1(config-if)#exit
```

You should see a log message showing the OSPF neighbor come up.

```
R1(config)#
*May 23 17:45:26.587: %OSPF-5-ADJCHG: Process 2, Nbr 4.4.4.4 on FastEthernet2/0 from LOADING to
FULL, Loading Done
R1(config)#+
```

If we now check the routes in the VRF RED routing table you should see 4.4.4.4 in there as well.

```
R1#sh ip route vrf RED

Routing Table: RED
Codes: L - local, C - connected, S - static, R - RIP, M - mobile, B - BGP
      D - EIGRP, EX - EIGRP external, O - OSPF, IA - OSPF inter area
      N1 - OSPF NSSA external type 1, N2 - OSPF NSSA external type 2
      E1 - OSPF external type 1, E2 - OSPF external type 2
      i - IS-IS, su - IS-IS summary, L1 - IS-IS level-1, L2 - IS-IS level-2
      ia - IS-IS inter area, * - candidate default, U - per-user static route
      o - ODR, P - periodic downloaded static route, H - NHRP, l - LISP
      + - replicated route, % - next hop override

Gateway of last resort is not set

      4.0.0.0/32 is subnetted, 1 subnets
O            4.4.4.4 [110/2] via 192.168.1.4, 00:01:57, FastEthernet2/0
      192.168.1.0/24 is variably subnetted, 2 subnets, 2 masks
C            192.168.1.0/24 is directly connected, FastEthernet2/0
L            192.168.1.1/32 is directly connected, FastEthernet2/0
R1#
```

We now need to repeat this process for R3 & R5

Router 6 will peer OSPF using process number 2 to a VRF configured on R3. It will use the local site addressing of 192.168.2.0/24.

## R5

```
R5#conf t
Enter configuration commands, one per line. End with CNTL/Z.
R5(config)#int lo0
R5(config-if)#ip add 6.6.6.6 255.255.255.255
R5(config-if)#ip ospf 2 area 2
R5(config-if)#int f0/0
R5(config-if)#ip add 192.168.2.6 255.255.255.0
R5(config-if)#ip ospf 2 area 2
R5(config-if)#no shut
R5(config-if)#
*May 23 17:53:52.083: %LINK-3-UPDOWN: Interface FastEthernet0/0, changed state to up
*May 23 17:53:53.083: %LINEPROTO-5-UPDOWN: Line protocol on Interface FastEthernet0/0, changed state to up
R5(config-if)#[REDACTED]
```

## R3

```
R3#conf t
Enter configuration commands, one per line. End with CNTL/Z.
R3(config)#int f2/0
R3(config-if)#no shut
R3(config-if)#ip add 192.168.2.3 255.255.255.0
R3(config-if)#exit
R3(config)#[REDACTED]
```

We also need to configure a VRF onto R3 as well.

```
R3(config)#ip vrf RED
R3(config-vrf)#rd 4:4
R3(config-vrf)#route-target both 4:4
R3(config-vrf)#exit
R3(config)#[REDACTED]
```

So now we have configured the VRF on R3 we need to move the interface F0/1 into that VRF.

Now notice what happens when you do that – the IP address is removed.

```
R3(config)#int f2/0
R3(config-if)#ip vrf forwarding RED
% Interface FastEthernet2/0 IPv4 disabled and address(es) removed due to enabling VRF RED
R3(config-if)#[REDACTED]
```

You just need to re-apply it.

```
R3(config)#int f2/0
R3(config-if)#ip address 192.168.2.1 255.255.255.0
R3(config-if)#end
R3#
```

Now if we view the config on R3 int f2/0 you can see the VRF configured.

```
R3#sh run int f2/0
Building configuration...

Current configuration : 107 bytes
!
interface FastEthernet2/0
  ip vrf forwarding RED
  ip address 192.168.2.1 255.255.255.0
  duplex full
end

R3#
```

Finally we just need to enable OSPF on that interface and verify the routes are in the RED routing table.

```
R3#conf t
Enter configuration commands, one per line. End with CNTL/Z.
R3(config)#int f2/0
R3(config-if)#ip ospf 2 area 2
R3(config-if)#exit
R3(config)#exit
R3#
*May 23 17:59:52.387: %SYS-5-CONFIG_I: Configured from console by console
R3#
*May 23 17:59:54.499: %OSPF-5-ADJCHG: Process 2, Nbr 6.6.6.6 on FastEthernet2/0 from LOADING to FULL, L
oading Done
R3#
```

Check the routes in vrf RED

```
R3#sh ip route vrf RED

Routing Table: RED
Codes: L - local, C - connected, S - static, R - RIP, M - mobile, B - BGP
      D - EIGRP, EX - EIGRP external, O - OSPF, IA - OSPF inter area
      N1 - OSPF NSSA external type 1, N2 - OSPF NSSA external type 2
      E1 - OSPF external type 1, E2 - OSPF external type 2
      i - IS-IS, su - IS-IS summary, L1 - IS-IS level-1, L2 - IS-IS level-2
      ia - IS-IS inter area, * - candidate default, U - per-user static route
      o - ODR, P - periodic downloaded static route, H - NHRP, l - LISP
      + - replicated route, % - next hop override

Gateway of last resort is not set

          6.0.0.0/32 is subnetted, 1 subnets
O        6.6.6.6 [110/2] via 192.168.2.6, 00:00:44, FastEthernet2/0
          192.168.2.0/24 is variably subnetted, 2 subnets, 2 masks
C        192.168.2.0/24 is directly connected, FastEthernet2/0
L        192.168.2.1/32 is directly connected, FastEthernet2/0
R3#
```

The final step to get full connectivity across the MPLS core is to redistribute the routes in OSPF on R1 and R3 into MP-BGP and MP-BGP into OSPF, this is what we are going to do now.

We need to redistribute the OSPF routes from R4 into BGP in the VRF on R1, the OSPF routes from R6 into MP-BGP in the VRF on R3 and then the routes in MP-BGP in R1 and R3 back out to OSPF.

### Before we start lets do some verifications

Check the routes on R4

```
R4#sh ip route
Codes: L - local, C - connected, S - static, R - RIP, M - mobile, B - BGP
      D - EIGRP, EX - EIGRP external, O - OSPF, IA - OSPF inter area
      N1 - OSPF NSSA external type 1, N2 - OSPF NSSA external type 2
      E1 - OSPF external type 1, E2 - OSPF external type 2
      i - IS-IS, su - IS-IS summary, L1 - IS-IS level-1, L2 - IS-IS level-2
      ia - IS-IS inter area, * - candidate default, U - per-user static route
      o - ODR, P - periodic downloaded static route, H - NHRP, l - LISP
      + - replicated route, % - next hop override

Gateway of last resort is not set

        4.0.0.0/32 is subnetted, 1 subnets
C            4.4.4.4 is directly connected, Loopback0
        192.168.1.0/24 is variably subnetted, 2 subnets, 2 masks
C            192.168.1.0/24 is directly connected, FastEthernet0/0
L            192.168.1.4/32 is directly connected, FastEthernet0/0
R4#
```

As expected we have the local interface and the loopback address.

When we are done we want to see 6.6.6.6 in there so we can ping across the MPLS

### Check the routes on R1

```
R1#sh ip route
Codes: L - local, C - connected, S - static, R - RIP, M - mobile, B - BGP
      D - EIGRP, EX - EIGRP external, O - OSPF, IA - OSPF inter area
      N1 - OSPF NSSA external type 1, N2 - OSPF NSSA external type 2
      E1 - OSPF external type 1, E2 - OSPF external type 2
      i - IS-IS, su - IS-IS summary, L1 - IS-IS level-1, L2 - IS-IS level-2
      ia - IS-IS inter area, * - candidate default, U - per-user static route
      o - ODR, P - periodic downloaded static route, H - NHRP, l - LISP
      + - replicated route, % - next hop override

Gateway of last resort is not set

        1.0.0.0/32 is subnetted, 1 subnets
C            1.1.1.1 is directly connected, Loopback0
        2.0.0.0/32 is subnetted, 1 subnets
O            2.2.2.2 [110/2] via 10.0.0.2, 00:55:27, FastEthernet0/0
        3.0.0.0/32 is subnetted, 1 subnets
O            3.3.3.3 [110/3] via 10.0.0.2, 00:49:04, FastEthernet0/0
        10.0.0.0/8 is variably subnetted, 3 subnets, 2 masks
C            10.0.0.0/24 is directly connected, FastEthernet0/0
L            10.0.0.1/32 is directly connected, FastEthernet0/0
O            10.0.1.0/24 [110/2] via 10.0.0.2, 00:54:19, FastEthernet0/0
R1#
```

Remember we have a VRF configured on this router so this command will show routes in the global routing table (the MPLS Core) and it will not show the 192.168.1.0/24 route as that is in VRF RED – to see that we run the following command.

```
R1#sh ip route vrf RED

Routing Table: RED
Codes: L - local, C - connected, S - static, R - RIP, M - mobile, B - BGP
      D - EIGRP, EX - EIGRP external, O - OSPF, IA - OSPF inter area
      N1 - OSPF NSSA external type 1, N2 - OSPF NSSA external type 2
      E1 - OSPF external type 1, E2 - OSPF external type 2
      i - IS-IS, su - IS-IS summary, L1 - IS-IS level-1, L2 - IS-IS level-2
      ia - IS-IS inter area, * - candidate default, U - per-user static route
      o - ODR, P - periodic downloaded static route, H - NHRP, l - LISP
      + - replicated route, % - next hop override

Gateway of last resort is not set

        4.0.0.0/32 is subnetted, 1 subnets
O          4.4.4.4 [110/2] via 192.168.1.4, 00:17:23, FastEthernet2/0
              192.168.1.0/24 is variably subnetted, 2 subnets, 2 masks
C            192.168.1.0/24 is directly connected, FastEthernet2/0
L            192.168.1.1/32 is directly connected, FastEthernet2/0
R1#
```

Here you can see Routing Table: RED is shown and the routes to R4 are now visible with 4.4.4.4 being in OSPF.

So we need to do the following;

- Redistribute OSPF into MP-BGP on R1
- Redistribute MP-BGP into OSPF on R1
- Redistribute OSPF into MP-BGP on R3
- Redistribute MP-BGP into OSPF on R3

### **Redistribute OSPF into MP-BGP on R1**

```
R1#conf t
Enter configuration commands, one per line.  End with CNTL/Z.
R1(config)#router bgp 1
R1(config-router)#address-family ipv4 vrf RED
R1(config-router-af)#redistribute ospf 2
R1(config-router-af)#exit
R1(config-router)#exit
R1(config)#+
```

### **Redistribute OSPF into MP-BGP on R3**

```
R3#conf t
Enter configuration commands, one per line.  End with CNTL/Z.
R3(config)#router bgp 1
R3(config-router)#address-family ipv4 vrf RED
R3(config-router-af)#redistribute ospf 2
R3(config-router-af)#exit
R3(config-router)#exit
R3(config)#+
```

This has enabled redistribution of the OSPF routes into BGP. We can check the routes from R4 and R6 are now showing in the BGP table for their VRF with this command

```
sh ip bgp vpng4 vrf RED
```

```
R1#sh ip bgp vpng4 vrf RED
BGP table version is 7, local router ID is 1.1.1.1
Status codes: s suppressed, d damped, h history, * valid, > best, i - internal,
               r RIB-failure, S Stale, m multipath, b backup-path, f RT-Filter,
               x best-external, a additional-path, c RIB-compressed,
Origin codes: i - IGP, e - EGP, ? - incomplete
RPKI validation codes: V valid, I invalid, N Not found

      Network          Next Hop            Metric LocPrf Weight Path
Route Distinguisher: 4:4 (default for vrf RED)
 *>  4.4.4.4/32      192.168.1.4          2        32768  ?
 *>i 6.6.6.6/32      3.3.3.3             2       100      0  ?
 *>  192.168.1.0     0.0.0.0             0        32768  ?
 *>i 192.168.2.0     3.3.3.3             0       100      0  ?
R1#
```

Here we can see that 4.4.4.4 is now in the BGP table in VRF RED on R1 with a next hop of 192.168.1.4 (R4) and also 6.6.6.6 is in there as well with a next hop of 3.3.3.3 (which is the loopback of R3 – showing that it is going over the MPLS and R1 is not in the picture)

The same should be true on R3.

```
R3#sh ip bgp vpng4 vrf RED
BGP table version is 7, local router ID is 3.3.3.3
Status codes: s suppressed, d damped, h history, * valid, > best, i - internal,
               r RIB-failure, S Stale, m multipath, b backup-path, f RT-Filter,
               x best-external, a additional-path, c RIB-compressed,
Origin codes: i - IGP, e - EGP, ? - incomplete
RPKI validation codes: V valid, I invalid, N Not found

      Network          Next Hop            Metric LocPrf Weight Path
Route Distinguisher: 4:4 (default for vrf RED)
 *>i 4.4.4.4/32      1.1.1.1             2       100      0  ?
 *>  6.6.6.6/32      192.168.2.6          2        32768  ?
 *>i 192.168.1.0    1.1.1.1             0       100      0  ?
 *>  192.168.2.0    0.0.0.0             0        32768  ?
R3#
```

Which it is! 6.6.6.6 is now in the BGP table in VRF RED on R3 with a next hop of 192.168.2.6 (R6) and also 4.4.4 is in there as well with a next hop of 1.1.1.1 (which is the loopback of R1 – showing that it is going over the MPLS and R2 is not in the picture)  
The final step is to get the routes that have come across the MPLS back into OSPF and then we can get end to end connectivity.

R1

```
R1#conf t
Enter configuration commands, one per line. End with CNTL/Z.
R1(config)#router ospf 2
R1(config-router)#redistribute bgp 1 subnets
R1(config-router)#exit
R1(config)#exit
```

### R3

```
R3#conf t
Enter configuration commands, one per line. End with CNTL/Z.
R3(config)#router ospf 2
R3(config-router)#redistribute bgp 1 subnets
R3(config-router)#exit
R3(config)#exit
```

If all has worked we should be now able to ping 6.6.6.6 from R4

**Before we do let's see what the routing table looks like on R4.**

```
R4#sh ip route
Codes: L - local, C - connected, S - static, R - RIP, M - mobile, B - BGP
      D - EIGRP, EX - EIGRP external, O - OSPF, IA - OSPF inter area
      N1 - OSPF NSSA external type 1, N2 - OSPF NSSA external type 2
      E1 - OSPF external type 1, E2 - OSPF external type 2
      i - IS-IS, su - IS-IS summary, L1 - IS-IS level-1, L2 - IS-IS level-2
      ia - IS-IS inter area, * - candidate default, U - per-user static route
      o - ODR, P - periodic downloaded static route, H - NHRP, l - LIS
      + - replicated route, % - next hop override

Gateway of last resort is not set

        4.0.0.0/32 is subnetted, 1 subnets
C          4.4.4.4 is directly connected, Loopback0
        6.0.0.0/32 is subnetted, 1 subnets
O IA    6.6.6.6 [110/3] via 192.168.1.1, 00:01:35, FastEthernet0/0
              192.168.1.0/24 is variably subnetted, 2 subnets, 2 masks
C          192.168.1.0/24 is directly connected, FastEthernet0/0
L          192.168.1.4/32 is directly connected, FastEthernet0/0
O IA   192.168.2.0/24 [110/2] via 192.168.1.1, 00:01:35, FastEthernet0/0
R4#
```

Great we have 6.6.6.6 in there

**Also check the routing table on R5.**

```
R5#sh ip route
Codes: L - local, C - connected, S - static, R - RIP, M - mobile, B - BGP
      D - EIGRP, EX - EIGRP external, O - OSPF, IA - OSPF inter area
      N1 - OSPF NSSA external type 1, N2 - OSPF NSSA external type 2
      E1 - OSPF external type 1, E2 - OSPF external type 2
      i - IS-IS, su - IS-IS summary, L1 - IS-IS level-1, L2 - IS-IS level-2
      ia - IS-IS inter area, * - candidate default, U - per-user static route
      o - ODR, P - periodic downloaded static route, H - NHRP, l - LIS
      + - replicated route, % - next hop override

Gateway of last resort is not set

        4.0.0.0/32 is subnetted, 1 subnets
O IA    4.4.4.4 [110/3] via 192.168.2.1, 00:01:29, FastEthernet0/0
              6.0.0.0/32 is subnetted, 1 subnets
C          6.6.6.6 is directly connected, Loopback0
O IA   192.168.1.0/24 [110/2] via 192.168.2.1, 00:01:29, FastEthernet0/0
              192.168.2.0/24 is variably subnetted, 2 subnets, 2 masks
C          192.168.2.0/24 is directly connected, FastEthernet0/0
L          192.168.2.6/32 is directly connected, FastEthernet0/0
R5#
```

Brilliant we have 4.4.4.4 in there so we should be able to ping across the MPLS.

```
R4#ping 6.6.6.6
Type escape sequence to abort.
Sending 5, 100-byte ICMP Echos to 6.6.6.6, timeout is 2 seconds:
!!!!!
Success rate is 100 percent (5/5), round-trip min/avg/max = 332/387/428 ms
R4#
```

Which we can – to prove this is going over the MPLS and be label switched and not routed, lets do a trace.

```
R4#trace 6.6.6.6
Type escape sequence to abort.
Tracing the route to 6.6.6.6
VRF info: (vrf in name/id, vrf out name/id)
 1 192.168.1.1 148 msec 128 msec 136 msec
 2 10.0.0.2 [MPLS: Labels 17/19 Exp 0] 448 msec 584 msec 612 msec
 3 192.168.2.1 [MPLS: Label 19 Exp 0] 344 msec 324 msec 276 msec
 4 192.168.2.6 616 msec 672 msec 544 msec
R4#
```

# COMPUTER VISION



## JNAN VIKAS MANDAL'S

Mohanlal Raichand Mehta College of Commerce  
Diwali Maa College of Science  
Amritlal Raichand Mehta College of Arts  
Dr. R.T. Doshi College of Computer Science  
NAAC Re-Accredited Grade 'A+' (CGPA : 3.31) (3rd Cycle)

### DEPARTMENT OF INFORMATION TECHNOLOGY

#### CERTIFICATE

This is to certify that \_\_\_\_\_ bearing seat no. \_\_\_\_\_ has done the project work/journal work in the subject of Computer Vision of semester 2 Practical Examination during the academic year 2024-25 under the guidance of **Mrs.Vinaya Mangnale** being the partial requirement for the fulfilment of the curriculum of Degree in Master of Science in Information Technology under University of Mumbai.

Place: Airoli

Date:

Sign of Subject in- Charge

Sign of External Examiner

Sign of Coordinator

# INDEX

Practical No.	Title	Date	Sign
1	Perform Geometric Transformation		
2	Perform Image Stitching.		
3	Perform Camera Calibration.		
4	Perform the following: a. Face detection b. Object detection c. Pedestrian detection d. Face recognition		
5	Construct 3D model from images.		
6	Implement object detection and tracking from video.		
7	Perform Feature extraction using RANSAC		
8	Perform Colorization		
9	Perform Text detection and recognition		
10	Perform Image matting and composting		

## PRACTICAL NO.1

Aim- Perform geometric transformation.

CODE-

```
import cv2
import numpy as np
image = cv2.imread("C:/Users/Admin/Downloads/image1.jpg")

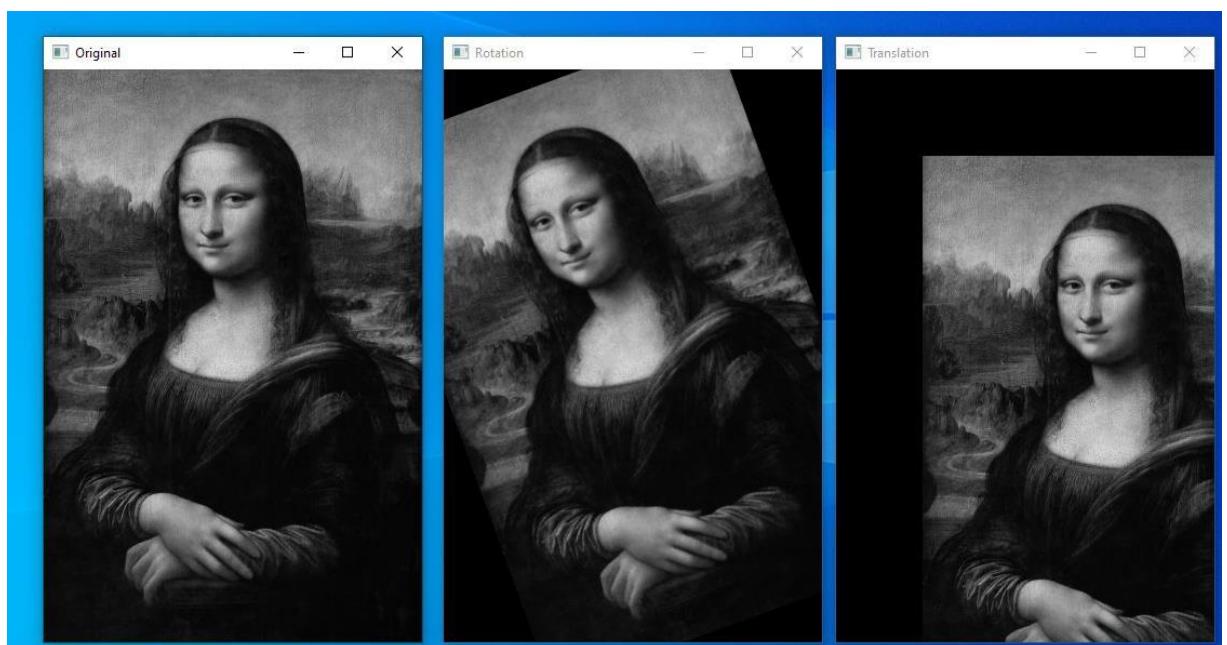
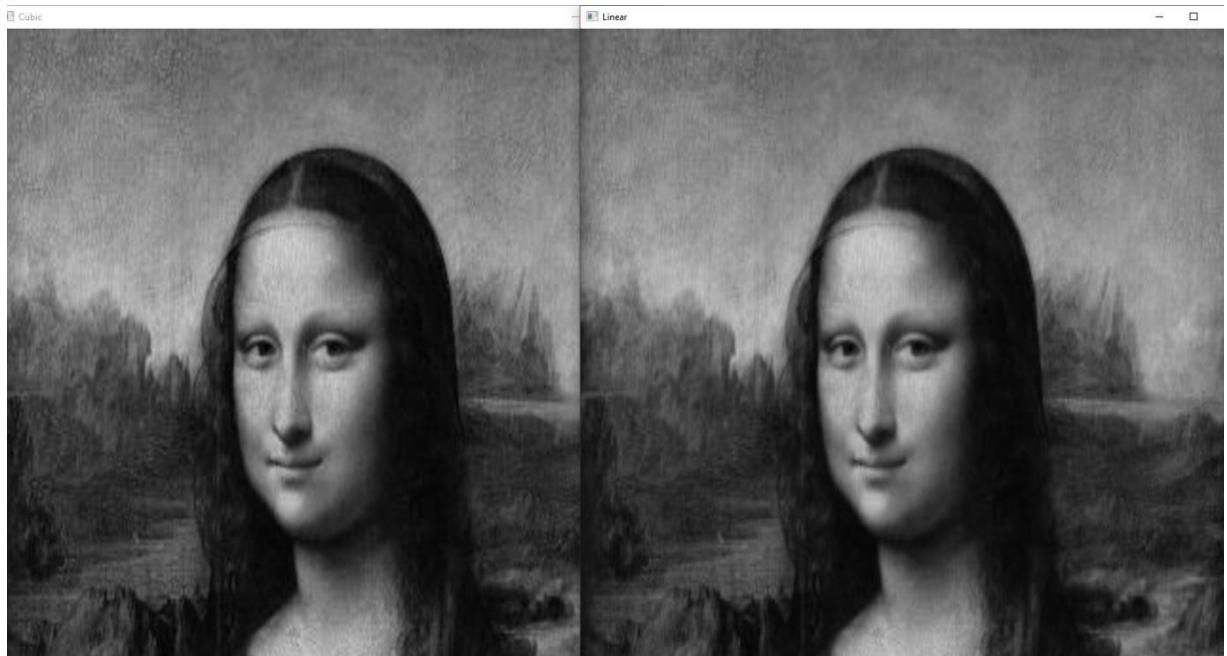
#perform linear and cubic
image_re_linear = cv2.resize(image, None, fx=2.5,fy=3.5, interpolation=cv2.INTER_LINEAR)
image_re_cubic = cv2.resize(image , None, fx=2.5,fy=3.5, interpolation=cv2.INTER_CUBIC)

# Perform the translation
matrix = np.float32([[1, 0, 80], [0, 1, 80]])
translated = cv2.warpAffine(image, matrix, (image.shape[1], image.shape[0]))

#Rotation
height, width = image.shape[:2]
matrix1 = cv2.getRotationMatrix2D((width/2,height/2) , 20 , 1)
Rotation = cv2.warpAffine(image, matrix1 , (width , height))

# Display the translated image
cv2.imshow('Rotation' , Rotation)
cv2.imshow('Translation', translated)
cv2.imshow('Cubic' , image_re_cubic)
cv2.imshow('Linear', image_re_linear)
cv2.imshow('Original' , image)
cv2.waitKey(100000)
cv2.destroyAllWindows()
```

## OUTPUT-



## PRACTICAL NO.2

Aim- Perform Image Stitching.

CODE-

```
# python -m pip install pillow
from PIL import Image

# Read the three images
image1 = Image.open('C:\\Users\\Admin\\Downloads\\shizu.jpg')
image2 = Image.open('C:\\Users\\Admin\\Downloads\\dora.png')
image3 = Image.open('C:\\Users\\Admin\\Downloads\\nobi.jpg')

# Show the original images
image1.show()
image2.show()
image3.show()

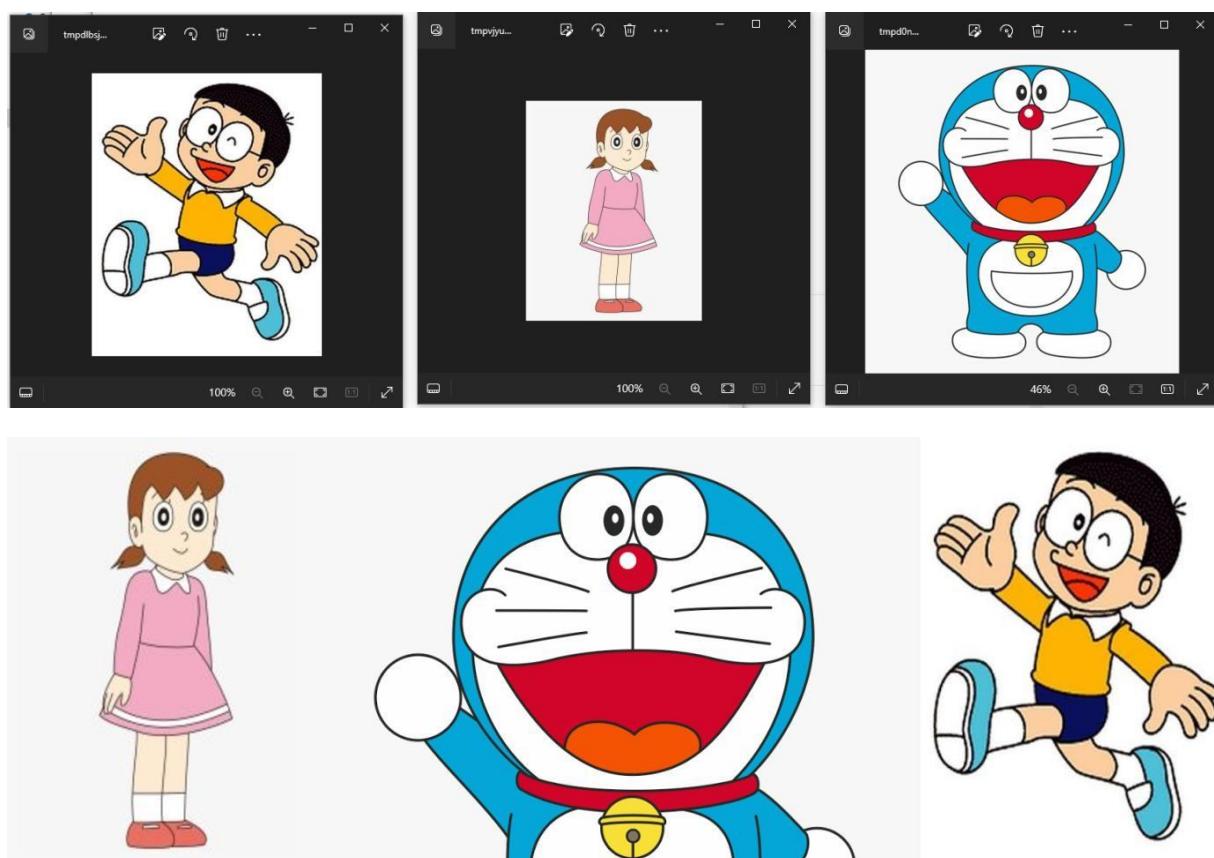
# Resize the first image
image1 = image1.resize((426, 600))
image1_size = image1.size

# Resize the third image to match the height of the first image
image3 = image3.resize((426, 600))

# Get the sizes of the images
image2_size = image2.size
image3_size = image3.size
```

```
# Create a new image with enough space for all three images side by side  
new_image = Image.new('RGB', (image1_size[0] + image2_size[0] + image3_size[0],  
image1_size[1]), (250, 250, 250))  
  
# Paste the images into the new image  
new_image.paste(image1, (0, 0))  
new_image.paste(image2, (image1_size[0], 0))  
new_image.paste(image3, (image1_size[0] + image2_size[0], 0))  
  
# Show the final combined image  
new_image.show()
```

OUTPUT-



## PRACTICAL NO.3

Aim- Perform Camera Calibration.

CODE-

```
# Import required modules
import cv2 #pip install opencv
import numpy as np #pip install numpy
import os
import glob

# Define the dimensions of the checkerboard
CHECKERBOARD = (6, 9)
# Termination criteria for corner refinement
criteria = (cv2.TERM_CRITERIA_EPS + cv2.TERM_CRITERIA_MAX_ITER, 30, 0.001)

# Vector for 3D points
threepoints = []
# Vector for 2D points
twodpoints = []

# 3D points real-world coordinates
objectp3d = np.zeros((1, CHECKERBOARD[0] * CHECKERBOARD[1], 3), np.float32)
objectp3d[0, :, :2] = np.mgrid[0:CHECKERBOARD[0],
0:CHECKERBOARD[1]].T.reshape(-1, 2)
prev_img_shape = None

# Specify the path to the folder containing checkerboard images
image_folder_path = r'C:\Users\Janhavi\Downloads\images'

# Extracting path of individual images stored in the specified directory
images = glob.glob(os.path.join(image_folder_path, '*.jpg'))

# Check if images were found in the specified path
if not images:
    print(f"No images found in the directory: {image_folder_path}")
else:
    for filename in images:
        image = cv2.imread(filename)
        grayColor = cv2.cvtColor(image, cv2.COLOR_BGR2GRAY)
        ret, corners = cv2.findChessboardCorners(grayColor, CHECKERBOARD,
        cv2.CALIB_CB_ADAPTIVE_THRESH + cv2.CALIB_CB_FAST_CHECK +
        cv2.CALIB_CB_NORMALIZE_IMAGE)
```

```

threepoints.append(objectp3d)
corners2 = cv2.cornerSubPix(grayColor, corners, (11, 11), (-1, -1), criteria)
twodpoints.append(corners2)

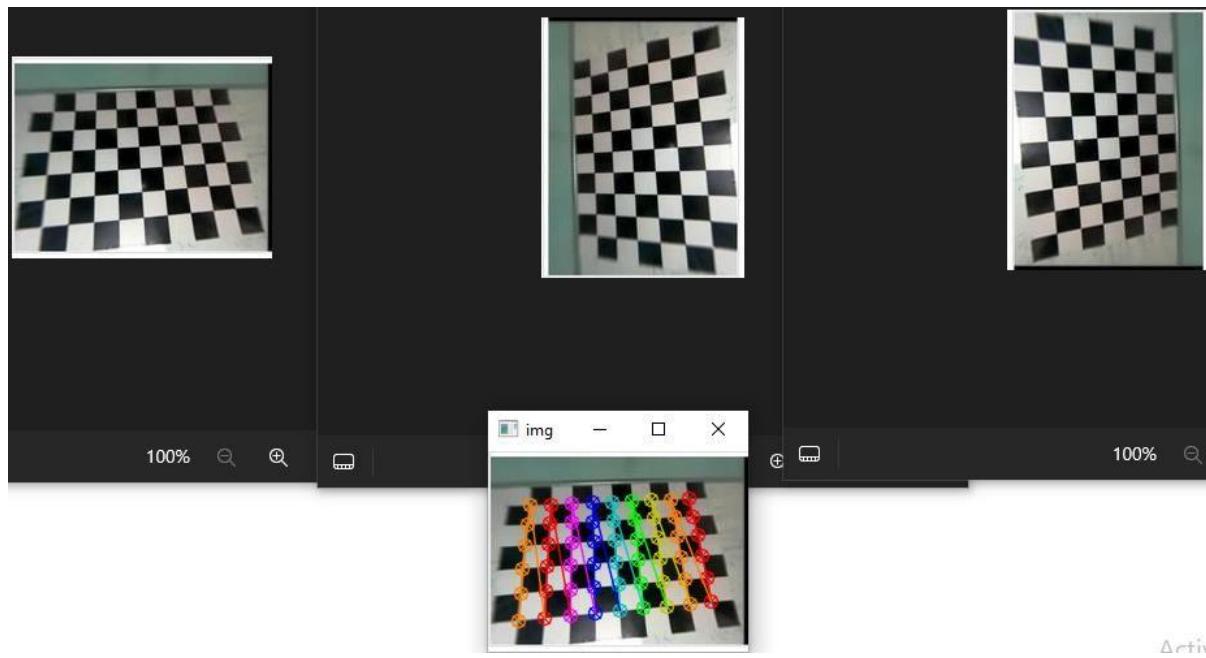
# Draw and display the corners
image = cv2.drawChessboardCorners(image, CHECKERBOARD, corners2, ret)
cv2.imshow('img', image)
cv2.waitKey(0)
cv2.destroyAllWindows()
h, w = image.shape[:2]

# Perform camera calibration
ret, matrix, distortion, r_vecs, t_vecs = cv2.calibrateCamera(
    threepoints, twodpoints, grayColor.shape[::-1], None, None)

# Displaying the calibration results
print("Camera matrix:")
print(matrix)
print("\nDistortion coefficient:")
print(distortion)
print("\nRotation Vectors:")
print(r_vecs)
print("\nTranslation Vectors:")
print(t_vecs)

```

## OUTPUT-



## PRACTICAL NO.4

Perform the following:

### Face detection

CODE-

```
from mtcnn import MTCNN
import cv2

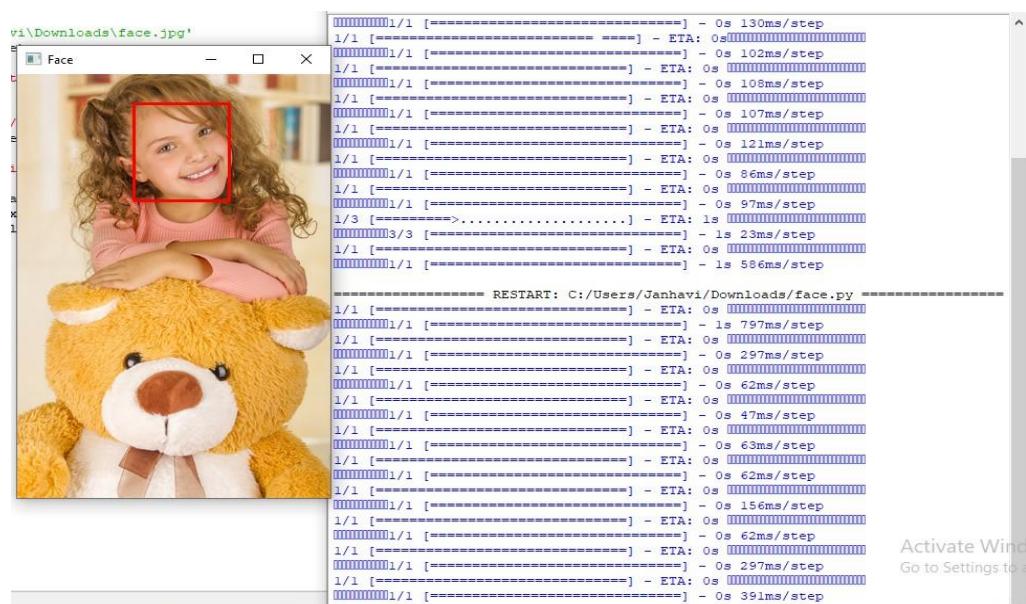
#load image from file //
filename = r'C:\Users\Janhavi\Downloads\face.jpg'
pixels = cv2.imread(filename)

# we need to create a detector by default weights //
detector = MTCNN()

# to find faces in an img //
faces = detector.detect_faces(pixels)

# display faces on the original image //
for face in faces:
    x, y, width, height = face['box']
    cv2.rectangle(pixels, (x, y), (x+width, y+height), (0, 0, 255), 2)
    cv2.imshow('Face', pixels)
    cv2.waitKey(0)
```

OUTPUT-



## Object detection

CODE-

```
import cv2
from matplotlib import pyplot as plt

# Paths
image_path = r"C:\Users\Janhavi\Downloads\stop.jpg"
cascade_path = 'stop_data.xml'

# Load image
img = cv2.imread(image_path)
if img is None:
    raise FileNotFoundError("Image file could not be loaded. Check the path.")

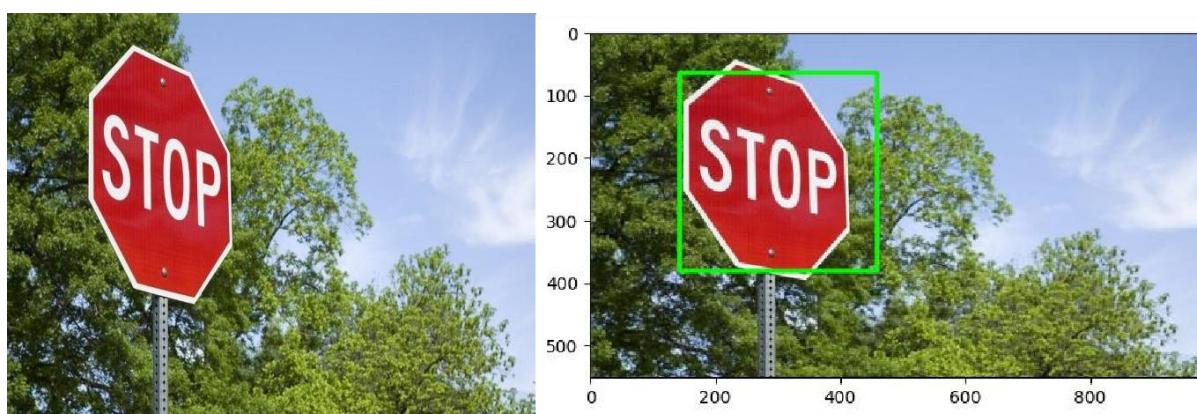
# Convert image
img_gray = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)
img_rgb = cv2.cvtColor(img, cv2.COLOR_BGR2RGB)

# Load cascade
stop_data = cv2.CascadeClassifier(cascade_path)
if stop_data.empty():
    raise FileNotFoundError("Cascade file could not be loaded. Check the path.")

# Detect and draw rectangles
for (x, y, w, h) in stop_data.detectMultiScale(img_gray, minSize=(20, 20)):
    cv2.rectangle(img_rgb, (x, y), (x + h, y + w), (0, 255, 0), 5)

# Display image
plt.imshow(img_rgb)
plt.show()
```

OUTPUT-



## Pedestrian detection

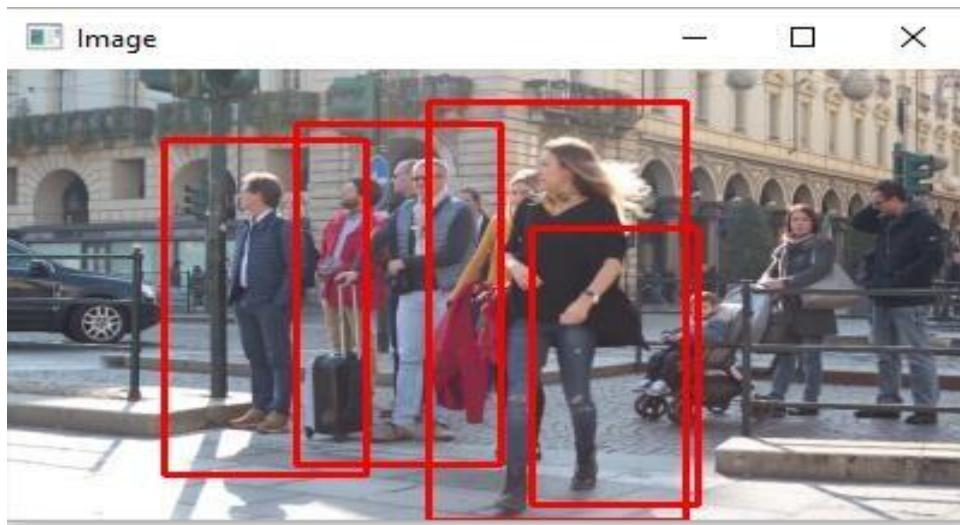
CODE-

```
# pip install imutils
import cv2
import imutils

# Initializing the HOG person detector
hog = cv2.HOGDescriptor()
hog.setSVMDetector(cv2.HOGDescriptor_getDefaultPeopleDetector())

# Video capture from file
cap = cv2.VideoCapture(r'C:\Users\Janhavi\Downloads\pd.mp4')
while cap.isOpened():
    ret, image = cap.read()
    if not ret:
        break
    image = imutils.resize(image, width=min(400, image.shape[1]))
    regions, _ = hog.detectMultiScale(image, winStride=(4, 4), padding=(4, 4), scale=1.05)
    for (x, y, w, h) in regions:
        cv2.rectangle(image, (x, y), (x + w, y + h), (0, 0, 255), 2)
    cv2.imshow("Image", image)
    if cv2.waitKey(25) & 0xFF == ord('q'):
        break
cap.release()
cv2.destroyAllWindows()
```

OUTPUT-



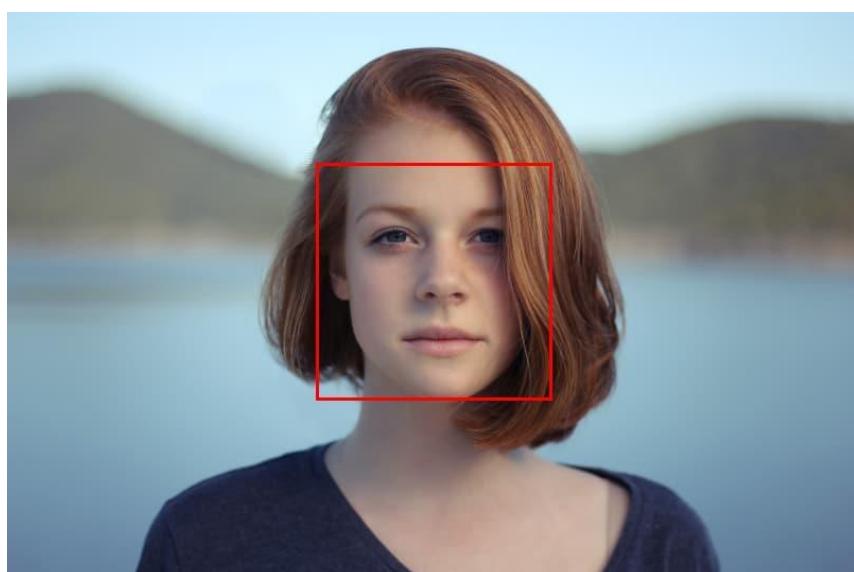
## Face Recognition

CODE-

```
# pip install imutils
import cv2
import imutils
hog = cv2.HOGDescriptor()
hog.setSVMDetector(cv2.HOGDescriptor_getDefaultPeopleDetector())
cap = cv2.VideoCapture(r'C:\Users\Janhavi\Downloads\pd.mp4')
while cap.isOpened():
    ret, image = cap.read()
    if not ret:
        break
    image = imutils.resize(image, width=min(400, image.shape[1]))
    regions, _ = hog.detectMultiScale(image, winStride=(4, 4), padding=(4, 4), scale=1.05)
    for (x, y, w, h) in regions:
        cv2.rectangle(image, (x, y), (x + w, y + h), (0, 0, 255), 2)
    cv2.imshow("Image", image)
    if cv2.waitKey(25) & 0xFF == ord('q'):
        break

# Release video capture and close all windows
cap.release()
cv2.destroyAllWindows()
```

OUTPUT-



## PRACTICAL NO.5

Aim- Construct 3D model from images.

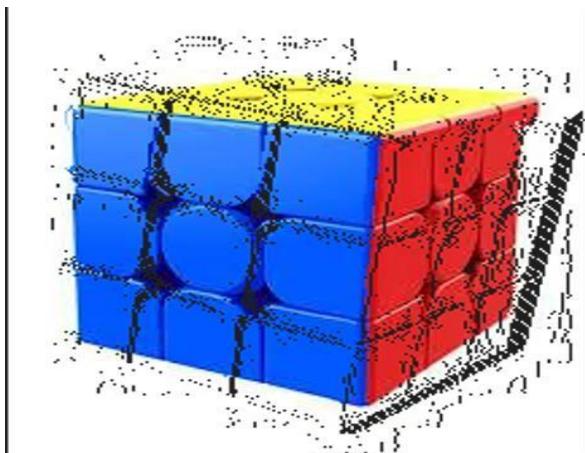
CODE-

```
from PIL import Image
import numpy as np

def shift_image(img, depth_img, shift_amount=10):
    # Ensure base image has alpha
    img = img.convert("RGBA")
    data = np.array(img)
    depth_img = depth_img.convert("L")
    depth_data = np.array(depth_img)
    deltas = ((depth_data / 255.0) * float(shift_amount)).astype(int)
    shifted_data = np.zeros_like(data)
    height, width, _ = data.shape
    for y, row in enumerate(deltas):
        for x, dx in enumerate(row):
            if x + dx < width and x + dx >= 0:
                shifted_data[y, x + dx] = data[y, x]
    shifted_image = Image.fromarray(shifted_data.astype(np.uint8))
    return shifted_image

img = Image.open(r"C:\Users\Janhavi\Downloads\cube1.jpeg")
depth_img = Image.open(r"C:\Users\Janhavi\Downloads\cube3.jpeg")
shifted_img = shift_image(img, depth_img, shift_amount=10)
shifted_img.show()
```

OUTPUT-



## PRACTICAL NO.6

Aim- Implement object detection and tracking from video.

CODE-

```
import cv2
import numpy as np
import imutils

# Video capture from file
cap = cv2.VideoCapture(r'C:\Users\Janhavi\Downloads\football2.mp4')

if not cap.isOpened():
    print("Error: Could not open video.")
    exit()

object_detected = False
object_bbox = None

def find_initial_football_position(frame):
    hsv = cv2.cvtColor(frame, cv2.COLOR_BGR2HSV)
    lower_white = np.array([0, 0, 150])
    upper_white = np.array([180, 25, 255])
    mask = cv2.inRange(hsv, lower_white, upper_white)
    mask = cv2.erode(mask, None, iterations=2)
    mask = cv2.dilate(mask, None, iterations=2)
    contours, _ = cv2.findContours(mask.copy(), cv2.RETR_EXTERNAL,
                                   cv2.CHAIN_APPROX_SIMPLE)
    if contours:
        largest_contour = max(contours, key=cv2.contourArea)
        (x, y, w, h) = cv2.boundingRect(largest_contour)
        return (x, y, w, h)
    else:
        return None

ret, frame = cap.read()
if ret:
    frame = imutils.resize(frame, width=800)
    object_bbox = find_initial_football_position(frame)
    if object_bbox:
        object_detected = True

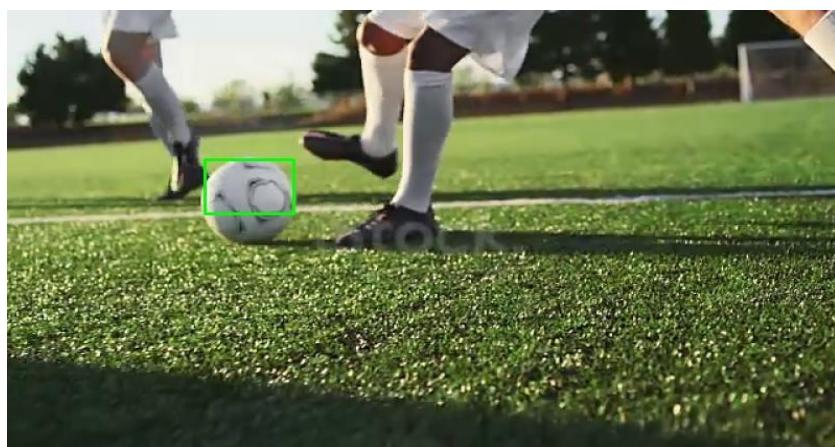
while cap.isOpened():
    ret, frame = cap.read()
```

```

if not ret:
    break
frame = imutils.resize(frame, width=800)
hsv = cv2.cvtColor(frame, cv2.COLOR_BGR2HSV)
lower_white = np.array([0, 0, 150])
upper_white = np.array([180, 25, 255])
mask = cv2.inRange(hsv, lower_white, upper_white)
mask = cv2.erode(mask, None, iterations=2)
mask = cv2.dilate(mask, None, iterations=2)
contours, _ = cv2.findContours(mask.copy(), cv2.RETR_EXTERNAL,
                               cv2.CHAIN_APPROX_SIMPLE)
object_detected = False
object_bbox = None
for contour in contours:
    # Compute bounding box and area
    (x, y, w, h) = cv2.boundingRect(contour)
    area = cv2.contourArea(contour)
    if area > 1000 and w > 30 and h > 30:
        # Update object bbox
        object_bbox = (x, y, w, h)
        object_detected = True
        break
if object_detected:
    (x, y, w, h) = object_bbox
    cv2.rectangle(frame, (x, y), (x + w, y + h), (0, 255, 0), 2)
cv2.imshow("Football Detection", frame)
if cv2.waitKey(25) & 0xFF == ord('q'):
    break
cap.release()
cv2.destroyAllWindows()

```

## OUTPUT-



## PRACTICAL NO.7

Aim-Perform Feature extraction using RANSAC

CODE-

```
import cv2
import numpy as np
img1 = cv2.imread(r'C:\Users\Janhavi\Downloads\milk.jpg', 0)
img2 = cv2.imread(r'C:\Users\Janhavi\Downloads\mocha.png', 0)

orb = cv2.ORB_create()

keypoints1, descriptors1 = orb.detectAndCompute(img1, None)
keypoints2, descriptors2 = orb.detectAndCompute(img2, None)

bf = cv2.BFMatcher(cv2.NORM_HAMMING, crossCheck=True)
matches = bf.match(descriptors1, descriptors2)
matches = sorted(matches, key=lambda x: x.distance)

print("Number of keypoints in image 1:", len(keypoints1))
print("Number of keypoints in image 2:", len(keypoints2))
print("Number of matches found:", len(matches))

src_pts = np.float32([keypoints1[m.queryIdx].pt for m in matches]).reshape(-1, 1, 2)
dst_pts = np.float32([keypoints2[m.trainIdx].pt for m in matches]).reshape(-1, 1, 2)

# Compute the homography matrix using RANSAC
H, mask = cv2.findHomography(src_pts, dst_pts, cv2.RANSAC, 5.0)
print("Homography Matrix (H):")
print(H)
```

OUTPUT-

```
Python 3.8.3 Shell
File Edit Shell Debug Options Window Help
Python 3.8.3 (tags/v3.8.3:6f8c832, May 13 2020, 22:37:02) [MSC v.1924 64 bit (AMD64)] on win32
Type "help", "copyright", "credits" or "license()" for more information.
>>>
===== RESTART: C:/Users/Janhavi/Downloads/feature-ext.py =====
Number of keypoints in image 1: 500
Number of keypoints in image 2: 500
Number of matches found: 84
Homography Matrix (H):
[[ -1.06165354e+00 -5.98587766e-02 2.20239084e+02]
 [ -1.16448342e+00 -8.25116005e-02 2.45538196e+02]
 [ -4.77827955e-03 -3.02996075e-04 1.00000000e+00]]
```



## PRACTICAL NO.8

Aim-Perform Colorization

CODE-

```
import numpy as np
import cv2
from cv2 import dnn

# Model file paths
proto_file = r'C:\Users\Admin\Downloads\Model\colorization_deploy_v2.prototxt'
model_file = r'C:\Users\Admin\Downloads\Model\colorization_release_v2.caffemodel'
hull_pts = r'C:\Users\Admin\Downloads\Model\pts_in_hull.npy'

img_path = r'C:\Users\Admin\Downloads\flower.jpg'

img = cv2.imread(img_path)
if img is None:
    print(f'Error: Unable to load image at {img_path}')
    exit()
scaled = img.astype("float32") / 255.0
lab_img = cv2.cvtColor(scaled, cv2.COLOR_BGR2LAB)

net = dnn.readNetFromCaffe(proto_file, model_file)
kernel = np.load(hull_pts)

class8 = net.getLayerId("class8_ab")
conv8 = net.getLayerId("conv8_313_rh")
pts = kernel.transpose().reshape(2, 313, 1, 1)
net.getLayer(class8).blobs = [pts.astype("float32")]
net.getLayer(conv8).blobs = [np.full([1, 313], 2.606, dtype="float32")]

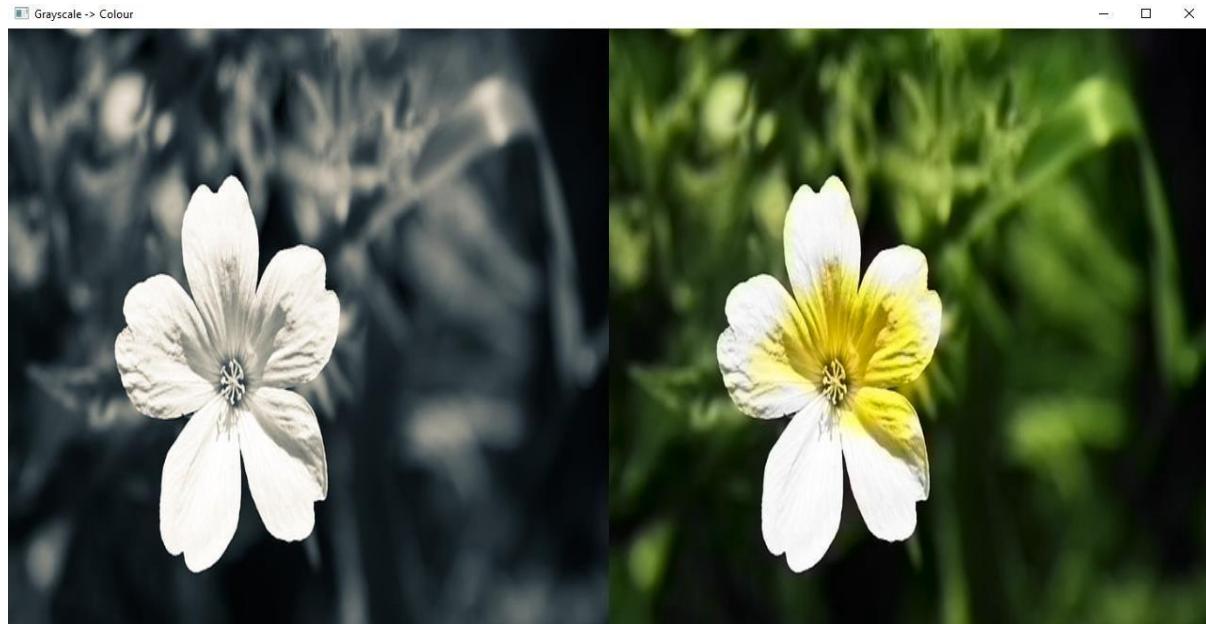
resized = cv2.resize(lab_img, (224, 224))
L = cv2.split(resized)[0]
L -= 50

net.setInput(cv2.dnn.blobFromImage(L))
ab_channel = net.forward()[0, :, :, :].transpose((1, 2, 0))
ab_channel = cv2.resize(ab_channel, (img.shape[1], img.shape[0]))

# Convert to colorized image in BGR
L = cv2.split(lab_img)[0]
```

```
colorized = cv2.cvtColor(np.concatenate((L[:, :, np.newaxis], ab_channel), axis=2),  
cv2.COLOR_LAB2BGR)  
colorized = np.clip((255 * colorized).astype("uint8"), 0, 255)  
  
# Resize images for display  
img_display = cv2.resize(img, (640, 640))  
colorized_display = cv2.resize(colorized, (640, 640))  
  
# Concatenate and display images  
result = np.hstack([img_display, colorized_display])  
cv2.imshow("Grayscale -> Colour", result)  
cv2.waitKey(0)  
cv2.destroyAllWindows()
```

## OUTPUT-



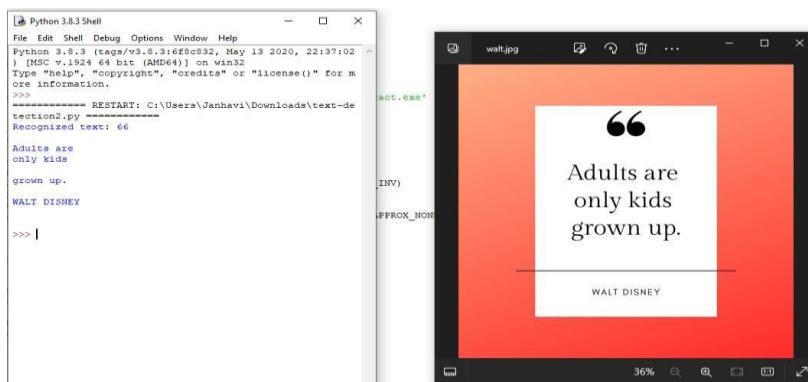
## PRACTICAL NO.9

Aim-Perform Text detection and recognition

CODE-

```
#pip install pytesseract
#Download the tesseract executable file from https://sourceforge.net/projects/tesseract-ocr-
alt/
import cv2
import pytesseract
import os
os.environ['TESSDATA_PREFIX'] = r'C:\Program Files\Tesseract-OCR'
pytesseract.pytesseract.tesseract_cmd = r'C:\Program Files\Tesseract-OCR\tesseract.exe'
image_path = r'C:\Users\Janhavi\Downloads\walt.jpg'
img = cv2.imread(image_path)
gray = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)
ret, thresh1 = cv2.threshold(gray, 0, 255, cv2.THRESH_OTSU |
cv2.THRESH_BINARY_INV)
rect_kernel = cv2.getStructuringElement(cv2.MORPH_RECT, (18, 18))
dilation = cv2.dilate(thresh1, rect_kernel, iterations=1)
contours, hierarchy = cv2.findContours(dilation, cv2.RETR_EXTERNAL,
cv2.CHAIN_APPROX_NONE)
im2 = img.copy()
output_path = r'C:\Users\Janhavi\Downloads\recognized.txt'
with open(output_path, "w+") as file:
    for cnt in contours:
        x, y, w, h = cv2.boundingRect(cnt)
        cropped = im2[y:y + h, x:x + w]
        text = pytesseract.image_to_string(cropped)
        print(f'Recognized text: {text}')
        file.write(text + "\n")
```

OUTPUT-



## PRACTICAL NO.10

Aim-Perform Image matting and compositing

### Image matting

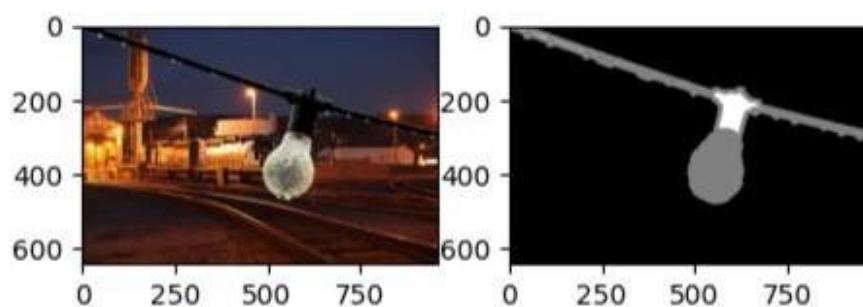
CODE-

```
# Import necessary libraries
import matplotlib.pyplot as plt
from PIL import Image
import requests

# Load the image and trimap
url = "https://github.com/hustvl/ViTMatte/blob/main/demo/bulb_rgb.png?raw=true"
image = Image.open(requests.get(url, stream=True).raw).convert("RGB")
url = "https://github.com/hustvl/ViTMatte/blob/main/demo/bulb_trimap.png?raw=true"
trimap = Image.open(requests.get(url, stream=True).raw)

# Display the image and trimap
plt.figure(figsize=(15, 15))
plt.subplot(1, 2, 1)
plt.imshow(image)
plt.subplot(1, 2, 2)
plt.imshow(trimap)
plt.show()
```

OUTPUT-



## Image compositing

### CODE-

```
from PIL import Image

# Load images and convert to mode 'L'
im1 = Image.open(r'C:\Users\Admin\Downloads\im1.png').convert('L')
im2 = Image.open(r'C:\Users\Admin\Downloads\im2.webp').convert('L')
mask = Image.open(r'C:\Users\Admin\Downloads\im3.webp').convert('L')

# Ensure all images are the same size
size = im1.size
im2 = im2.resize(size)
mask = mask.resize(size)

# Composite the images
im3 = Image.composite(im1, im2, mask)

# Show the result
im3.show()
```

### OUTPUT-

